

# A General-purpose Distributed and Parallel Programming System at Scale

Dr. Tsung-Wei Huang

Department of Electrical and Computer Engineering  
University of Illinois at Urbana-Champaign, IL, USA



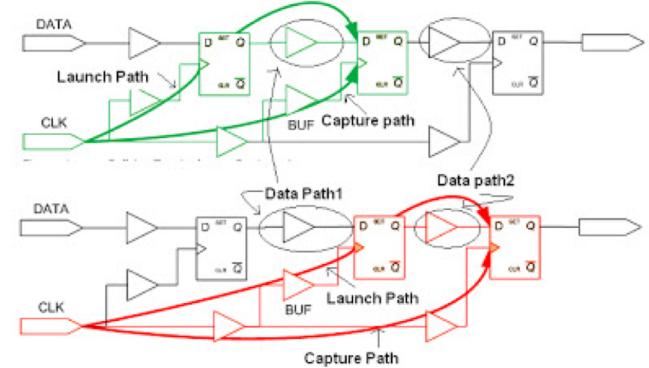
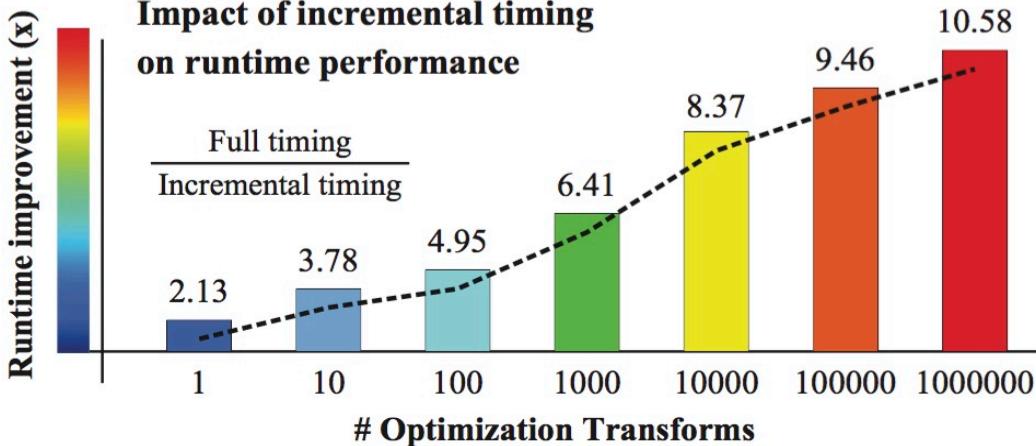
# Outline

---

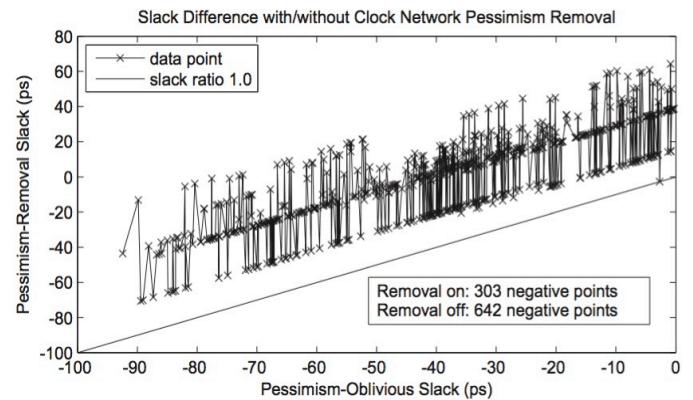
- My early PhD research
  - OpenTimer: A VLSI timing analysis tool
  - Express parallelism in the right way
- A general-purpose distributed programming system
  - DtCraft and its ecosystem
  - Machine learning, graph algorithms, chip design
  - Cpp-Taskflow: Modern C++ parallel task programming
- Conclusion and future work

# My Early PhD Research

- Static timing analysis (STA)
  - Key component in VLSI design
  - Verify the circuit timing
- Many challenges
  - Incremental timing
  - Parallelization, scalability



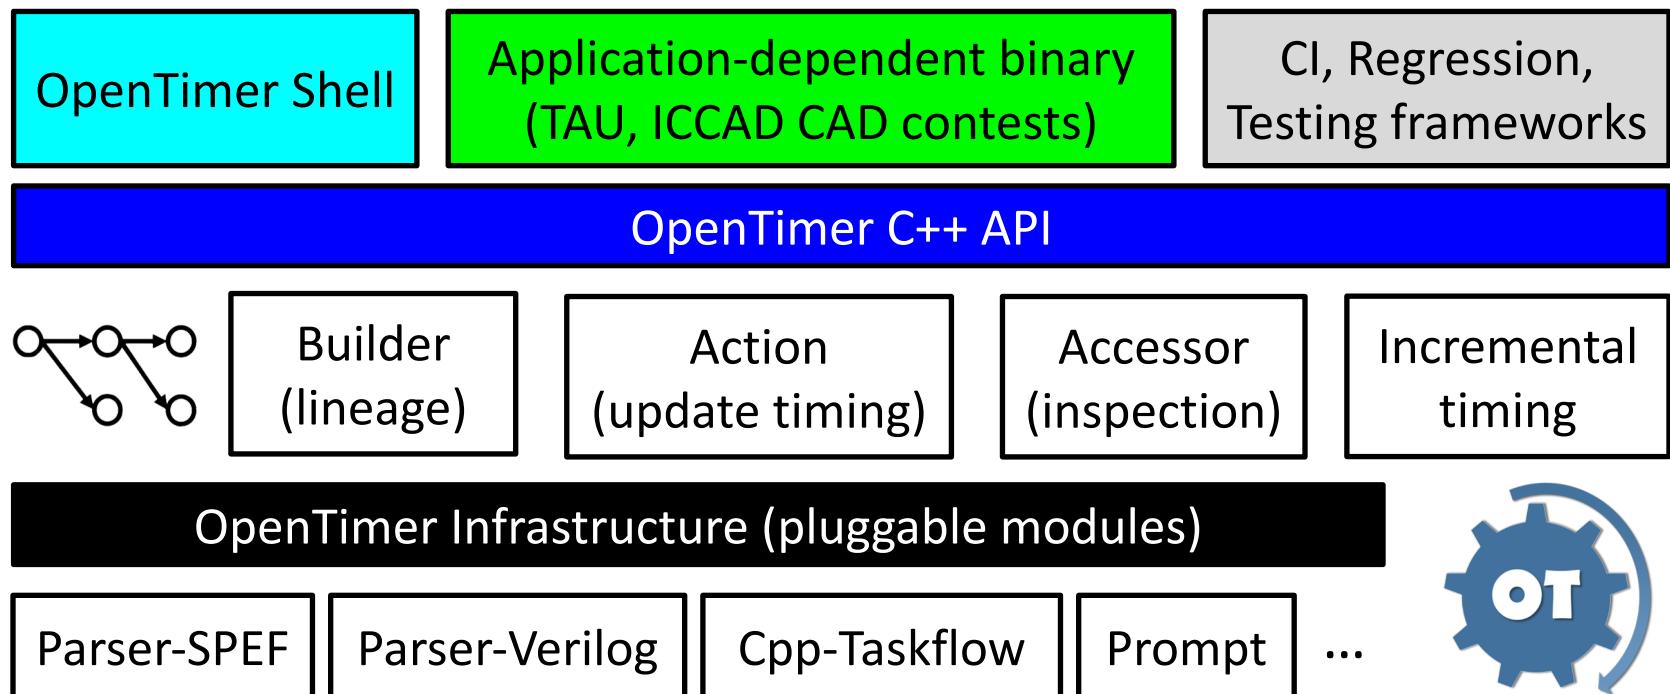
Setup/hold timing checks



Path-based timing analysis to reduce pessimism

# OpenTimer: An Open-source STA Tool

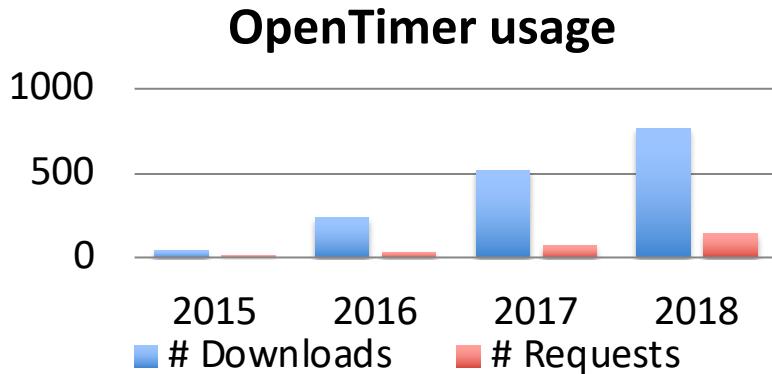
- TAU14 (1<sup>st</sup> place), TAU15 (2<sup>nd</sup> place), TAU16 (1<sup>st</sup> place)
- Best tool awards (WOSET ICCAD18, LibreCore16)
- Golden timers in VLSI and CAD communities



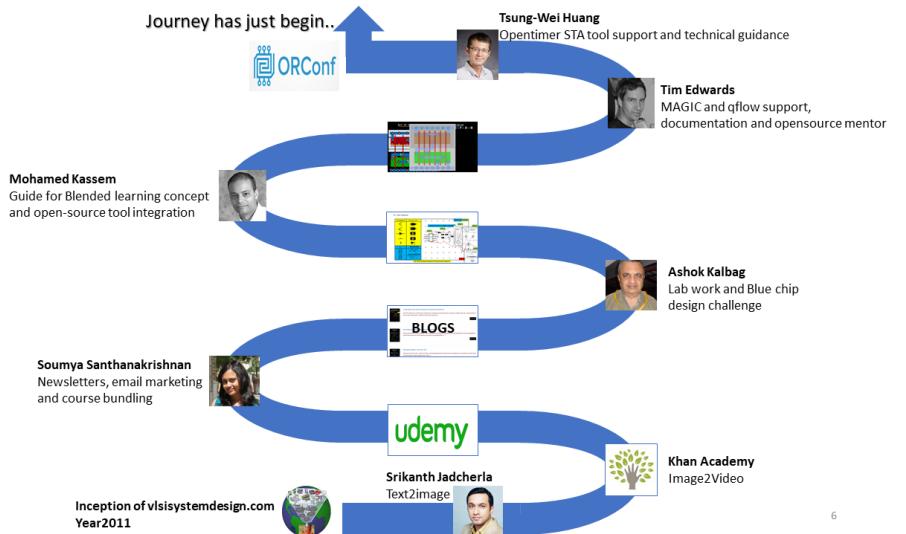
# OpenTimer Community – Thank You!



OpenTimer is begin used in many projects



We are growing! > 500 downloads in 2018



Collaboration with VSD (start-up) and Udemy online education platform



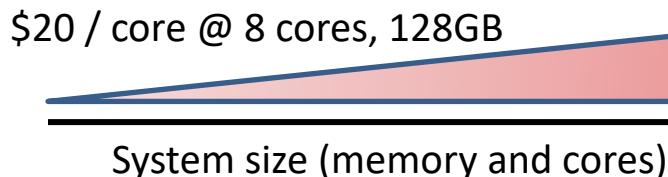
Acknowledged by ACM TAU16-19 Contests

# Collaboration with IBM

- “*We need a new **distributed** timing analysis method to deal with the large design complexity,*” IBM Timing Group, Fishkill, NY, 2015

- Existing tools are architecturally limited by one machine
- Can take **800GB RAM** and **14 days** on a single machine
- Maintaining “bare-metal” machines is not cost-efficient

Hourly Bare Metal		Monthly Bare Metal	
Single Processor	Dual Processor	Quad Processor	
Starting at \$159/mo	Starting at \$379/mo	Starting at \$1,299/mo	
2.66GHz to 2.9GHz	2.0GHz to 2.9GHz	2.0GHz to 2.7GHz	
Up to 8 cores	Up to 16 cores	Up to 40 cores	
Up to 128GB RAM	Up to 256GB RAM	Up to 512GB RAM	
Up to 16TB local storage	Up to 144TB local storage	Up to 96TB local storage	
20TB public outbound	20TB public outbound	20TB public outbound	
<a href="#">Order Now</a>	<a href="#">Order Now</a>	<a href="#">Order Now</a>	



*Objective: Execute the same # computations, distributed across an array of less expensive per unit cost cloud systems*

Cost per unit of computation

*Leverage economy of scale by decreasing the per unit cost of computation “hill”*

# Existing Tools in Cloud Computing

---

- **Hadoop**

- Distributed MapReduce platform on HDFS



- **Zookeeper**

- Coordination service for distributed application



- **Mesos**

- A high-performance cluster manager



- **Spark**

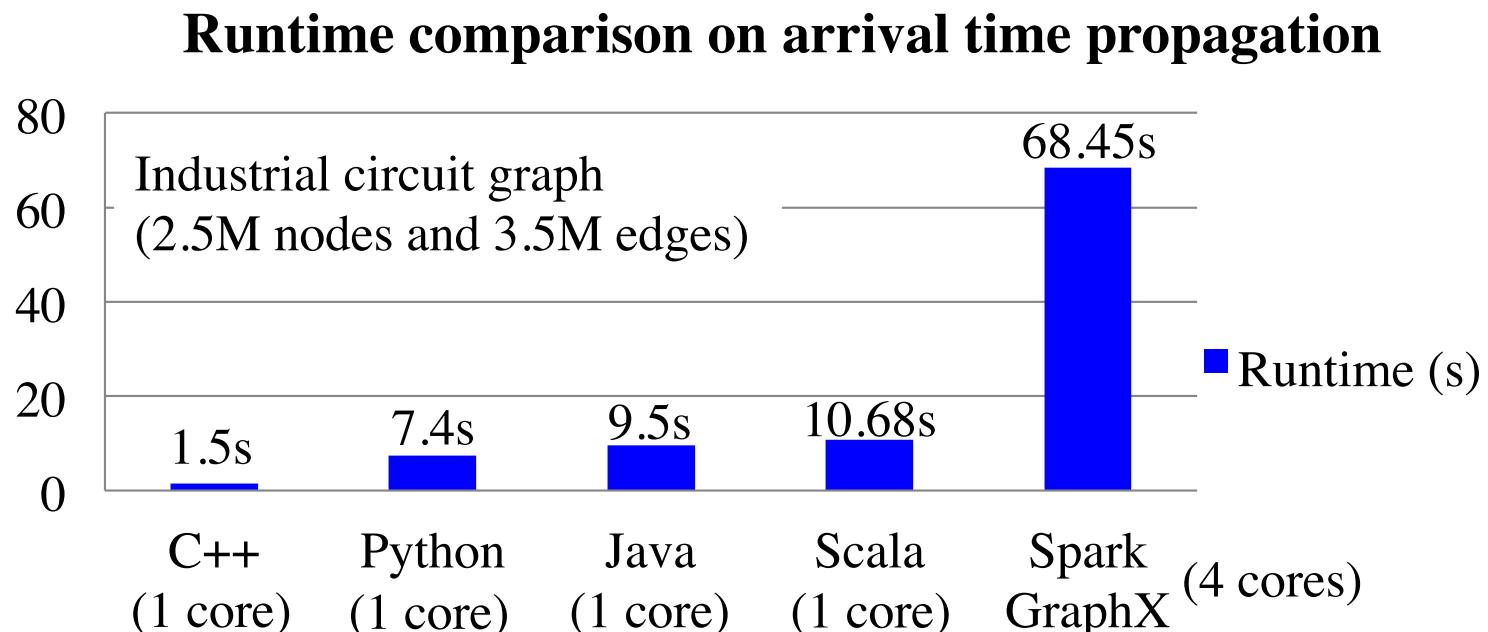
- A general computing engine for big-data analytics



- ...

***Are these tools suitable for our application?***

# Big-data Tools is NOT an Easy Fit ...

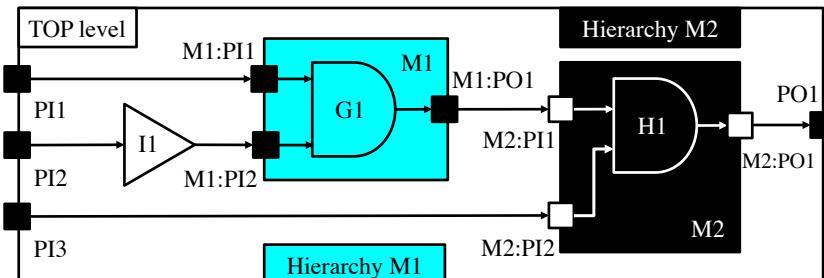


Method	Spark 2.0 (RDD + GraphX Pregel)	Java (SP)	C++ (SP)
Runtime (s)	68.45	9.5	1.50

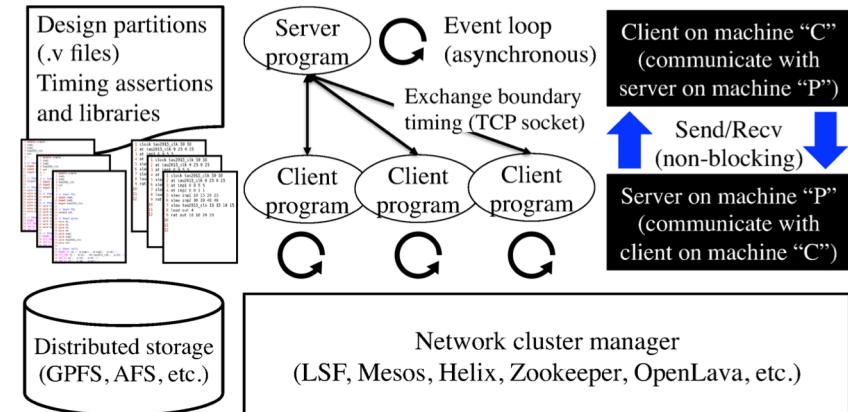
*MapReduce overhead*      *Language overhead*

# If it is Hard, let me Hard-code it ...

- A “specialized” distributed timer
  - Low-level Linux sockets and message passing code
  - Explicit resource partition and mapping
- Single-server multiple-client model
  - Server coordinates with clients to exchange data



Three partitions, top-level, M1, and M2  
(given by design teams)



T.-W. Huang, et al, “A Distributed Timing Analysis Framework for Large Designs,” IEEE/ACM DAC16  
T.-W. Huang, et al, US Patent US20170242945A1, assignee IBM, 2017  
T.-W. Huang, et al, US Patent US9836572B2, assignee IBM, 2017

# Observations

---

- Existing big data tools have large room to improve
- Fast prototype but hard long-term maintenance
- Non-unified tooling environment
- Cumbersome software dependencies
- Cannot afford hard-coded solution all the time
- Difficult low-level concurrency controls
- Hand-written message passing code
- ...

***We want a general-purpose programming system to streamline the building of parallel/distributed systems!***

# DARPA's Electronic Resurgence Initiative

- ❑ \$1.5B investment to reduce the design complexity
  - ❑ People without chip design experience can design chips



Chip layout army



Massive cloud computing!



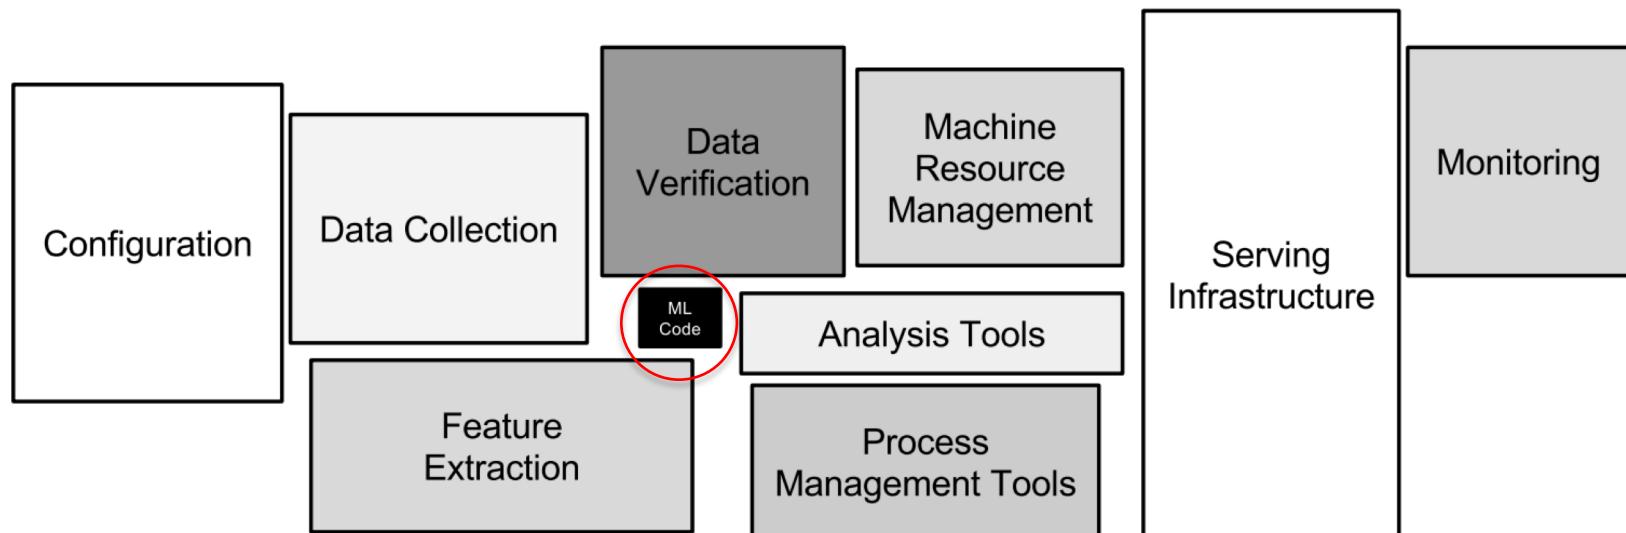
No-touch foundry

- ❑ A general-purpose silicon compiler (IDEA & POSH)
  - ❑ fully-automated layout generator, no human in the loop
  - ❑ Massive software efforts and tool integrations
  - ❑ *Parallel/distributed computing and machine learning\**

\*DARPA IDEA/POSH Kick-off Meeting, Boston, June, 2018

# Hidden Technical Debts in ML Systems

- ❑ Only a **TINY** fractions of ML systems is ML code
  - ❑ Writing ML code is cheap; system maintenance is\$\$\$\$
- ❑ ML innovations are moving to “**system innovations**”
  - ❑ Spark MLflow, TensorFlow 2.0, PyTorch, MindsDB
  - ❑ *Parallel/distributed computing, transparency, scalability*



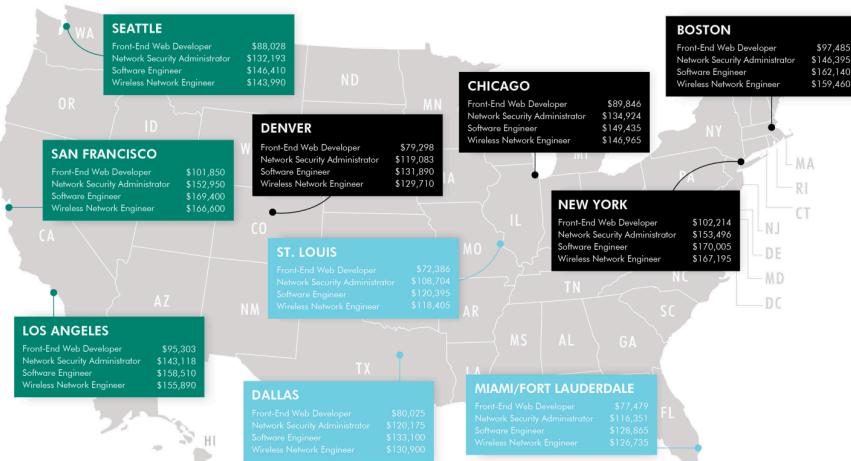
# Keep Programmability in Mind

- In the cloud era ...
  - Hardware is just a commodity
  - Building a cluster is cheap
  - Coding takes people and time



## TECH SALARIES ACROSS THE U.S.

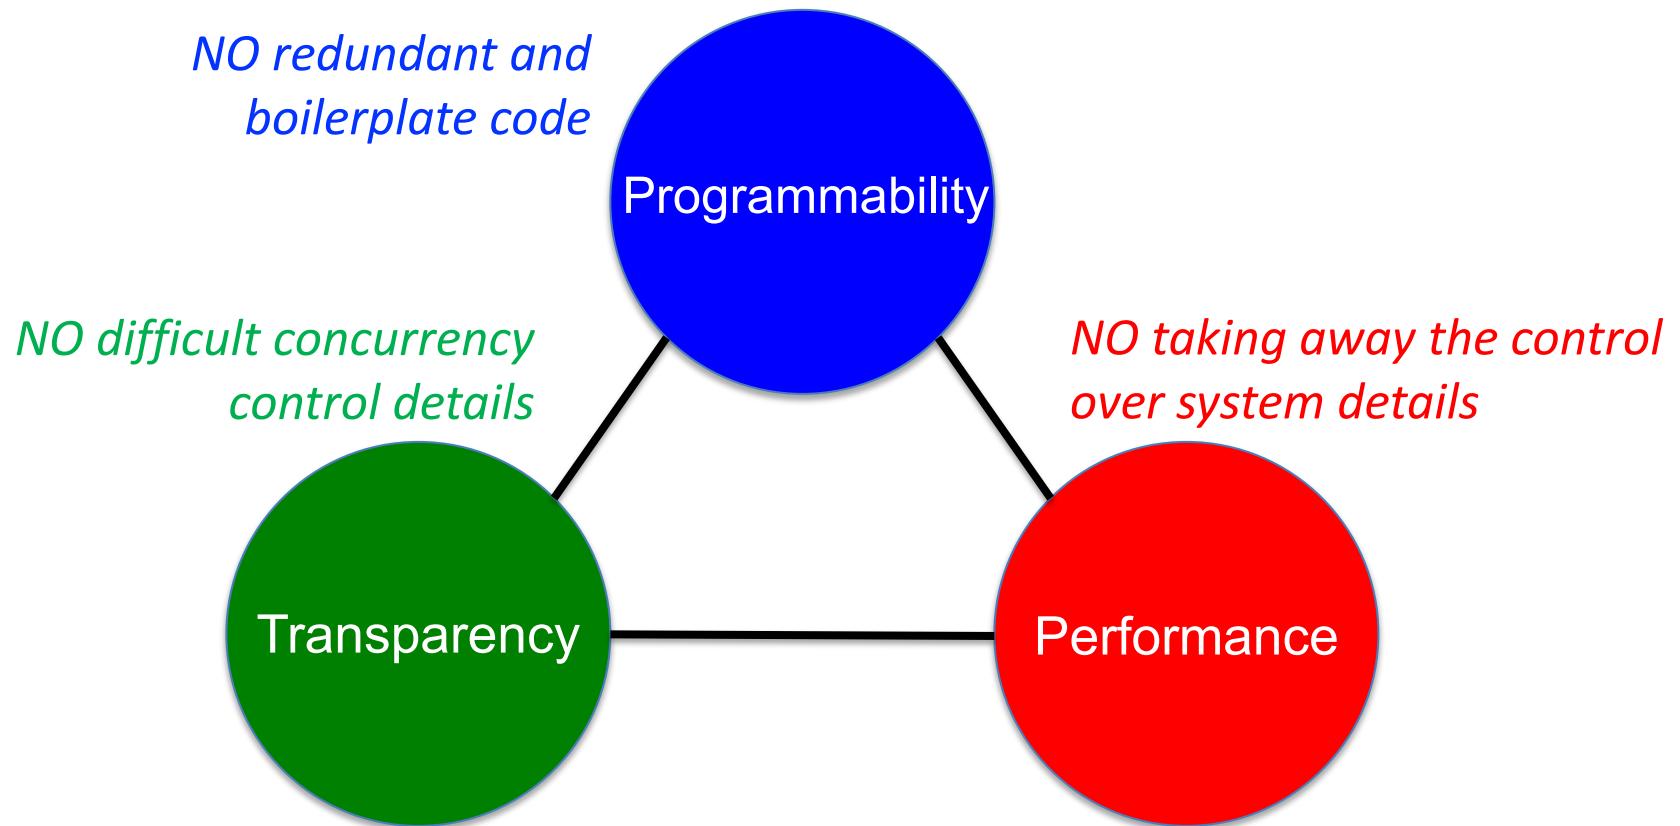
Compare midpoint starting salaries for technology positions in different cities.\*



***Programmability can affect the “performance” and “productivity” in all aspects (details, styles, high-level decisions, etc)!***

2018 Avg Software Engineer salary (NY) > \$170K

# Want a Unified Programming System



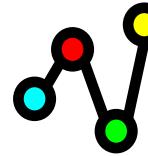
*“We want to let users easily express their parallel/distributed computing workload without taking away the control over system details to achieve high performance, using our expressive API written in modern C++”*

# Outline

---

- ❑ My early PhD research
  - ❑ OpenTimer: A VLSI timing analysis tool
  - ❑ Express parallelism in the right way
- ❑ A general-purpose distributed programming system
  - ❑ DtCraft and its ecosystem
  - ❑ Machine learning, graph algorithms, chip design
  - ❑ Cpp-Taskflow: Modern C++ parallel task programming
- ❑ Conclusion and future work

# A New Solution: DtCraft



Application pipeline (domain-specific applications)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



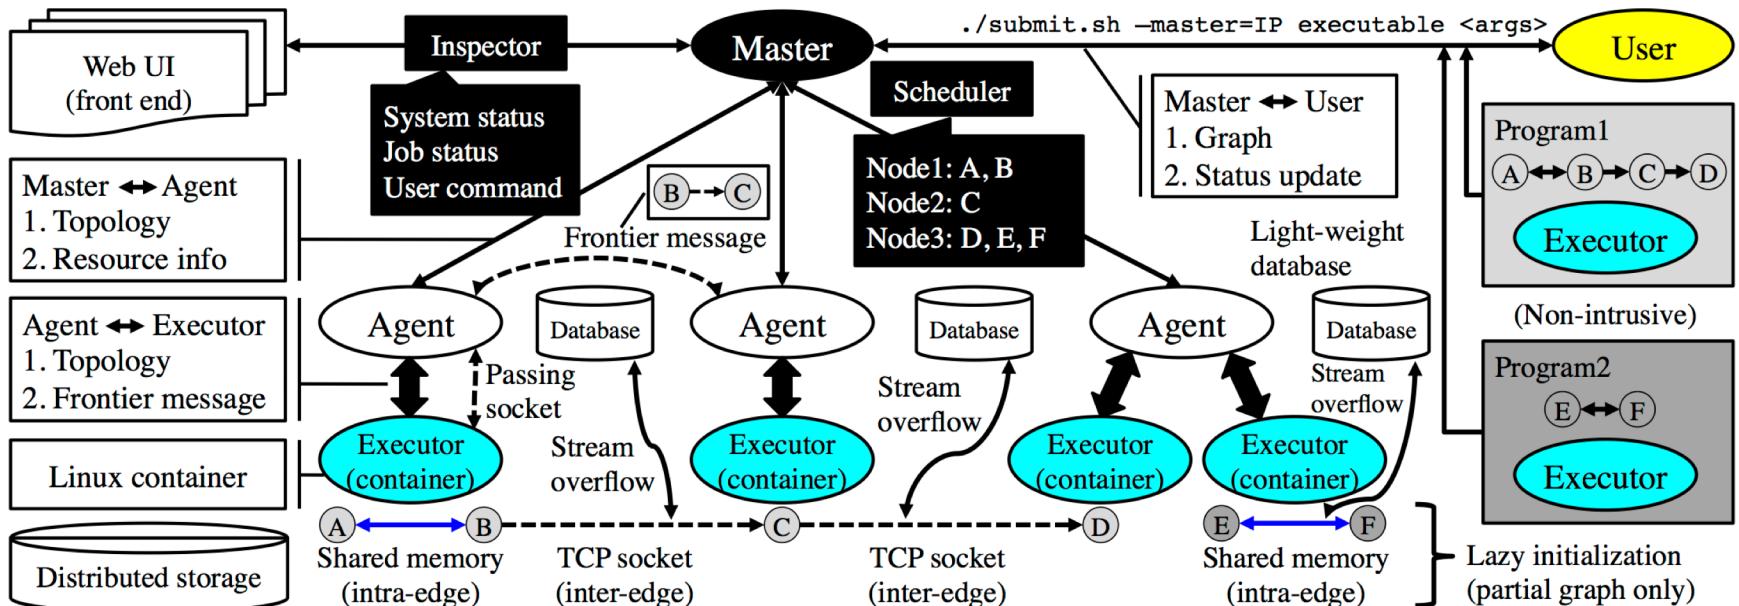
...



*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

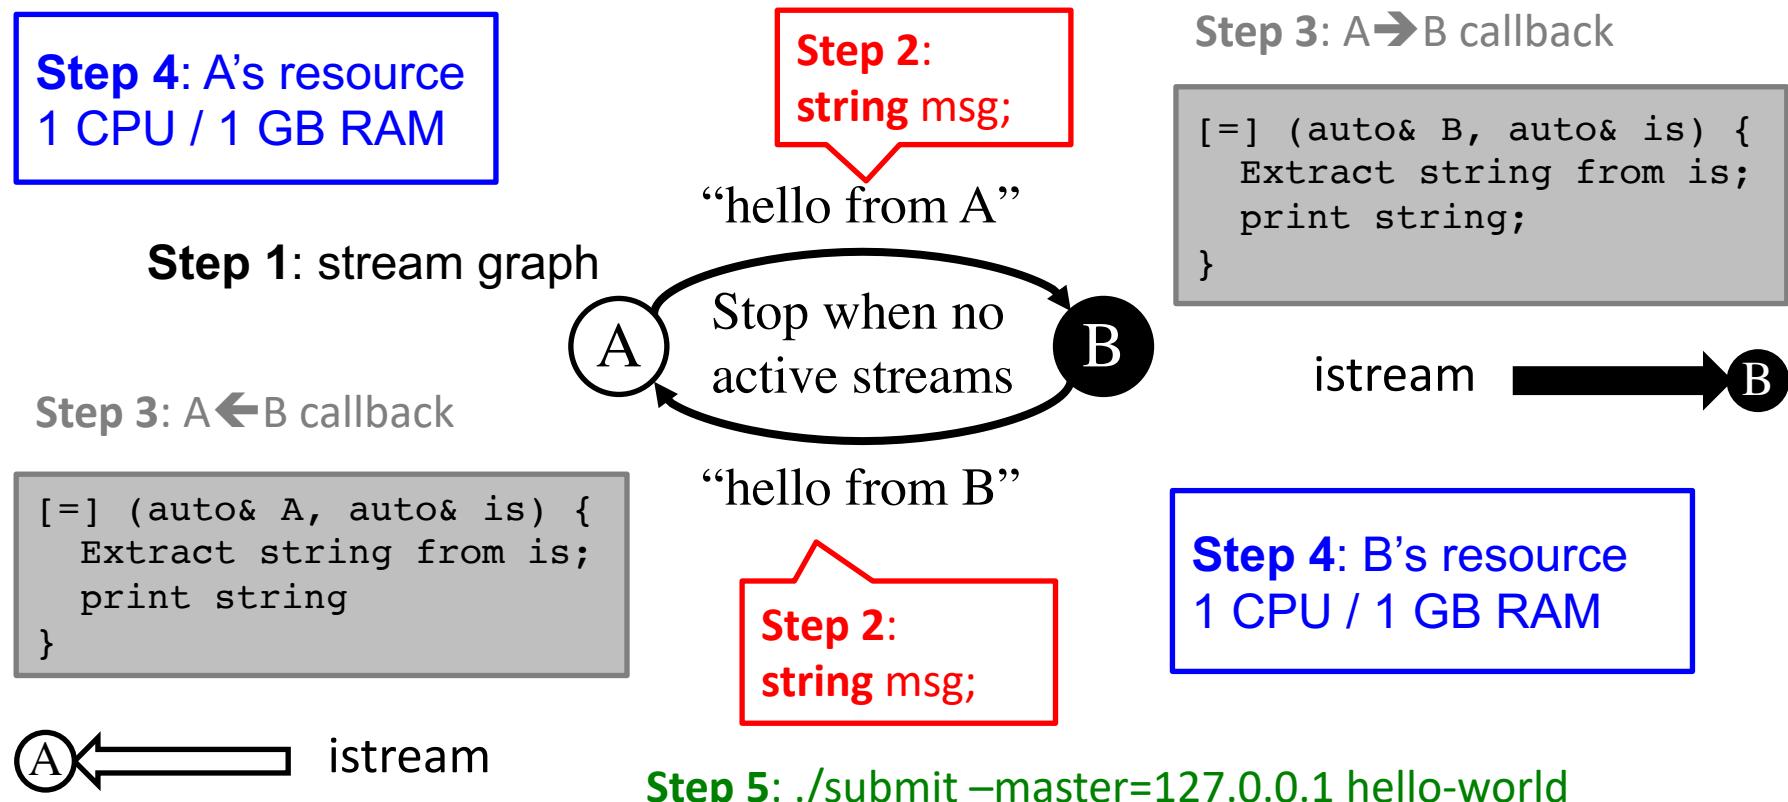
# DtCraft C++ Programming Runtime

- You describe an application in a *stream graph*
  - “No arm wrestling” with difficult concurrency details
- Everything is by default distributed
  - Same program runs on one and many machines



# A Hello-World Example

- An iterative and incremental flow
  - Two vertices + two streams



# DtCraft Code of Hello World

```
dtc::Graph G;  
  
auto A = G.vertex();  
auto B = G.vertex();  
  
G.container().add(A).cpu(1).memory(1_GB);  
  
auto AB = G.stream(A, B).on(  
[] (dtc::Vertex& B, dtc::InputStream& is) {  
    if(std::string b; is(b) != -1) {  
        dtc::cout("Received: ", b, '\n');  
        return dtc::Event::REMOVE;  
    }  
    return dtc::Event::DEFAULT;  
});  
  
auto BA = G.stream(B, A);  
  
A.on([&AB] (dtc::Vertex& v) {  
    (*v.ostream(AB))("hello world from A's");  
    dtc::cout("Sent 'hello world from A' to stream ", AB, "\n");  
});  
  
G.container().add(B).cpu(1).memory(1_GB);  
  
B.on([&BA] (dtc::Vertex& v) {  
    (*v.ostream(BA))("hello world from B's");  
    dtc::cout("Sent 'hello world from B' to stream ", BA, "\n");  
});  
  
BA.on([] (dtc::Vertex& A, dtc::InputStream& is) {  
    if(std::string a; is(a) != -1) {  
        dtc::cout("Received: ", a, '\n');  
        return dtc::Event::REMOVE;  
    }  
    return dtc::Event::DEFAULT;  
});  
  
dtc::Executor(G).run();
```

- Only a couple lines of code
- Single sequential program
- No explicit data management
- Distributed across computers
- Easy-to-use interface
- Asynchronous by default
- Scalable to many threads
- Scalable to many machines
- In-context resource controls
- Transparent concurrency
- Automatic Linux container
- ... and more

# Without DtCraft ...

```
auto count_A = 0;
auto count_B = 0;

// Send a random binary data to fd and add the
// received data to the counter.
auto pinpong(int fd, int& count) {
    auto data = random<bool>()
    auto w = write(fd, &data, sizeof(data));
    if(w == -1 && errno != EAGAIN) {
        throw system_error("Failed on write");
    }
    data = 0;
    auto r = read(fds, &data, sizeof(data));
    if(r == -1 && errno != EAGAIN) {
        throw system_error("Failed on read");
    }
    count += data;
}

int fd = -1;
std::error_code errc;
```

*server.cpp*

```
if(getenv("MODE") == "SERVER") {
    fd = make_socket_server_fd("9999", errc);
}
else {
    fd = make_socket_client_fd("127.0.0.1", "9999", errc);
}

if(fd == -1) {
    throw system_error("Failed to make socket");
}
```

*client.cpp*

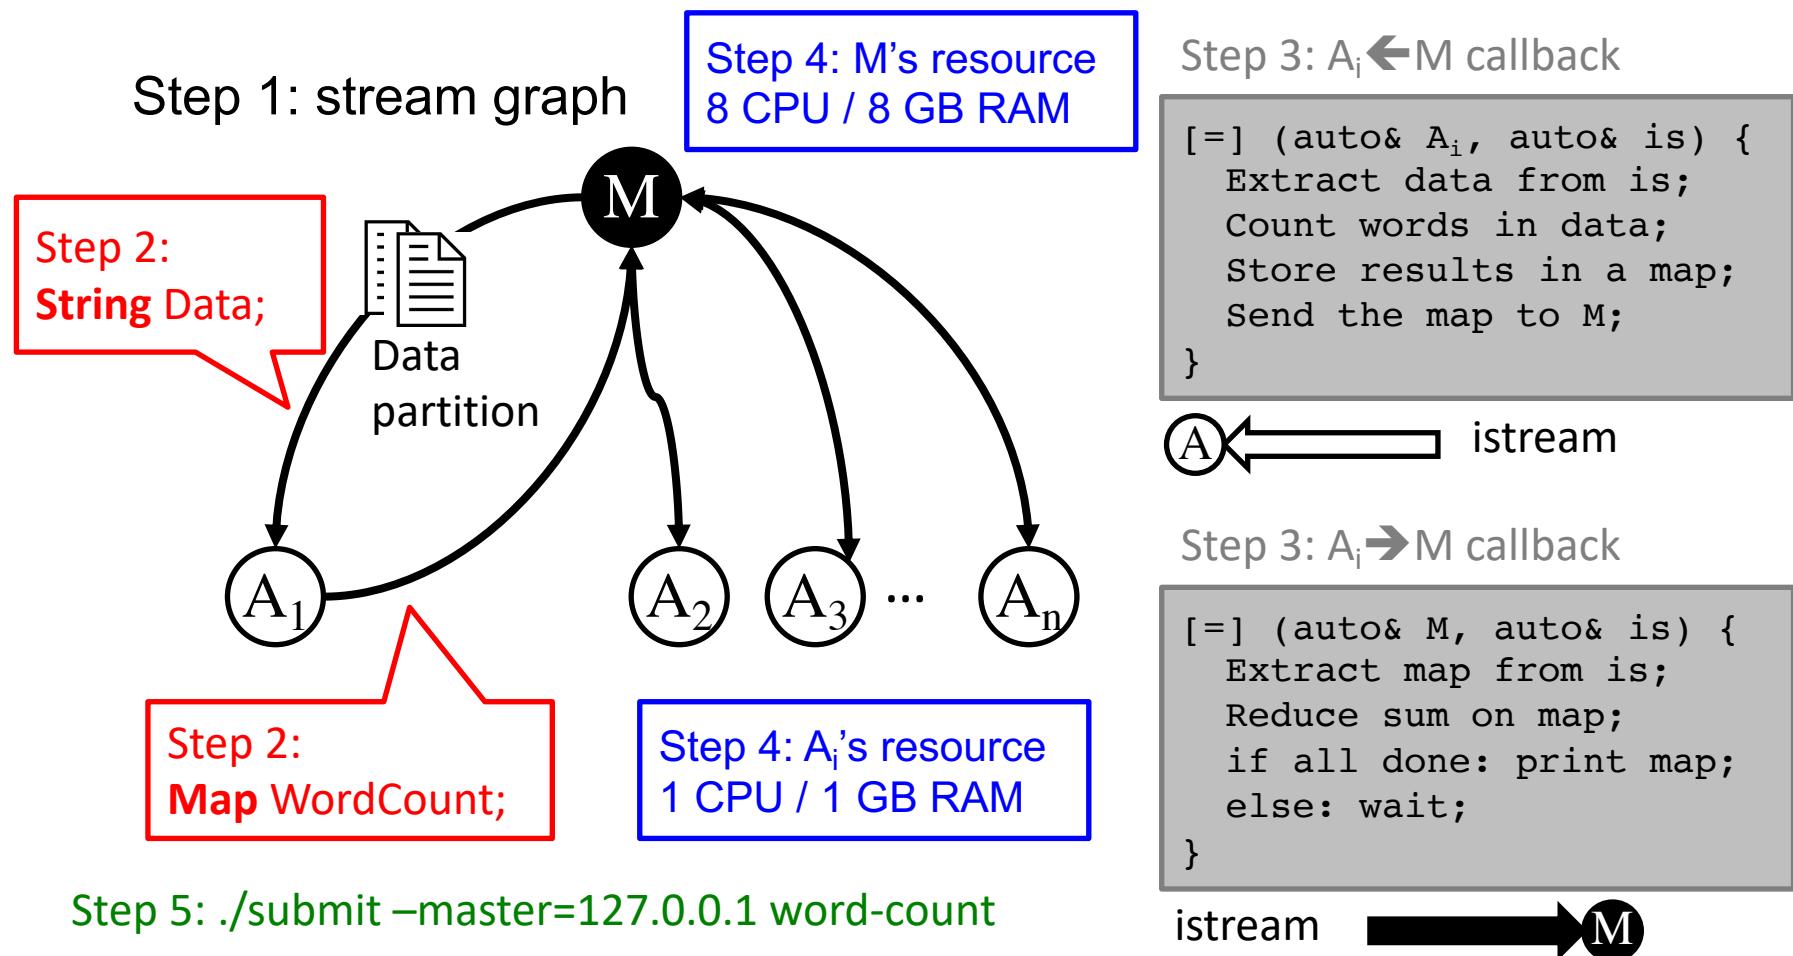
*Branch your code to server and client  
for distributed computation!*  
*simple.cpp → server.cpp + client.cpp*

```
int make_socket_server_fd(
    std::string_view port,
    std::error_code errc
) {
    int fd {-1};
    if(fd != -1) {
        ::close(fd);
        fd = -1;
    }
    make_fd_close_on_exec(fd);
    tries = 3;
    issue_connect:
    ret = ::connect(fd, ptr->ai_addr, ptr->ai_addrlen);
    if(ret == -1) {
        if(errno == EINTR) {
            goto issue_connect;
        }
        else if(errno == EAGAIN & tries--) {
            std::this_thread::sleep_for(std::chrono::milliseconds(500));
            goto issue_connect;
        }
        else if(errno != EINPROGRESS) {
            goto try_next;
        }
        errc = make_posix_error_code(errno);
    }
    // Poll the socket. Note that writable return doesn't mean it is connected
    if(select_on_write(fd, 5, errc) && !errc) {
        int optval = -1;
        socklen_t optlen = sizeof(optval);
        if(::getsockopt(fd, SOL_SOCKET, SO_ERROR, &optval, &optlen) < 0) {
            errc = make_posix_error_code(errno);
            goto try_next;
        }
        if(optval != 0) {
            errc = make_posix_error_code(optval);
            goto try_next;
        }
        // Try to connect again
        for(auto i {0}; i < tries; ++i) {
            // Ignore EINPROGRESS
            if(ptr->ai_family == AF_INET) {
                struct sockaddr_in* paddr =
                    reinterpret_cast<struct sockaddr_in*>(ptr->ai_addr);
                if(paddr->sin_port == 0) {
                    paddr->sin_port = htons(ntohs(htons(0)) + i);
                }
            }
            else if(ptr->ai_family == AF_INET6) {
                struct sockaddr_in6* paddr =
                    reinterpret_cast<struct sockaddr_in6*>(ptr->ai_addr);
                if(paddr->sin6_port == 0) {
                    paddr->sin6_port = htons(ntohs(htons(0)) + i);
                }
            }
            ::setsockopt(fd, SOL_SOCKET, SO_REUSEADDR, &i, sizeof(i));
            if(::bind(fd, ptr->ai_addr, ptr->ai_addrlen) != -1) {
                errc = std::error_code();
                goto done;
            }
            if(::listen(fd, 5) != -1) {
                errc = std::error_code();
                goto done;
            }
            if(fd != -1) {
                ::close(fd);
                fd = -1;
            }
        }
    }
    ::freeaddrinfo(res);
}
return fd;
```

*A lot of boilerplate code  
plus hundred lines of  
scripts to enable  
distributed flow...*

# Word Counting MapReduce Example

- A master manages several workers



# System Implementation Details

## ❑ Kernel – Master, Agent, and Executor

- ❑ Master: global scheduler, deployment
- ❑ Agent: local scheduler
- ❑ Executor: task executor

## ❑ Event-driven environment

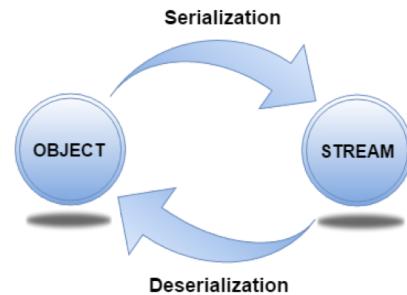
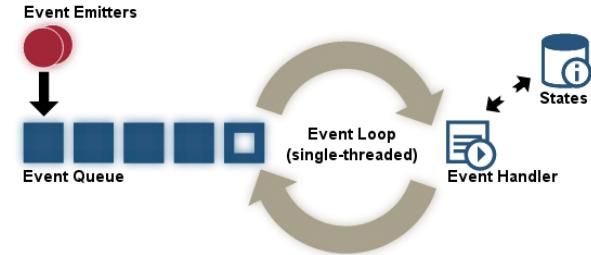
- ❑ A new C++ reactor library
- ❑ Thread-safe, lock-free, non-blocking IO

## ❑ Streaming interface

- ❑ Built-in serialization library

## ❑ Linux container

- ❑ Secure, safe, and robust resource control



# DtCraft to Handle Machine Learning

Application pipeline (domain-specific applications)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



...

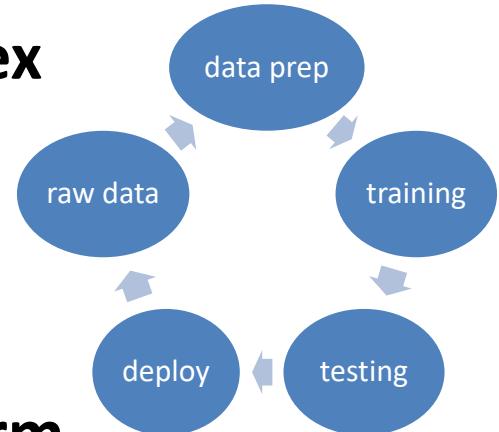


*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# MLCraft – Machine Learning at Scale

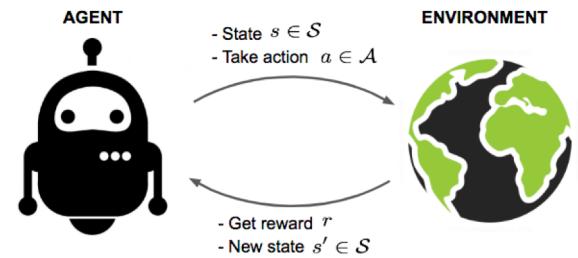
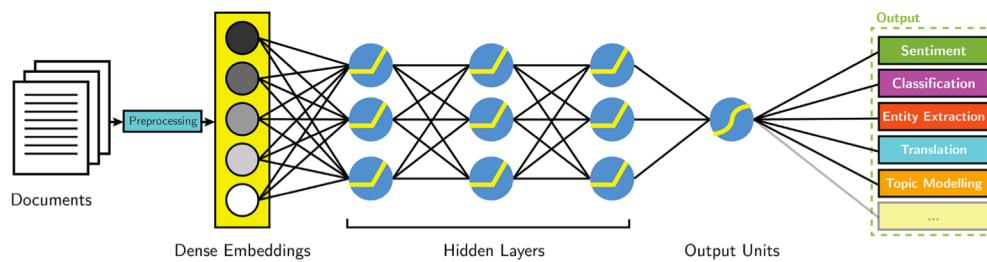
## □ ML system development is very complex

- Hard to track experiments
- Hard to reproduce results
- Hard to deploy and maintain



## □ MLCraft aims to be a unified ML platform

- Automate ML workflows and developments
- Open API and interface to work with existing libraries
- More robust, predictable, easier to maintain



```

dtc::Graph G;

auto src = G.insert<dtc::cell::MnistStreamFeeder>(
    mnist_image_file,
    mnist_label_file,
    [] (Eigen::MatrixXf& images, Eigen::VectorXi& labels) {
        images /= 255.0f;
        return std::make_tuple(images, labels);
    }
);

src.window(1ms).frequency(1000);

auto dnn = G.insert<dtc::cell::Visitor1x1>(
    [] (dtc::ml::DnnClassifier& c) {
        dtc::cout("Creating a dnn classifier [784x30x10]\n").flush();
        c.fully_connected_layer(784, 30, dtc::ml::Activation::RELU)
            .fully_connected_layer(30, 10, dtc::ml::Activation::NONE)
            .optimizer<dtc::ml::AdamOptimizer>();
    },
    [i=0] (dtc::ml::DnnClassifier& c, std::tuple<Eigen::MatrixXf, Eigen::VectorXi>& data) mutable {
        auto& [images, labels] = data;
        dtc::cout(
            "Accuracy at training cycle ", i++, " [", images.rows(), " images]: ",
            ((labels-c.infer(images)).array() == 0).count() / static_cast<float>(images.rows()), "%"
        ).flush();
        c.train(images, labels, 10, 64, 0.01f, [](){});
    }
);

dnn.in(src.out());

G.container().add(dnn);
G.container().add(src);

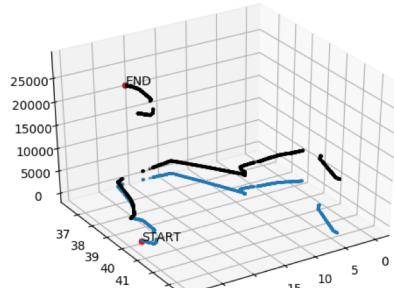
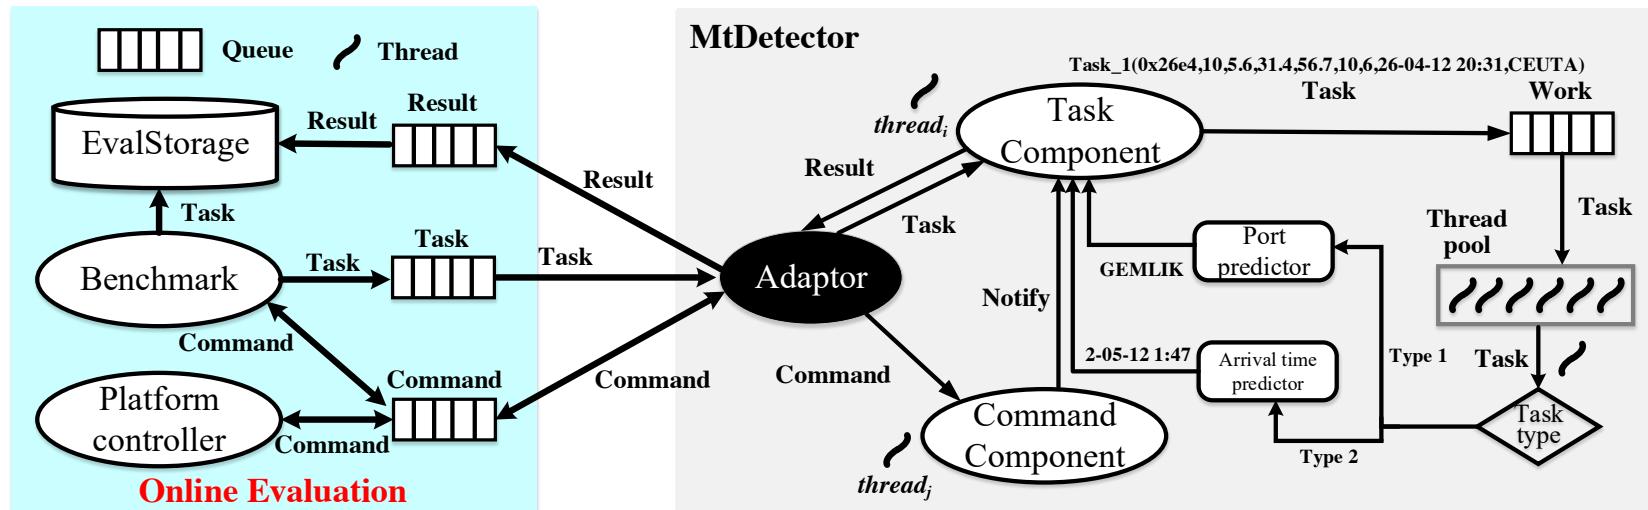
dtc::Executor(G).run();

```

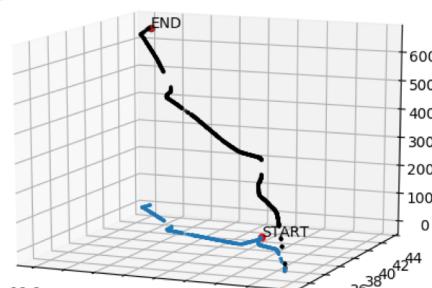
Only “60-line” code to  
automate a distributed  
ML workflow!

# Marine Traffic Prediction at Stream Scale

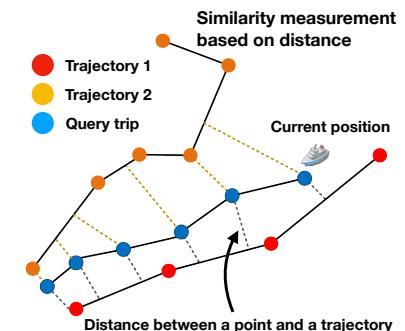
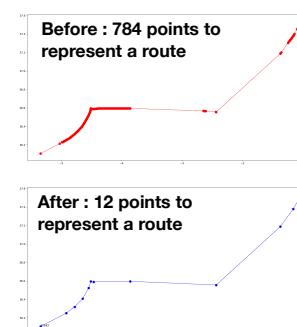
- MtDetector: A stream-based online ML system
  - Top-5 in ACM DEBS 2018 Grand Challenge (out of 50+)



MAE = 120 minutes

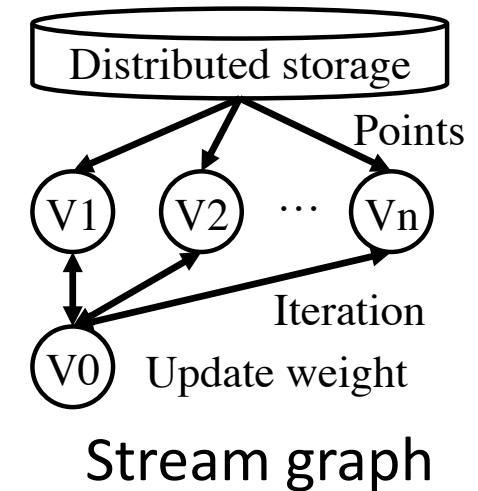
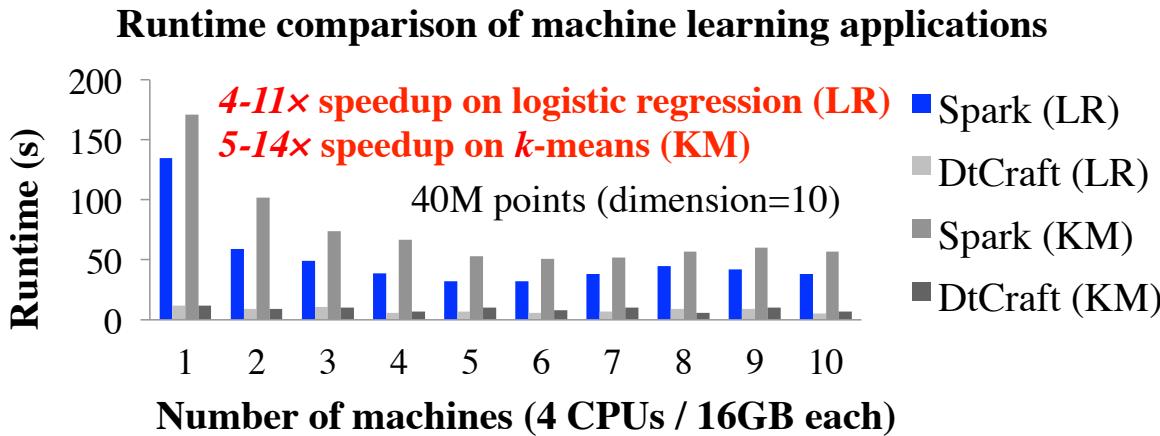


MAE = 267 minutes

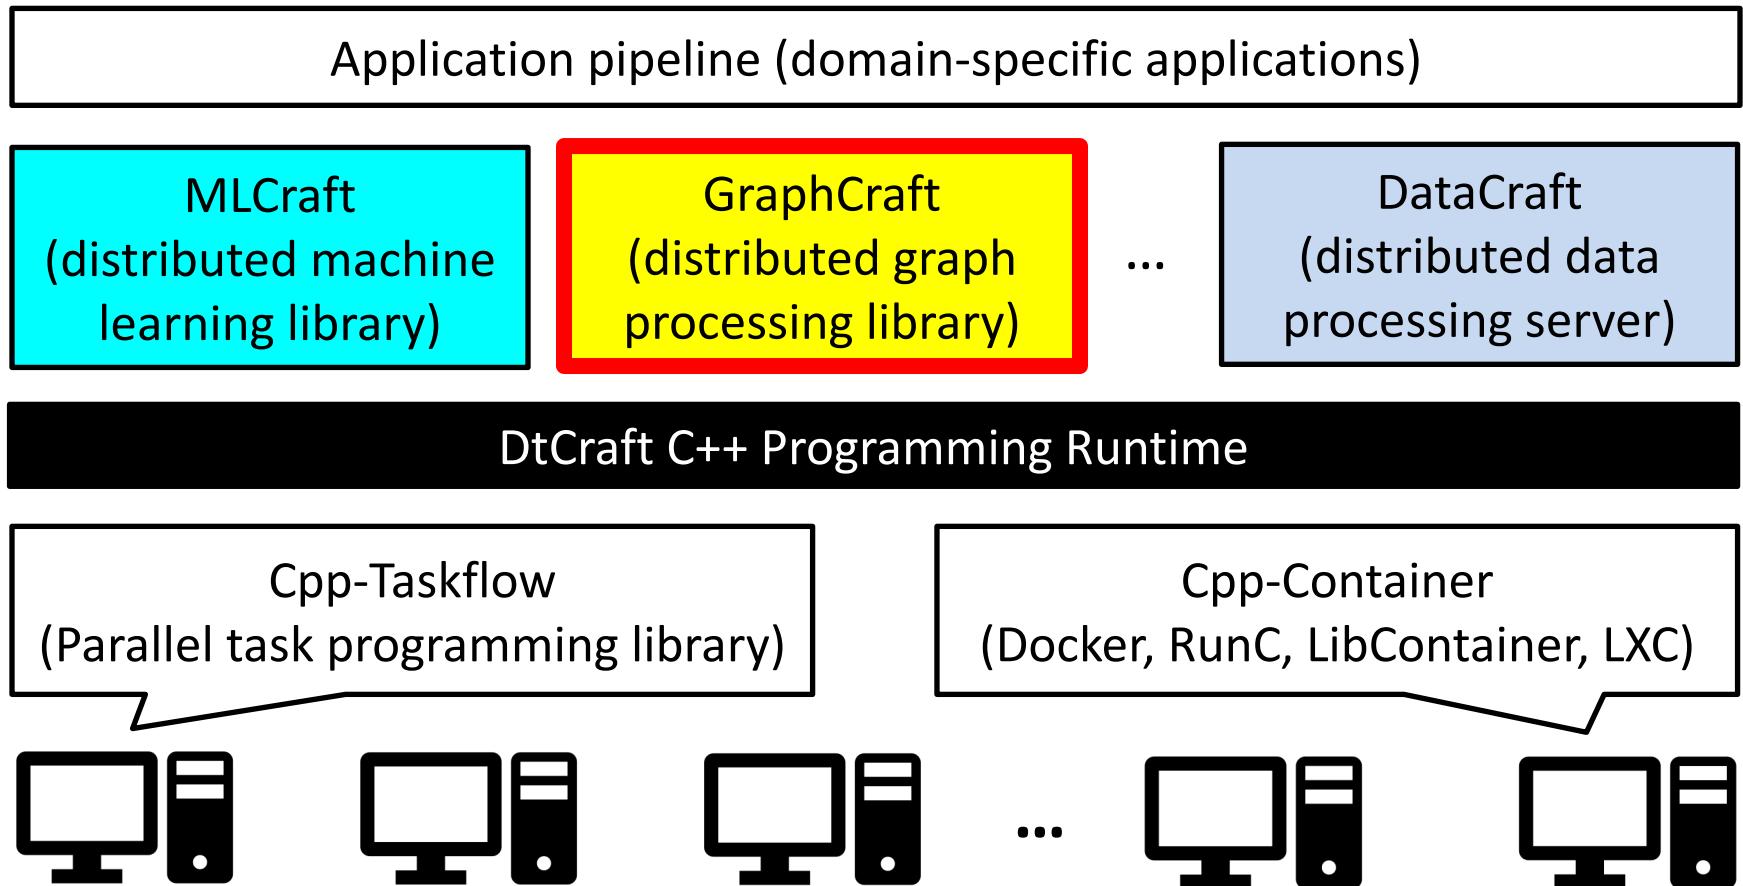


# Machine Learning Performance

- ❑ Logistics regression and *k*-means clustering
  - ❑ 100 training iterations
- ❑ Compared with Spark ML library
  - ❑ 15x speed-up with 2-5x less memory
  - ❑ No extra overhead to cache the data for reuse



# DtCraft to Handle Graph Algorithms

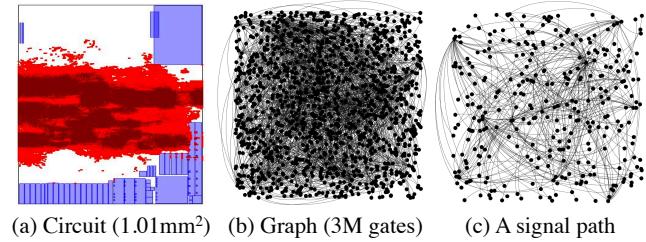


*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# Graph Algorithms Performance

## □ Shortest path finding

- A circuit graph of 10B transistors
- Pregel algorithms

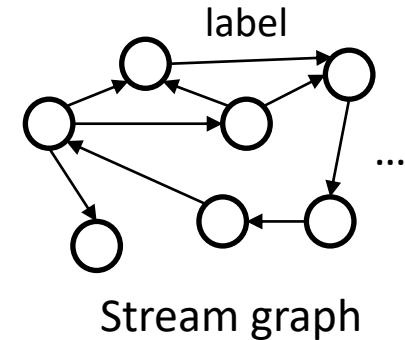


(a) Circuit (1.01mm<sup>2</sup>) (b) Graph (3M gates)

(c) A signal path

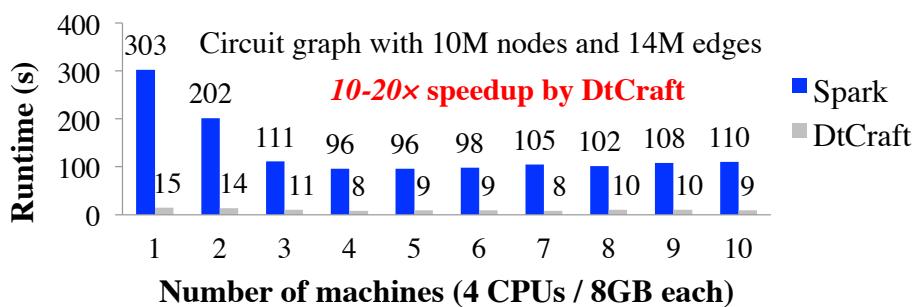
## □ Compared with Spark GraphX

- 10-20x speed-up
- Better scalability in increasing graph sizes
- Less communication cost

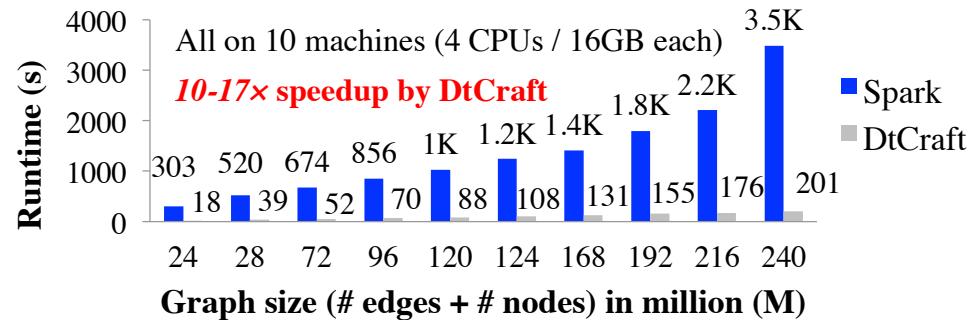


Stream graph

Runtime comparison of shortest path finding



Performance scalability (runtime vs graph size)



# DtCraft to Handle Distributed Timing

Application pipeline (Distributed Timing Analysis)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



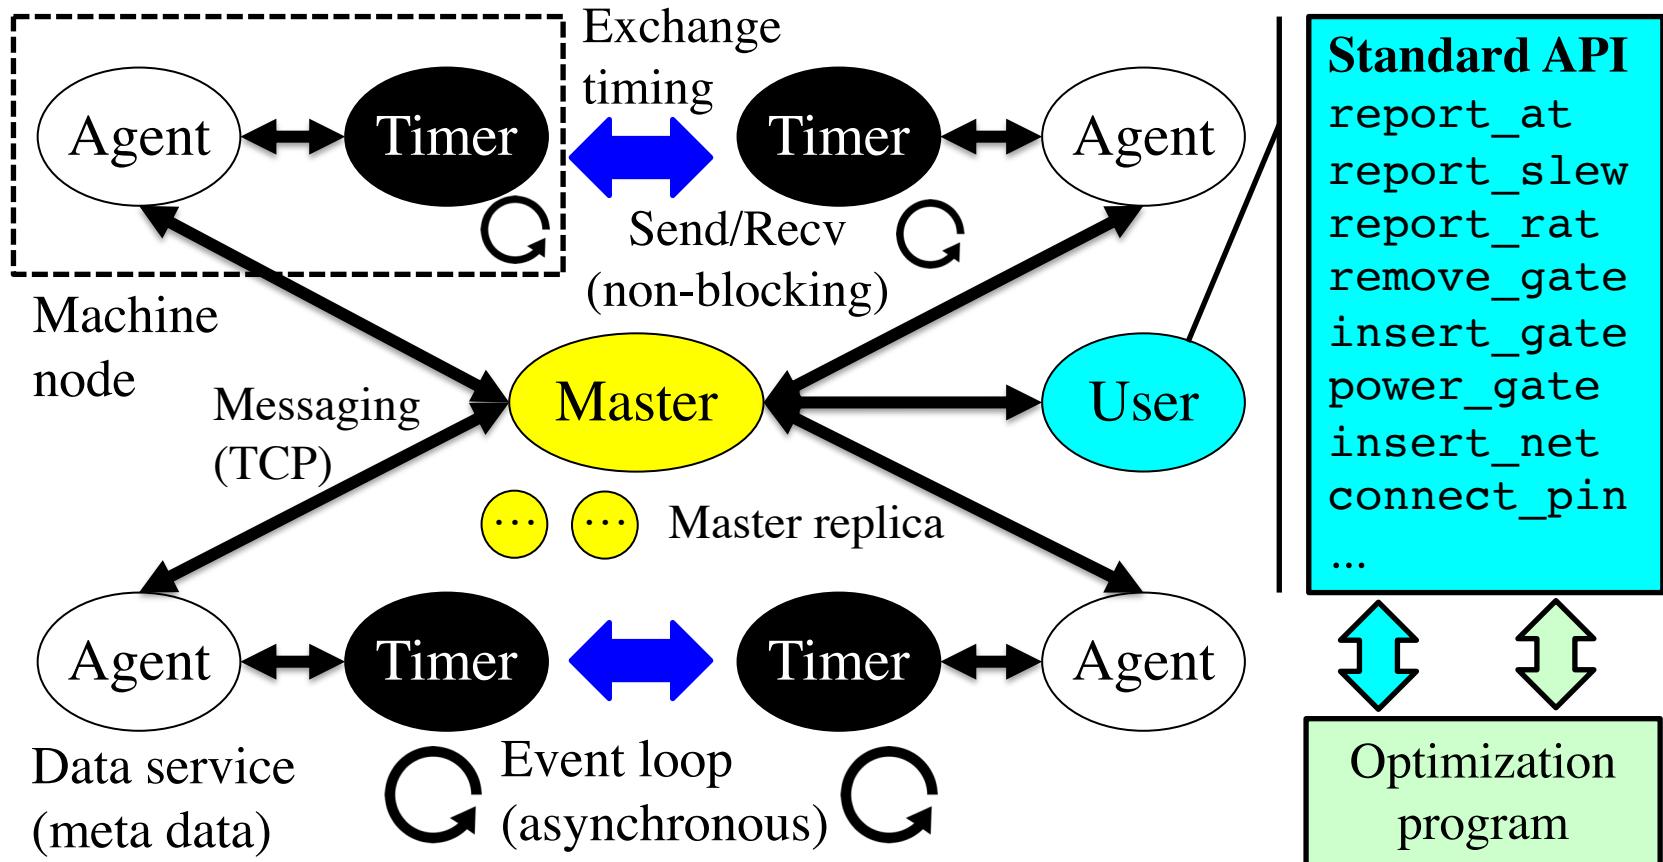
...



*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# Distributed Timing Analysis

## □ Master coordinator and worker (a timer per partition)



# Exchange Timing Values

## DtCraft

```
// Timing data (early/late rise/fall)
struct Timing {
    string pin;
    array<float, 4> value;

    template <typename T>
    auto archive(T& ar) {
        return ar(pin, value);
    }
};

// Timing path
struct Path {
    float slack;
    vector<string> pins;

    template <typename T>
    auto archive(T& ar) {
        return ar(slack, pins);
    }
};

// Exchange timing through DtCraft stream
stream.on(
    [] (Vertex& v, InputStream& is) {
        if (Timing timing; is(timing) != -1) {
            // Call OpenTimer to run incremental timing
        }
    }
);
```

*In-context streaming with < 30 lines*



## Existing framework

Hard-code your message format

```
// OpenTimer.proto
package OpenTimer;

// Message format for timing
message Timing {
    required string pin = 0;
    required float er = 1;
    required float ef = 2;
    required float lr = 3;
    required float lf = 4;
}

// Message format for path
message Path {
    required float slack = 0;
    repeated string pins = 1;
}
```

*Many extra stuff ↧*  
**Extra.pb.h**  
**Extra.pb.cpp**  
...  
Source.cpp

Google Protocol Buffer  
(open-source compiler)

C++/Java/Python  
source code generator

**.cpp/.h class methods**  
ParseFromArray(void\*, size\_t)  
SerializeToArray(void\*, size\_t)

Message wrapper

**Derived packet struct**  
header\_t header  
void\* buffer



*Out-of-context streaming takes > 300 lines*

# Deploy the Distributed Timer in One line

## DtCraft

```
// Create a timer vertex for Top
auto Top = G.Vertex().on(
    [=] () {
        OpenTimer timer ("Top.v");
    }
);

// Create a timer vertex for Macro 1
auto M1 = G.Vertex().on(
    [=] () {
        OpenTimer timer ("M1.v");
    }
);

// Create a timer vertex for Macro 2
auto M2 = G.Vertex().on(
    [=] () {
        OpenTimer timer ("M2.v");
    }
);

// Create streams ...
...

// Distribute timers to machines.
G.container().add(Top).num_cpus(4).memory_(4_GB);
G.container().add(M1).num_cpus(1).memory(8_GB);
G.container().add(M2).num_cpus(2).memory(6_GB);

~$ ./submit --master=127.0.0.1 binary
```



## Existing framework

*Duplicate the code for each partition*

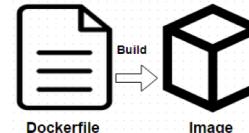
Top.cpp



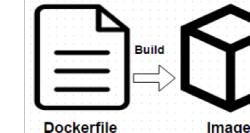
M1.cpp



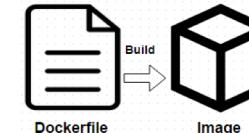
M2.cpp



Container 1



Container 2



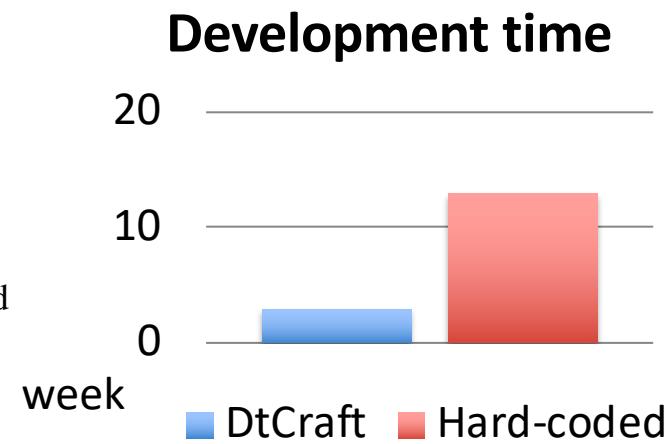
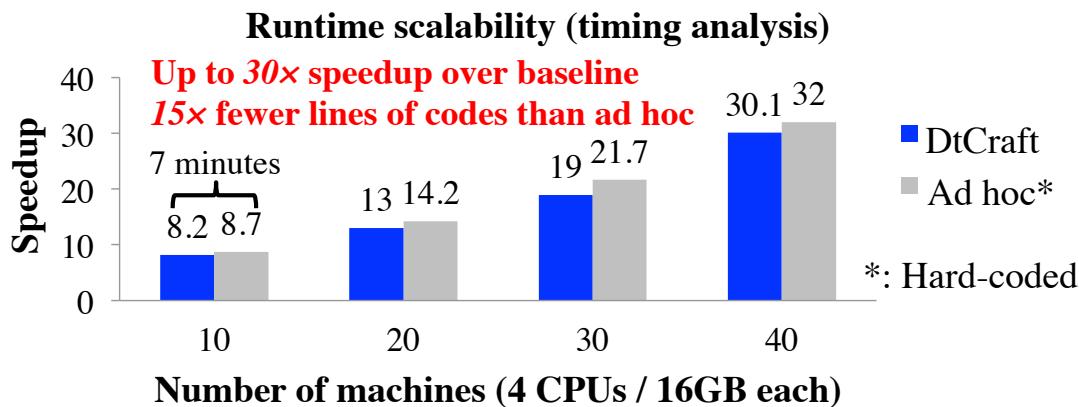
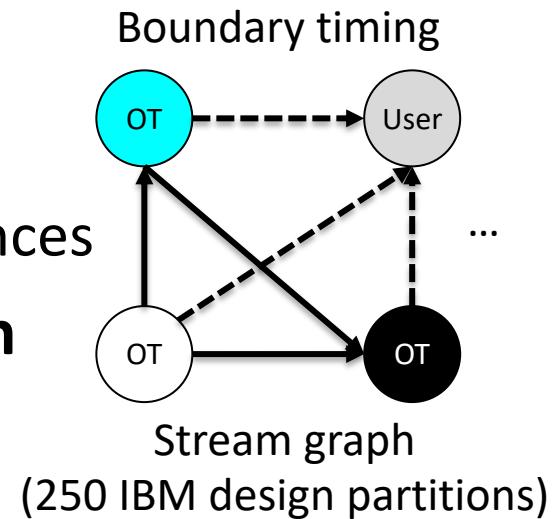
Container 3

*Wrap up with submission scripts*



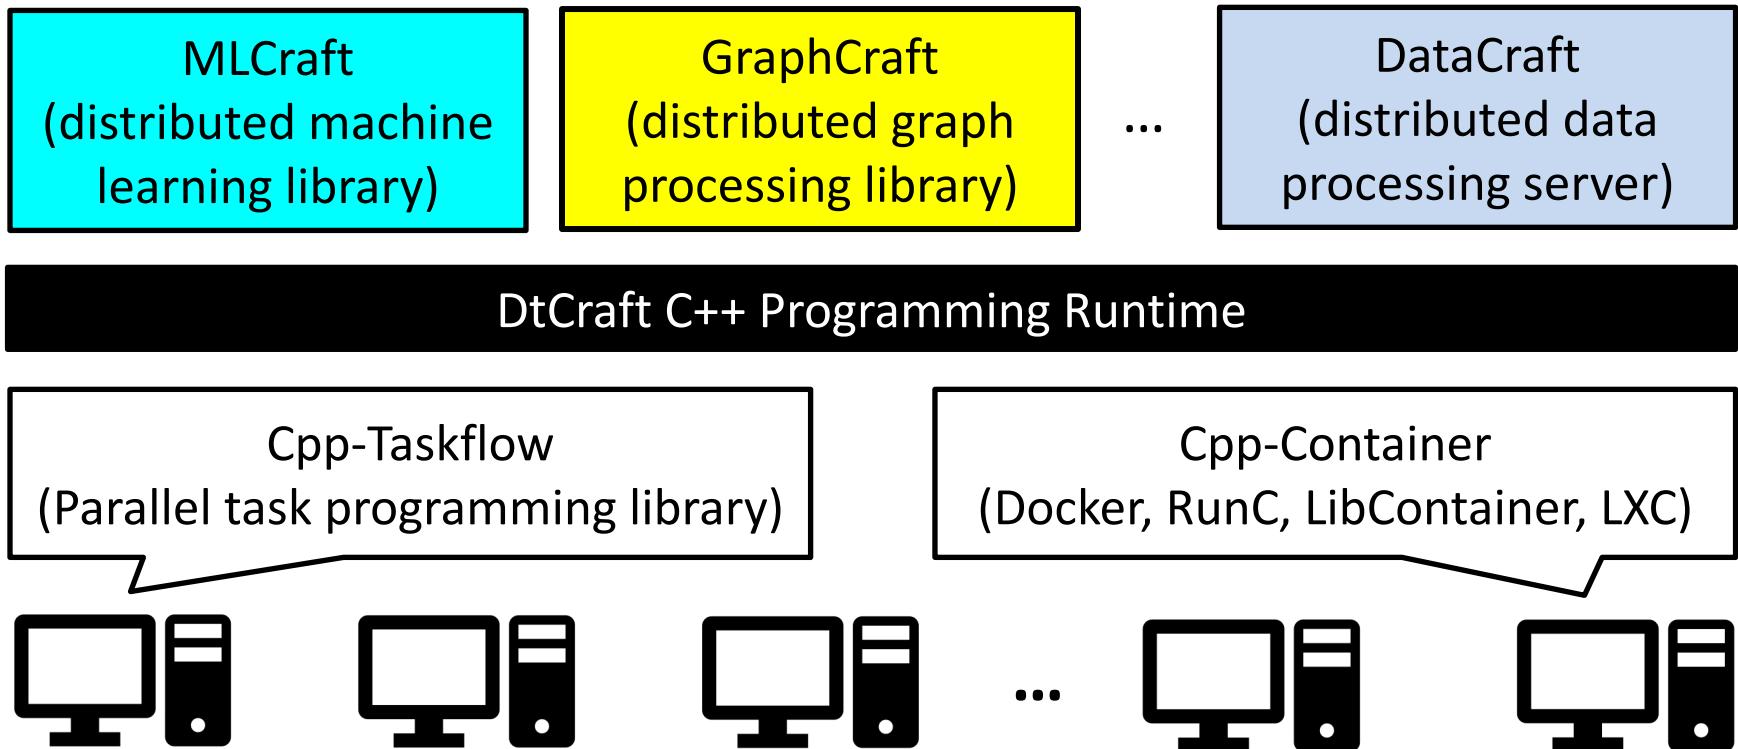
# Distributed Timing Performance

- ❑ IBM design of 250 partitions
  - ❑ OpenTimer as the timing analyzer
  - ❑ On a cluster of 40 Amazon EC2 instances
- ❑ Compared with hard-coded solution
  - ❑ 15x fewer lines of code
  - ❑ Only 7% performance loss



# DtCraft to Handle a Layout Generator

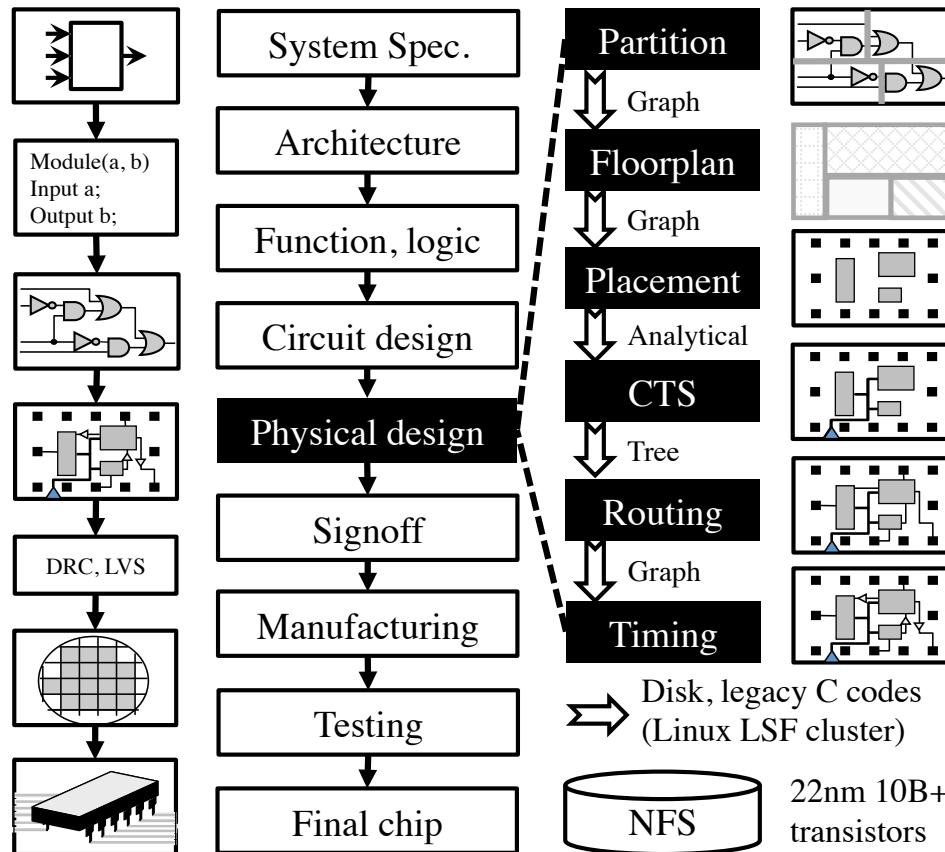
Application pipeline (**End-to-end distributed CAD flow**)



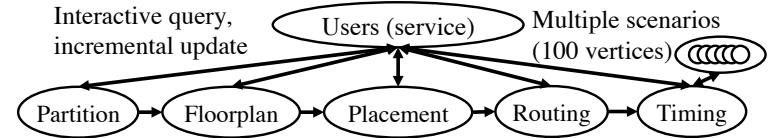
*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# A Fully Automated Layout Generator

## □ DARPA's vision: A GCC-like silicon compiler



Stream graph  
(106 vertices and 220 streams)



## □ Tool collections

- Partitioner & Floorplanner
- Placer & router
- Timer
- ...

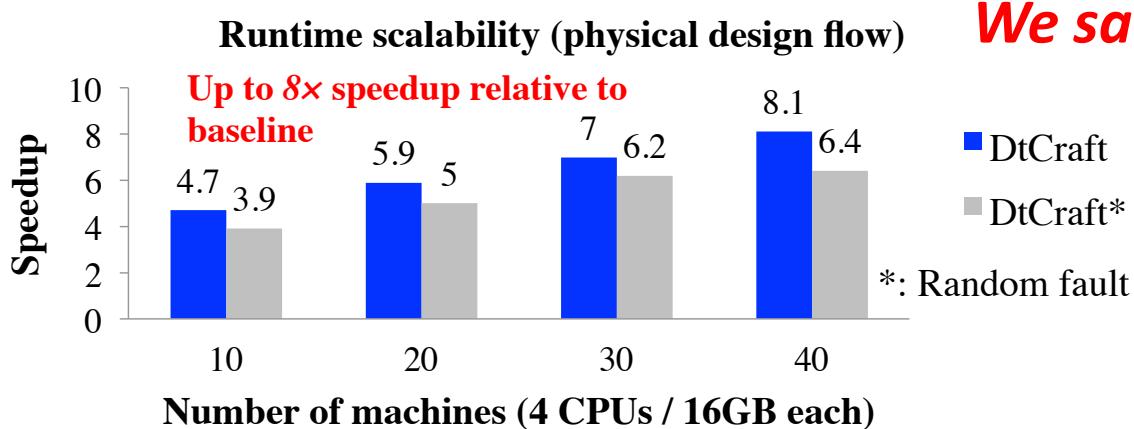
## □ IDEA/POSH end state ...

```
$ git clone https://github.com/darpa/idea
$ git clone https://github.com/darpa/posh
$ cd posh
$ make soc42
```

T.-W. Huang, et al, "DtCraft: A high-performance distributed execution engine at scale," IEEE TCAD, 2018

# Performance of Distributed Integration

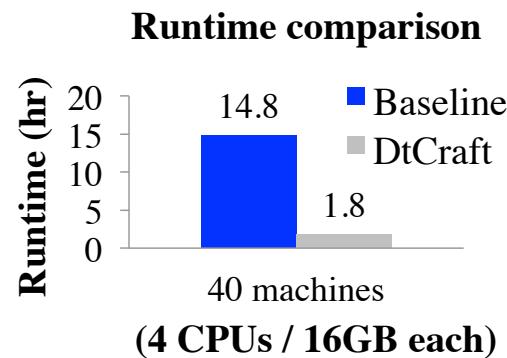
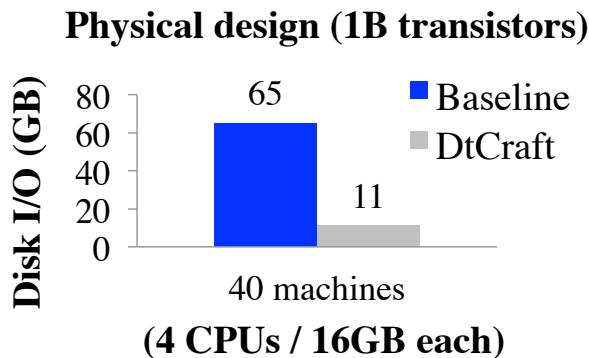
- ❑ Compared with existing “script-based” solutions
- ❑ Run on 40 Amazon EC2 nodes



**We saved 20x code in general!**

\* Up to 8x speed-up compared to the baseline

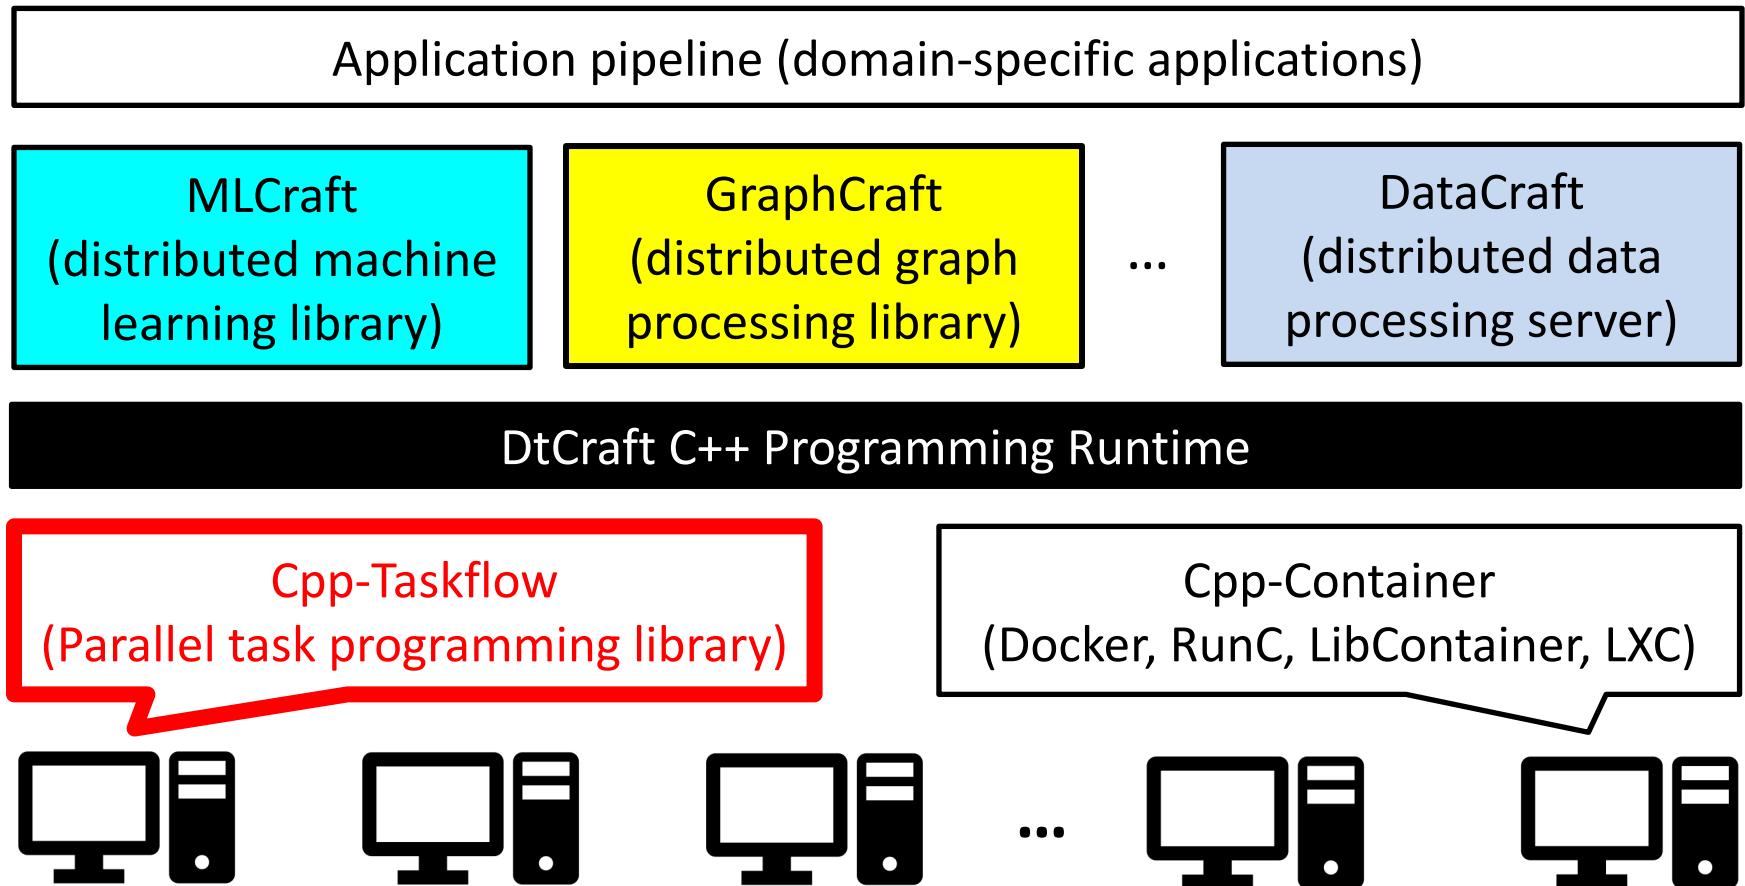
\* Design completion time is reduced from by **13 hours**



\* In-memory streaming reduced the I/O overhead by **~55GB**

\* Less I/O overhead translates to faster runtime

# DtCraft to Handle Low-level Concurrency



*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# Cpp-Taskflow: Parallel Programming Library

We aim to help developers **quickly** write **efficient** parallel programs using task-based approaches with **Modern C++**

## □ We want parallel code

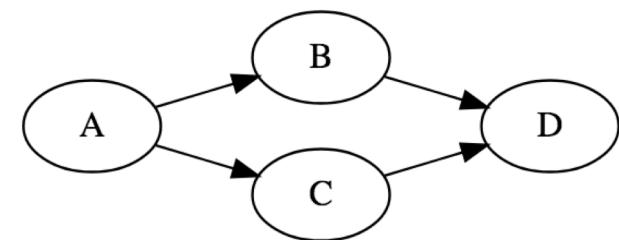
- Simple
- Expressive
- Transparent

## □ We don't want hard details

- Explicit thread management
- Difficult lock controls
- Daunting class declaration

## Cpp-Taskflow's Project Mantra

```
tf::Taskflow tf;  
  
auto [A, B, C, D] = tf.silent_emplace(  
    []() { std::cout << "Task A\n"; },  
    []() { std::cout << "Task B\n"; },  
    []() { std::cout << "Task C\n"; },  
    []() { std::cout << "Task D\n"; }  
);  
  
A.precede(B, C); // A runs before B and C  
B.precede(D); // B runs before D  
C.precede(D); // C runs before D  
  
tf.wait_for_all(); // block until finish
```

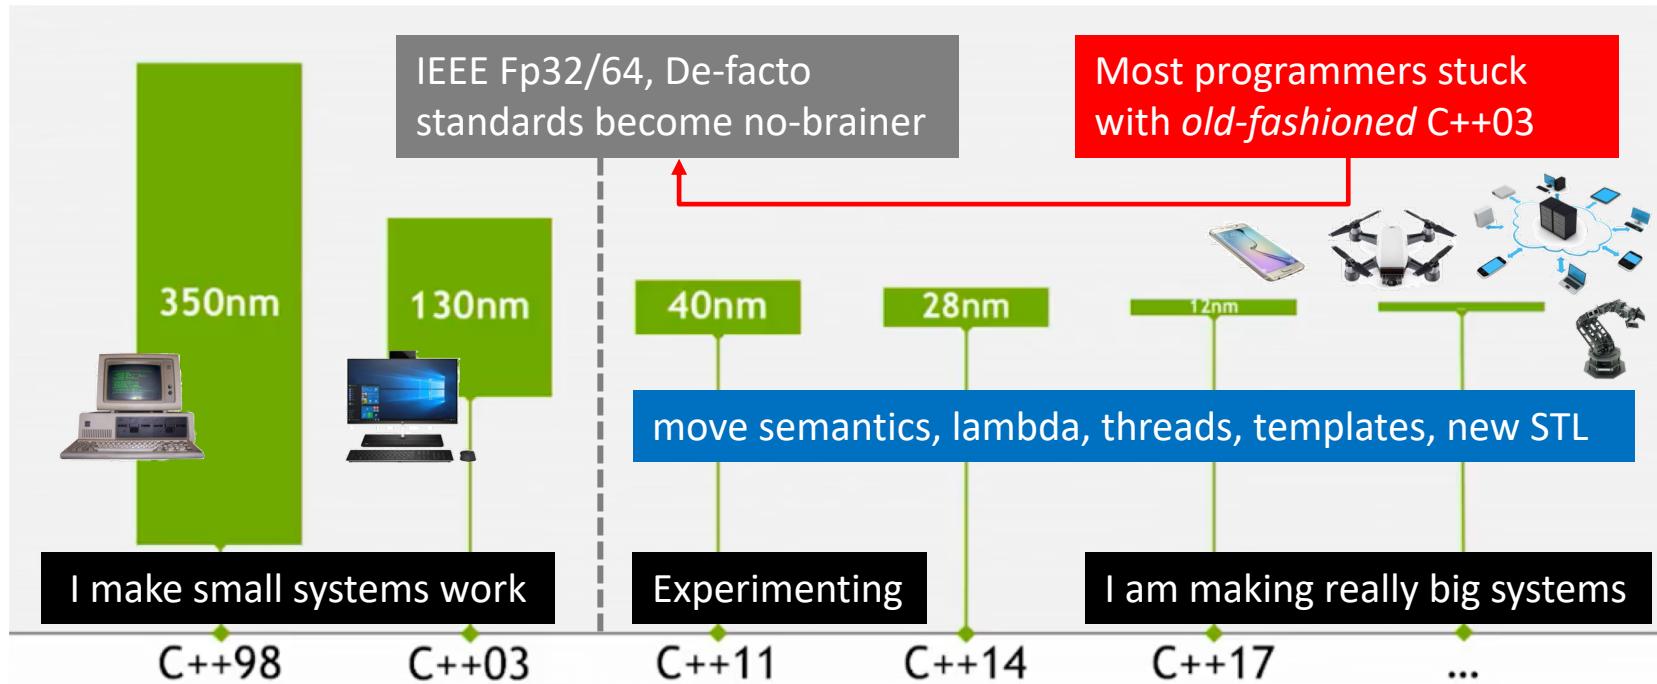


T.-W. Huang, et al, “Cpp-Taskflow: Fast Parallel Task-based Programming using Modern C++,” IEEE IPDPS, 2019

Only **10 lines** of code to enable parallel task execution!

# “Modern C++” Enables New Technology

- ❑ If you were able to tape out C++ ...

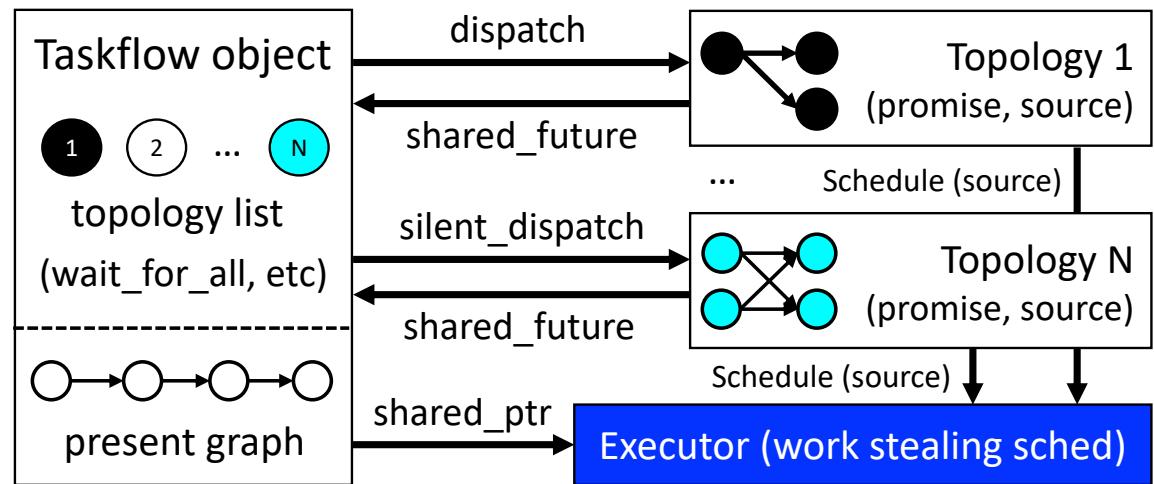
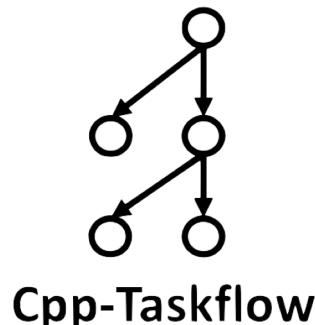


- ❑ It's much more than just being modern

- ❑ Must “rethink” the way we used to design a program
- ❑ Achieve the performance previously not possible

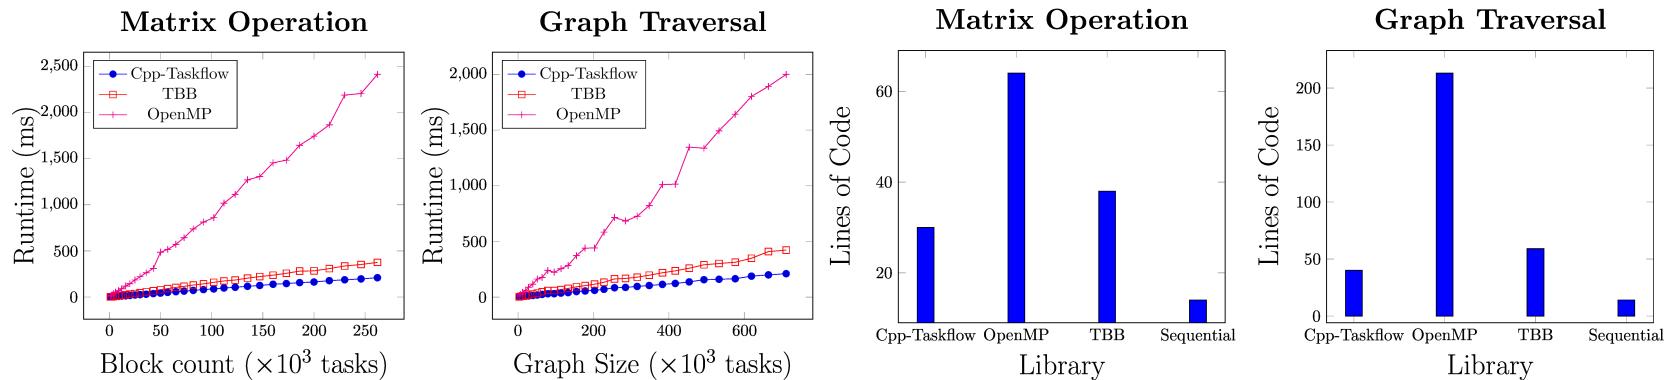
# Implementation Details of Cpp-Taskflow

- A new parallel task programming model
  - Simple yet efficient
- Unified interface for both static and dynamic tasking
  - No need to learn a different API set
- Pluggable scheduler and executor interface
  - Modular and customizable
  - Sharable



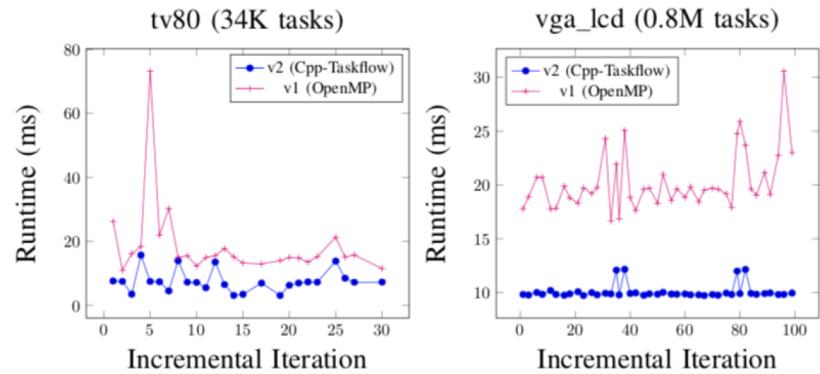
# Cpp-Taskflow's Performance

- Compared with OpenMP and Intel TBB
  - Faster, more scalable, and fewer lines of code



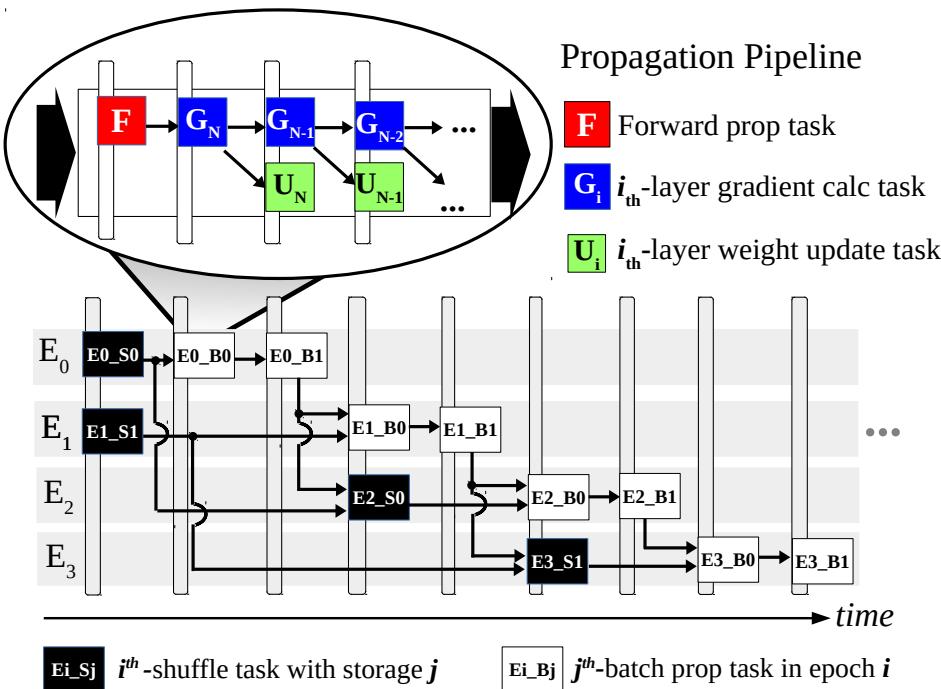
- Large VLSI timing analysis (OpenMP vs Cpp-Taskflow)
  - 2x speed-up over OpenMP
  - 9123 vs 4482 lines of code

*Cost to develop is \$275K with OpenMP  
vs \$130K with Cpp-Taskflow!  
(<https://dwheeler.com/sloccount/>)*



# Simplify the Development of Parallel ML

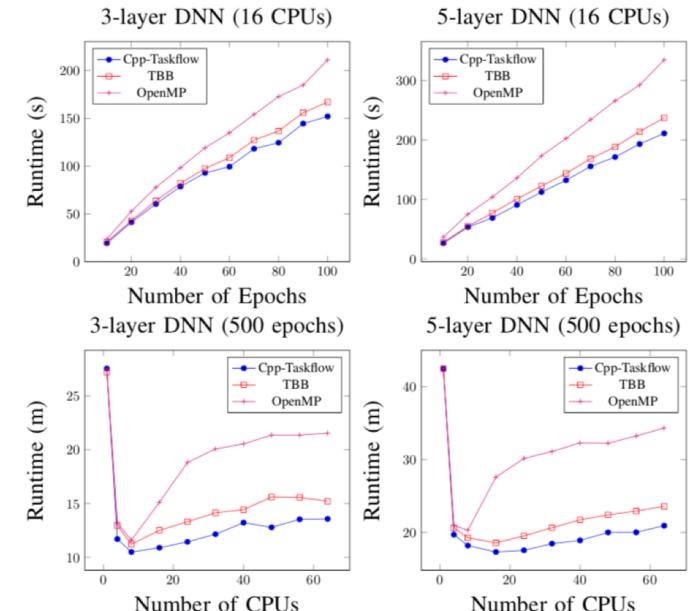
- “OMG, never do so in TensorFlow,” G\_\_\_\_ Research
  - RunQueue, Python threads are nightmare in TensorFlow
- Cpp-Taskflow makes developers’ lives easier
  - 2x fewer LoC and 2x faster



Cpp-Taskflow			OpenMP			TBB			Sequential		
LOC	CC	T	LOC	CC	T	LOC	CC	T	LOC	CC	T
59	11	3	162	23	9	90	12	3	33	9	2

CC: cyclomatic complexity of the implementation

T: development time (in hours) by an experienced programmer



# What Cpp-Taskflow Users Say ...

---

*“Cpp-Taskflow is the cleanest Task API I have ever seen,”*  
*Damien Hocking*

*“Cpp-Taskflow has a very simple and elegant tasking interface; the performance also scales very well,” Totalgee*

*“Best poster award for open-source parallel programming library,” Cpp-Conference (voted by 1000+ developers)*

**People are using Cpp-Taskflow to speed up TensorFlow’s kernel!**



lifeiteng commented on Dec 11, 2018



**Is your feature request related to a problem? Please describe.**

I want to use TaskFlow in a C++ project which depends on TensorFlow. When compiling the Op & Kernel with `c++17`, there is a problem [tensorflow/tensorflow#23561](#)

# Outline

---

- My early PhD research
  - OpenTimer: A VLSI timing analysis tool
  - Express parallelism in the right way
- A general-purpose distributed programming system
  - DtCraft and its ecosystem
  - Machine learning, graph algorithms, chip design
  - Cpp-Taskflow: Modern C++ parallel task programming
- Conclusion and future work

# Future Work (marked red)

- ❑ Enhance both app- and system-level capabilities
  - ❑ New workloads, new libraries, new API, etc.

Application pipeline (+ VLSI, HPC, Simulation, Optimization, etc)

MLCraft  
(+ reinforcement ML)

GraphCraft  
(+ algorithms)

DataCraft  
(+ DFS, Ceph)

Blockchain

...

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(+ ML compiler)

Cpp-Container  
(+ FPGA, GPU)

Security  
(+ OpenSSL protocols)

...



...



# Collaboration Plan

---

- With faculty members working on (not limited to):**
  - HPC, FPGA, networking, machine learning, AI, computer systems, VLSI, OS virtualization, etc.
- With students (anyone interested in computing!)**
  - Undergraduate and graduate research advising
  - Existing courses (EE, CE, CS)
    - Intro, VLSI, Digital Design, Architecture, Programming, OS, etc.
  - New course ideas (at all levels)* 
    - Big-data cloud computing, scalable AI systems, virtualization, modern C++ programming, principles of open source, etc.
  - Professional activities (contests, conferences, etc.)
    - Programming contests, research competition, conferences, etc.
- PI of DARPA research grant “OpenTimer and DtCraft”**
  - Brings collaborations with faculty members & students

# Acknowledgment (Users & Sponsors)



SYNOPSYS®

