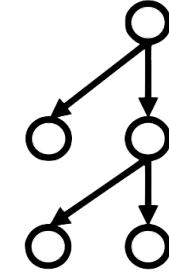


A Modern C++ Parallel Task Programming Library



GitHub: <https://github.com/cpp-taskflow>



Docs: <https://cpp-taskflow.github.io/cpp-taskflow/>

C.-X. Lin, Tsung-Wei Huang, G. Guo, and M. Wong

University of Utah, Salt Lake City, UT, USA

University of Illinois at Urbana-Champaign, IL, USA



Cpp-Taskflow's Project Mantra

A programming library helps developers **quickly** write **efficient** parallel programs on a manycore architecture using **task-based** models in **modern C++**

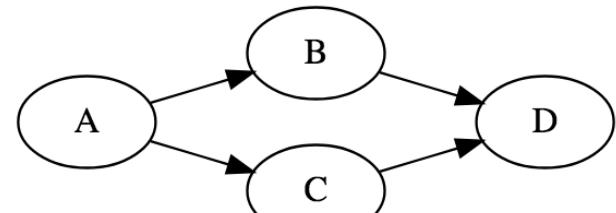
- Parallel computing is important in modern software
 - Multimedia, machine learning, scientific computing, etc.
- Task-based approach scales best with manycore arch
 - We should write tasks NOT threads
 - Not trivial due to dependencies (race, lock, bugs, etc.)
- We want developers to write parallel code that is:
 - Simple, expressive, and transparent
- We don't want developers to manage:
 - Threads, concurrency controls, scheduling

Hello-World in Cpp-Taskflow

```
#include <taskflow/taskflow.hpp> // Cpp-Taskflow is header-only
int main(){
    tf::Taskflow tf;
    auto [A, B, C, D] = tf.emplace(
        [] () { std::cout << "TaskA\n"; }
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B); // A runs before B
    A.precede(C); // A runs before C
    B.precede(D); // B runs before D
    C.precede(D); // C runs before D
    tf::Executor().run(tf); // create an executor to run the taskflow
    return 0;
}
```

Only **15 lines** of code to get a parallel task execution!

- ✓ No hardcode threads
- ✓ No concurrency controls
- ✓ No explicit task scheduling
- ✓ No extra library dependency



Hello-World in OpenMP

```
#include <omp.h> // OpenMP is a lang ext to describe parallelism in compiler directives
int main(){
    #omp parallel num_threads(std::thread::hardware_concurrency())
    {
        int A_B, A_C, B_D, C_D;
        #pragma omp task depend(out: A_B, A_C) ← Task dependency clauses
        {
            std::cout << "TaskA\n";
        }
        #pragma omp task depend(in: A_B; out: B_D) ← Task dependency clauses
        {
            std::cout << " TaskB\n";
        }
        #pragma omp task depend(in: A_C; out: C_D) ← Task dependency clauses
        {
            std::cout << " TaskC\n";
        }
        #pragma omp task depend(in: B_D, C_D) ← Task dependency clauses
        {
            std::cout << "TaskD\n";
        }
    }
    return 0;
}
```

*OpenMP task clauses are **static** and **explicit**;
Programmers are responsible a **proper order of writing tasks** consistent with sequential execution*

Hello-World in Intel's TBB Library

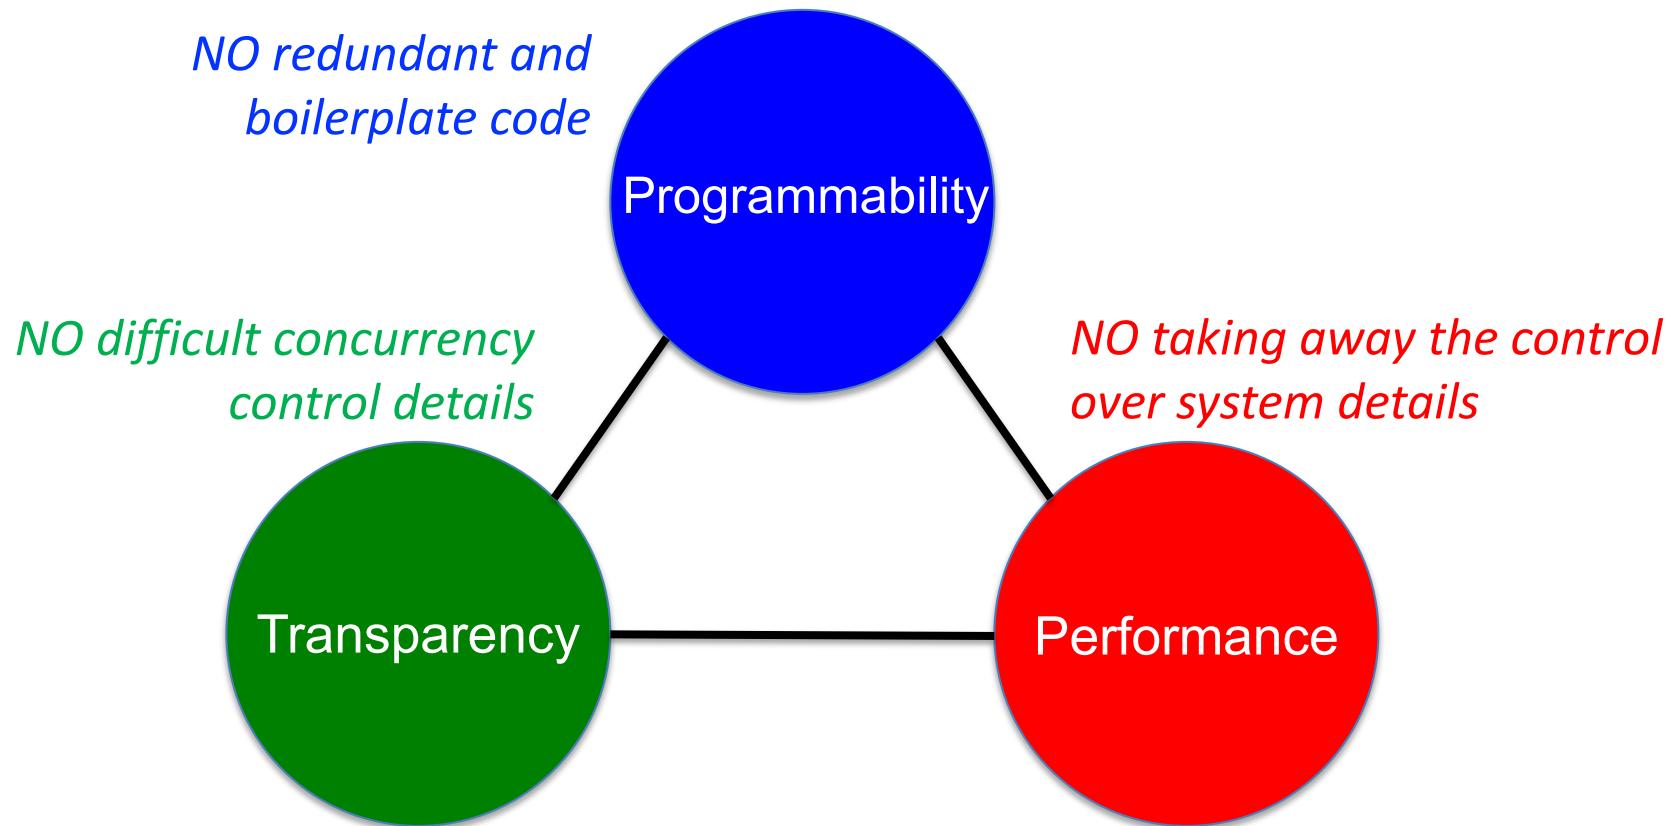
```
#include <tbb.h> // Intel's TBB is a general-purpose parallel programming library in C++
int main(){
    using namespace tbb;
    using namespace tbb::flow;
    int n = task_scheduler_init::default_num_threads();
    task_scheduler_init init(n);
    graph g;
    continue_node<continue_msg> A(g, [] (const continue msg &) {
        std::cout << "TaskA";
    });
    continue_node<continue_msg> B(g, [] (const continue msg &) {
        std::cout << "TaskB";
    });
    continue_node<continue_msg> C(g, [] (const continue msg &) {
        std::cout << "TaskC";
    });
    continue_node<continue_msg> D(g, [] (const continue msg &) {
        std::cout << "TaskD";
    });
    make_edge(A, B);
    make_edge(A, C);
    make_edge(B, D);
    make_edge(C, D);
    A.try_put(continue_msg());
    g.wait_for_all();
}
```

*Use TBB's FlowGraph
for task parallelism*

*Declare a task as a
continue_node*

TBB has excellent performance in generic parallel computing. Its drawback is mostly in the ease-of-use standpoint (simplicity, expressivity, and programmability).

Our Goal of Parallel Task Programming



“We want to let users easily express their parallel computing workload without taking away the control over system details to achieve high performance, using our expressive API in modern C++”

Accelerating DNN Training

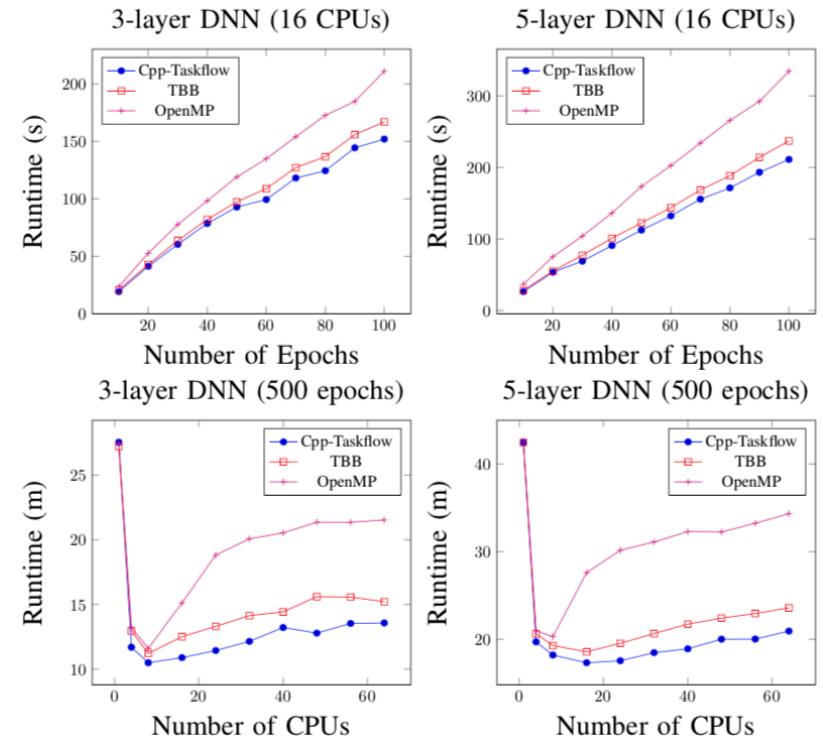
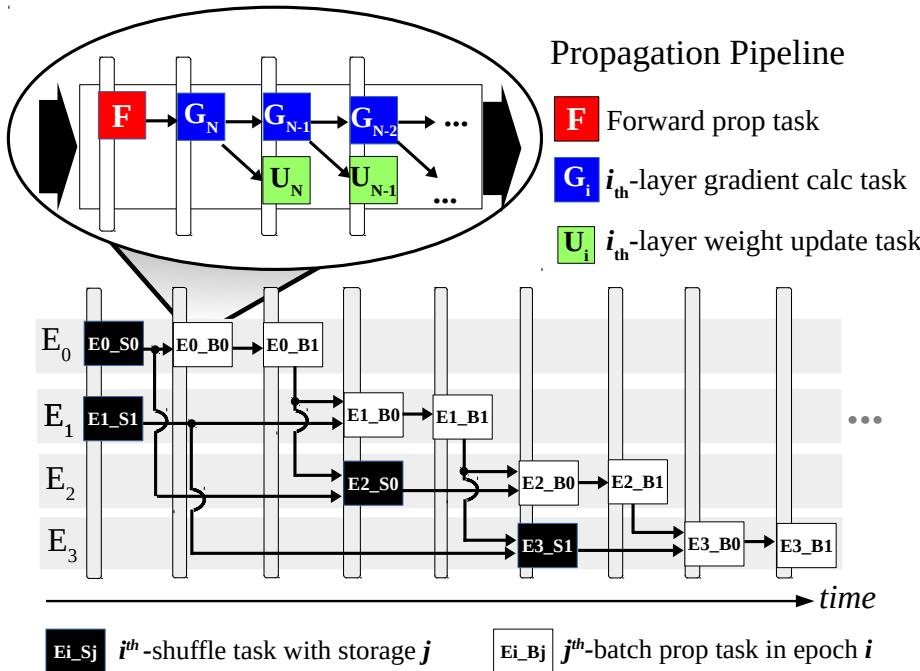
□ 3-layer DNN and 5-layer DNN image classifier

Cpp-Taskflow			OpenMP			TBB			Sequential		
LOC	CC	T	LOC	CC	T	LOC	CC	T	LOC	CC	T
59	11	3	162	23	9	90	12	3	33	9	2

CC: cyclomatic complexity of the implementation

T: development time (in hours) by an experienced programmer

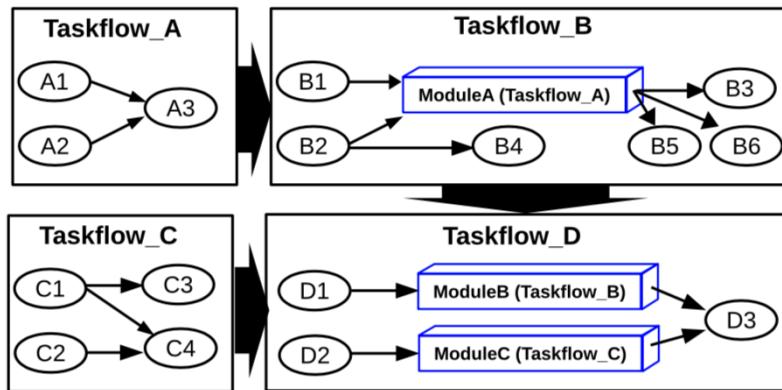
Dev time (hrs): 3 (Cpp-Taskflow) vs 9 (OpenMP)



Cpp-Taskflow is about 10%-17% faster than OpenMP and Intel TBB in avg, using the least amount of source code

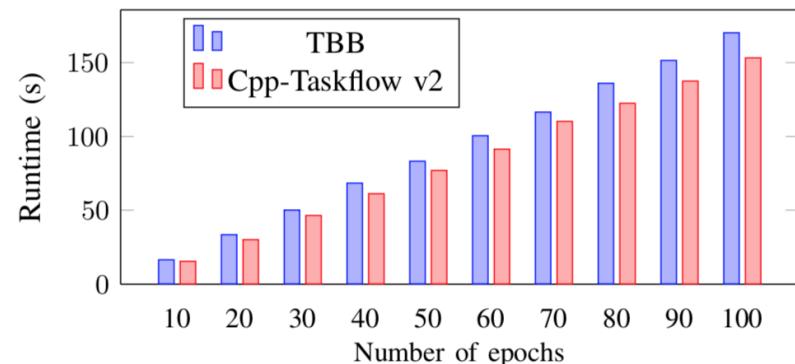
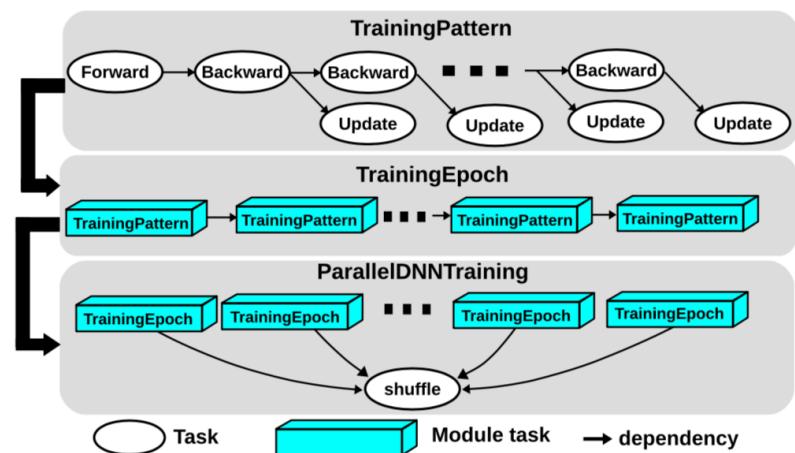
Cpp-Taskflow is Composable

- Large parallel graphs from small parallel patterns
 - Key to improving programming productivity



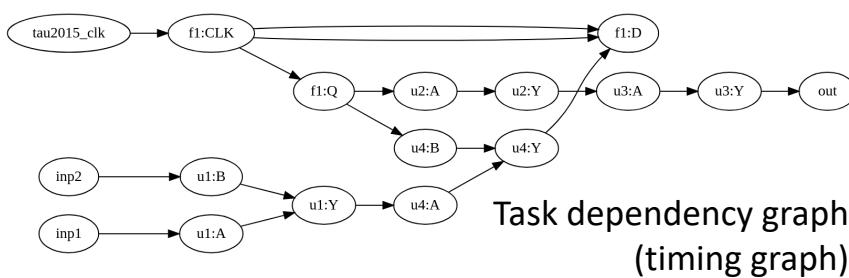
*Describes **end-to-end parallelisms** both **inside** and **outside** a machine learning workflow -> less code, more powerful, and better runtime*

22% less coding complexity and up to **40%** faster than Intel TBB in Neural Architecture Search (NAS) applications

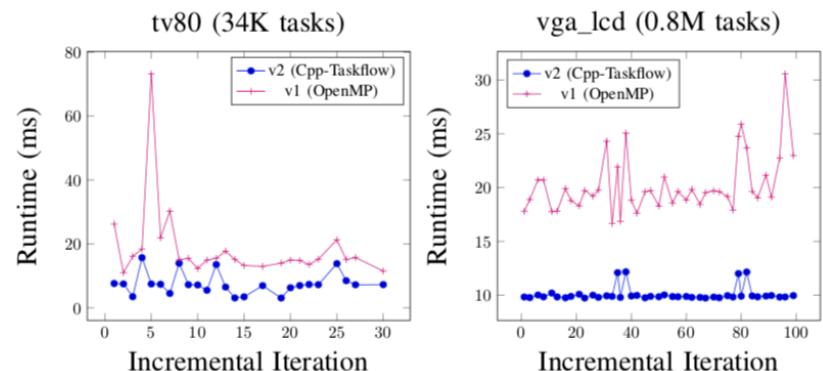


Large-Scale Graph Analytics

- OpenTimer v1: A VLSI Static Timing Analysis Tool
 - v1 first released in 2015 (open-source under GPL)
 - Loop-based parallelism using OpenMP 4.0
- OpenTimer v2: A New Parallel Incremental Timer
 - v2 first released in 2018 (open-source under MIT)
 - Task-based parallel decomposition using Cpp-Taskflow



Cpp-Taskflow saved 4K+ lines of parallel code (<https://dwheeler.com/sloccount/>)



v2 (Cpp-Taskflow) is **1.4-2x** faster than v1 (OpenMP)

Community

[Unwatch](#) ▾

157

[Unstar](#)

2.6k

[Fork](#)

286

❑ GitHub: <https://github.com/cpp-taskflow> (MIT)

- ❑ README to start with Cpp-Taskflow in just a few mins
- ❑ Doxygen-based C++ API and step-by-step tutorials
 - <https://cpp-taskflow.github.io/cpp-taskflow/index.html>
- ❑ Showcase presentation: <https://cpp-taskflow.github.io/>

Cpp-Taskflow 2.2.0

Quick Start

Modern C++ Parallel Task Programming Library

Cpp-Taskflow is a fast C++ header-only library to help you quickly write parallel programs with complex task dependencies. Cpp-Taskflow is by far faster, more expressive, fewer lines of code, and easier for drop-in integration than existing parallel task programming libraries such as OpenMP Tasking and Intel Thread Building Block (TBB) FlowGraph.

Matrix Operation Graph Traversal Matrix Operation Graph Traversal

Runtime (ms) Runtime (ms) Lines of Code Lines of Code

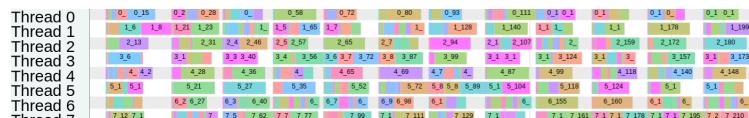
Block count ($\times 10^3$) tasks Graph Size ($\times 10^3$ tasks) Cpp-Taskflow Library Library

Generated by doxygen 1.8.13

Cpp-Taskflow is committed to support both academic and industry research projects, making it reliable and cost-effective for long-term and large-scale developments. Our users say:

↳ Quick Start
 ↳ Releases
 ↳ Master Branch (GitHub)
 ↳ Release 2.2.0 (2019/06/15)
 ↳ Release 2.1.0 (2019/02/15)
 ↳ Release 2.0.0 (2018/08/28)
 ↳ Cookbook
 ↳ C0: Project Motivation
 ↳ C1: Create a Taskflow
 ↳ C2: Execute a Taskflow
 ↳ C3: Create a Parallel For-loop Graph
 ↳ C4: Create a Parallel Reduction Graph
 ↳ C5: Spawn Task Dependency Graph
 ↳ Frequently Asked Questions
 ↳ General Questions
 ↳ Programming Questions
 ↳ Reference

Cpp-Taskflow API documentation



Cpp-Taskflow thread observer
 (profiling, debugging, testing)

*“Cpp-Taskflow has the cleanest C++ Task API I have ever seen,”
 Damien Hocking*

*“Best poster award for open-source parallel programming library,” 2018 Cpp-Conference
 (voted by 1K+ professional developers)*

Beyond Cpp-Taskflow: Heteroflow



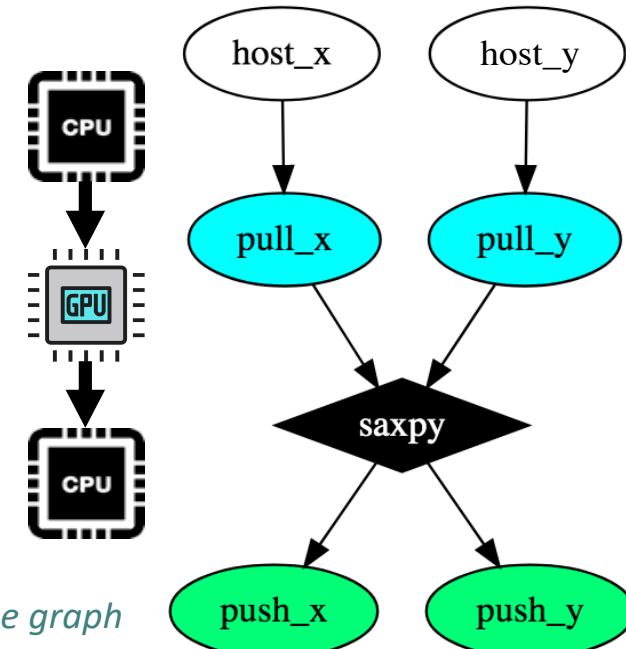
□ Concurrent CPU-GPU task programming library

```
#include <heteroflow/heteroflow.hpp>
__global__ void saxpy(int n, float a, float *x, float *y) {
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

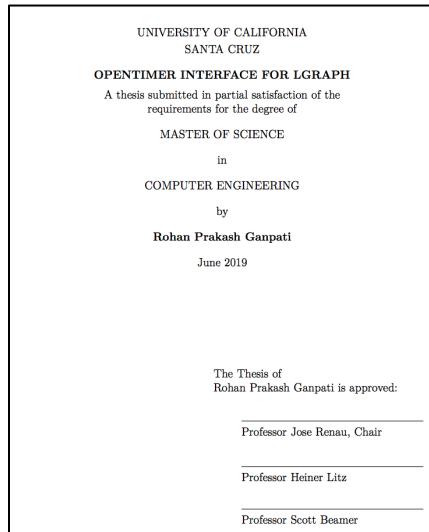
int main(void) {
    const int N = 1<<20;
    std::vector<float> x, y;
    hf::Heteroflow hf;           // create a heteroflow object
    auto host_x = hf.host([&](){ x.resize(N, 1.0f); });
    auto host_y = hf.host([&](){ y.resize(N, 2.0f); });
    auto pull_x = hf.pull(x);
    auto pull_y = hf.pull(y);
    auto kernel = hf.kernel(saxpy, N, 2.0f, pull_x, pull_y)
                  .shape((N+255)/256, 256);
    auto push_x = hf.push(pull_x, x);
    auto push_y = hf.push(pull_y, y);
    host_x.precede(pull_x);     // host_x to run before pull_x
    host_y.precede(pull_y);     // host_y to run before pull_y
    kernel.precede(push_x, push_y).succeed(pull_x, pull_y);
    hf::Executor().run(hf).wait(); // create an executor to run the graph
}
```

Only **20 lines** of code to enable parallel CPU-GPU task execution!

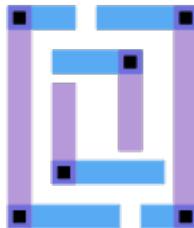
- ✓ No device memory controls
- ✓ No manual device offloading
- ✓ No explicit CPU-GPU synchronization
- ✓ No hardcoded scheduling



Thank You All (Users + Sponsors) ☺



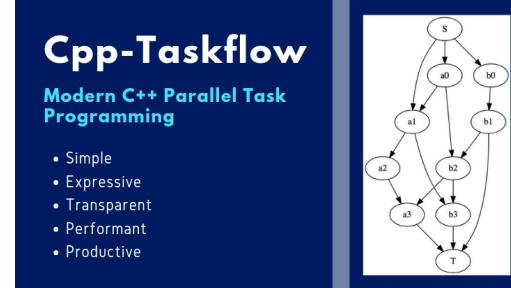
Cpp-Taskflow integration with LGraph
(master thesis by R. Ganpati @ UCSC)



Qflow Placement &
Route



VSD open-source flow



Cpp-learning's highlight
(written by Hayabusa)



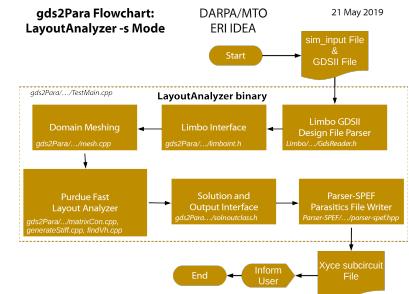
IDEA grant



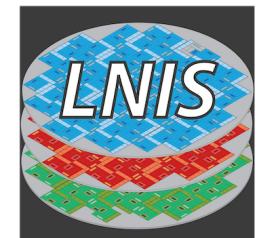
Golden timer in ACM
TAU contests



NovusCore's World of Warcraft emulator



Purdue's gds2Para



Galois
Parallel Graph Processing Systems