

# General-purpose Parallel and Distributed Programming Systems at Scale

---

Dr. Tsung-Wei Huang

Department of Electrical and Computer Engineering

University of Utah, Salt Lake City, UT



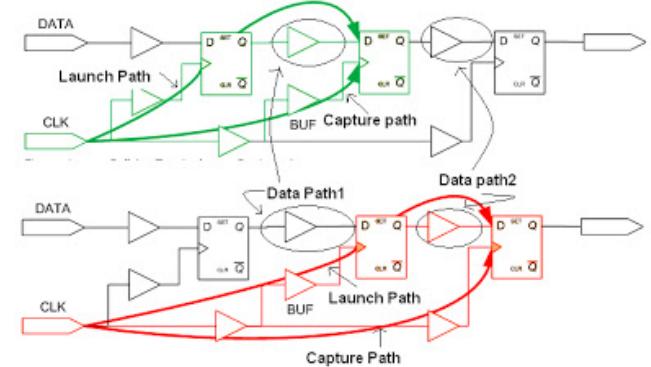
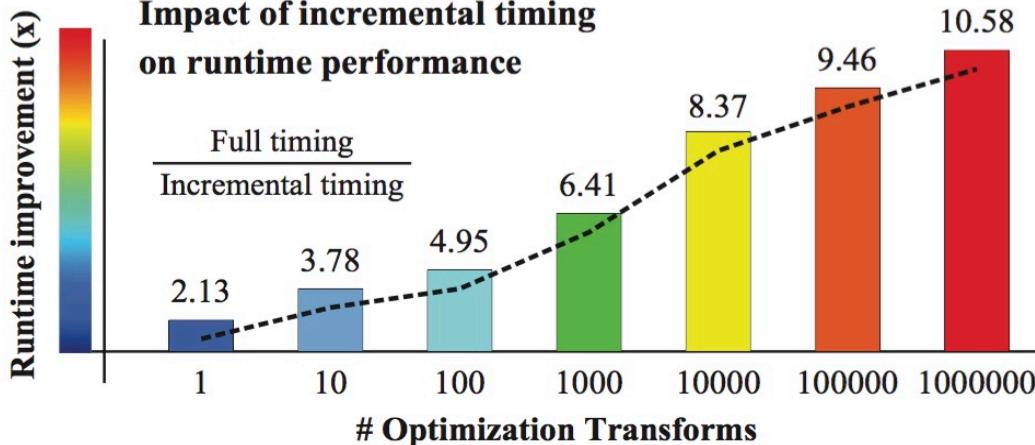
# Outline

---

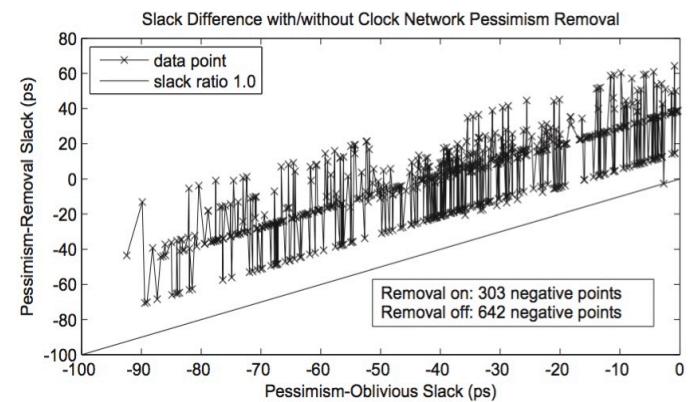
- **My early research at UIUC**
  - OpenTimer: A VLSI timing analysis tool
  - Express parallelism in the right way
- **A general-purpose distributed programming system**
  - DtCraft and its ecosystem
  - Machine learning, graph algorithms, chip design
  - Cpp-Taskflow: Modern C++ parallel task programming
- **Conclusion and future work**

# My Early Research at UIUC

- Static timing analysis (STA)
  - Key component in VLSI design
  - Verify the circuit timing
- Many challenges
  - Incremental timing
  - Parallelization, scalability



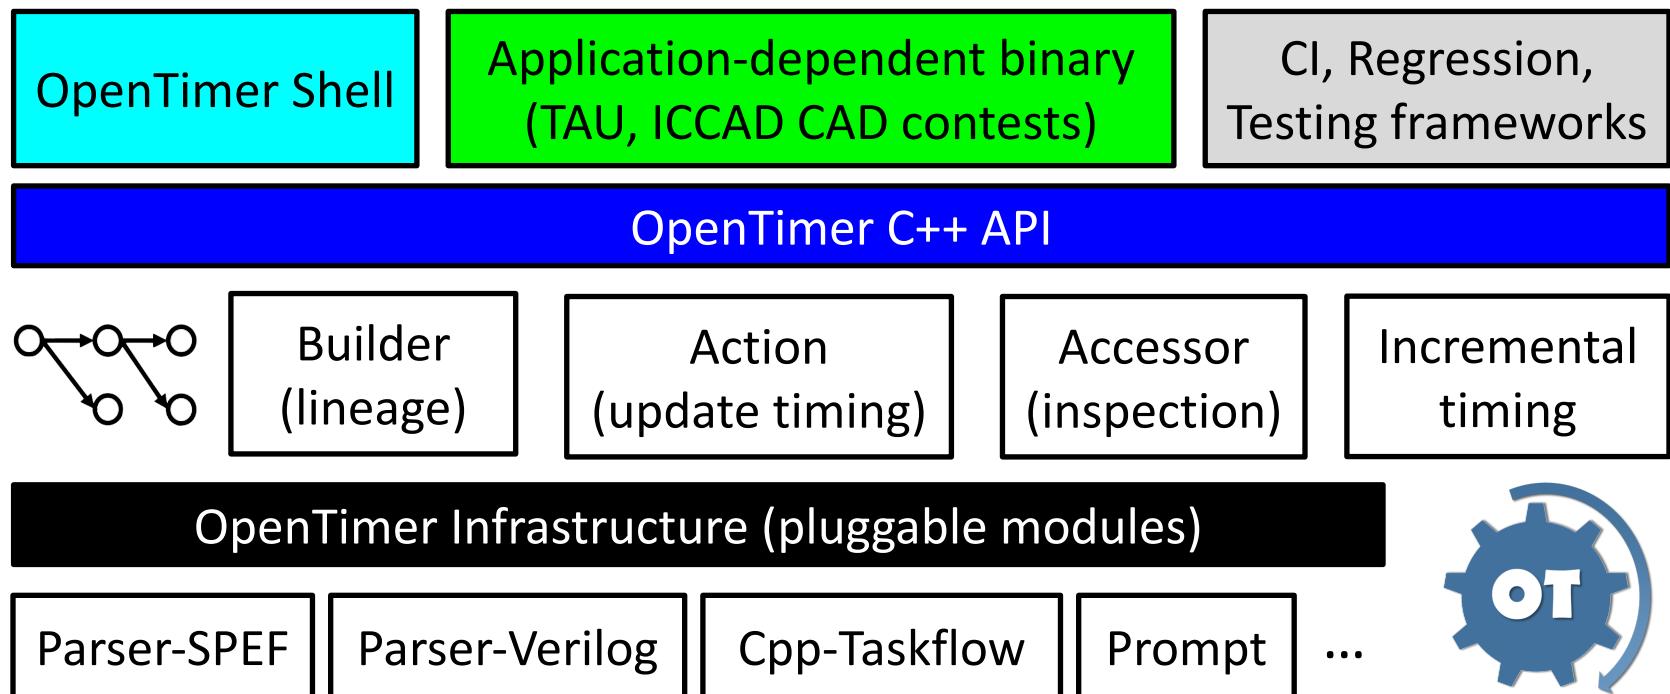
Setup/hold timing checks



Path-based timing analysis to reduce pessimism

# OpenTimer: An Open-source STA Tool

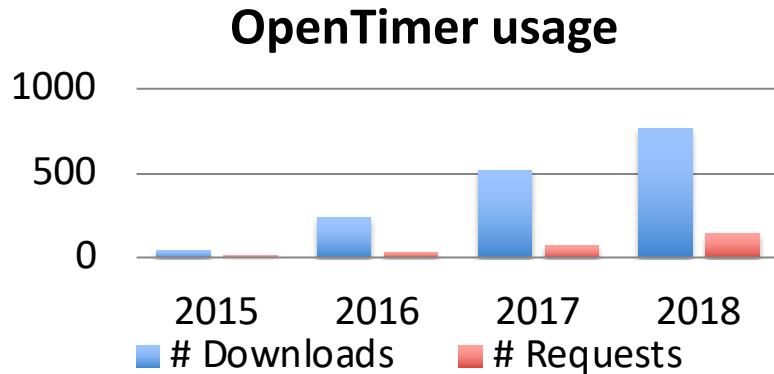
- TAU14 (1<sup>st</sup> place), TAU15 (2<sup>nd</sup> place), TAU16 (1<sup>st</sup> place)
- Best tool awards (WOSET ICCAD18, LibreCore16)
- Golden timers in VLSI and CAD communities



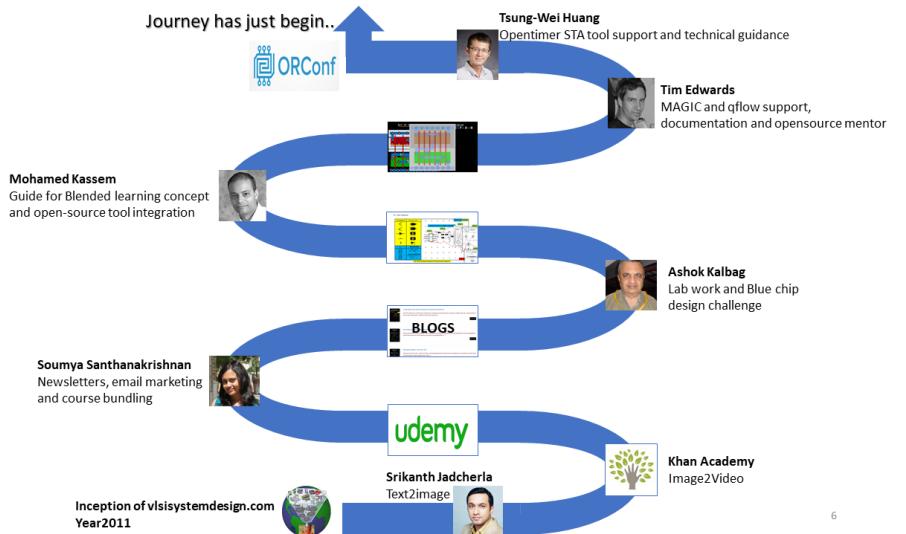
# OpenTimer Community – Thank You!



OpenTimer is begin used in many projects



We are growing! > 500 downloads in 2018



Collaboration with VSD (start-up) and Udemy online education platform



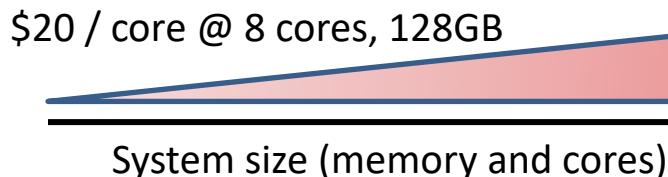
Acknowledged by ACM TAU16-19 Contests

# Collaboration with IBM

- “*We need a new **distributed** timing analysis method to deal with the large design complexity,*” IBM Timing Group, Fishkill, NY, 2015

- Existing tools are architecturally limited by one machine
- Can take **800GB RAM** and **14 days** on a single machine
- Maintaining “bare-metal” machines is not cost-efficient

Hourly Bare Metal		Monthly Bare Metal	
Single Processor	Dual Processor	Quad Processor	
Starting at \$159/mo	Starting at \$379/mo	Starting at \$1,299/mo	
2.66GHz to 2.9GHz	2.0GHz to 2.9GHz	2.0GHz to 2.7GHz	
Up to 8 cores	Up to 16 cores	Up to 40 cores	
Up to 128GB RAM	Up to 256GB RAM	Up to 512GB RAM	
Up to 16TB local storage	Up to 144TB local storage	Up to 96TB local storage	
20TB public outbound	20TB public outbound	20TB public outbound	
<a href="#">Order Now</a>	<a href="#">Order Now</a>	<a href="#">Order Now</a>	



*Objective: Execute the same # computations, distributed across an array of less expensive per unit cost cloud systems*

Cost per unit of computation

*Leverage economy of scale by decreasing the per unit cost of computation “hill”*

# Existing Tools in Cloud Computing

---

- **Hadoop**

- Distributed MapReduce platform on HDFS



- **Zookeeper**

- Coordination service for distributed application



- **Mesos**

- A high-performance cluster manager



- **Spark**

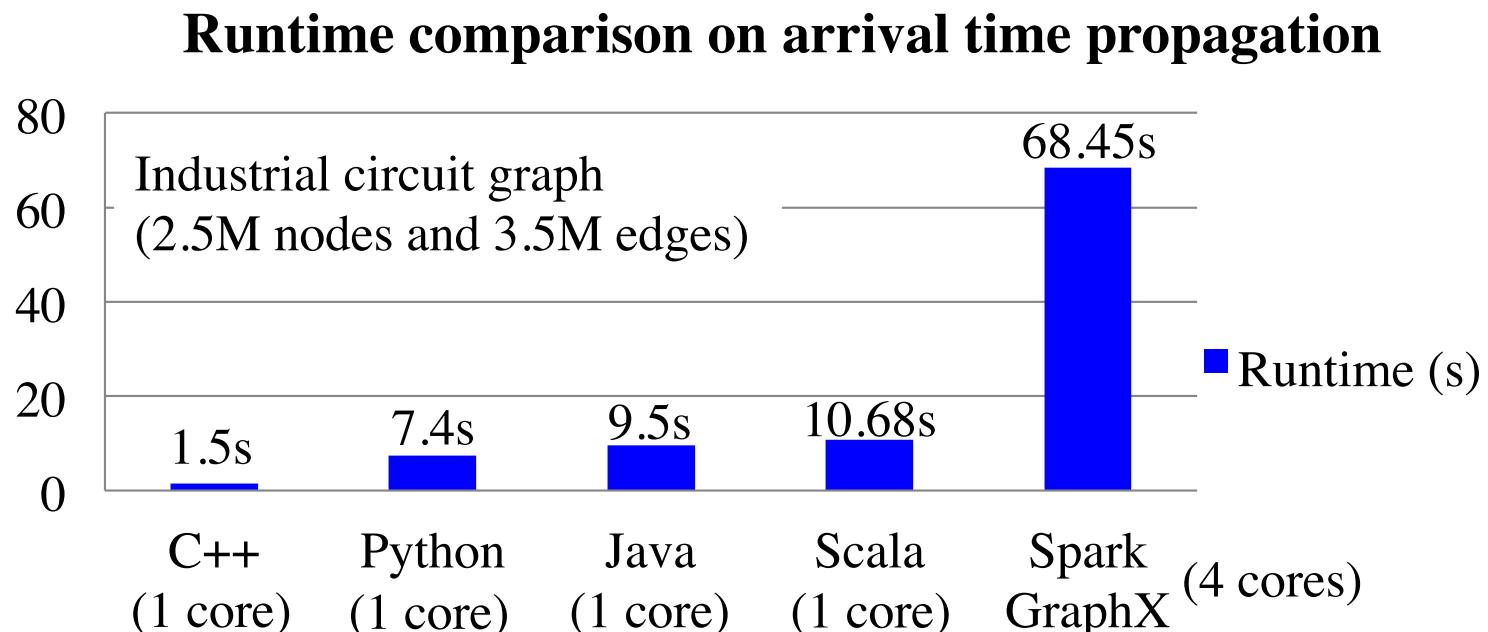
- A general computing engine for big-data analytics



- ...

***Are these tools suitable for our application?***

# Big-data Tools is NOT an Easy Fit ...

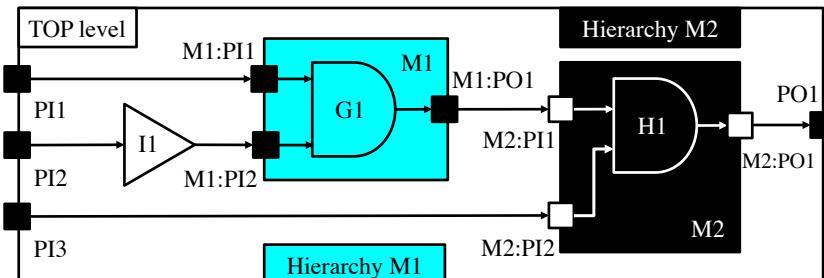


Method	Spark 2.0 (RDD + GraphX Pregel)	Java (SP)	C++ (SP)
Runtime (s)	68.45	9.5	1.50

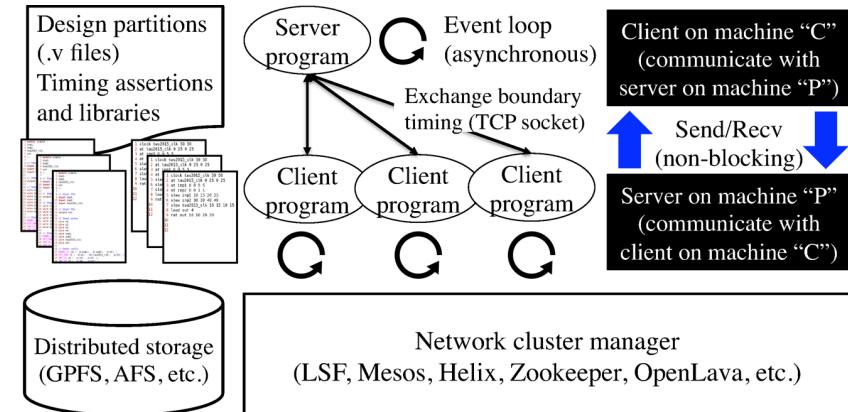
*MapReduce overhead*      *Language overhead*

# If it is Hard, let me Hard-code it ...

- A “specialized” distributed timer
  - Low-level Linux sockets and message passing code
  - Explicit resource partition and mapping
- Single-server multiple-client model
  - Server coordinates with clients to exchange data



Three partitions, top-level, M1, and M2  
(given by design teams)



T.-W. Huang, et al, “A Distributed Timing Analysis Framework for Large Designs,” IEEE/ACM DAC16  
T.-W. Huang, et al, US Patent US20170242945A1, assignee IBM, 2017  
T.-W. Huang, et al, US Patent US9836572B2, assignee IBM, 2017

# Observations

---

- Existing big-data tools have large room to improve
- Fast prototype but hard long-term maintenance
- Non-unified tooling environment
- Cumbersome software dependencies
- Cannot afford hard-coded solution all the time
- Difficult low-level concurrency controls
- Hand-written message passing code
- ...



*We want a general-purpose  
programming system to streamline  
the building of parallel/distributed  
systems!*

# DARPA's Electronic Resurgence Initiative

- ❑ \$1.5B investment to reduce the design complexity
  - ❑ People without chip design experience can design chips



Chip layout army



Massive cloud computing!



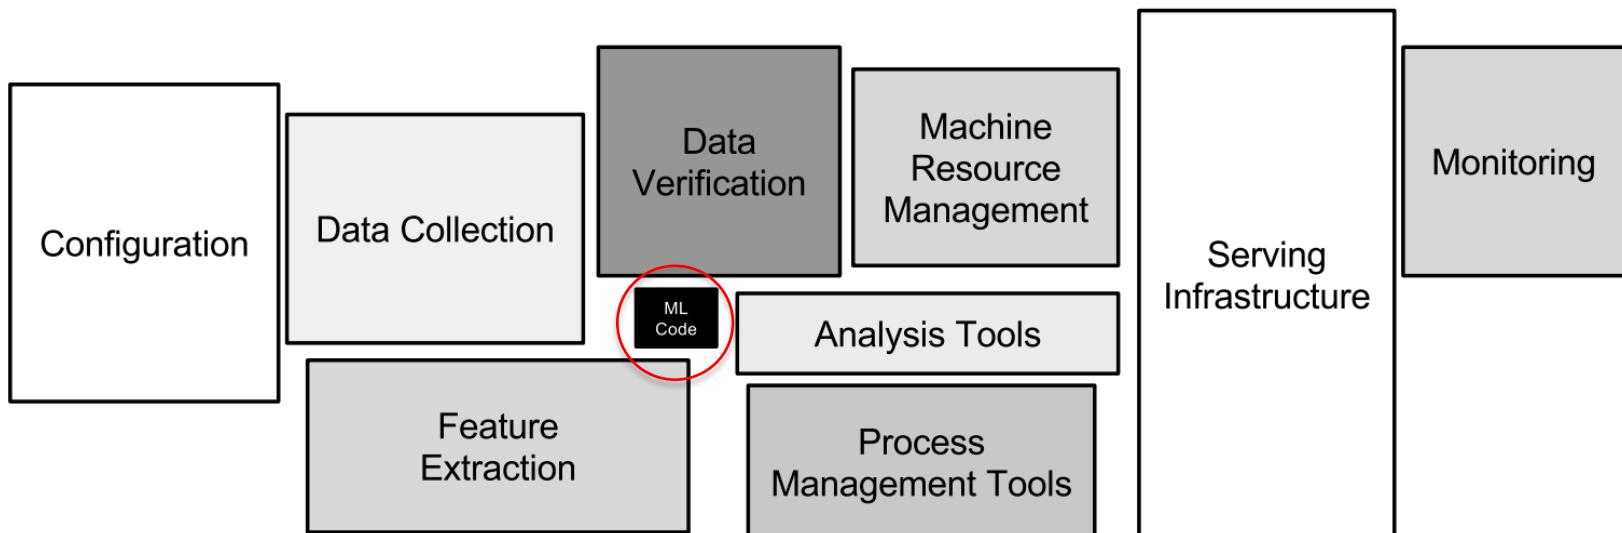
No-touch foundry

- ❑ A general-purpose silicon compiler (IDEA & POSH)
  - ❑ fully-automated layout generator, no human in the loop
  - ❑ Massive software efforts and tool integrations
  - ❑ *Parallel/distributed computing and machine learning\**

\*DARPA IDEA/POSH Kick-off Meeting, Boston, June, 2018

# Hidden Technical Debts in ML Systems

- ❑ Only a **TINY** fractions of ML systems is ML code
  - ❑ Writing ML code is cheap; system maintenance is\$\$\$\$
- ❑ ML innovations are moving to “**system innovations**”
  - ❑ Spark MLflow, TensorFlow 2.0, PyTorch, MindsDB
  - ❑ *Parallel/distributed computing, transparency, scalability*



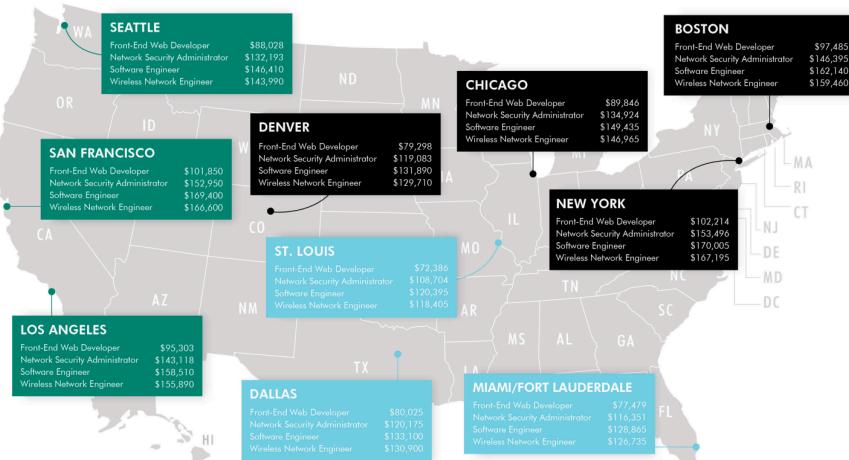
# Keep Programmability in Mind

- In the cloud era ...
  - Hardware is just a commodity
  - Building a cluster is cheap
  - Coding takes people and time



## TECH SALARIES ACROSS THE U.S.

Compare midpoint starting salaries for technology positions in different cities.\*



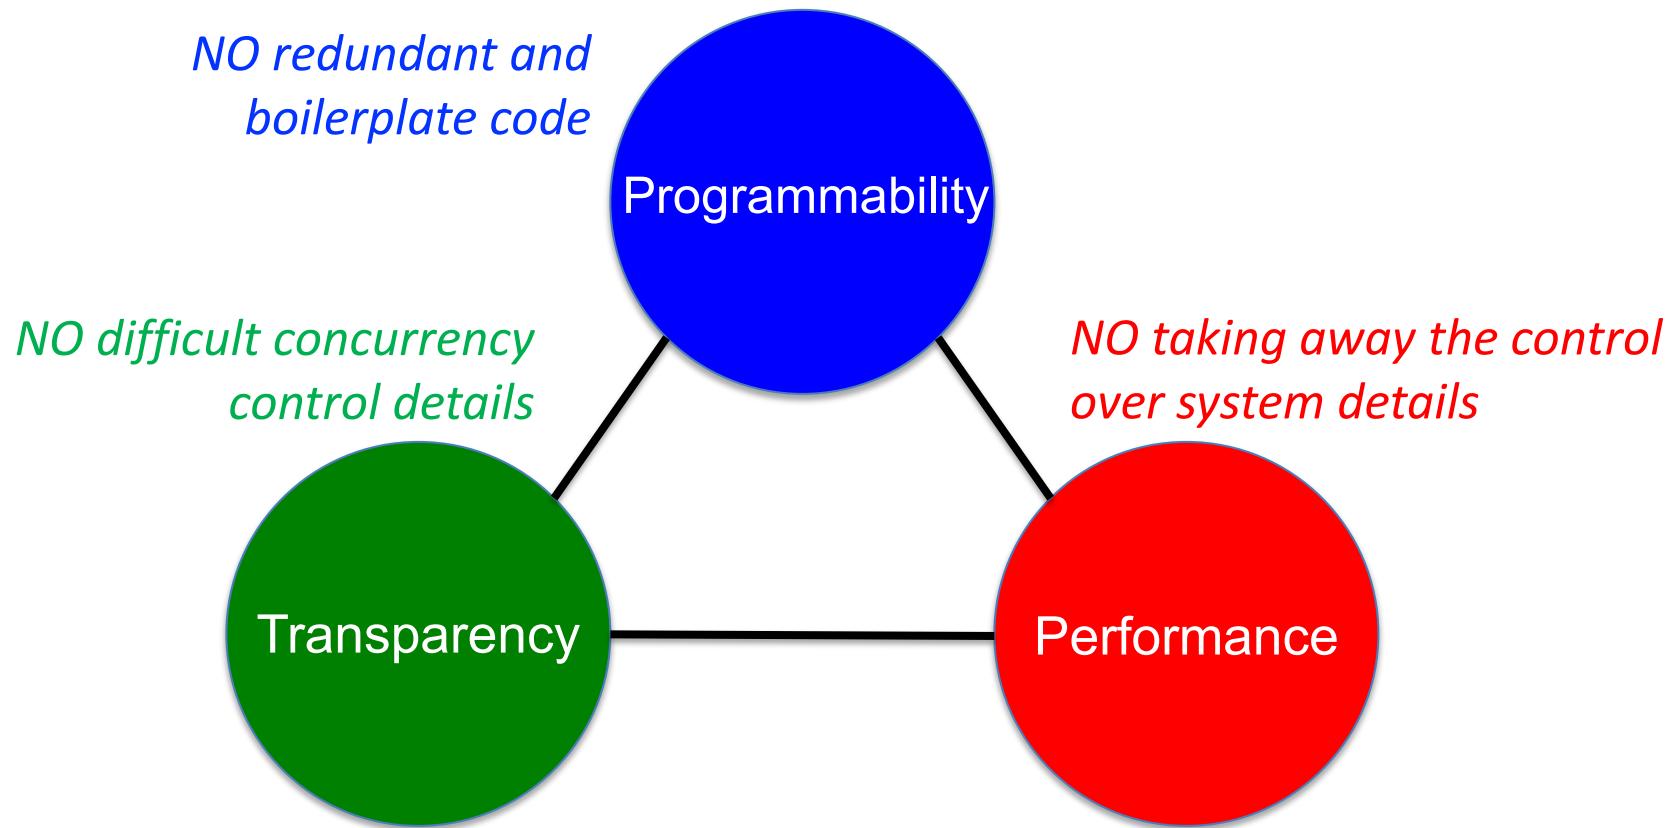
Google Cloud



*Programmability can affect the “performance” and “productivity” in all aspects (details, styles, high-level decisions, etc)!*

2018 Avg Software Engineer salary (NY) > \$170K

# Want a Unified Programming System



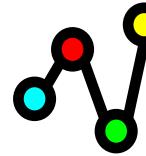
*“We want to let users easily express their parallel/distributed computing workload without taking away the control over system details to achieve high performance, using our expressive API written in modern C++”*

# Outline

---

- ❑ My early research at UIUC
  - ❑ OpenTimer: A VLSI timing analysis tool
  - ❑ Express parallelism in the right way
- ❑ A general-purpose distributed programming system
  - ❑ DtCraft and its ecosystem
  - ❑ Machine learning, graph algorithms, chip design
  - ❑ Cpp-Taskflow: Modern C++ parallel task programming
- ❑ Conclusion and future work

# A New Solution: DtCraft



Application pipeline (domain-specific applications)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



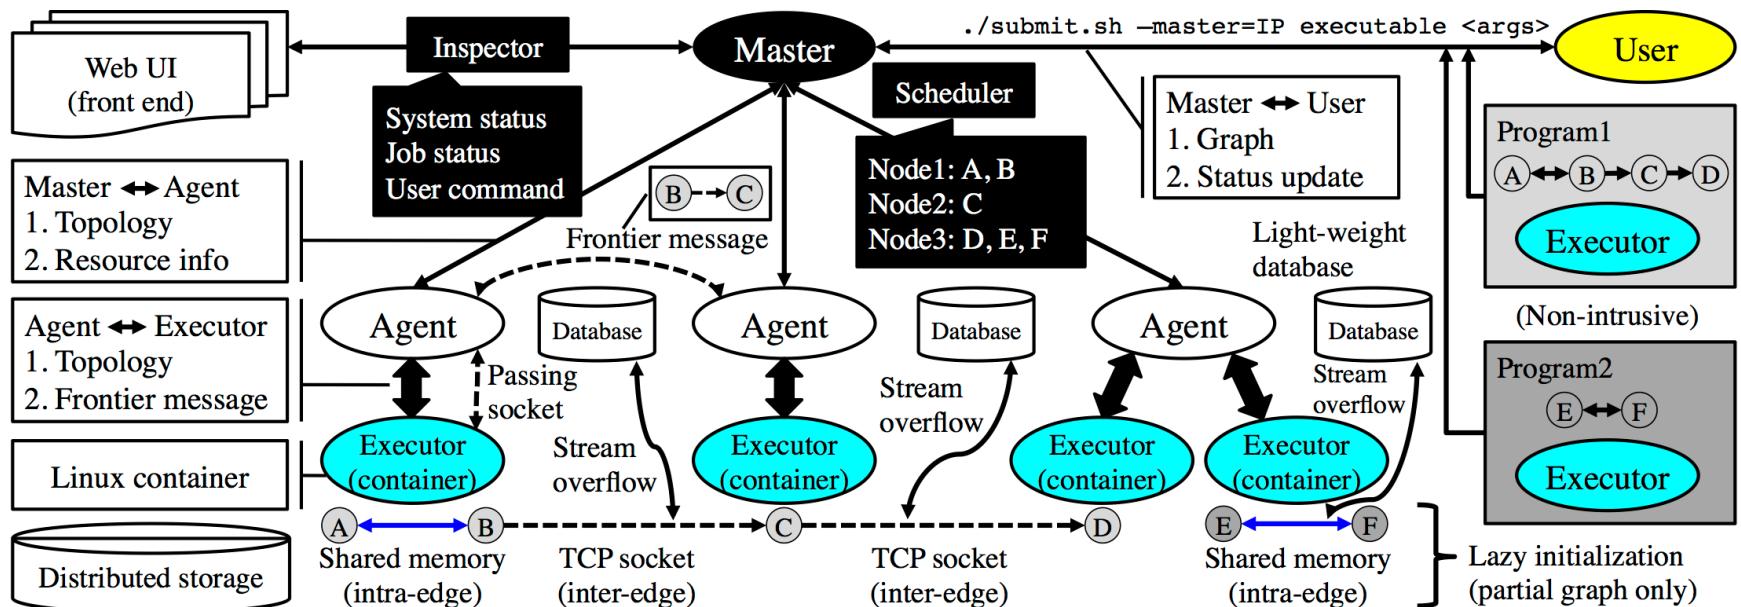
...



*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

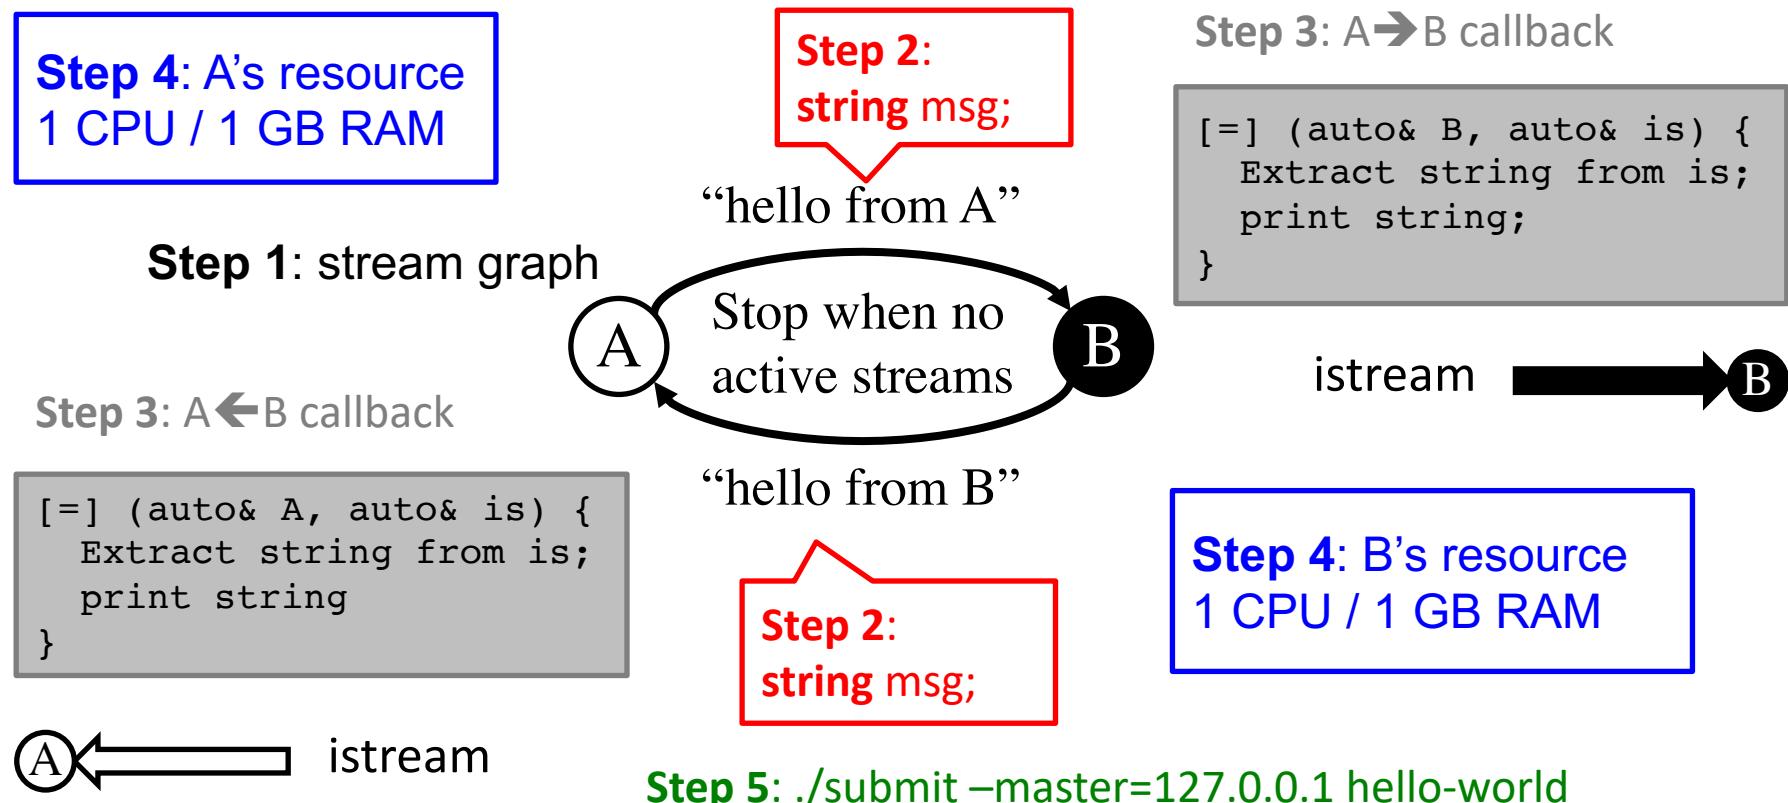
# DtCraft C++ Programming Runtime

- You describe an application in a *stream graph*
  - “No arm wrestling” with difficult concurrency details
- Everything is by default distributed
  - Same program runs on one and many machines



# A Hello-World Example

- An iterative and incremental flow
  - Two vertices + two streams



# DtCraft Code of Hello World

```
dtc::Graph G;  
  
auto A = G.vertex();  
auto B = G.vertex();  
  
G.container().add(A).cpu(1).memory(1_GB);  
  
auto AB = G.stream(A, B).on(  
[] (dtc::Vertex& B, dtc::InputStream& is) {  
    if(std::string b; is(b) != -1) {  
        dtc::cout("Received: ", b, '\n');  
        return dtc::Event::REMOVE;  
    }  
    return dtc::Event::DEFAULT;  
});  
  
auto BA = G.stream(B, A);  
  
A.on([&AB] (dtc::Vertex& v) {  
    (*v.ostream(AB))("hello world from A's");  
    dtc::cout("Sent 'hello world from A' to stream ", AB, "\n");  
});  
  
G.container().add(B).cpu(1).memory(1_GB);  
  
B.on([&BA] (dtc::Vertex& v) {  
    (*v.ostream(BA))("hello world from B's");  
    dtc::cout("Sent 'hello world from B' to stream ", BA, "\n");  
});  
  
BA.on([] (dtc::Vertex& A, dtc::InputStream& is) {  
    if(std::string a; is(a) != -1) {  
        dtc::cout("Received: ", a, '\n');  
        return dtc::Event::REMOVE;  
    }  
    return dtc::Event::DEFAULT;  
});  
  
dtc::Executor(G).run();
```

- Only a couple lines of code
- Single sequential program
- No explicit data management
- Distributed across computers
- Easy-to-use interface
- Asynchronous by default
- Scalable to many threads
- Scalable to many machines
- In-context resource controls
- Transparent concurrency
- Automatic Linux container
- ... and more

# Without DtCraft ...

```
auto count_A = 0;
auto count_B = 0;

// Send a random binary data to fd and add the
// received data to the counter.
auto pinpong(int fd, int& count) {
    auto data = random<bool>()
    auto w = write(fd, &data, sizeof(data));
    if(w == -1 && errno != EAGAIN) {
        throw system_error("Failed on write");
    }
    data = 0;
    auto r = read(fds, &data, sizeof(data));
    if(r == -1 && errno != EAGAIN) {
        throw system_error("Failed on read");
    }
    count += data;
}

int fd = -1;
std::error_code errc;                                server.cpp

if(getenv("MODE") == "SERVER") {
    fd = make_socket_server_fd("9999", errc);
}
else {
    fd = make_socket_client_fd("127.0.0.1", "9999", errc);
}

if(fd == -1) {
    throw system_error("Failed to make socket");
}
```

client.cpp

```
int make_socket_server_fd(
    std::string_view port,
    std::error_code errc
) {
    int fd {-1}
    if(fd != -1) {
        ::close(fd);
        fd = -1;
    }
    make_fd_close_on_exec(fd);
    tries = 3;
    issue_connect:
    ret = ::connect(fd, ptr->ai_addr, ptr->ai_addrlen);
    if(ret == -1) {
        if(errno == EINTR) {
            goto issue_connect;
        }
        else if(errno == EAGAIN & tries--) {
            std::this_thread::sleep_for(std::chrono::milliseconds(500));
            goto issue_connect;
        }
        else if(errno != EINPROGRESS) {
            goto try_next;
        }
        errc = make_posix_error_code(errno);
    }
    // Poll the socket. Note that writable return doesn't mean it is connected
    if(select_on_write(fd, 5, errc) && !errc) {
        int optval = -1;
        socklen_t optlen = sizeof(optval);
        if(::getsockopt(fd, SOL_SOCKET, SO_ERROR, &optval, &optlen) < 0) {
            errc = make_posix_error_code(errno);
            goto try_next;
        }
        if(optval != 0) {
            errc = make_posix_error_code(optval);
            goto try_next;
        }
    }
    // Try each address
    for(auto& ptr : addrs) {
        // Ignore if fd is already connected
        if(ptr->ai_flags & AI_CONNECTED) {
            goto try_next;
        }
        if(::bind(fd, ptr->ai_addr, ptr->ai_addrlen) == -1) {
            errc = make_posix_error_code(errno);
            goto try_next;
        }
        ::setsockopt(fd, SOL_SOCKET, SO_REUSEADDR, &optval, sizeof(optval));
        if(::bind(fd, ptr->ai_addr, ptr->ai_addrlen) == -1) {
            errc = make_posix_error_code(errno);
            goto try_next;
        }
        if(fd != -1) {
            ::close(fd);
            fd = -1;
        }
        make_fd_close_on_exec(fd);
    }
    ::freeaddrinfo(res);
    return fd;
}
```

*Branch your code to server and client  
for distributed computation!  
simple.cpp → server.cpp + client.cpp*

*A lot of boilerplate code  
plus hundred lines of  
scripts to enable  
distributed flow...*

# DtCraft to Handle Machine Learning

Application pipeline (domain-specific applications)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



...

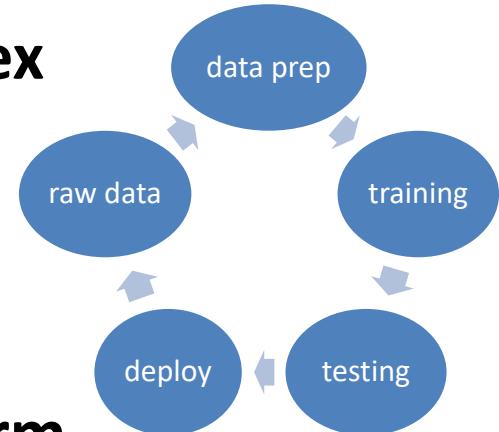


*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# MLCraft – Machine Learning at Scale

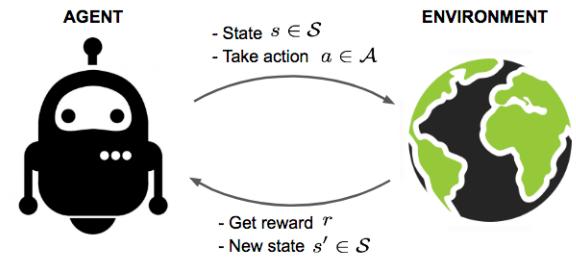
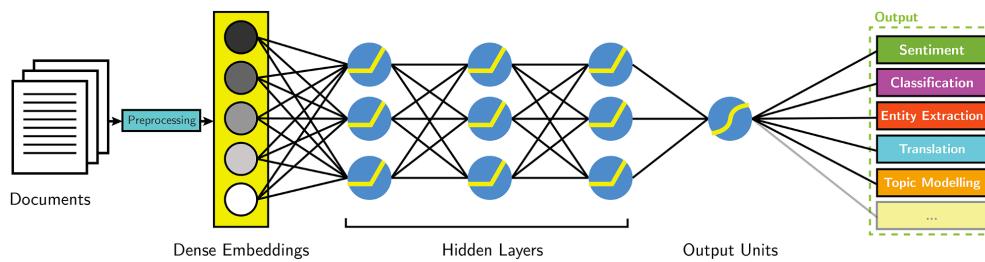
## □ ML system development is very complex

- Hard to track experiments
- Hard to reproduce results
- Hard to deploy and maintain



## □ MLCraft aims to be a unified ML platform

- Automate ML workflows and developments
- Open API and interface to work with existing libraries
- More robust, predictable, easier to maintain



```

dtc::Graph G;

auto src = G.insert<dtc::cell::MnistStreamFeeder>(
    mnist_image_file,
    mnist_label_file,
    [] (Eigen::MatrixXf& images, Eigen::VectorXi& labels) {
        images /= 255.0f;
        return std::make_tuple(images, labels);
    }
);

src.window(1ms).frequency(1000);

auto dnn = G.insert<dtc::cell::Visitor1x1>(
    [] (dtc::ml::DnnClassifier& c) {
        dtc::cout("Creating a dnn classifier [784x30x10]\n").flush();
        c.fully_connected_layer(784, 30, dtc::ml::Activation::RELU)
            .fully_connected_layer(30, 10, dtc::ml::Activation::NONE)
            .optimizer<dtc::ml::AdamOptimizer>();
    },
    [i=0] (dtc::ml::DnnClassifier& c, std::tuple<Eigen::MatrixXf, Eigen::VectorXi>& data) mutable {
        auto& [images, labels] = data;
        dtc::cout(
            "Accuracy at training cycle ", i++, " [", images.rows(), " images]: ",
            ((labels-c.infer(images)).array() == 0).count() / static_cast<float>(images.rows()), "%"
        ).flush();
        c.train(images, labels, 10, 64, 0.01f, [](){});
    }
);

dnn.in(src.out());

G.container().add(dnn);
G.container().add(src);

dtc::Executor(G).run();

```

Only “60-line” code to  
automate a distributed  
ML workflow!

# DtCraft to Handle Distributed Timing

Application pipeline (Distributed Timing Analysis)

MLCraft  
(distributed machine learning library)

GraphCraft  
(distributed graph processing library)

...

DataCraft  
(distributed data processing server)

DtCraft C++ Programming Runtime

Cpp-Taskflow  
(Parallel task programming library)

Cpp-Container  
(Docker, RunC, LibContainer, LXC)



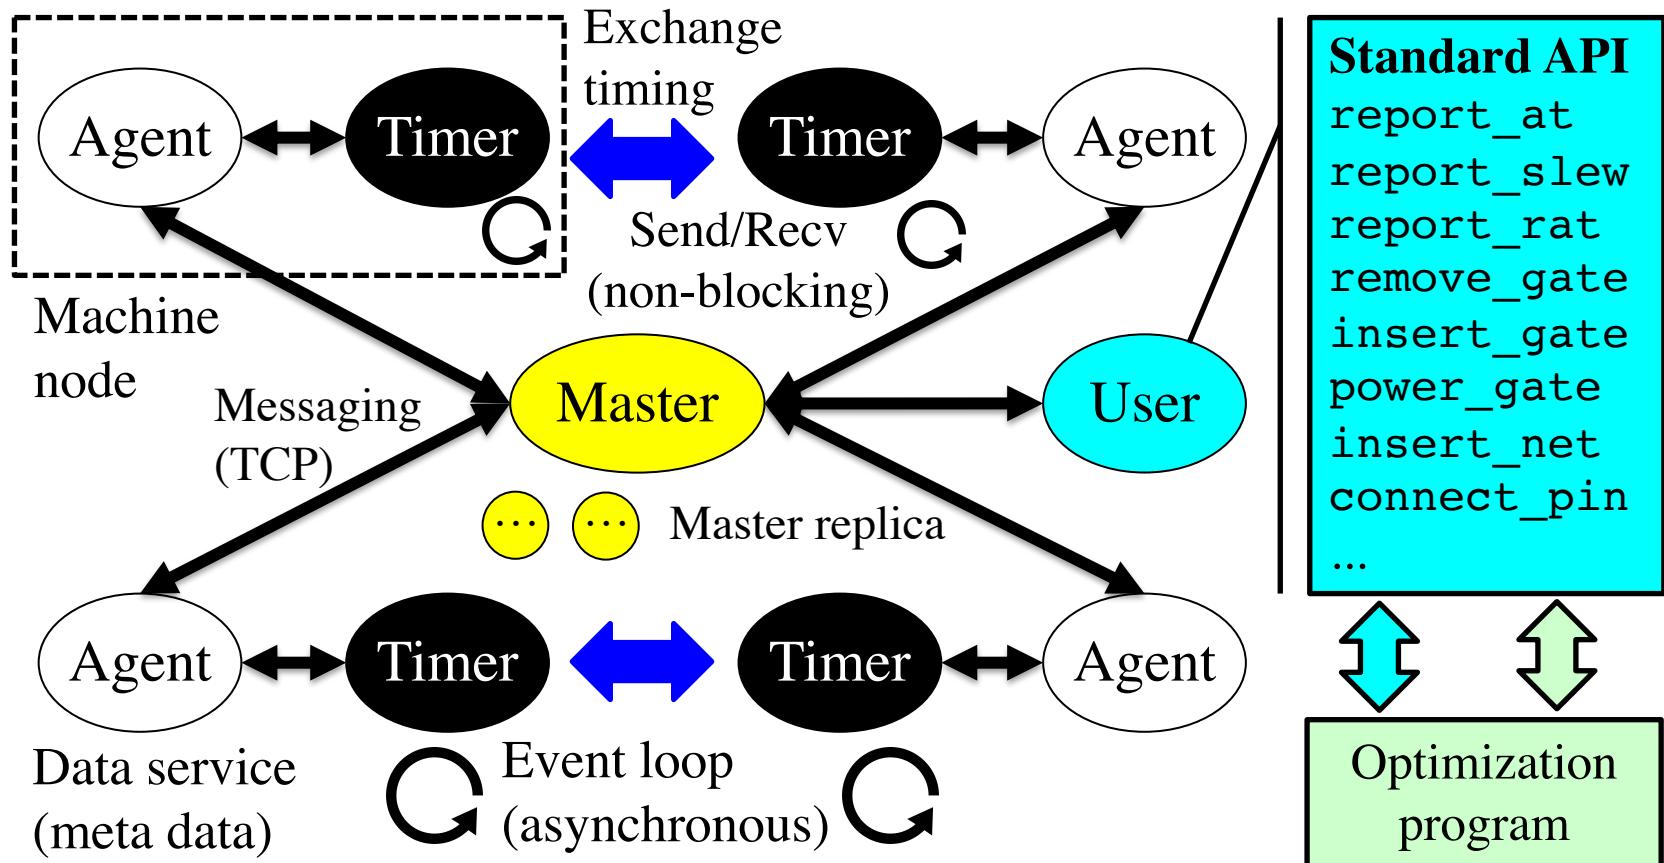
...



DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)

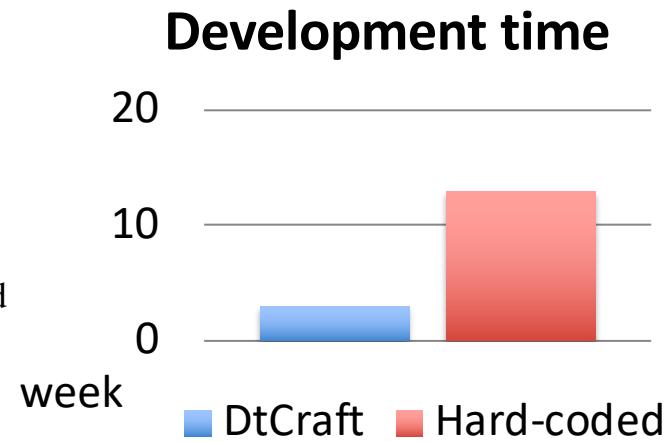
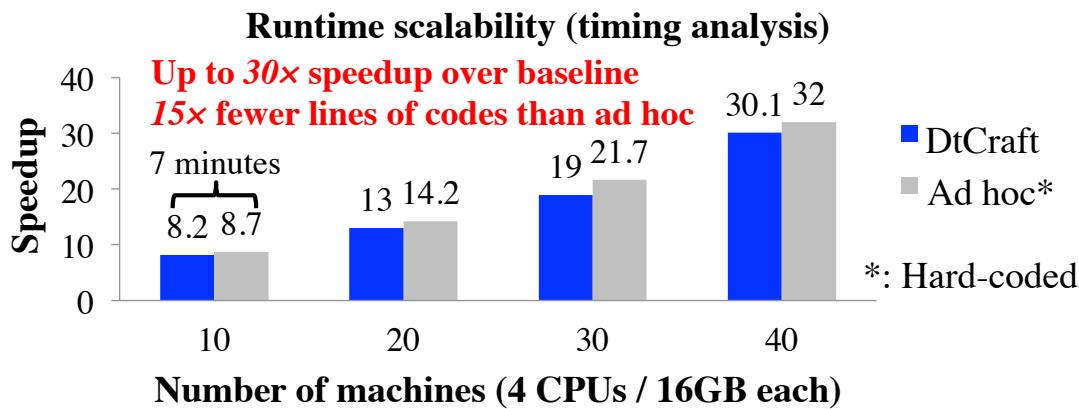
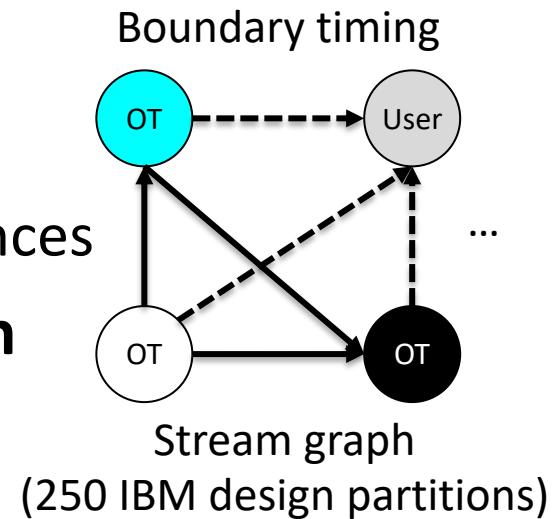
# Distributed Timing Analysis

## □ Master coordinator and worker (a timer per partition)



# Distributed Timing Performance

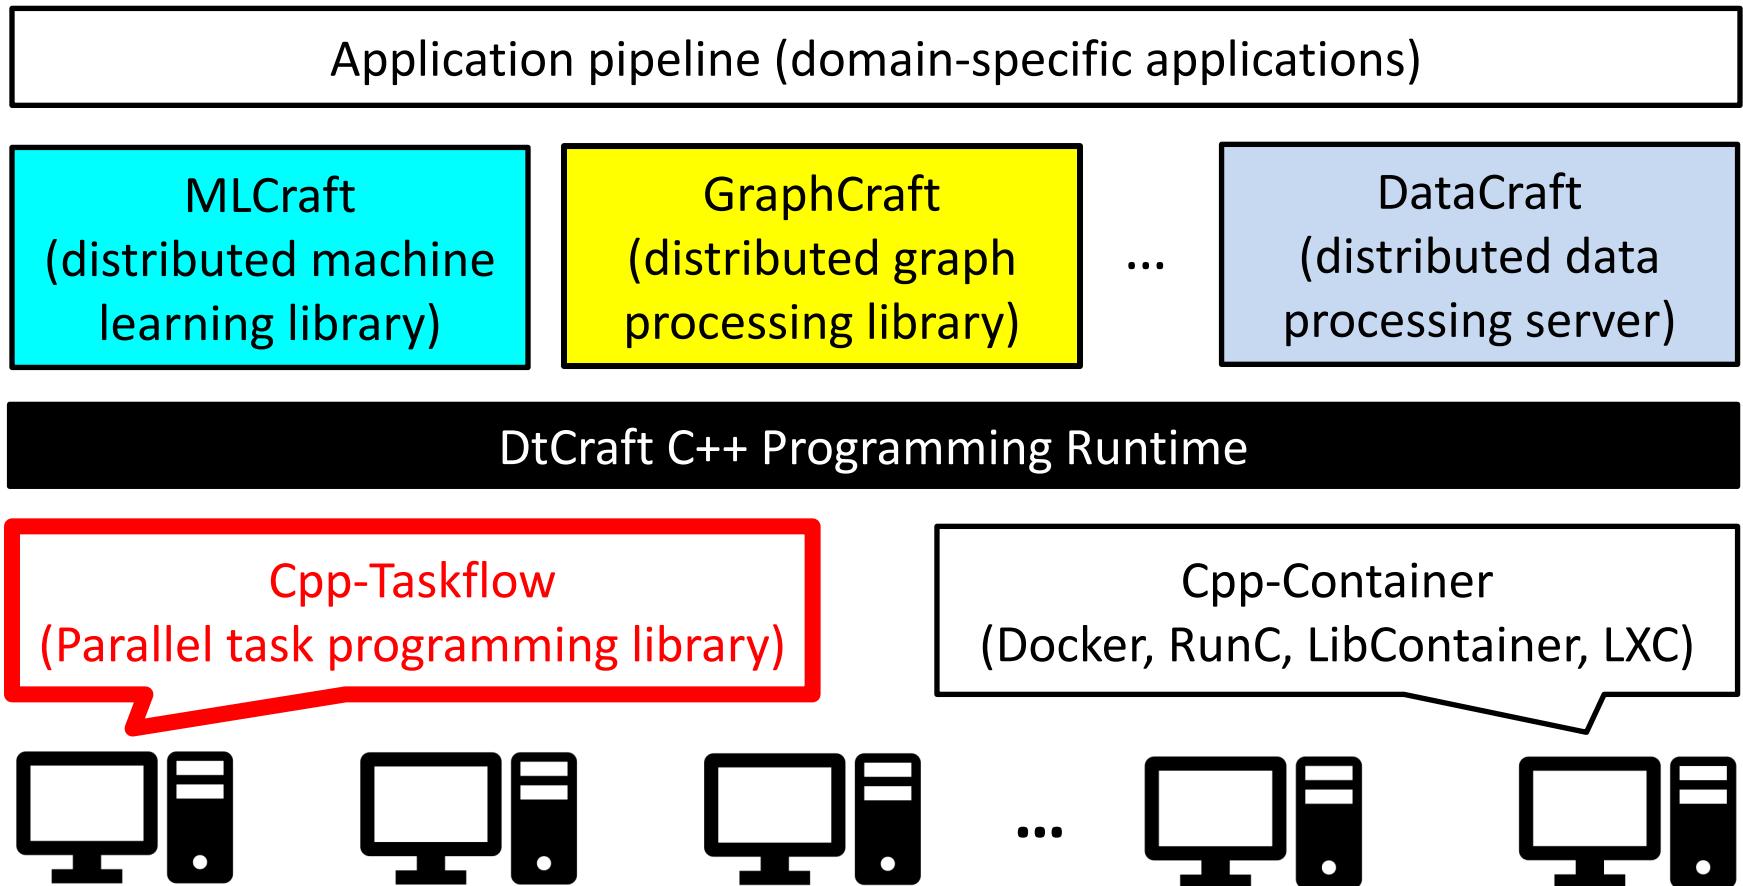
- ❑ IBM design of 250 partitions
  - ❑ OpenTimer as the timing analyzer
  - ❑ On a cluster of 40 Amazon EC2 instances
- ❑ Compared with hard-coded solution
  - ❑ 15x fewer lines of code
  - ❑ Only 7% performance loss





*Improving the programmability of a system is not only making programming easier and fast, but is enabling a new innovation to solve problems more efficiently*

# DtCraft to Handle Low-level Concurrency



*DARPA IDEA Grant, NSF Grant, Best Open-source Tool Awards (IEEE IPDPS, ACM MM, WOSET, CPPConf, ACM DEBS, ACM/IEEE DAC, IEEE/ACM ICCAD, IEEE TCAD, ACM TAU, US Patents, etc.)*

# Cpp-Taskflow's Project Mantra

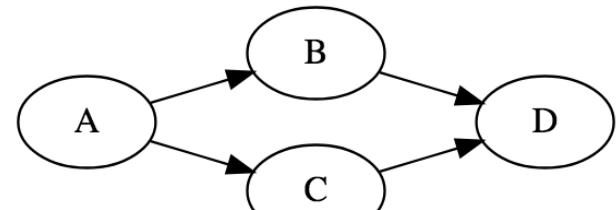
A programming library helps developers **quickly** write **efficient** parallel programs on a **manycore machine** using **task-based** approaches in **modern C++**

- Task-based approach scales best with multicore arch
  - We should write tasks instead of threads
  - Not trivial due to dependencies (race, lock, bugs, etc)
- We want developers to write parallel code that is:
  - Simple, expressive, and transparent
- We don't want developers to manage:
  - Explicit thread management
  - Difficult concurrency controls and daunting class objects

# Hello-World in Cpp-Taskflow

```
#include <taskflow/taskflow.hpp> // Cpp-Taskflow is header-only
int main(){
    tf::Taskflow tf;
    auto [A, B, C, D] = tf.emplace(
        [] () { std::cout << "TaskA\n"; }
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B);           // A runs before B
    A.precede(C);           // A runs before C
    B.precede(D);           // B runs before D
    C.precede(D);           // C runs before D
    tf::Executor().run(tf); // create an executor to run the taskflow
    return 0;
}
```

Only **15 lines** of code to get a parallel task execution!



# Hello-World in OpenMP

```
#include <omp.h> // OpenMP is a lang ext to describe parallelism in compiler directives
int main(){
    #omp parallel num_threads(std::thread::hardware_concurrency())
    {
        int A_B, A_C, B_D, C_D;
        #pragma omp task depend(out: A_B, A_C) ← Task dependency clauses
        {
            std::cout << "TaskA\n";
        }
        #pragma omp task depend(in: A_B; out: B_D) ← Task dependency clauses
        {
            std::cout << " TaskB\n";
        }
        #pragma omp task depend(in: A_C; out: C_D) ← Task dependency clauses
        {
            std::cout << " TaskC\n";
        }
        #pragma omp task depend(in: B_D, C_D) ← Task dependency clauses
        {
            std::cout << "TaskD\n";
        }
    }
    return 0;
}
```

*OpenMP task clauses are **static** and **explicit**;  
Programmers are responsible a **proper order of writing tasks** consistent with sequential execution*

# Hello-World in Intel's TBB Library

```
#include <tbb.h> // Intel's TBB is a general-purpose parallel programming library in C++
int main(){
    using namespace tbb;
    using namespace tbb::flow;
    int n = task_scheduler_init::default_num_threads();
    task_scheduler_init init(n);
    graph g;
    continue_node<continue_msg> A(g, [] (const continue msg &) {
        std::cout << "TaskA";
    });
    continue_node<continue_msg> B(g, [] (const continue msg &) {
        std::cout << "TaskB";
    });
    continue_node<continue_msg> C(g, [] (const continue msg &) {
        std::cout << "TaskC";
    });
    continue_node<continue_msg> D(g, [] (const continue msg &) {
        std::cout << "TaskD";
    });
    make_edge(A, B);
    make_edge(A, C);
    make_edge(B, D);
    make_edge(C, D);
    A.try_put(continue_msg());
    g.wait_for_all();
}
```

*Use TBB's FlowGraph  
for task parallelism*

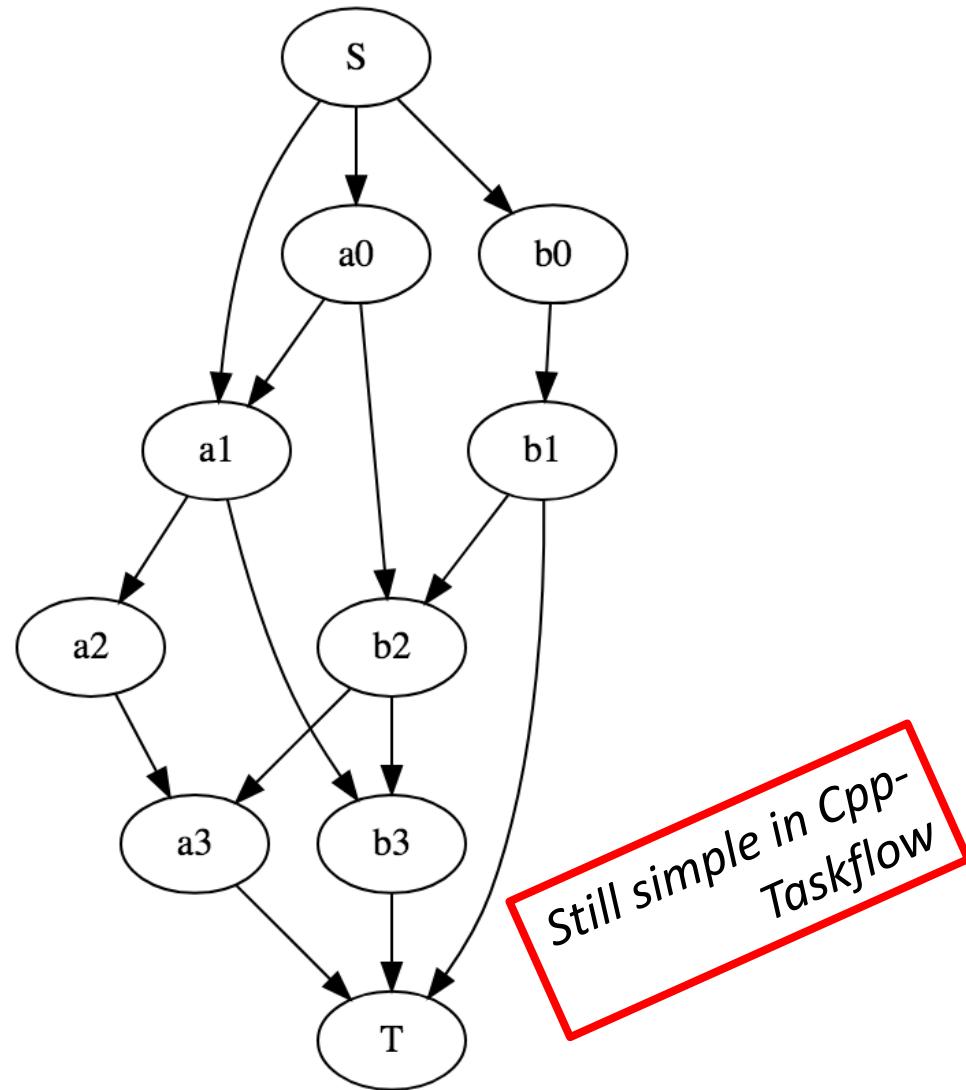
*Declare a task as a  
continue\_node*

*TBB has excellent performance in generic parallel computing. Its drawback is mostly in the ease-of-use standpoint (simplicity, expressivity, and programmability).*

*Somehow, this looks more like “hello universe” ...*

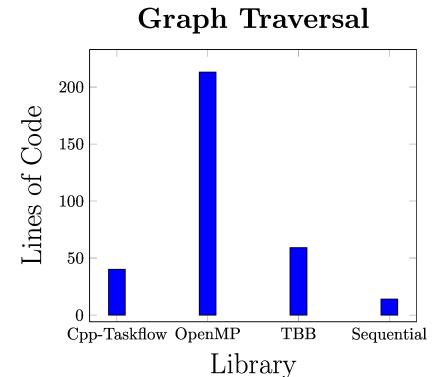
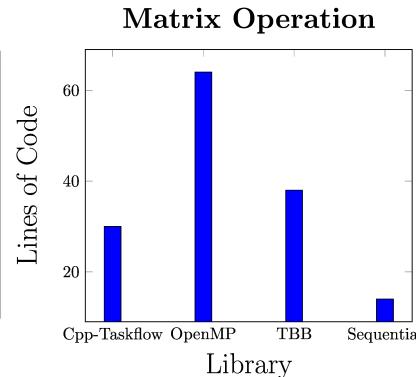
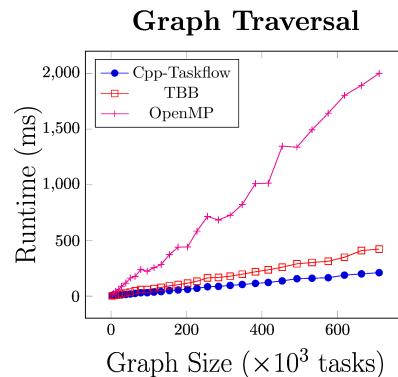
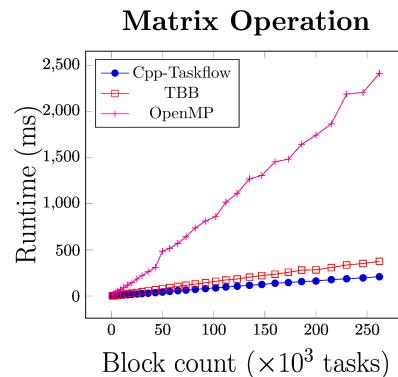
# A Slightly More Complicated Example

```
// source dependencies
S.precede(a0);    // S runs before a0
S.precede(b0);    // S runs before b0
S.precede(a1);    // S runs before a1
// a_ -> others
a0.precede(a1);   // a0 runs before a1
a0.precede(b2);   // a0 runs before b2
a1.precede(a2);   // a1 runs before a2
a1.precede(b3);   // a1 runs before b3
a2.precede(a3);   // a2 runs before a3
// b_ -> others
b0.precede(b1);   // b0 runs before b1
b1.precede(b2);   // b1 runs before b2
b2.precede(b3);   // b2 runs before b3
b2.precede(a3);   // b2 runs before a3
// target dependencies
a3.precede(T);    // a3 runs before T
b1.precede(T);    // b1 runs before T
b3.precede(T);    // b3 runs before T
```



# Micro-benchmark Performance

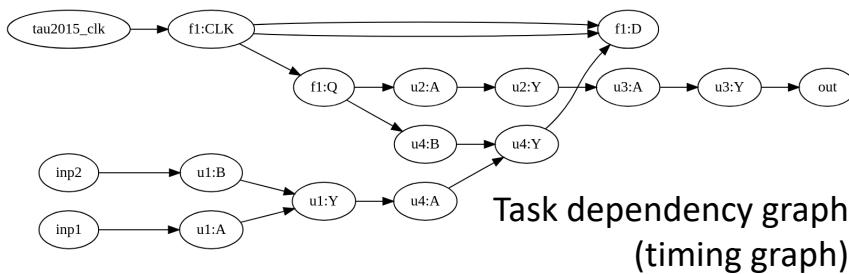
- Measured the “pure” tasking performance
  - Wavefront computing (regular compute pattern)
  - Graph traversal (irregular compute pattern)
  - Compared with OpenMP 4.5 and Intel TBB FlowGraph
    - C++ v8 with -fomp -O2 -std=c++17
    - Evaluated on a 4-core AMD CPU machine



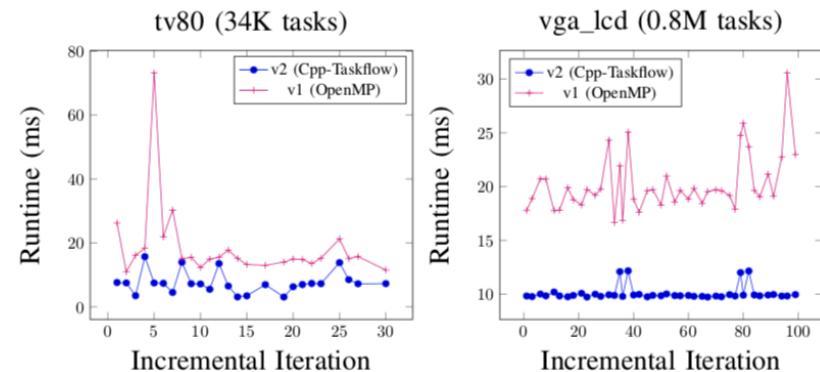
Cpp-Taskflow scales the best when task counts (problem size) increases, using the least amount of code

# Large-Scale VLSI Timing Analysis

- OpenTimer v1: A VLSI Static Timing Analysis Tool
  - v1 first released in 2015 (open-source under GPL)
  - Loop-based parallelism using OpenMP 4.0
- OpenTimer v2: A New Parallel Incremental Timer
  - v2 first released in 2018 (open-source under MIT)
  - Task-based parallel decomposition using Cpp-Taskflow



**Cost to develop is \$275K with OpenMP  
vs \$130K with Cpp-Taskflow!**  
(<https://dwheeler.com/sloccount/>)



v2 (Cpp-Taskflow) is 1.4-2x faster than  
v1 (OpenMP)

# Deep Learning Model Training

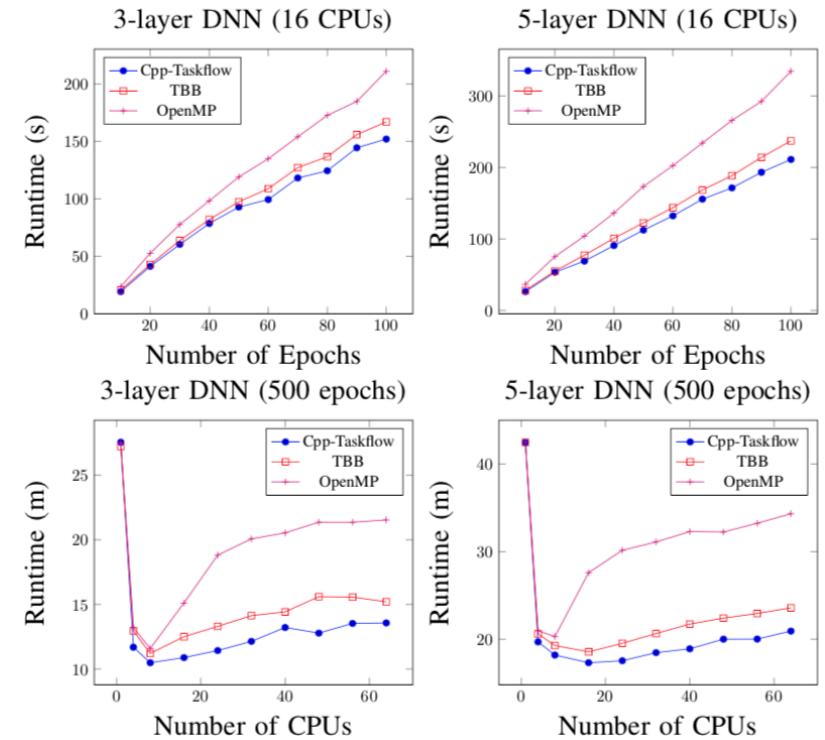
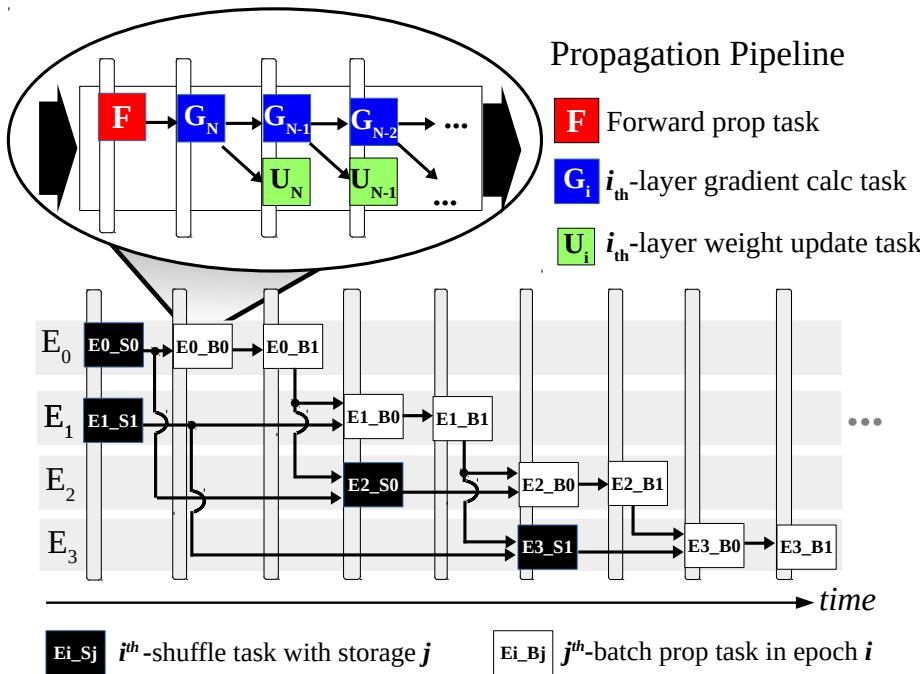
## □ 3-layer DNN and 5-layer DNN image classifier

Cpp-Taskflow			OpenMP			TBB			Sequential		
LOC	CC	T	LOC	CC	T	LOC	CC	T	LOC	CC	T
59	11	3	162	23	9	90	12	3	33	9	2

CC: cyclomatic complexity of the implementation

T: development time (in hours) by an experienced programmer

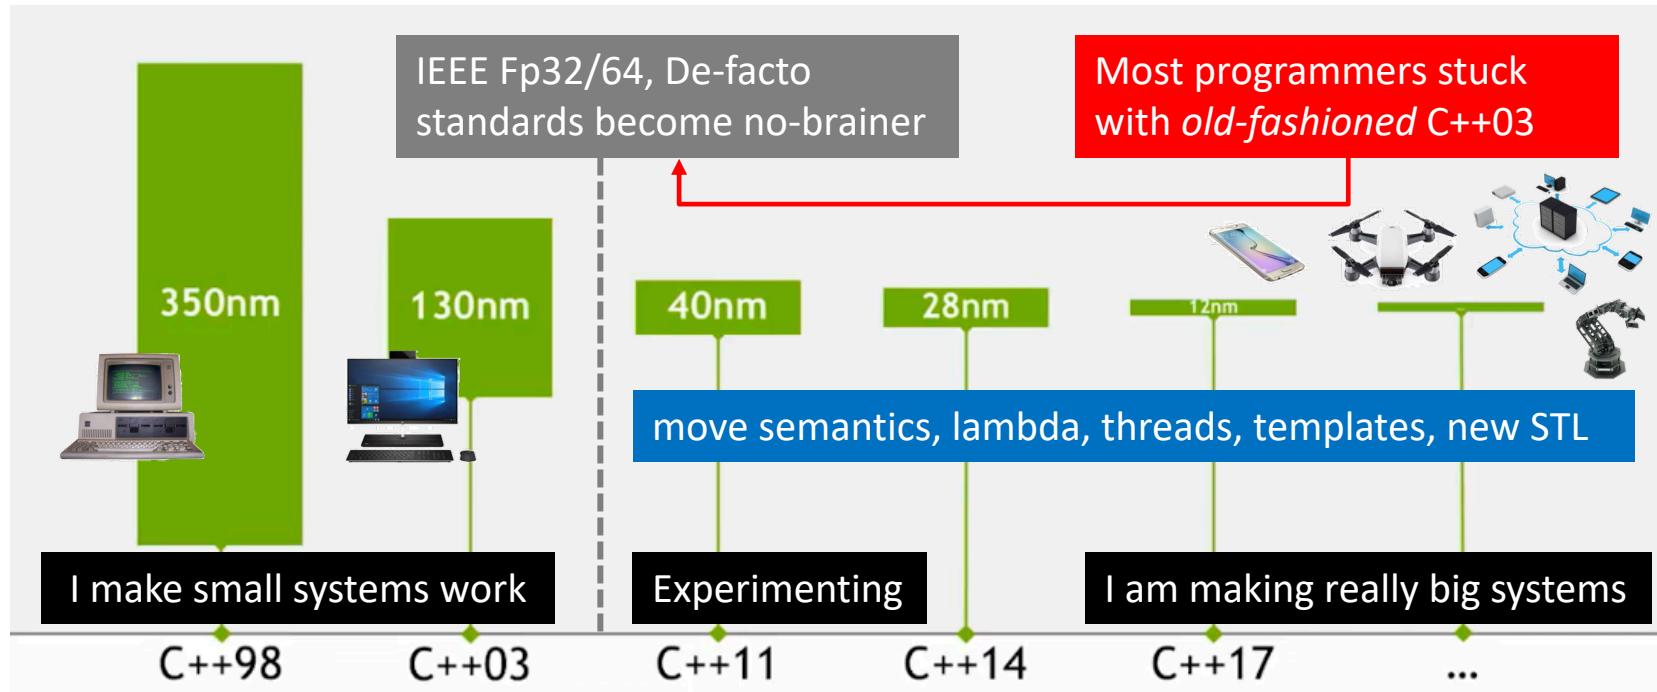
Dev time (hrs): 3 (Cpp-Taskflow) vs 9 (OpenMP)



Cpp-Taskflow is about 10%-17% faster than OpenMP and Intel TBB in avg, using the least amount of source code

# “Modern C++” Enables New Technology

- If you were able to tape out C++ ...



- It's much more than just being modern

- Must “rethink” the way we used to design a program
- Achieve the performance previously not possible



*Modern programming languages  
allow us to achieve the performance  
scalability and programming  
productivity that were previously out  
of reach*

# What Cpp-Taskflow Users Say ...

---

*“Cpp-Taskflow is the cleanest Task API I have ever seen,”*  
*Damien Hocking*

*“Cpp-Taskflow has a very simple and elegant tasking interface; the performance also scales very well,” Totalgee*

*“Best poster award for open-source parallel programming library,” Cpp-Conference (voted by 1000+ developers)*

**People are using Cpp-Taskflow to speed up TensorFlow’s kernel!**



lifeiteng commented on Dec 11, 2018



Is your feature request related to a problem? Please describe.

I want to use TaskFlow in a C++ project which depends on TensorFlow. When compiling the Op & Kernel with `c++17`, there is a problem [tensorflow/tensorflow#23561](#)

# Outline

---

- My early research at UIUC
  - OpenTimer: A VLSI timing analysis tool
  - Express parallelism in the right way
- A general-purpose distributed programming system
  - DtCraft and its ecosystem
  - Machine learning, graph algorithms, chip design
  - Cpp-Taskflow: Modern C++ parallel task programming
- Conclusion and future work

# Conclusion

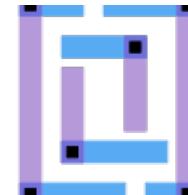
---

- **Parallel task programming systems**
  - DtCraft: distributed computing
  - Cpp-Taskflow: multicore task programming
  - Solution at programming level matters a lot
    - Performance is top priority but never underestimate productivity
- **Three takeaways**
  - Need high-level abstractions for parallel programming
  - Need descent programming productivity
  - Need modern software technologies
- **CE students have unique strength in SE**
  - We understand both hardware and software
  - We understand both science and engineering

# Acknowledgment (Users & Sponsors)



**SYNOPSYS**<sup>®</sup>



UNIVERSITY OF CALIFORNIA  
**SANTA CRUZ**



國立臺灣大學  
National Taiwan University



國立清華大學  
NATIONAL TSING HUA UNIVERSITY



國立交通大學  
National Chiao Tung University

*Thank you!*

<https://tsung-wei-huang.github.io/>

**Please contact me if you are interested in doing  
research for building software to solve real-world  
computer engineering problems**