

A General-purpose Parallel and Heterogeneous Task Programming System at Scale



Dr. Tsung-Wei Huang


Department of Electrical and Computer Engineering

University of Utah, Salt Lake City, UT





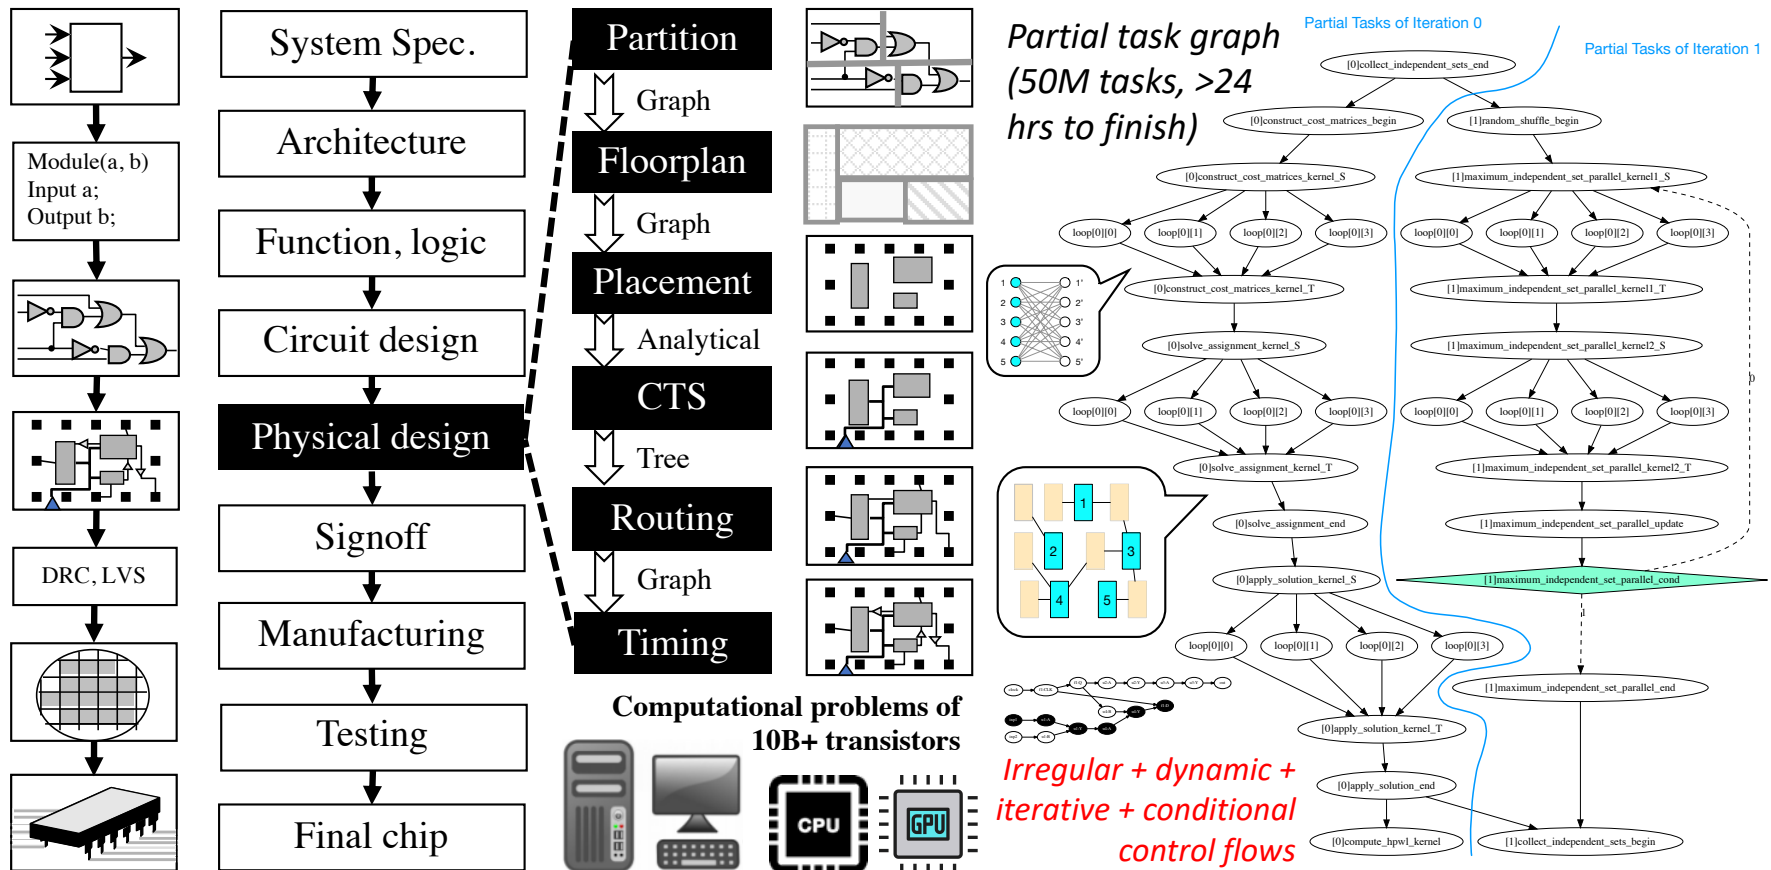
*How can we make it easier for scientific software developers to program large parallel and heterogeneous resources with **high performance scalability** and **simultaneous high productivity**?*



Real Experience: Parallelizing VLSI CAD Software

❑ CAD has many of the most difficult CS problems

❑ multidisciplinary, irregular, dynamic, billions of tasks, etc



CAD is Demanding New Parallelism!

☐ Yet, most existing solutions are incremental ...

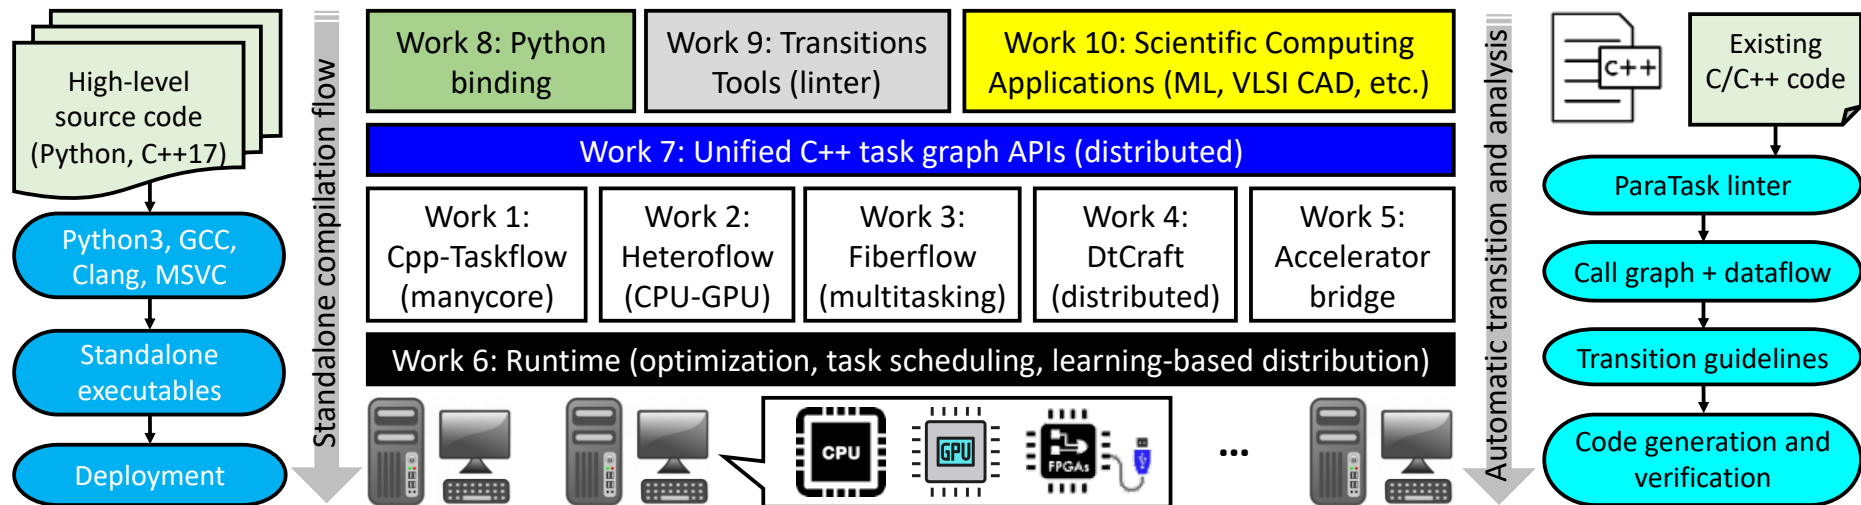
- ☐ Augment codebase with OpenMP, Pthreads, MPI, Intel TBB, and CUDA to introduce incremental parallelism
- ☐ Solutions are heavily hard-coded by resorting everything to “heroic programmers”

☐ Why are we sluggishly changing this?

- ☐ Existing programming tools (HPX, SYCL, Kokkos, RAJA, PaRSEC, StarPU, etc.) are perfect in data-parallel programs but short in:
 - Steep learning curve of new models (syntax + semantic)
 - Dynamic/conditional/cyclic control flows
 - Composition to handle complex heterogeneous workflows
- ☐ Lack of automatic transition and verification tools

A General-purpose Parallel Task Programming System

- ❑ Task-based approach scales best with parallel arch
- ❑ We need to handle various computational needs
 - ❑ Dynamic controls, cyclic flows, composition, irregularity
- ❑ Transition tools to ease the adoption of new models
 - ❑ Never acceptable if everything is done manually




This is an on-going large system project with many components under construction!

“Hello World” in Cpp-Taskflow [IPDPS19]

```
#include <taskflow/taskflow.hpp> // Cpp-Taskflow is header-only
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B);           // A runs before B
    A.precede(C);           // A runs before C
    B.precede(D);           // B runs before D
    C.precede(D);           // C runs before D
    executor.run(taskflow); // create an executor to run the taskflow
    return 0;
}
```

Only **17 lines** of code to get a parallel task execution!



“Hello World” in OpenMP

```
#include <omp.h> // OpenMP is a lang ext to describe parallelism in compiler directives
```

```
int main(){
```

```
    #omp parallel num_threads(std::thread::hardware_concurrency())
```

```
{
```

```
    int A_B, A_C, B_D, C_D;
```

```
    #pragma omp task depend(out: A_B, A_C) ←
```

Task dependency clauses

```
{
```

```
        std::cout << "TaskA\n";
```

```
}
```

```
    #pragma omp task depend(in: A_B; out: B_D) ←
```

Task dependency clauses

```
{
```

```
        std::cout << "TaskB\n";
```

```
}
```

```
    #pragma omp task depend(in: A_C; out: C_D) ←
```

Task dependency clauses

```
{
```

```
        std::cout << "TaskC\n";
```

```
}
```

```
    #pragma omp task depend(in: B_D, C_D) ←
```

Task dependency clauses

```
{
```

```
        std::cout << "TaskD\n";
```

```
}
```

```
}
```



```
    return 0;
```

```
}
```

*OpenMP task clauses are **static** and **explicit**;
Programmers are responsible for a **proper order of
writing tasks** consistent with sequential execution*

“Hello World” in Intel’s TBB Library

```
#include <tbb.h> // Intel’s TBB is a general-purpose parallel programming library in C++
```

```
int main(){  
    using namespace tbb;  
    using namespace tbb::flow;  
    int n = task_scheduler_init::default_num_threads ();  
    task_scheduler_init init(n);  
    graph g;   
    continue_node<continue_msg> A(g, [] (const continue_msg &) {  
        std::cout << "TaskA";  
    });  
    continue_node<continue_msg> B(g, [] (const continue_msg &) {  
        std::cout << "TaskB";  
    });  
    continue_node<continue_msg> C(g, [] (const continue_msg &) {   
        std::cout << "TaskC";  
    });  
    continue_node<continue_msg> C(g, [] (const continue_msg &) {  
        std::cout << "TaskD";  
    });  
    make_edge(A, B);  
    make_edge(A, C);  
    make_edge(B, D);  
    make_edge(C, D);  
    A.try_put(continue_msg());  
    g.wait_for_all();  
}
```


*Use TBB’s FlowGraph
for task parallelism*

*Declare a task as a
continue_node*

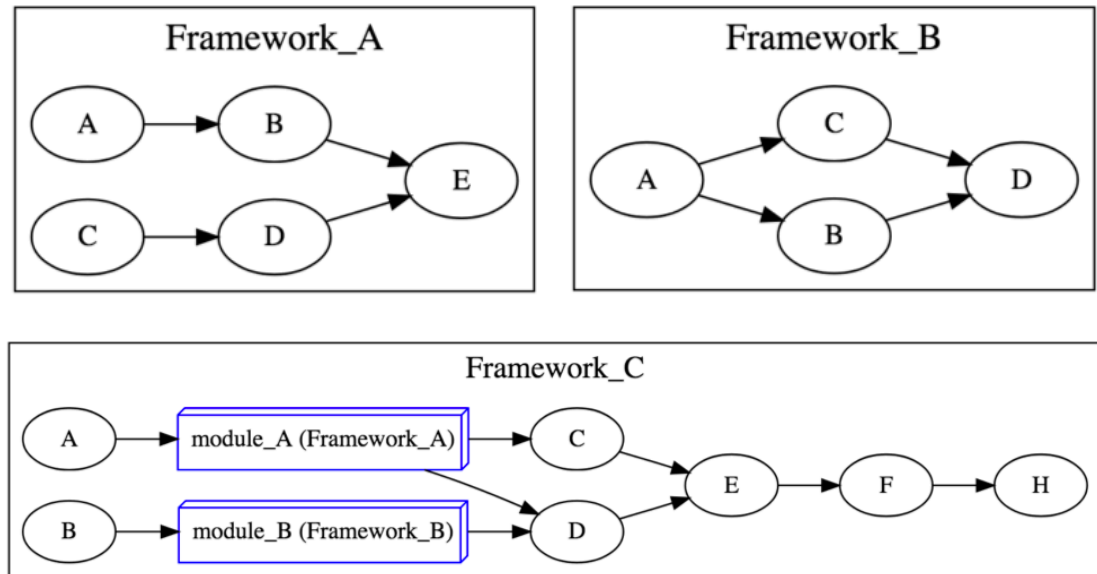
*TBB has excellent performance in generic parallel computing. Its drawback is mostly in the **ease-of-use** standpoint (simplicity, expressivity, and programmability).*

“Hello Universe” in Cpp-Taskflow


Dynamic tasking



Composable tasking for *complex workflows*



Conditional tasking for *cyclic* and *dynamic* control flows



More on: <https://github.com/cpp-taskflow/cpp-taskflow>

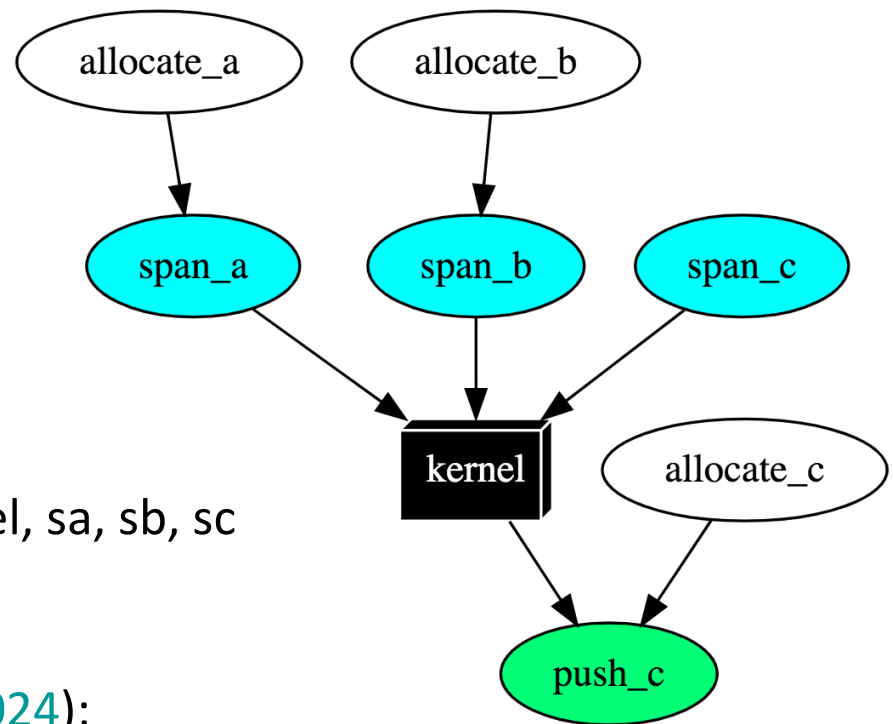
Concurrent CPU-GPU Task Programming

```
auto ha = hf.host([](){});  
auto hb = hf.host([](){});  
auto hc = hf.host([](){});  
auto sa = hf.span(1024);  
auto sb = hf.span(1024);  
auto sc = hf.span(1024);
```

```
auto op = hf.kernel(  
    {(1024+32-1)/32}, 32, 0, fn_kernel, sa, sb, sc  
);
```

```
auto cc = hf.copy(host_data, sc, 1024);
```

```
ha.precede(sa);  
hb.precede(sb);  
op.succeed(sa, sb, sc).precede(cc);  
cc.succeed(hc);
```



kernel is non-trivial, but what makes heterogeneous programming difficult is the “surrounding tasks”


Real Use Case: VLSI Timing Analysis

❑ OpenTimer v1: A VLSI Static Timing Analysis Tool

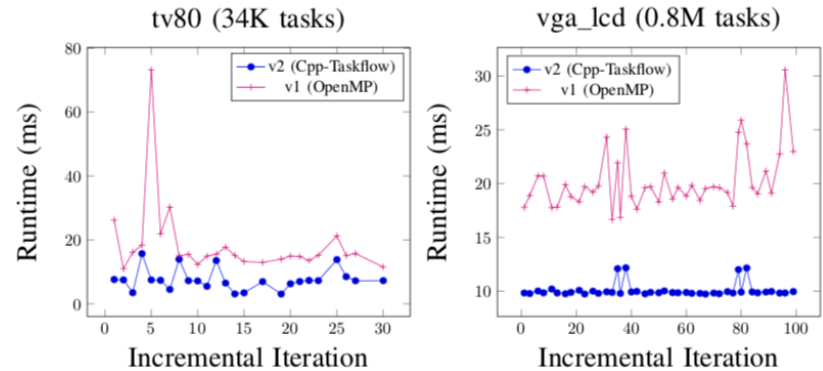
- ❑ v1 first released in 2015 (open-source under GPL)
- ❑ Loop-based parallelisms using OpenMP 4.0

❑ OpenTimer v2: A New Parallel Incremental Timer

- ❑ v2 first released in 2018 (open-source under MIT)
- ❑ Task-based parallel decomposition using Cpp-Taskflow



*Saved 4K lines of parallel code!
Cost to develop is \$275K with OpenMP vs
\$130K with Cpp-Taskflow! (measured by
sloccount)*



v2 (Cpp-Taskflow) is 1.4-2x faster than v1 (OpenMP)

Current and Future Work

- ❑ **A general-purpose parallel task programming system**
 - ❑ Simple, efficient, and transparent using modern C++
 - ❑ Multithreading and CPU-GPU tasking
 - ❑ Real case use in VLSI timing analysis with billion-tasking
- ❑ **On-going items**
 - ❑ Developing new task models using cudaGraph
 - ❑ Developing accelerator bridge to handle other devices
 - ❑ Developing transition tools to ease the adoption
- ❑ **Open to collaboration for more use cases!!!**
 - ❑ <https://github.com/cpp-taskflow/cpp-taskflow>
 - ❑ <https://github.com/tsung-wei-huang/DtCraft>