

# Distributed Timing Analysis in 100 Line Code

Tsung-Wei Huang

Research Assistant Professor

Department of Electrical and Computer Engineering

University of Illinois at Urbana-Champaign, IL, USA



# What we Need Today

---

## ❑ OpenTimer

- ❑ Static timing analysis (STA) tool
- ❑ Installation guide (by Kunal Ghosh):  
<https://www.udemy.com/vsd-a-complete-guide-to-install-open-source-eda-tools/>

## ❑ DtCraft

- ❑ Distributed programming system
- ❑ Website: <http://dtcraft.web.engr.illinois.edu/>
- ❑ GitHub: <https://github.com/twhuang-uiuc/DtCraft>

## ❑ A Linux machine (Ubuntu recommended)

- ❑ G++ 7.2 (for C++17)

## ❑ Demo code:

<http://web.engr.illinois.edu/~thuang19/webinar.tar.gz>

# Install DtCraft

---

## □ Download DtCraft

- <http://dtcraft.web.engr.illinois.edu/download.html>

```
~$ git clone https://github.com/twhuang-uiuc/DtCraft.git
```

```
~$ cd DtCraft
```

## □ Build DtCraft

- Disable shared library for simplicity (--disable-shared)

```
~$ ./configure --disable-shared
```

```
~$ make
```

## □ Make sure you have GCC/G++ 7 installed (C++17)

```
~$ sudo apt-get update
```

```
~$ sudo apt-get install gcc-7 g++-7
```

# Outline

---

- Express your parallelism in the right way
  - A “hard-coded” distributed timing analysis framework
- Boost your productivity in writing parallel code
  - DtCraft system
- Leverage your time to produce promising results
  - Demo 1: A vanilla example
  - Demo 2: Distributed timing using DtCraft
  - Lab

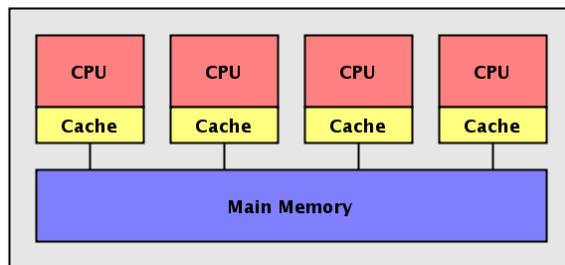
# Outline

---

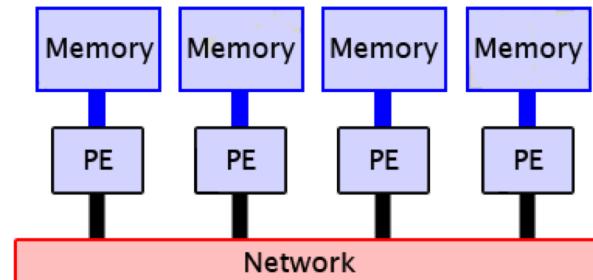
- ❑ Express your parallelism in the right way
  - ❑ A “hard-coded” distributed timing analysis framework
- ❑ Boost your productivity in writing parallel code
  - ❑ DtCraft system
- ❑ Leverage your time to produce promising results
  - ❑ Demo 1: A vanilla example
  - ❑ Demo 2: Distributed timing using DtCraft
  - ❑ Lab

# Distributed Timing

- Deal with the ever-increasing design complexity
  - Billions of transistors result in very large timing graphs
  - Analyze the timing under different conditions
  - Vertical scaling is not cost efficient
- Want to scale out our computations
  - Leverage the power of computer clusters (cloud computing)



Single node (threaded)

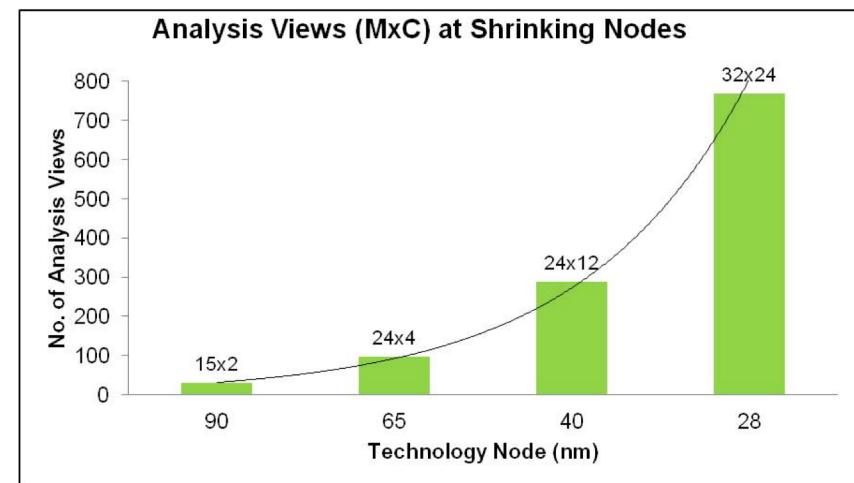
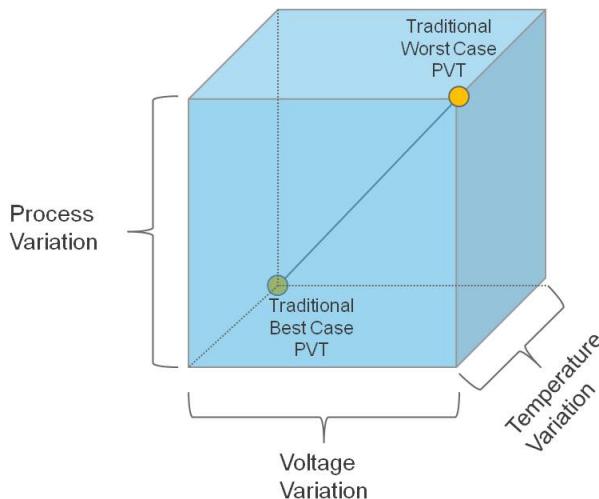


Datacenter (distributed)

# Motivation: Speed up Timing Closure

## ❑ Multi-mode multi-corner (MMMC) timing analysis

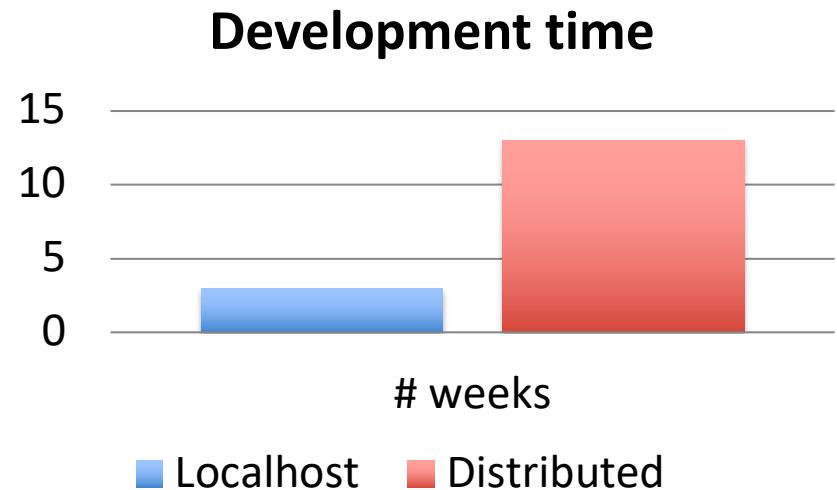
- ❑ Test modes, functional modes
- ❑ Process, voltage, temperature (PVT)
- ❑ Timing runs across all combinations
  - Temperature: Tmax/Tmin
  - Voltage: Vmin/Vmax
- ❑ Each combination is referred to as a *timing view*



# Good News and Bad News

---

- Each timing view is logically parallel to each other
- Developing a distributed program is very difficult
  - Several weeks more than a single-machine counterpart
  - Network programming, subtly buggy code, etc
- Scalability and transparency
  - Intend to focus on high level rather than low level
  - Want better productivity
  - Want better flexibility
  - Want better performance



# Distributed Systems in Big Data

---

- ❑ **Hadoop**

- ❑ Distributed MapReduce platform on HDFS

- ❑ **Cassandra**

- ❑ Scalable multi-master database

- ❑ **Chukwa**

- ❑ A distributed data collection system

- ❑ **Zookeeper**

- ❑ Coordination service for distributed application

- ❑ **Mesos**

- ❑ A high-performance cluster manager

- ❑ **Spark**

- ❑ A fast and general computing engine for big-data analytics

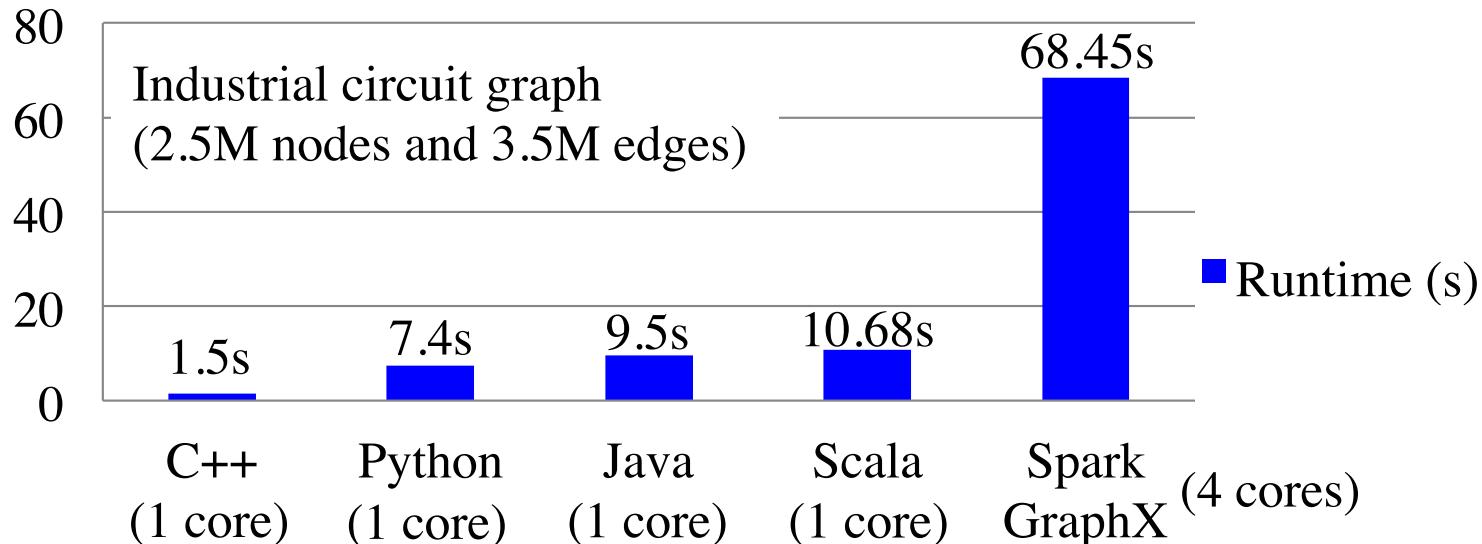
# The Questions are

---

- Are these packages suitable for timing?
- What are the potential hurdles for EDA to use them?
- How much code rewrite do I need?
- What is the significance of adopting new languages?
- Will I lose performance?

# Big-data Tool is Not an Easy Fit!

## Runtime comparison on arrival time propagation



Method	Spark (RDD + GraphX Pregel)	Java (SP)	C++ (SP)
Runtime (s)	68.45	9.5	1.50

Overhead of big data tools

Language difference

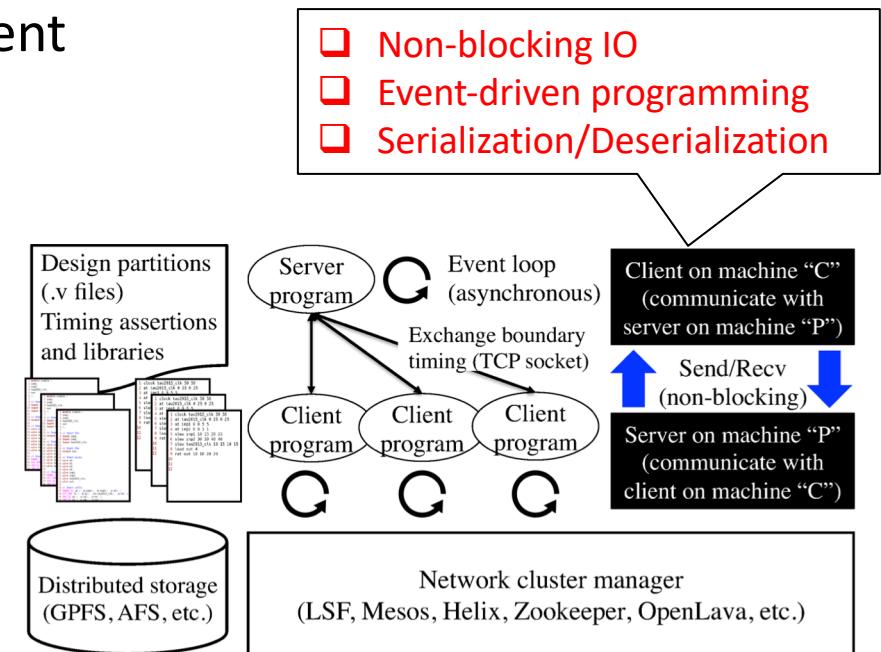
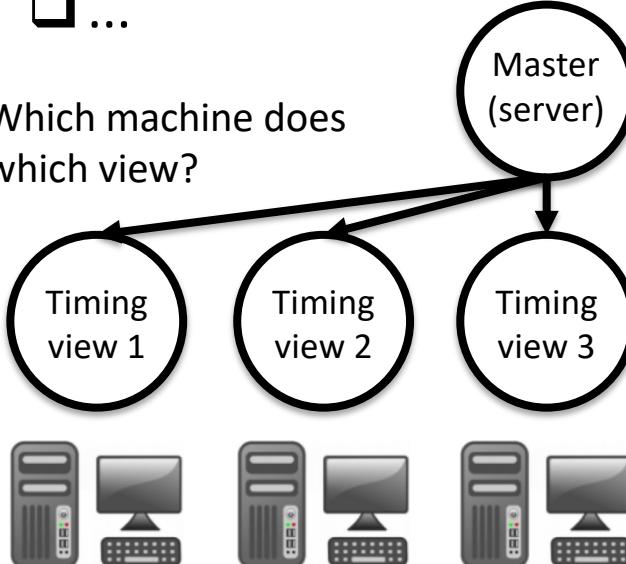
# A Hard-coded Distributed MMMC Framework

## □ Built from the scratch using raw Linux Socket

- Hard code using Linux sockets
- Explicit data movement
- Explicit job execution
- Explicit parallelism management
- ...

Difficult scalability ☹  
Large amount of code rewrites ☹

Which machine does which view?



# Observations

---

- ❑ Big data doesn't fit well in timing
  - ❑ IO-bound vs CPU-bound
  - ❑ Unstructured data vs structured data
  - ❑ JVM vs C/C++
- ❑ Life shouldn't be hard-coded
  - ❑ Deal with low-level socket programming
  - ❑ Move data explicitly between compute nodes
  - ❑ Manage cluster resources on your own
  - ❑ Result in a large amount of development efforts
- ❑ Want parallel programming *at scale more productive*

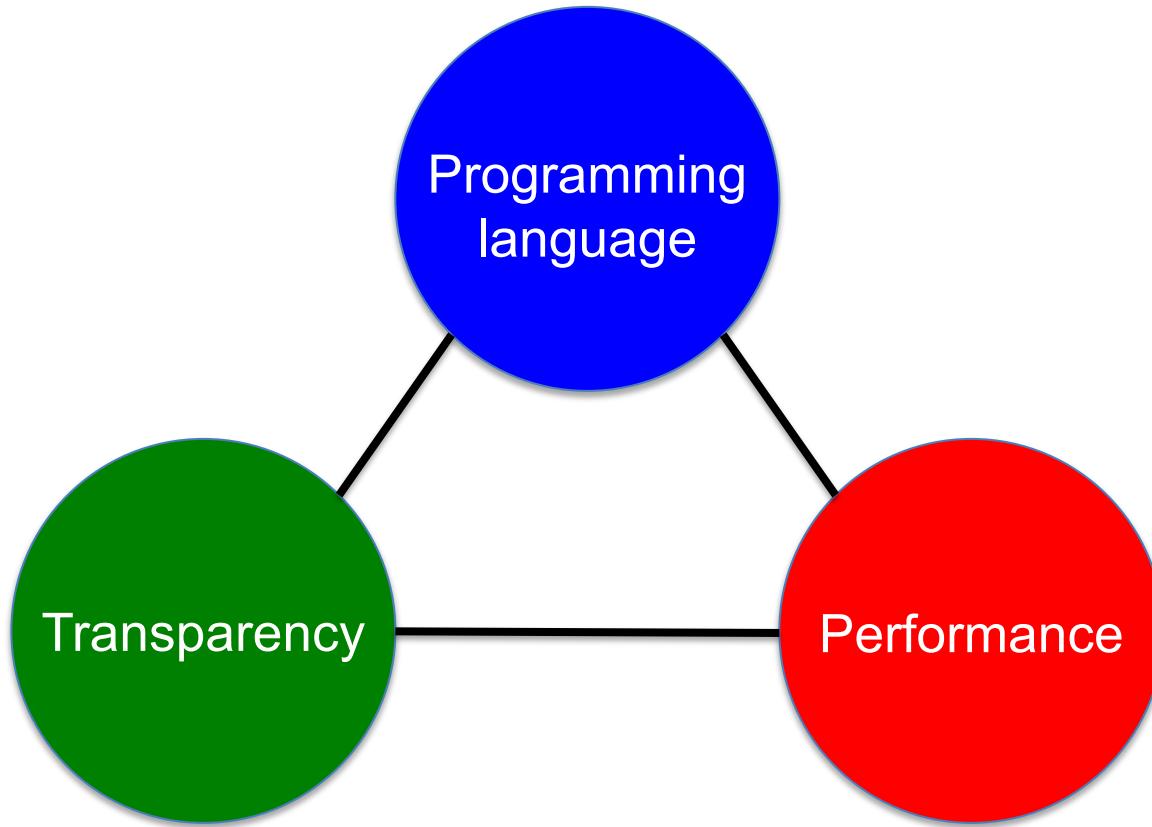
# Outline

---

- Express your parallelism in the right way
  - A “hard-coded” distributed timing analysis framework
- Boost your productivity in writing parallel code
  - DtCraft system
- Leverage your time to produce promising results
  - Demo 1: A vanilla example
  - Demo 2: Distributed timing using DtCraft
  - Lab

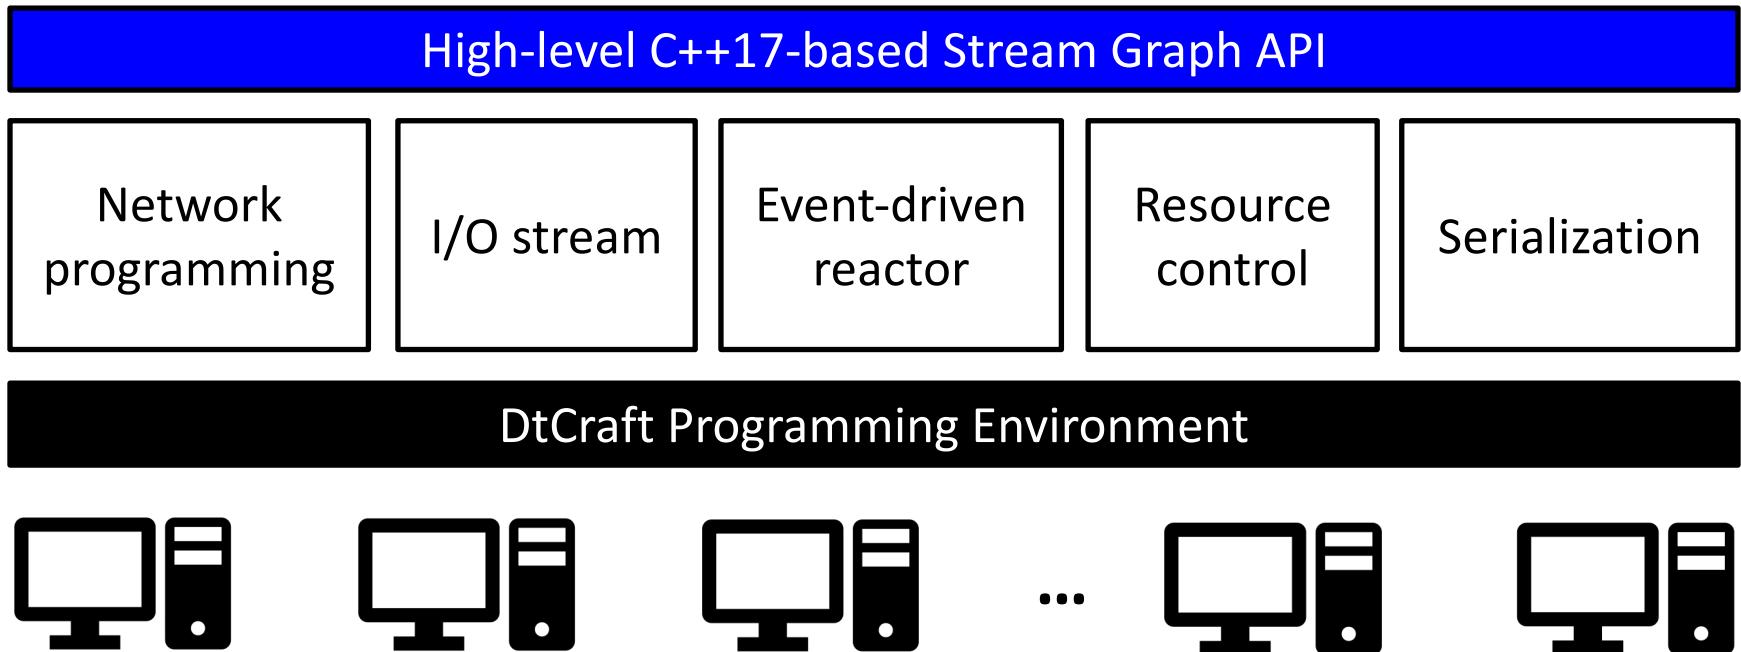
# What does “Productivity” really mean?

---



# A New Solution: DtCraft

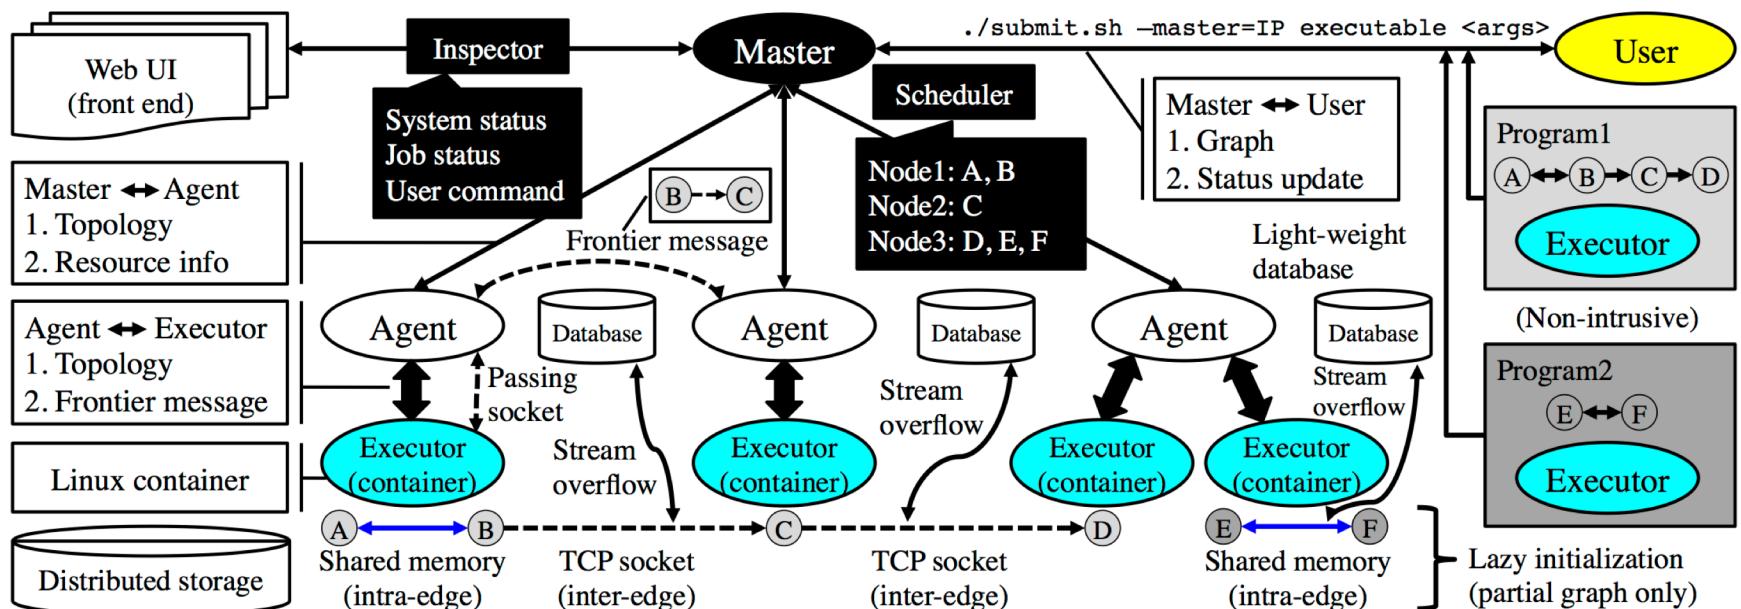
- ❑ A unified engine to simplify cluster programming
  - ❑ Completely built from the ground up using C++17
- ❑ Save your time away from the pain of DevOps



T.-W. Huang, C.-X. Lin, and M. D. F. Wong, “DtCraft: A high-performance distributed execution engine at scale,” IEEE TCAD, to appear, 2018

# System Architecture

- Express your parallelism in our *stream graph* model
  - Generic dataflow at any granularity
- Deliver transparent concurrency through the kernel
  - Automatic workload distribution and message passing

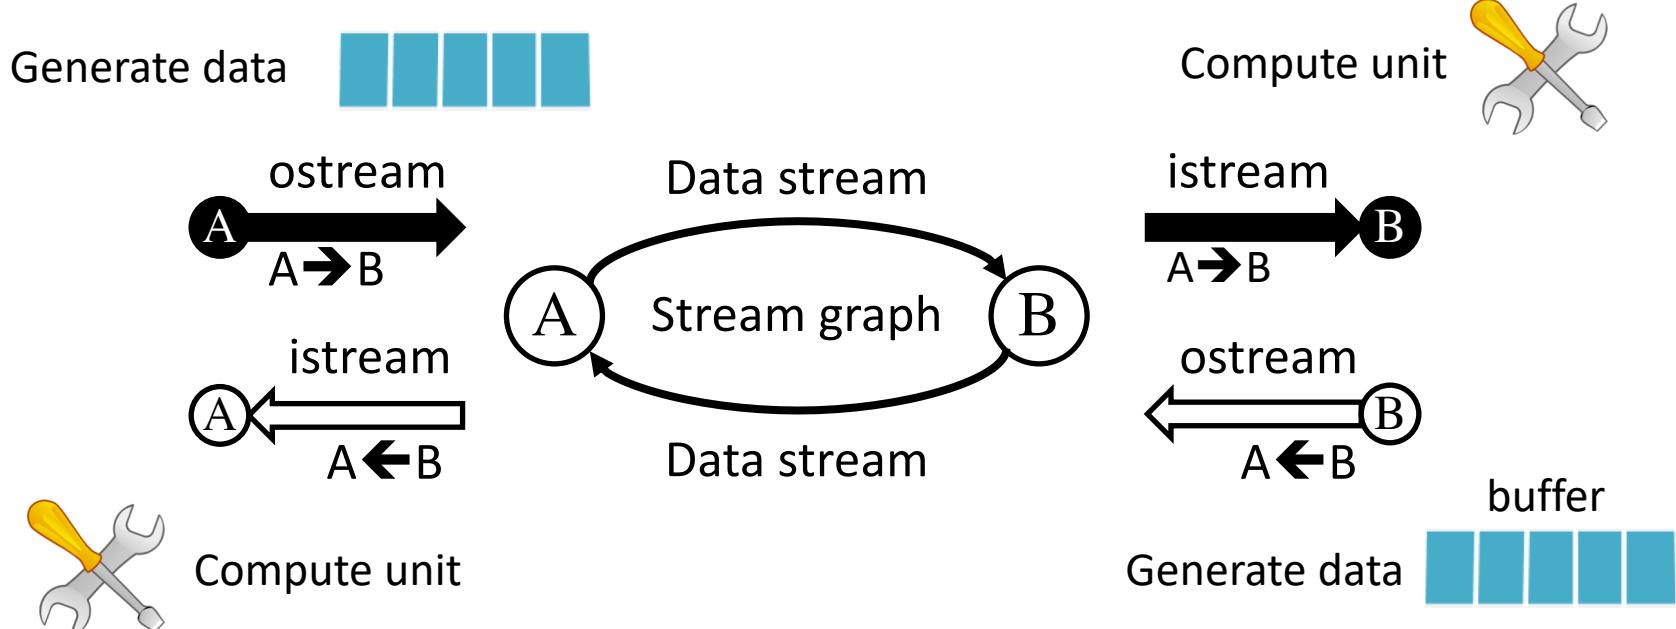


DtCraft website: <http://dtcraft.web.engr.illinois.edu/>

DtCraft github: <https://github.com/twhuang-uiuc/DtCraft>

# Stream Graph Programming Model

- A general representation of a dataflow
  - Abstraction over computation and communication
- Analogous to the assembly line model
  - Vertex storage → goods store
  - Stream processing unit → independent workers



# Outline

---

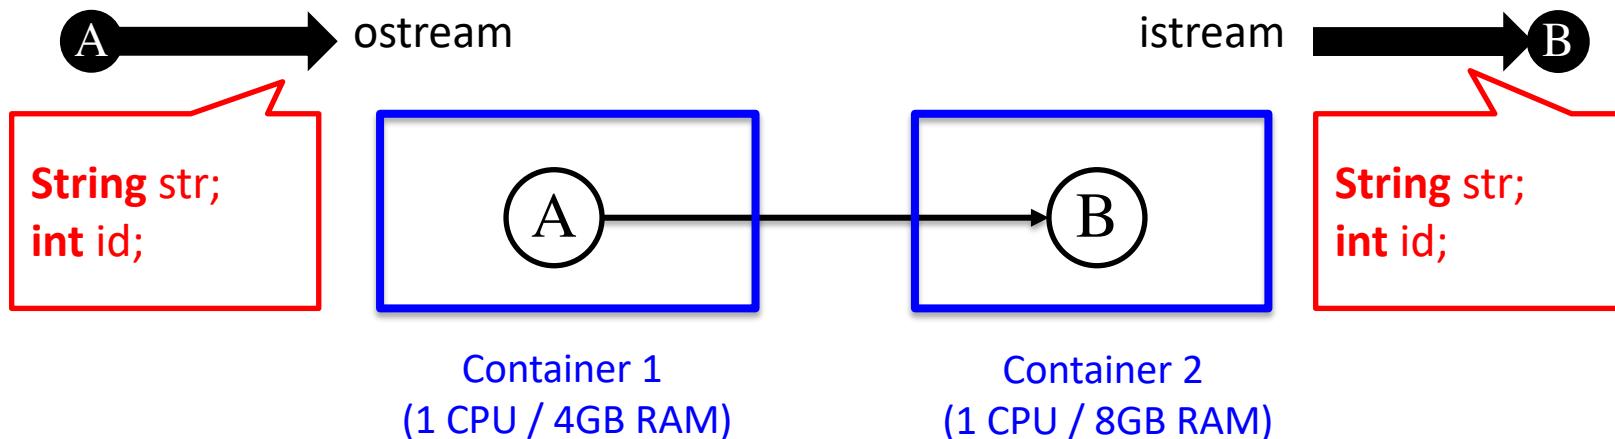
- ❑ Express your parallelism in the right way
  - ❑ A “hard-coded” distributed timing analysis framework
- ❑ Boost your productivity in writing parallel code
  - ❑ DtCraft system
- ❑ Leverage your time to produce promising results
  - ❑ Demo 1: A vanilla example
  - ❑ Demo 2: Distributed timing using DtCraft
  - ❑ Lab

# Write a DtCraft Application

- Step 1: Decide the stream graph of your application
- Step 2: Specify the data to stream between vertices
- Step 3: Define the stream computation callback
- Step 4: Attach resources on vertices (optional)
- Step 5: Submit

./submit –master=host hello-world

```
auto L = [=] (auto& vertex, auto& istream) {  
    if(string data; istream(data) != -1) {  
        // Your program control flow.  
    } else { ... }  
};
```

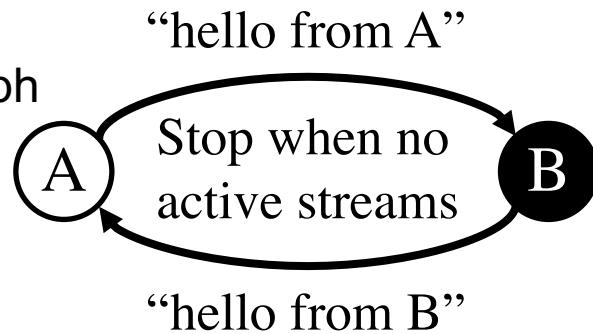


# A Vanilla Example

---

- A cycle of two vertices and two streams
  - Each vertex sends a hello message to the other
  - Closes the underlying stream channel

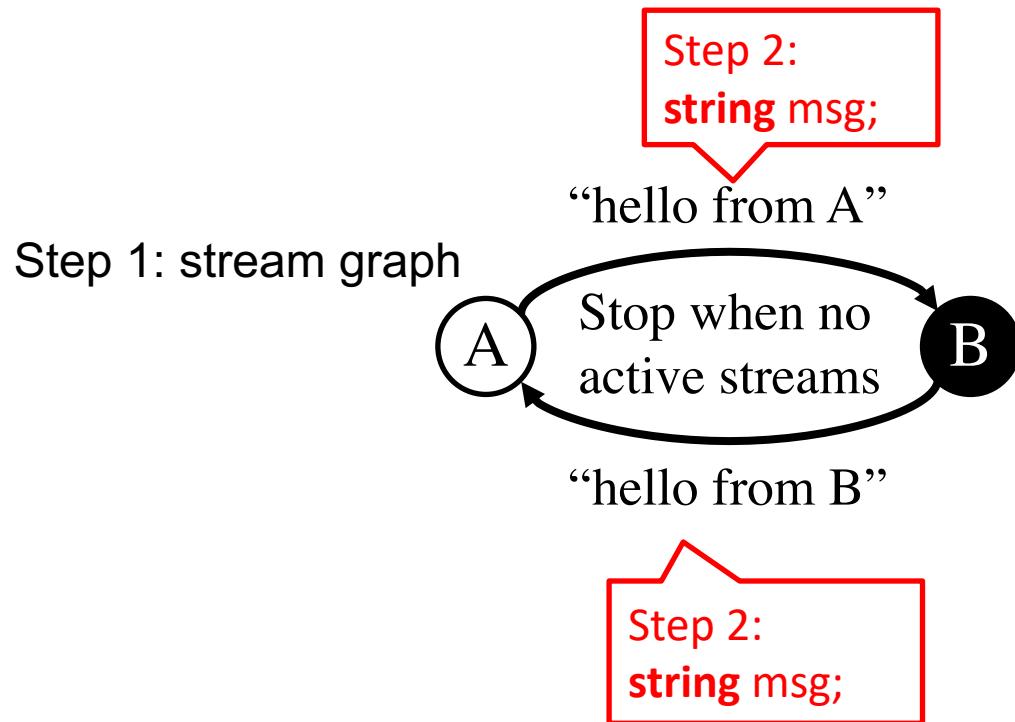
Step 1: stream graph



# A Vanilla Example

---

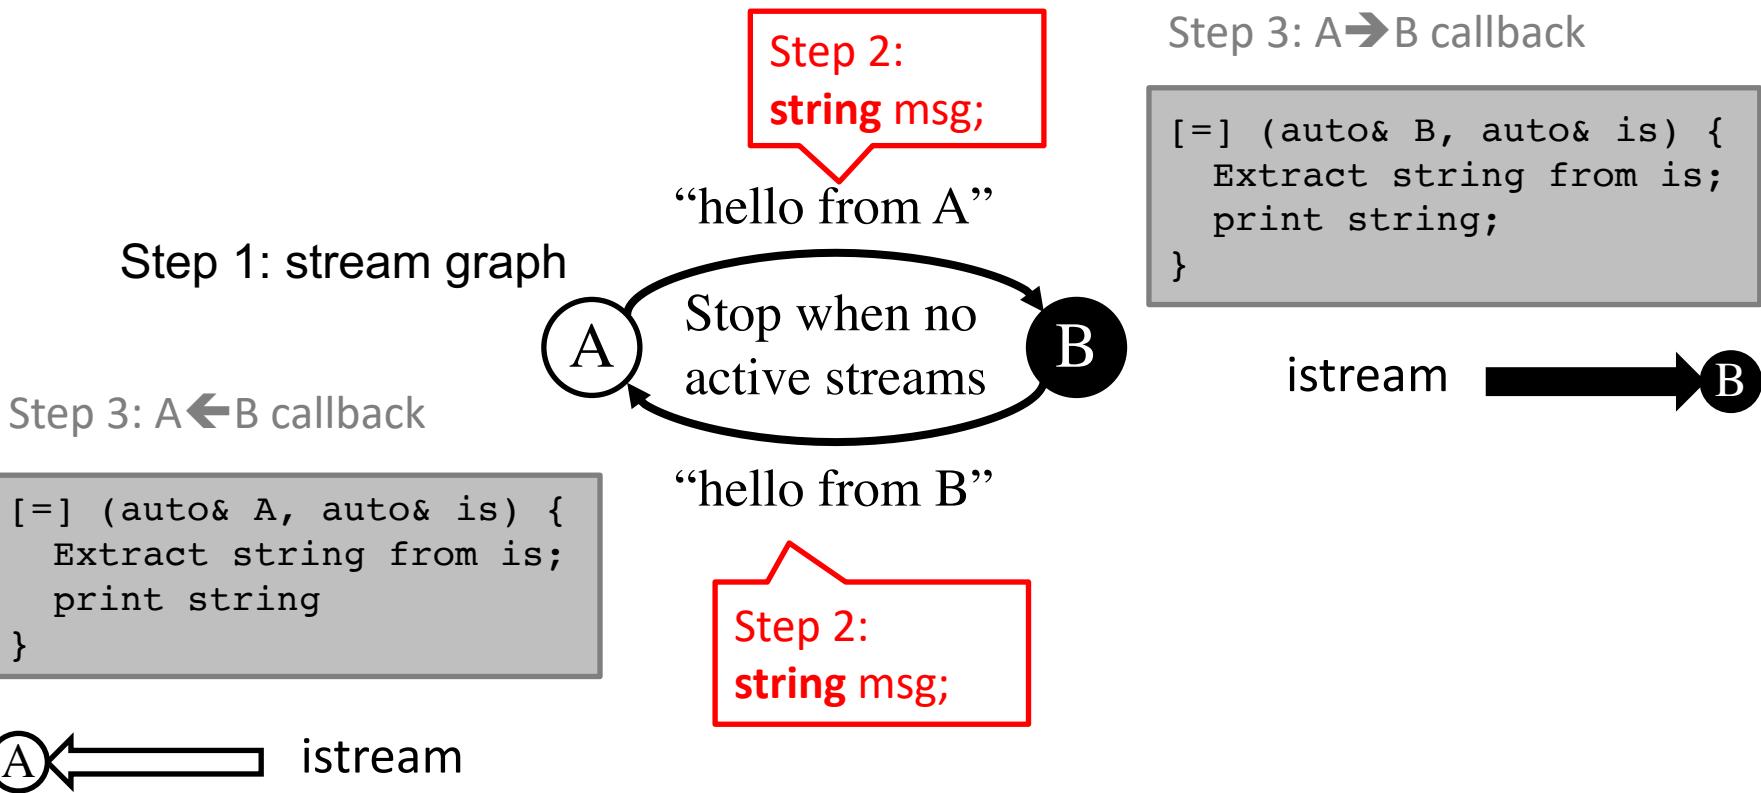
- ❑ A cycle of two vertices and two streams
  - ❑ Each vertex sends a hello message to the other
  - ❑ Closes the underlying stream channel



# A Vanilla Example

## □ A cycle of two vertices and two streams

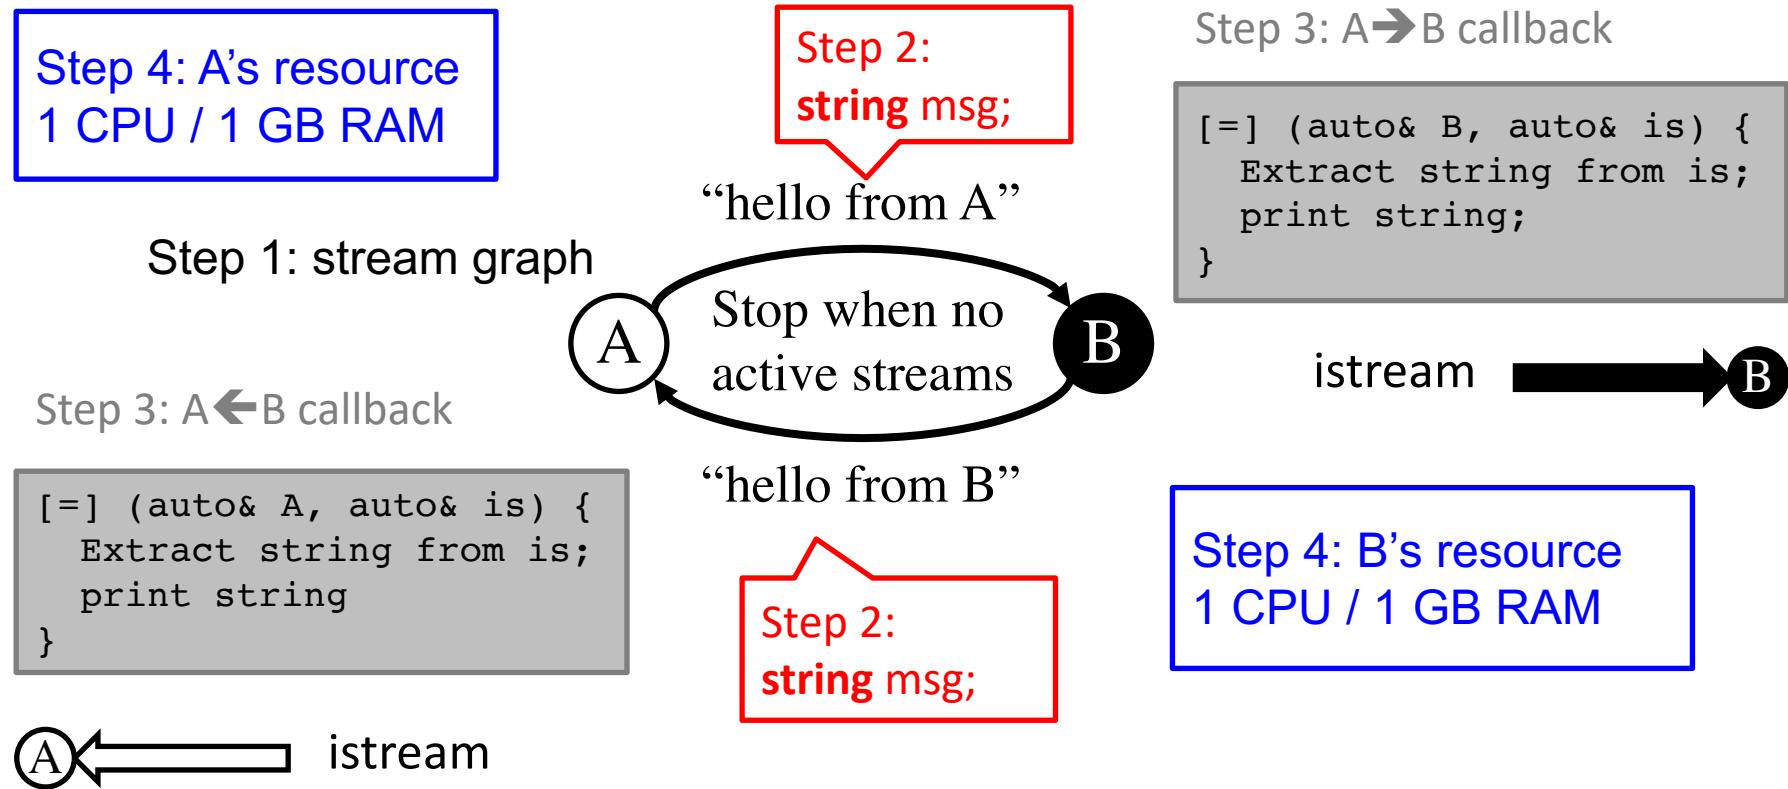
- Each vertex sends a hello message to the other
- Closes the underlying stream channel



# A Vanilla Example

## □ A cycle of two vertices and two streams

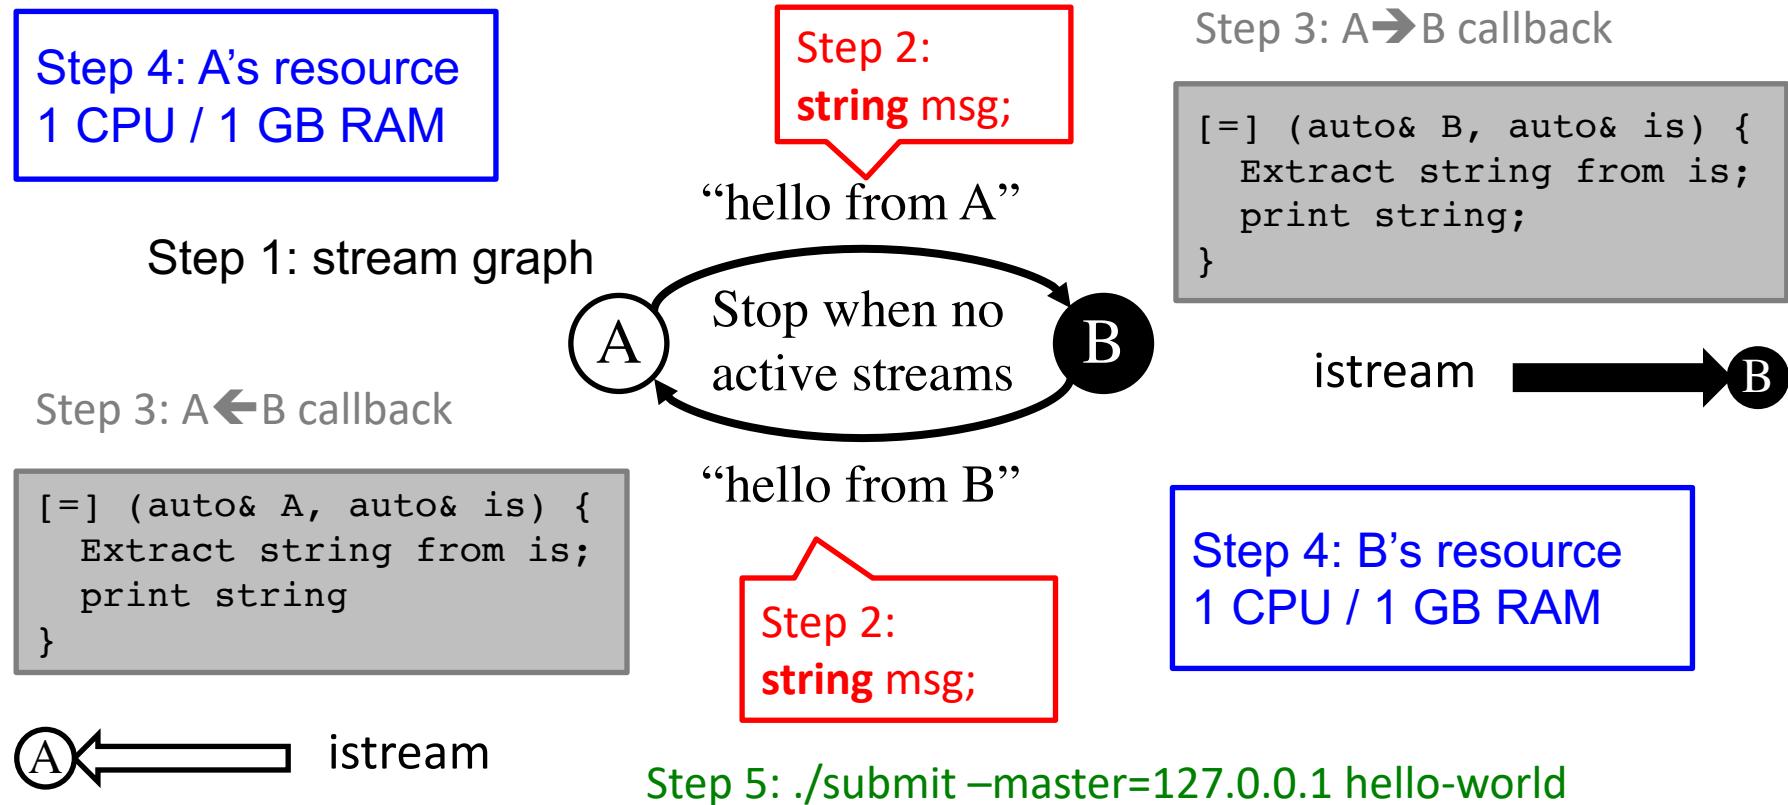
- Each vertex sends a hello message to the other
- Closes the underlying stream channel



# A Vanilla Example

## □ A cycle of two vertices and two streams

- Each vertex sends a hello message to the other
- Closes the underlying stream channel



# Demo (hello\_world.cpp)

---

## ❑ Hello world

- ❑ Create a stream graph of two vertices and two streams
- ❑ Use container interface to manage cluster resources

## ❑ Local mode execution

- ❑ Single process

## ❑ Distributed mode execution

- ❑ Launch master and agents to set up a DtCraft cluster
- ❑ Submit hello\_world to the cluster



Login user	host	DtCraft home
twhuang	csl-408-08.csl.Illinois.edu	/home/twhuang/DtCraft
twhuang	csl-408-14.csl.Illinois.edu	/home/twhuang/DtCraft



Notice: Replace with your own login/hosts/DtCraftHome.

Demo code: <http://web.engr illinois.edu/~thuang19/webinar.tar.gz>

# Debrief

```
dtc::Graph G;

auto A = G.vertex();
auto B = G.vertex();

G.container().add(A).cpu(1).memory(1_GB);

auto AB = G.stream(A, B).on(
[] (dtc::Vertex& B, dtc::InputStream& is) {
    if(std::string b; is(b) != -1) {
        dtc::cout("Received: ", b, '\n');
        return dtc::Event::REMOVE;
    }
    return dtc::Event::DEFAULT;
});

auto BA = G.stream(B, A);

A.on(
[&AB] (dtc::Vertex& v) {
    (*v.ostream(AB))("Hello world from A's");
    dtc::cout("Sent 'Hello world from A' to stream ", AB, "\n");
}
);

G.container().add(B).cpu(1).memory(1_GB);

B.on(
[&BA] (dtc::Vertex& v) {
    (*v.ostream(BA))("Hello world from B's");
    dtc::cout("Sent 'Hello world from B' to stream ", BA, "\n");
}
);

BA.on(
[] (dtc::Vertex& A, dtc::InputStream& is) {
    if(std::string a; is(a) != -1) {
        dtc::cout("Received: ", a, '\n');
        return dtc::Event::REMOVE;
    }
    return dtc::Event::DEFAULT;
}
);

dtc::Executor(G).run();
```

} Vertex

} Stream  
A→B

} Stream  
B→A

- Only a couple lines of code
- Single sequential program
- Distributed across computers
- No explicit data management
- Easy-to-use streaming interface
- Asynchronous by default
- Scalable to many threads
- Scalable to many machines
- In-context resource controls
- Scale out to heterogeneous devices
- Transparent concurrency controls
- Robust runtime via Linux container
- ... and more

# Distributed Hello-world without DtCraft ...

```
auto count_A = 0;
auto count_B = 0;

// Send a random binary data to fd and add the
// received data to the counter.
auto pinpong(int fd, int& count) {
    auto data = random<bool>()
    auto w = write(fd, &data, sizeof(data));
    if(w == -1 && errno != EAGAIN) {
        throw system_error("Failed on write");
    }
    data = 0;
    auto r = read(fds, &data, sizeof(data));
    if(r == -1 && errno != EAGAIN) {
        throw system_error("Failed on read");
    }
    count += data;
}

int fd = -1;
std::error_code errc;
```

*server.cpp*

```
if(getenv("MODE") == "SERVER") {
    fd = make_socket_server_fd("9999", errc);
}
else {
    fd = make_socket_client_fd("127.0.0.1", "9999", errc);
}

if(fd == -1) {
    throw system_error("Failed to make socket");
}
```

*client.cpp*

```
int make_socket_server_fd(
    std::string_view port,
    std::error_code errc
) {
    int fd {-1};
    if(fd != -1) {
        ::close(fd);
        fd = -1;
    }
    make_fd_close_on_exec(fd);
    tries = 3;
    issue_connect:
    ret = ::connect(fd, ptr->ai_addr, ptr->ai_addrlen);

    if(ret == -1) {
        if(errno == EINTR) {
            goto issue_connect;
        }
        else if(errno == EAGAIN & tries--) {
            std::this_thread::sleep_for(std::chrono::milliseconds(500));
            goto issue_connect;
        }
        else if(errno != EINPROGRESS) {
            goto try_next;
        }
        errc = make_posix_error_code(errno);
    }

    // Poll the socket. Note that writable return doesn't mean it is connected
    if(select_on_write(fd, 5, errc) && !errc) {
        int optval = -1;
        socklen_t optlen = sizeof(optval);
        if(::getsockopt(fd, SOL_SOCKET, SO_ERROR, &optval, &optlen) < 0) {
            errc = make_posix_error_code(errno);
            goto try_next;
        }
        if(optval != 0) {
            errc = make_posix_error_code(optval);
            goto try_next;
        }
        // Try for another
        for(aut
            if(::bind(fd, ptr->ai_addr, ptr->ai_addrlen) != -1) {
                errc = make_posix_error_code(errno);
                goto try_next;
            }
            if(::listen(fd, 5) != -1) {
                errc = make_posix_error_code(errno);
                goto try_next;
            }
            if(fd != -1) {
                ::close(fd);
                fd = -1;
            }
        }
        make_
        ::freeaddrinfo(res);
    }
    return fd;
}
```

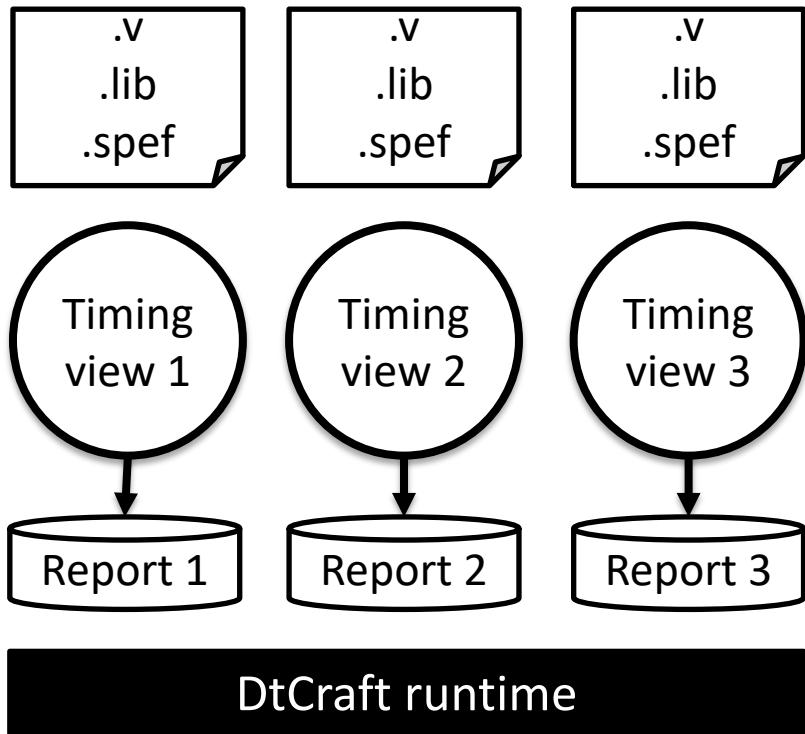
*A lot of boilerplate code  
for this trivial distributed  
program...*

*Branch your code to server and client for  
distributed computation!*

*simple.cpp → server.cpp + client.cpp*

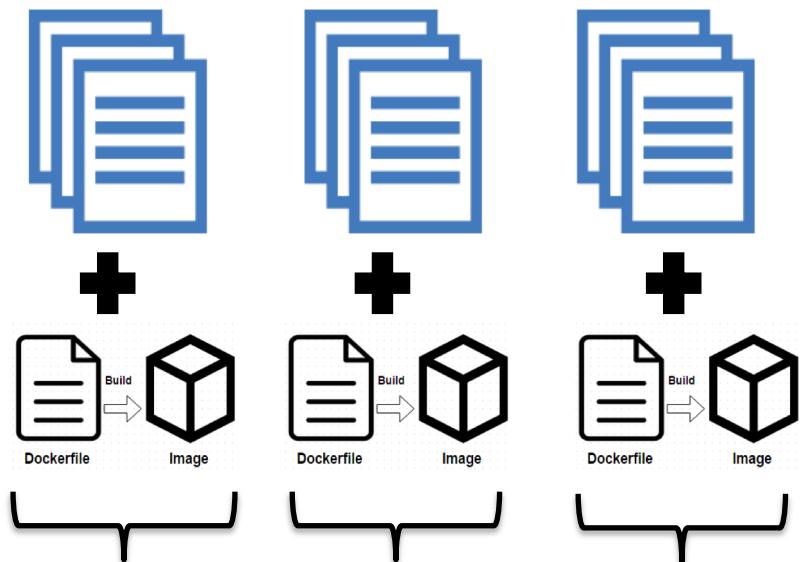
# Distributed Timing with DtCraft

## □ Three timing views



With DtCraft (<100 lines of code)

*Code duplication, separate control flows*



*Wrap up with scripts*

Without DtCraft (hard-coded)



# Demo (dta.cpp)

---

- Distributed timing analysis with three timing views
  - simple\_tv1 (P=1, V=0.7, T=70)
  - simple\_tv2 (P=0.5, V=0.95, T=85)
  - simple\_tv3 (P=0.9, V=0.5, T=60)
- OpenTimer how-to (by Kunal)
  - <https://www.udemy.com/vlsi-academy-sta-checks-2/>
- Local mode execution
- Distributed mode execution



Login user	host	DtCraft home
twhuang	csl-408-08.csl.Illinois.edu	/home/twhuang/DtCraft
twhuang	csl-408-14.csl.Illinois.edu	/home/twhuang/DtCraft

Notice: Replace with your own login/hosts/DtCraftHome.

Demo code: <http://web.engr illinois.edu/~thuang19/webinar.tar.gz>

# Debrief

---

## Transparency

- No low-level network programming details
- Automatic workload distribution

## Scalability

- Same code scales out automatically when new machines added
- Dynamic scaling
- Flexible partitions and in-context resource controls

## Programmability

- Can incorporate other programs together

## Productivity

- Less than 100 lines of code

# Exercise (lab.cpp)

---

- Distributed timing for two designs each with four views**
  - c17\_v1, c17\_v2, c17\_v3, c17\_v4
  - c499\_v1, c499\_v2, c499\_v3, c499\_v4
- Implement a stream graph**
  - Eight vertices each operating on one timing view
  - Two containers, one for c17\_v\* and another for c499\_v\*
  - Try different resource assignments for each container
- Submit the graph to your DtCraft cluster**
  - Local mode and distributed mode
- Report TNS and WNS for each view**
- Report elapsed time and peak memory for each container**
- Use at most 50 lines of code ☺**

# Example Solution (20 lines)

---

```
#include <dtc/dtc.hpp>

int main(int argc, char* argv[]) {
    using namespace dtc::literals;

    dtc::Graph G;
    auto c17tv1 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c17_tv1 tv1.report");
    auto c17tv2 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c17_tv2 tv2.report");
    auto c17tv3 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c17_tv3 tv3.report");
    auto c17tv4 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c17_tv4 tv4.report");

    auto c499tv1 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c499_tv1 tv1.report");
    auto c499tv2 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c499_tv2 tv2.report");
    auto c499tv3 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c499_tv3 tv3.report");
    auto c499tv4 = G.vertex().program("path_to_webinar/ot.sh path_to_webinar/c499_tv4 tv4.report");

    G.container().add(c17tv1).add(c17tv2).add(c17tv3).add(c17tv4).cpu(1).memory(1_GB);
    G.container().add(c499tv1).add(c499tv2).add(c499tv3).add(c499tv4).cpu(1).memory(1_GB);

    dtc::Executor(G).run();
    return 0;
}
```

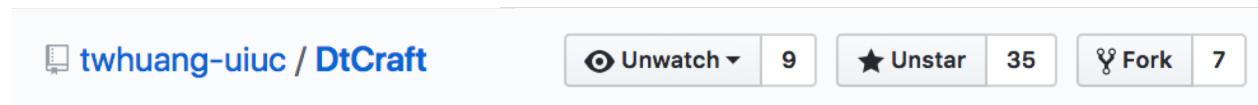
# Example Report

Design	TNS (ps)	WNS (ps)	Elapsed Time	Peak Memory
c17_tv1	-8.14469e+02	-2.29314e+01	0.23s	86M
c17_tv2	-2.79547e+03	-7.32566e+01		
c17_tv3	-7.15888e+02	-1.91890e+01		
c17_tv4	-1.69670e+03	-4.53052e+01		
c499_tv1	-4.97395e+05	-5.16786e+02	0.25s	90M
c499_tv2	-4.86524e+05	-5.05552e+02		
c499_tv3	-2.34709e+05	-2.44073e+02		
c499_tv4	-1.69093e+06	-1.75538e+03		

# Conclusion

---

- ❑ Express your parallelism in the right way
  - ❑ A “hard-coded” distributed timing analysis framework
- ❑ Boost your productivity in writing parallel code
  - ❑ DtCraft system
- ❑ Leverage your time to produce promising results
  - ❑ Demo 1: A vanilla example
  - ❑ Demo 2: Distributed timing using DtCraft
  - ❑ Lab



Please star and watch DtCraft at GitHub to receive updates!  
(<https://github.com/twhuang-uiuc/DtCraft>)

# Thank you!

Dr. Tsung-Wei Huang

[twh760812@gmail.com](mailto:twh760812@gmail.com)

Github: <https://github.com/twhuang-uiuc>

Website: <http://web.engr.illinois.edu/~thuang19/>

