

+ 25

Dynamic Asynchronous Tasking with Dependencies

TSUNG-WEI HUANG



20
25






Takeaways

- **Understand the importance of asynchronous tasking with dependencies**
- **Recognize the limitations of existing asynchronous tasking models**
- **Introduce a new dynamic task graph programming model called AsyncTask**
- **Overcome the scheduling challenges to support the model**
- **Demonstrate the efficiency of AsyncTask**
- **Conclude the talk**

Why Parallel Computing?

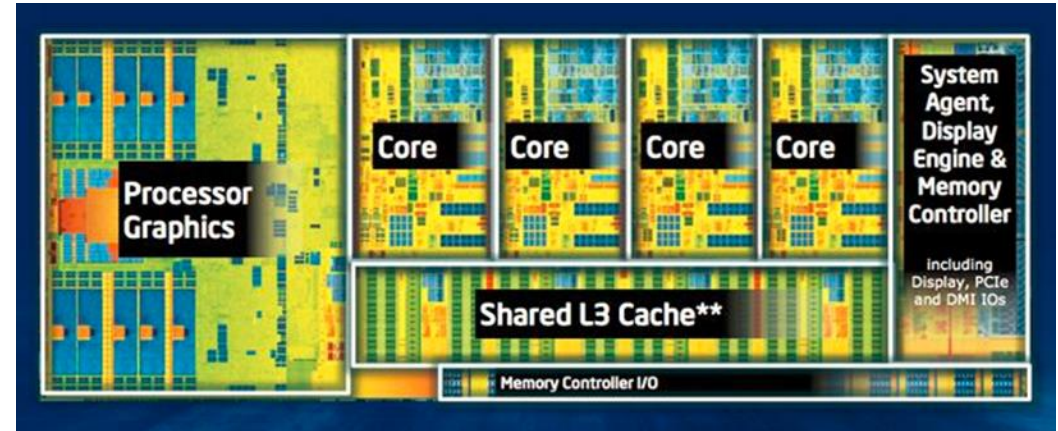
- Advances performance to a new level previously out of reach



Modern Hardware is Designed to Run in Parallel

- **Intel Haswell microarchitecture**

- Released in June 2013
- Typically comes with four cores
- Has an integrated GPU
- 1.4 B transistors with 22 nm technology
- Sophisticated design for ILP acceleration
- Deep pipeline – 16 stages

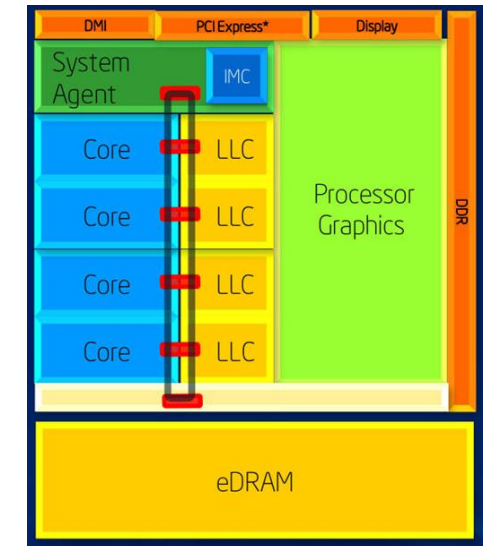


- **Superscalar architecture**

- Can issue and complete multiple independent instructions per cycle

- **Supports hyper-threading tech (HTT)**

- Allows a single physical CPU core to appear as two logical processors to the OS

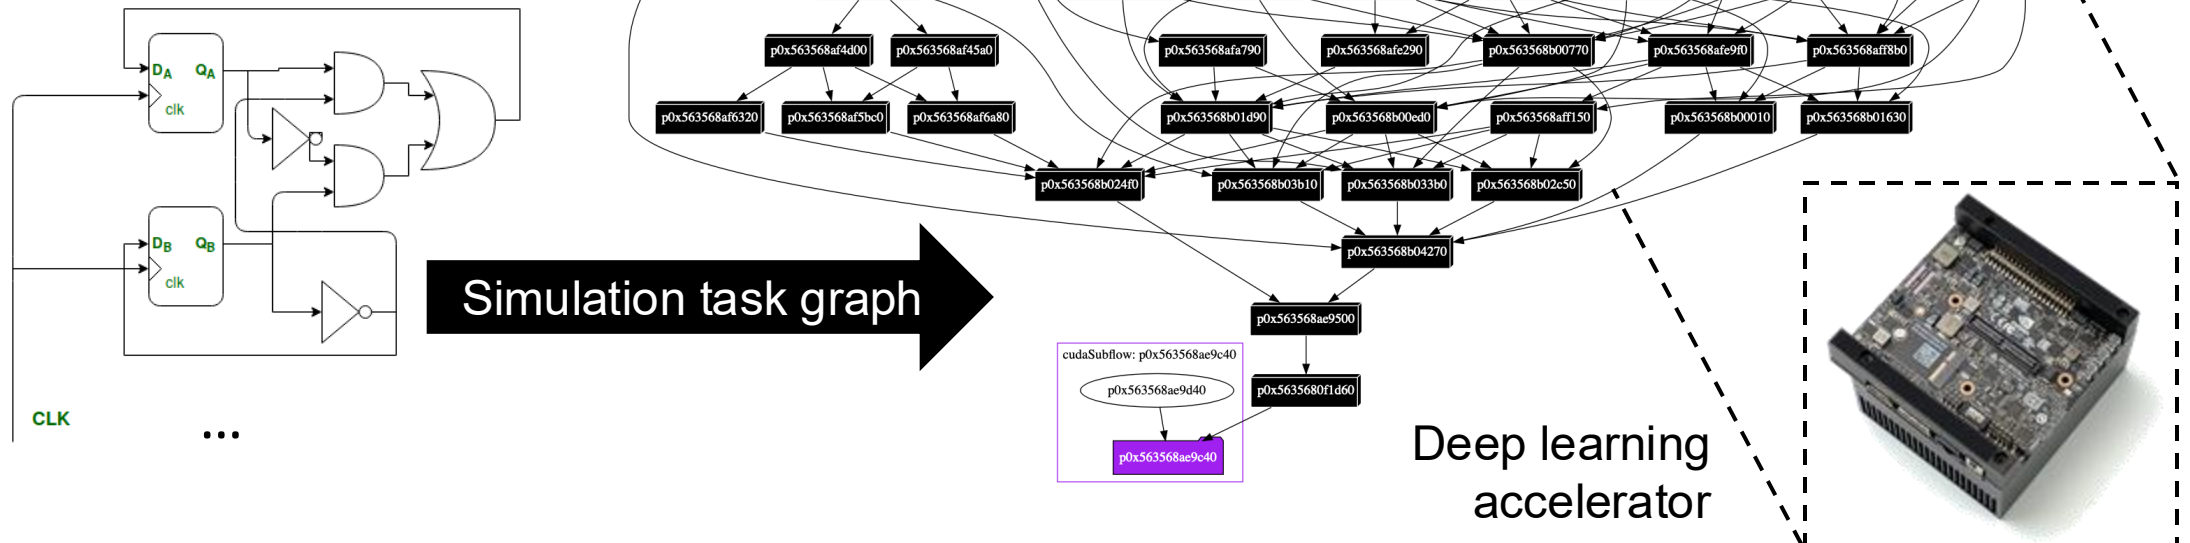


If you don't do parallel programming, you are not utilizing your hardware efficiently ...

Today's Parallel Computing Problem is Very Irregular


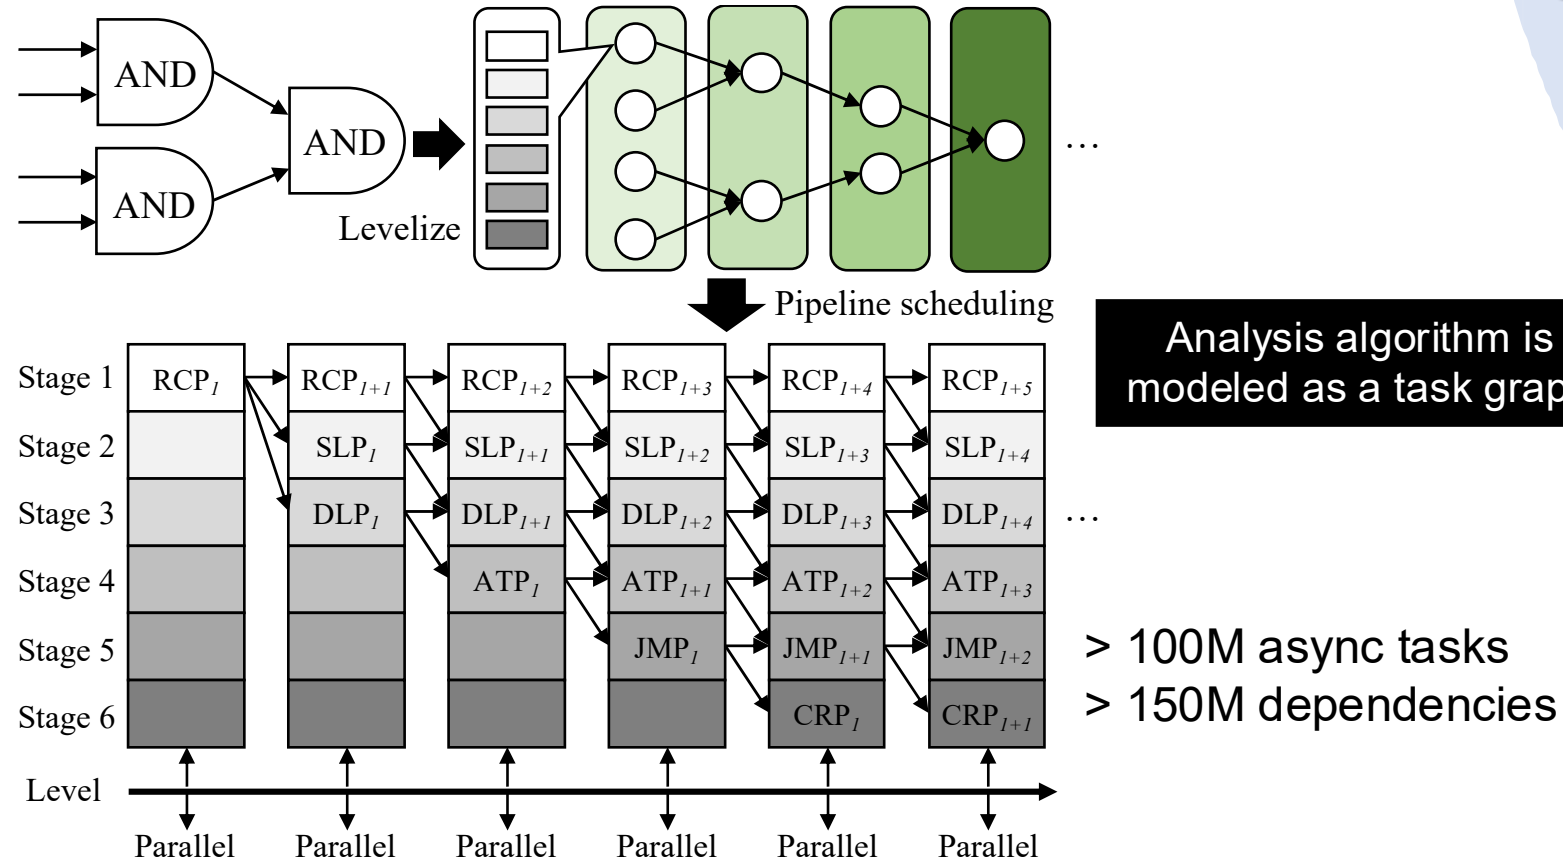
- **Computational task graph of a GPU-parallel circuit simulation workload¹**

- > 500M gates and nets
- > 1000 async kernel tasks
- > 1000 dependencies
- > hours to finish



Another Example of Irregular Parallel Workload

- CPU-parallel VLSI static timing analysis algorithm




Parallelizing such Irregular Problems is Not Easy ...

- **You need to deal with A LOT OF technical details**

- Parallelism abstraction (software + hardware)
- Concurrency control
- Synchronization
- Task and data race avoidance
- Dependency constraints
- Scheduling efficiencies (load balancing)
- Programming productivity
- Performance portability
- ...

- **And, don't forget about trade-offs**

- Performance vs Developer's intent

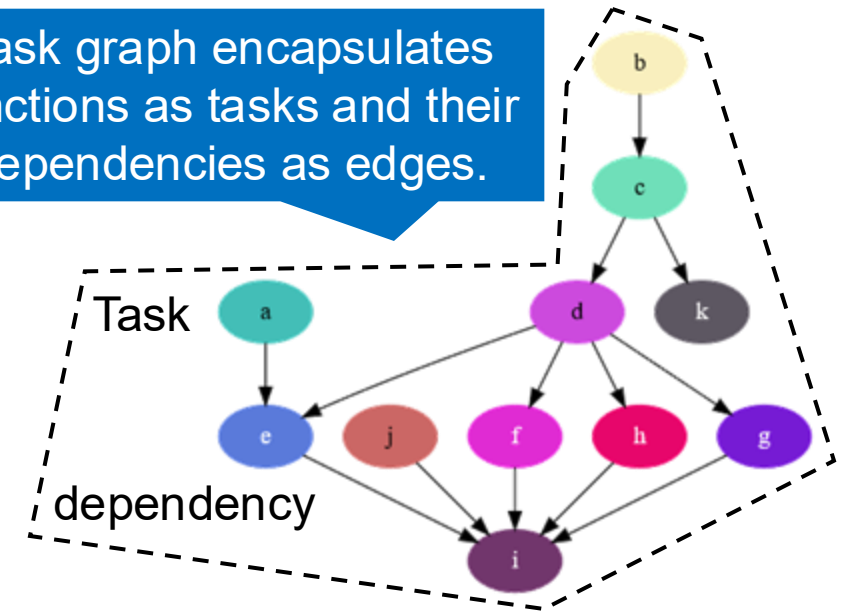


We want a solution that can sit on top to help programmers manage these details as much as possible because programmers care how fast (performance + productivity) they can get things done!

Why Task-parallel Programming (TPP)?

- **TPP is an effective solution for parallelizing irregular workloads**
 - Captures developers' intention in decomposing an algorithm into a *top-down* task graph
 - Delegates difficult scheduling details (e.g., load balancing) to an optimized runtime
- **Modern parallel programming libraries are moving towards task parallelism**
 - OpenMP 4.0 task dependency clauses (`omp depend`)
 - C++26 execution control library (`std::exec`)
 - TBB flow graph (`tbb::flow::graph`)
 - Taskflow control Taskflow graph (CTFG) model
 - ... (many others)

Task graph encapsulates functions as tasks and their dependencies as edges.



OpenMP

StarPU

kokkos



PaRSEC





Takeaways

- Understand the importance of asynchronous tasking with dependencies
- **Recognize the limitations of existing asynchronous tasking models**
- Introduce a new dynamic task graph programming model called AsyncTask
- Overcome the scheduling challenges to support the model
- Demonstrate the efficiency of AsyncTask
- Conclude the talk

Create an Asynchronous Task using `std::async`¹

- A high-level standard library facility to launch a task asynchronously

```
#include <future>
#include <iostream>
```

```
int compute(int v) {
    return v;
}
```

Use `std::async` to asynchronously run the function `compute(42)` on a new thread.

```
int main() {
    std::future<int> fu = std::async(std::launch::async, compute, 42);
    std::cout << fu.get() << std::endl; // prints 42
}
```

Return a `std::future` to wait for this asynchronous task to finish and access its result (i.e., 42)



An Example Implementation of `std::async`

```
template <typename F, typename... Args>
auto async(F&& func, Args&&... args) {
    using ReturnType = std::invoke_result_t<F, Args...>;
    // promise-future pair for intern-thread sync
    std::promise<ReturnType> prom;
    std::future<ReturnType> fu = prom.get_future();
    std::thread t([prom=std::move(prom),
        f=std::forward<F>(func), ...args=std::forward<Args>(args)] () mutable {
        if constexpr(std::is_void_v<ReturnType>) {
            f(std::move(args)...);
            prom.set_value();
        } else {
            prom.set_value(f(std::move(args)...));
        }
    });
    t.detach(); // mimic fire-and-forget behavior of std::async
    return fu;
}
```

I promise you that I will run your function, and you can access the result from the future object ...


We create a thread from a lambda function object that captures the function and its argument (with perfect forwarding¹) and invoke the function in the body.

Build a Task Graph w/ `std::async` and `std::future`

- `std::future` allows us to perform task-specific synchronization

```
auto A = std::async(std::launch::async,
    [](){ std::cout << "A\n"; }
);
A.wait();
auto B = std::async(std::launch::async,
    [](){ std::cout << "B\n"; }
);
auto C = std::async(std::launch::async,
    [](){ std::cout << "C\n"; }
);
B.wait();
C.wait();
auto D = std::async(std::launch::async,
    [](){ std::cout << "D\n"; }
);
D.wait();
```

We need to wait for A to finish before launching B and C asynchronously.



We need to wait for B and C to finish before launching D asynchronously

By properly synchronizing tasks using `future.wait`, we can dynamically create a task graph (i.e., dynamic task graph)

Sender-Receiver Version (with `std::exec`¹)

- A standardized abstraction for composing tasks and dependencies

```
exec::static_thread_pool pool;  
auto scheduler = pool.get_scheduler();
```

Schedule tasks on a pool of worker threads

```
// create a sender task for A  
auto sa = exec::then(exec::schedule(scheduler), []{ std::cout<<"A\n"; });  
exec::sync_wait(sa); // wait for A
```

```
// create two parallel sender tasks for B and C  
auto sb = exec::then(exec::schedule(scheduler), []{ std::cout<<"B\n"; });  
auto sc = exec::then(exec::schedule(scheduler), []{ std::cout<<"C\n"; });  
exec::sync_wait(exec::when_all(sb, sc)); // wait for B and C
```

```
// create a sender task for D  
auto sd = exec::then(exec::schedule(scheduler), []{ std::cout<<"D\n"; });  
exec::sync_wait(sd); // wait for D
```

Intel's TBB Library with `tbb::task_group`¹

- A class to create asynchronous tasks and wait for their completion

```
tbb::task_group tg;
```

A class in TBB to create asynchronous tasks and wait for their completion

```
// A
tg.run([] { std::cout << "A\n"; });
tg.wait();
```

Need to `task_group::wait` on A before running B and C

```
// B and C in parallel
tg.run([] { std::cout << "B\n"; });
tg.run([] { std::cout << "C\n"; });
tg.wait();
```

Need to `task_group::wait` on B and C before running D

```
// D
tg.run([] { std::cout << "D\n"; });
tg.wait();
```


OpenMP Tasking Model with depend Clauses¹

- Leverages compiler directives to define tasks and dependencies

```
#omp parallel
```

```
{
```

```
  int A_B, A_C, B_D, C_D;
```

```
  #pragma omp task depend(out: A_B, A_C)
```

```
  std::cout << "TaskA\n";
```

```
  #pragma omp task depend(in: A_B; out: B_D)
```

```
  std::cout << "TaskB\n";
```

```
  #pragma omp task depend(in: A_C; out: C_D)
```

```
  std::cout << "TaskB\n";
```


```
  #pragma omp task depend(in: B_D, C_D)
```

```
  std::cout << "TaskB\n";
```

```
}
```

Define dependency handles

Specify task dependencies using in and out clauses when creating an OpenMP task



With these OpenMP directives, the compiler will insert parallel code that launches asynchronous tasks and enforces their dependencies.

OpenCilk Version

- **A fork-join programming model relying on compiler-generated parallel code**
 - With language extensions like `cilk_spawn` and `cilk_sync`

```
void A() { std::cout << "A\n"; }  
void B() { std::cout << "B\n"; }  
void C() { std::cout << "C\n"; }  
void D() { std::cout << "D\n"; }  
int main() {  
    A();  
  
    cilk_spawn B();  
    C();  
    cilk_sync;  
  
    D();  
}
```

You need a compiler that supports OpenCilk syntax to run this code.

Spawn a child task on B using `cilk_spawn` and continue with C in the main thread

Synchronize both B and C using `cilk_sync` before running task D

Limitations of Existing Async Tasking Models

- ❌ **Tasks and their dependencies are decoupled during task graph creation**
 - If dependencies are not expressed alongside the task creation logic, it's difficult to reason about the overall task graph structure
 - Without a clear dependency structure, the runtime loses opportunities to optimize task placement and load balancing when constructing an asynchronous task
- ❌ **Correct placement of `wait` calls is left to programmers**
 - Programmers must determine a correct synchronization order at a fine-grained level
 - In the worst case, the number of `waits` equals the number of dependencies
 - In practice, many applications only care about the completion of the entire task graph instead of intermediate tasks, making such fine-grained waiting unnecessary, costly, and buggy
- ❌ **Limited support for building highly dynamic task graphs**
 - Highly dynamic task graphs → those whose structures, dependencies, and task content are highly dependent on runtime variables or dynamic control-flow results
 - Ex: OpenMP is not a good fit for this scenario as it relies on static compiler directives
- ❌ **May require a non-standard C++ compiler to generate parallel code**




Takeaways

- Understand the importance of asynchronous tasking with dependencies
- Recognize the limitations of existing asynchronous tasking models
- **Introduce a new dynamic task graph programming model called AsyncTask**
- Overcome the scheduling challenges to support the model
- Demonstrate the efficiency of AsyncTask
- Conclude the talk

Static vs Dynamic Task Graph Programming (TGP)

- All examples we've discussed so far are dynamic TGP (DTGP)

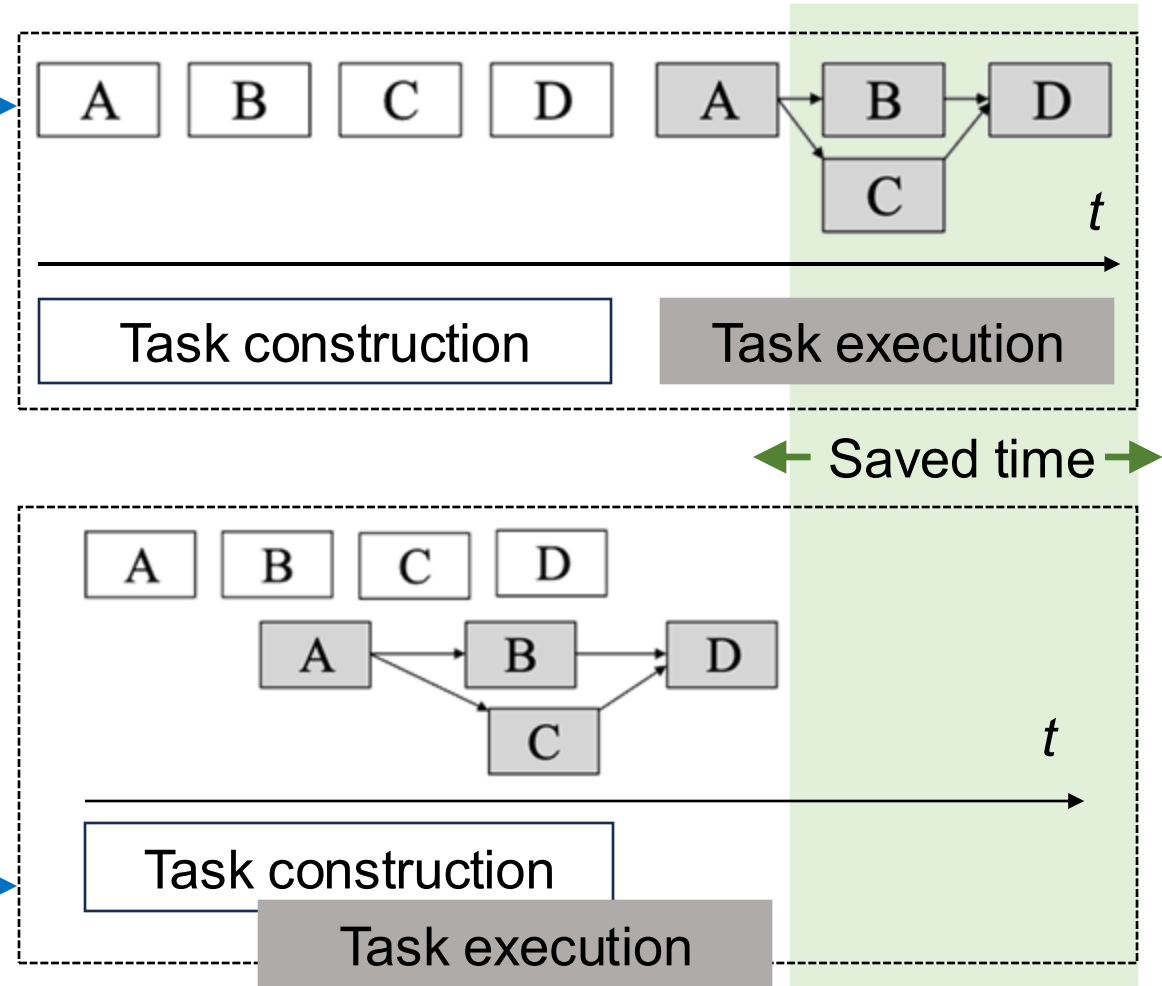
In static TGP (STGP), execution follows the *construct-and-run* model



STGP

DTGP

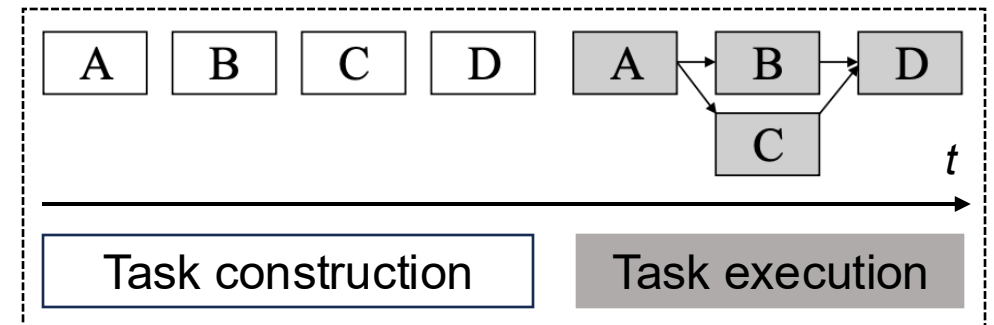

In DTGP, tasks can start as soon as their dependencies are met



Static Task Graph Programming in Taskflow¹

`#include <taskflow/taskflow.hpp> // Live: https://godbolt.org/z/j8hx3xnnx`

```
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B, C);
    D.succeed(B, C);
    executor.run(taskflow).wait();
    return 0;
}
```



AsyncTask: Dynamic Task Graph Programming in Taskflow

// Live: <https://godbolt.org/z/j76ThGbWK>

```
tf::Executor executor;

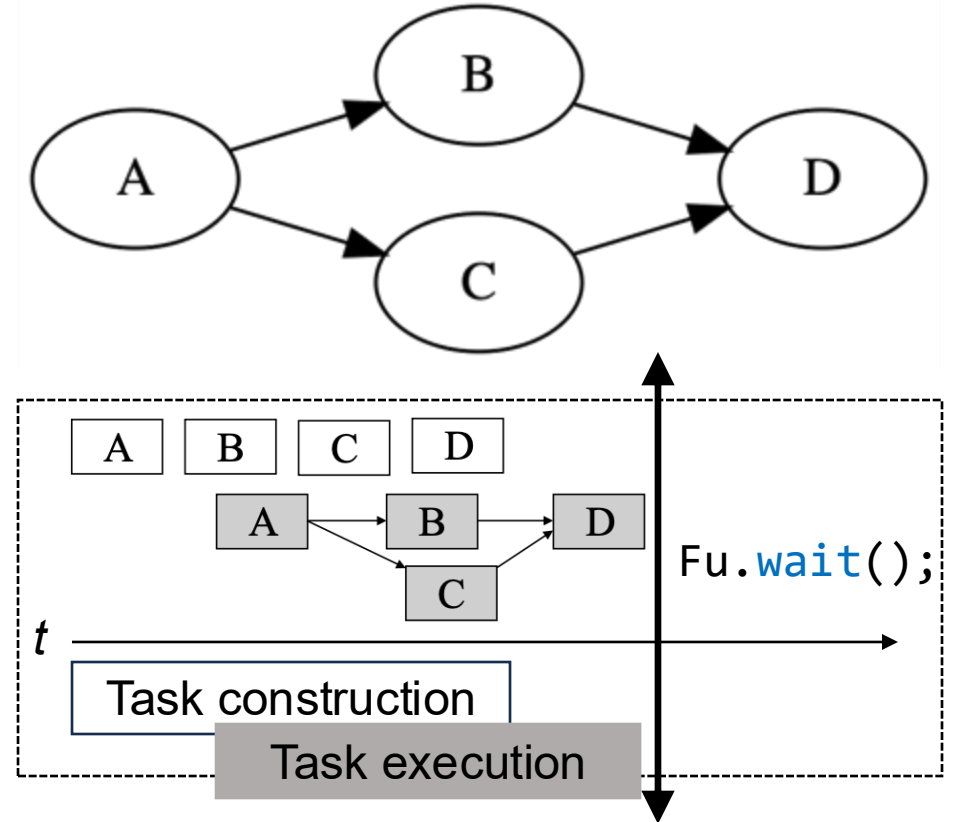
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});

auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);

auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);

auto [D, Fu] = executor.dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
```

```
Fu.wait();
```



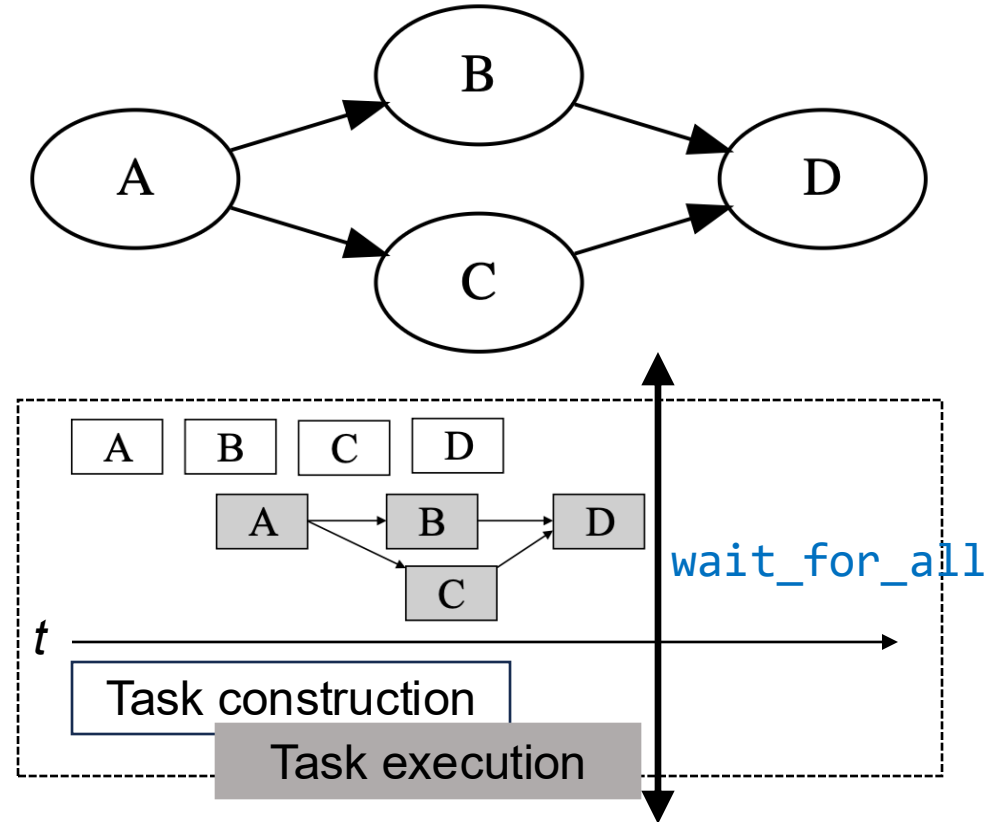
Specify variable task dependencies using C++ variadic parameter pack

Wait for All Tasks to Finish

// Live: <https://godbolt.org/z/T87PrTarx>

```
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
```

```
executor.wait_for_all();
```



Wait for the entire graph to finish.



Need a Correct Topological Order

```

auto A = executor.silent_dependent_async(
    [](){ std::cout << "TaskA\n"; }
);
auto B = executor.silent_dependent_async(
    [](){ std::cout << "TaskB\n"; }, A
);
auto C = executor.silent_dependent_async(
    [](){ std::cout << "TaskC\n"; }, A
);
auto D = executor.silent_dependent_async(
    [](){ std::cout << "TaskD\n"; }, B, C
);

```

Topological order #1: $A \rightarrow B \rightarrow C \rightarrow D$



Topological order #2: $A \rightarrow C \rightarrow B \rightarrow D$

```

auto A = executor.silent_dependent_async(
    [](){ std::cout << "TaskA\n"; }
);
auto C = executor.silent_dependent_async(
    [](){ std::cout << "TaskC\n"; }, A
);
auto B = executor.silent_dependent_async(
    [](){ std::cout << "TaskB\n"; }, A
);
auto D = executor.silent_dependent_async(
    [](){ std::cout << "TaskD\n"; }, B, C
);


```

Incorrect Topological Order ...

```
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B-is-unavailable-yet, C-is-unavailable-yet);

auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);

executor.wait_for_all();
```



An incorrect topological order (A→D→B→C) prevents you from expressing a correct dynamic task graph.

Variable Range of Task Dependencies

- **Both methods can take a variable range of dependent-async tasks**
 - Useful when the task dependencies come as a runtime variable (e.g., loaded from a file)

```
// Live: https://godbolt.org/z/6Pvco4KeE
std::vector<tf::AsyncTask> tasks = {
    executor.silent_dependent_async([](){ std::cout << "TaskA\n"; }),
    executor.silent_dependent_async([](){ std::cout << "TaskB\n"; }),
    executor.silent_dependent_async([](){ std::cout << "TaskC\n"; }),
    executor.silent_dependent_async([](){ std::cout << "TaskD\n"; })
};
// create a dependent-async tasks that depends on tasks, A, B, C, and D
executor.dependent_async([]() {}, tasks.begin(), tasks.end());
// create a silent-dependent-async task that depends on tasks, A, B, C, and D
executor.silent_dependent_async([]() {}, tasks.begin(), tasks.end());
```

While this feature may look trivial, I found it very difficult to achieve with existing asynchronous tasking libraries because their task creation and dependency expression are decoupled from each other ...

DTGP is Flexible for Runtime-driven Execution

- Assemble task graphs driven by runtime variables and control-flow results

```
if (a == true) {  
    G1 = build_task_graph1();  
    if (b == true) {  
        G2 = build_task_graph2();  
        G1.precede(G2);  
        if (c == true) {  
            ... // defined other TGPs  
        }  
    }  
    else {  
        G3 = build_task_graph3();  
        G1.precede(G3);  
    }  
}
```

```
G1 = build_task_graph1();  
G2 = build_task_graph2();  
if (G1.num_tasks() == 100) {  
    G1.precede(G2);  
}  
else {  
    G3 = build_task_graph3();  
    G1.precede(G2, G3);  
    if(G2.num_dependencies()>=10){  
        ... // define another TGP  
    } else {  
        ... // define another TGP  
    }  
}
```

This type of dynamic task graph is very difficult to achieve using static task graph programming ...

AsyncTask doesn't Touch Data Abstraction

- **Focus on coarse-grained task parallelism not fine-grained data parallelism**
 - Our goal is to have users describe tasks and their dependencies in an expressive language

```
template <typename F, typename... Tasks>  
auto dependent_async(F&& func, Tasks&&... tasks) {  
    ...  
}
```

This is how `std::async` is implemented
(e.g., args are captured with perfect forwarding)

- Users describe `func` as a lambda and capture necessary data or `func` arguments themselves
- **The advantage of this decision is twofold:**
 - Users retain full control over data layout and ownership, allowing them to optimize data structures and memory layout for their specific application domains
 - Letting users decide how and where to store data keeps `AsyncTask` lightweight and non-intrusive – no need to modify existing data structures to fit our framework
 - Ex: Models that count on data abstraction (e.g., Fastflow, TBB pipeline) require users to rewrite their code to library-specific data abstraction in order to gain parallelism

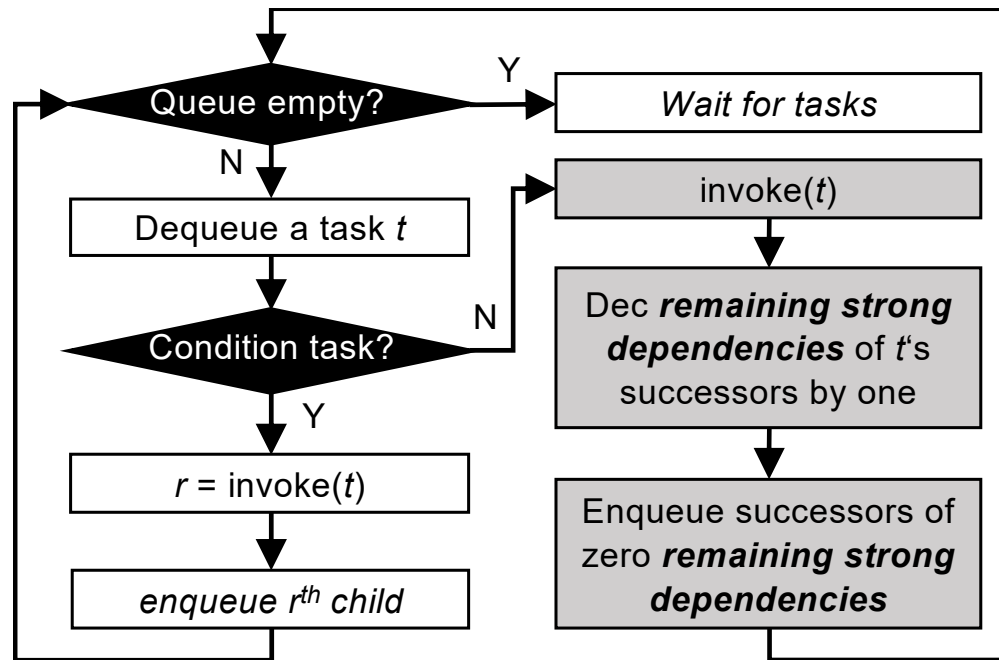


Takeaways


- Understand the importance of asynchronous tasking with dependencies
- Recognize the limitations of existing asynchronous tasking models
- Introduce a new dynamic task graph programming model called AsyncTask
- **Overcome the scheduling challenges to support the model**
- Demonstrate the efficiency of AsyncTask
- Conclude the talk

Overview: Taskflow's Work-stealing Scheduler¹

- Task-level scheduling

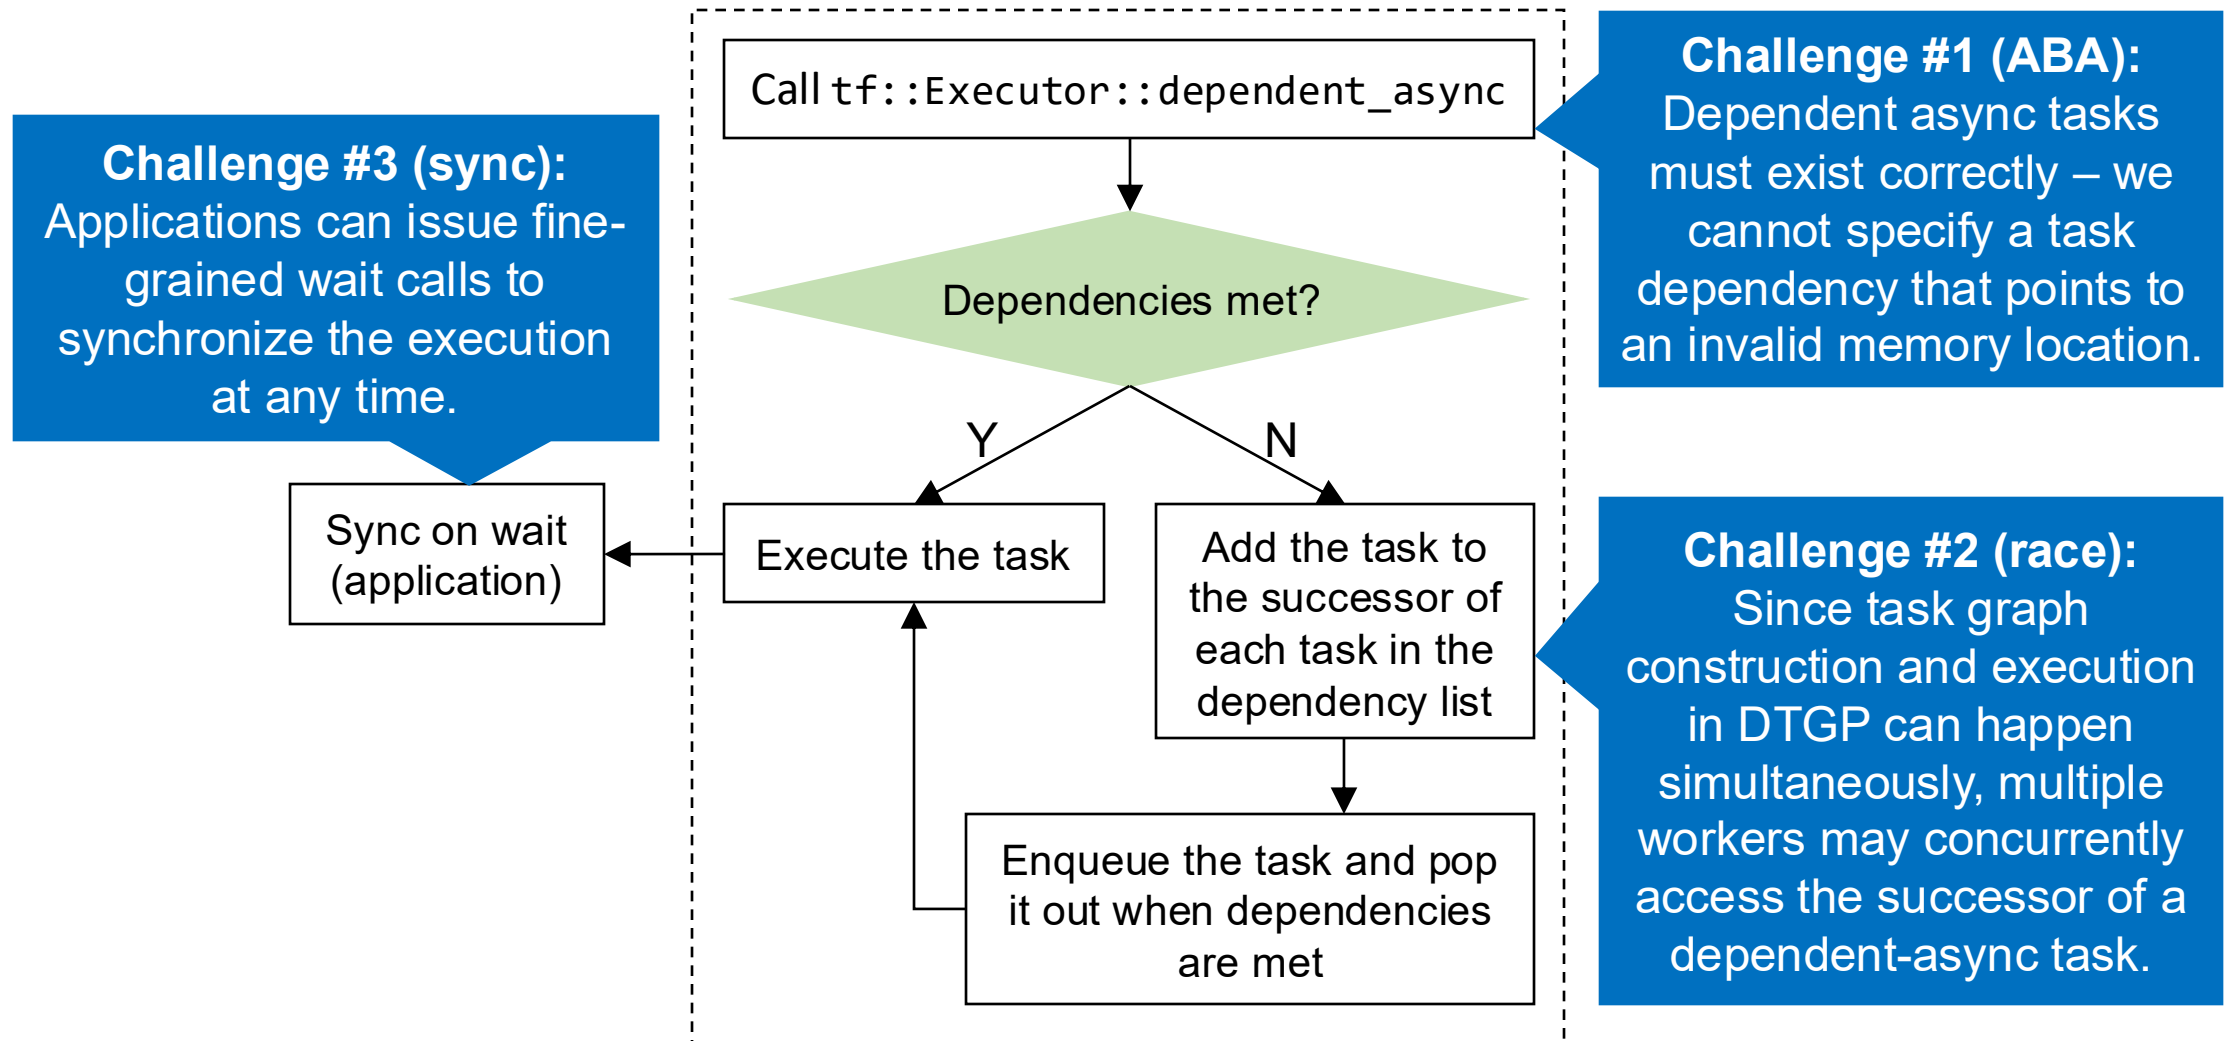


- Worker-level scheduling



Key results: schedule tasks with in-graph control flow with a strong balance between the number of active workers and dynamically generated tasks – *low latency, energy efficient, and high throughput*

Scheduling a Dynamic Task Graph

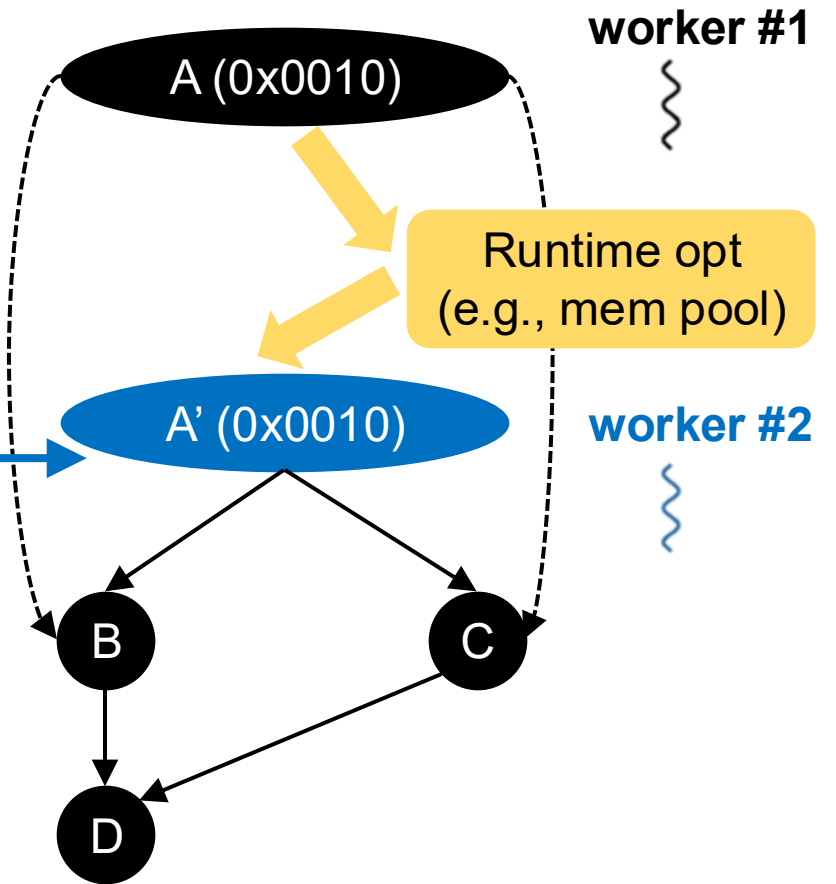


Solving Challenge #1: ABA Problem

```
tf::Executor executor;

auto A = executor.silent_dependent_async([]{
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([]{
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([]{
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async([]{
    std::cout << "TaskD\n";
}, B, C);

executor.wait_for_all();
```



Retain a Shared Ownership of Each Task Needed

```
tf::Executor executor;
```

```
tf::AsyncTask A = executor.silent_dependent_async([]{  
    std::cout << "TaskA\n";  
});
```

```
tf::AsyncTask B = executor.silent_dependent_async([]{  
    std::cout << "TaskB\n";  
}, A);
```

```
tf::AsyncTask C = executor.silent_dependent_async([]{  
    std::cout << "TaskC\n";  
}, A);
```

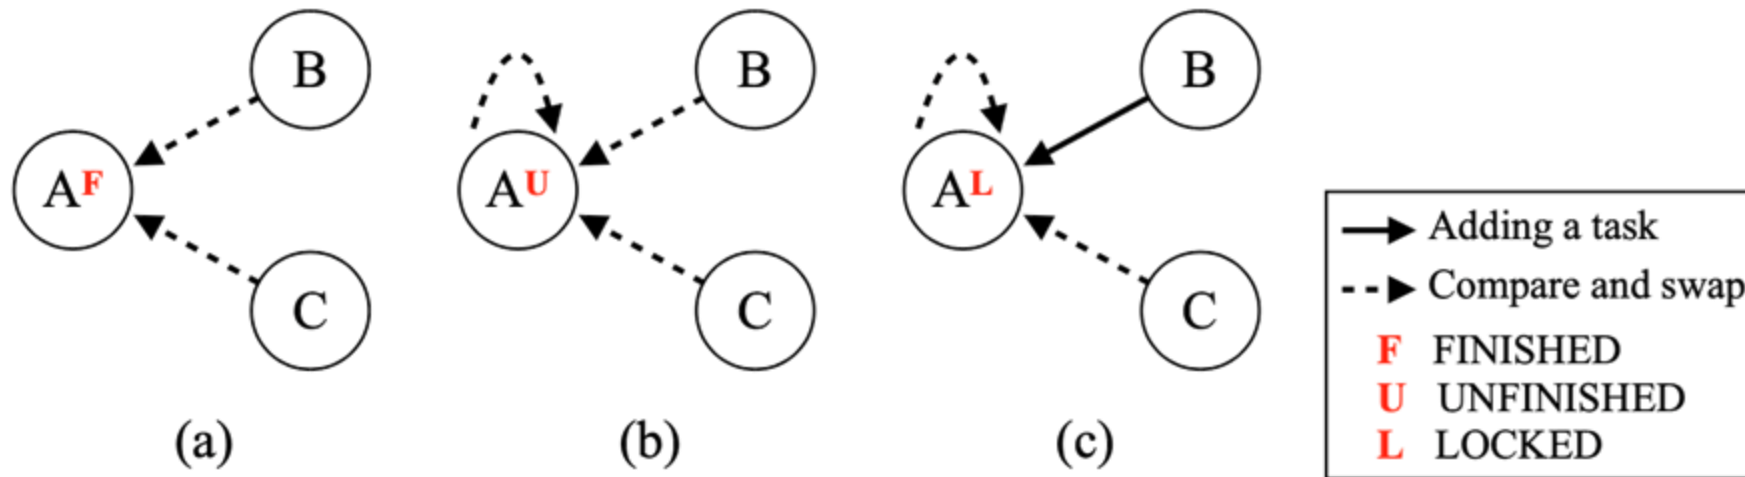
```
tf::AsyncTask D = executor.silent_dependent_async([]{  
    std::cout << "TaskD\n";  
}, B, C);
```

```
executor.wait_for_all();
```

tf::AsyncTask acts like
a std::shared_ptr to
ensure tasks stay alive
when they are used

Solving Challenge #2: Data Race

- **Both B and C want to add themselves to the successors of A**
 - Meanwhile, A may want to remove some of its successor when the task finishes



- **Use compare-and-swap (CAS) with spinning to enable exclusive access**
 - Spinning does not incur much overhead because most task graphs are sparse
 - If your task graph is very dense, probably DTGP is not the right solution to your application

Solving Challenge #3: Synchronization

- **Users can issue both coarse- and fine-grained synchronizations at any time**

- Coarse-grained sync: `executor.wait_for_all()`
- Fine-grained sync: `future.wait()`

```
tf::Executor executor;
auto A = executor.silent_dependent_async([]{});
auto B = executor.silent_dependent_async([], A);
executor.wait_for_all(); // wait for A and B
```

```
auto C = executor.silent_dependent_async([], A);
auto D = executor.silent_dependent_async([], B, C);
executor.wait_for_all(); // wait for C and D
```

```
// lock-based sync
std::unique_lock lock(mtx);
cv.wait(lock, [&]() {
    return num_tasks == 0;
});
```

```
// atomic wait-based sync
auto n = num_tasks.load();
while(n != 0) {
    num_tasks.wait(n);
    n = num_tasks.load();
};
```

We leverage C++20 atomic variables to perform waiting/notifying operations¹, which allows much of the synchronization to occur in user space rather than in the kernel space (~11% performance improvement).

Our Scheduling Algorithm is Lock-free¹

Algorithm 1 `dependent_async(callable, deps)`

```

1: Create a future
2:  $num\_deps \leftarrow \text{sizeof}(deps)$ 
3:  $task \leftarrow \text{initialize\_task}(callable, num\_deps, future)$ 
4: for all  $dep \in deps$  do
5:   process_dependent(task, dep, num_deps)
6: end for
7: if  $num\_deps == 0$  then
8:   schedule_async_task(task)
9: end if
10: return ( $task, future$ )

```

Algorithm 2 `process_dependent(task, dep, num_deps)`

```

1:  $dep\_state \leftarrow dep.state$ 
2:  $target\_state \leftarrow UNFINISHED$ 
3: if  $dep\_state.CAS(target\_state, LOCKED)$  then
4:    $dep.successors.push(task)$ 
5:    $dep\_state \leftarrow UNFINISHED$ 
6: else if  $target\_state == FINISHED$  then
7:    $num\_deps \leftarrow \text{AtomDec}(task.join\_counter)$ 
8: else
9:   goto line 2
10: end if

```

Algorithm 3 `schedule_async_task(task)`

```

1:  $target\_state \leftarrow UNFINISHED$ 
2: while not  $task.state.CAS(target\_state, FINISHED)$  do
3:    $target\_state \leftarrow UNFINISHED$ 
4: end while
5: Invoke(task.callable)
6: for all  $successor \in task.successors$  do
7:   if  $\text{AtomDec}(successor.join\_counter) == 0$  then
8:     schedule_async_task(successor)
9:   end if
10: end for
11: if  $\text{AtomDec}(task.ref\_count) == 0$  then
12:   Delete task
13: end if

```

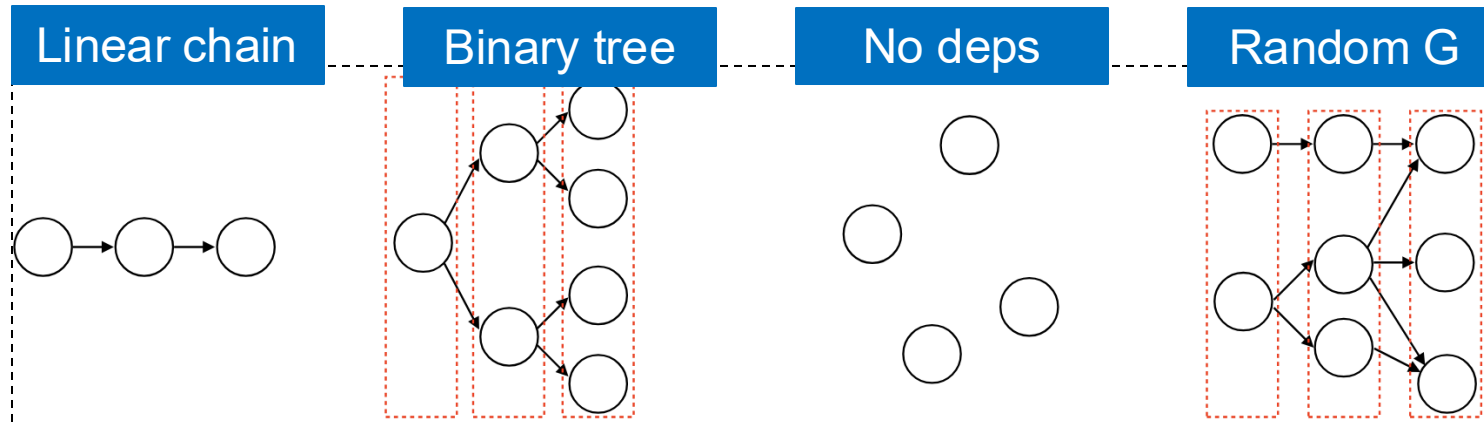


Takeaways

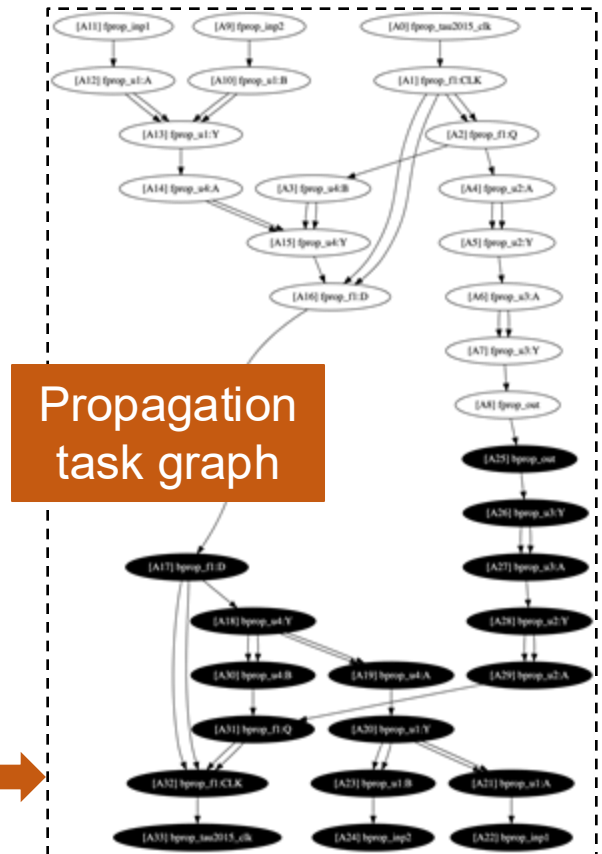
- Understand the importance of asynchronous tasking with dependencies
- Recognize the limitations of existing asynchronous tasking models
- Introduce a new dynamic task graph programming model called AsyncTask
- Overcome the scheduling challenges to support the model
- **Demonstrate the efficiency of AsyncTask**
- Conclude the talk

Evaluation of AsyncTask (1/2)

- **Evaluated on both microbenchmarks and a real-world application**
 - Study the performance of AsyncTask w/o and w/ the impact of application tasks
- **Microbenchmarks**
 - Measure the performance on four commonly used graph patterns



- **Real-world application: VLSI Static Timing Analysis¹**
 - Parallelize the timing propagation algorithm using AsyncTask

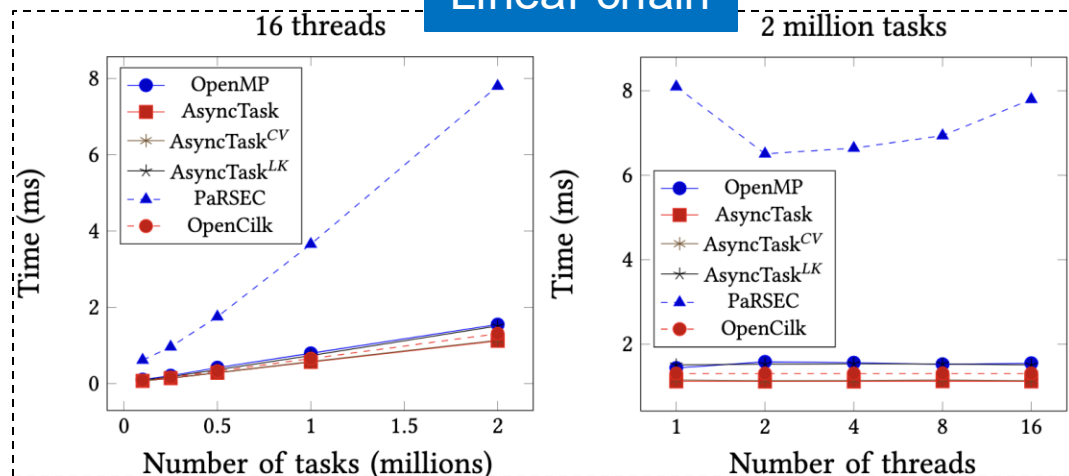


Evaluation of AsyncTask (2/2)

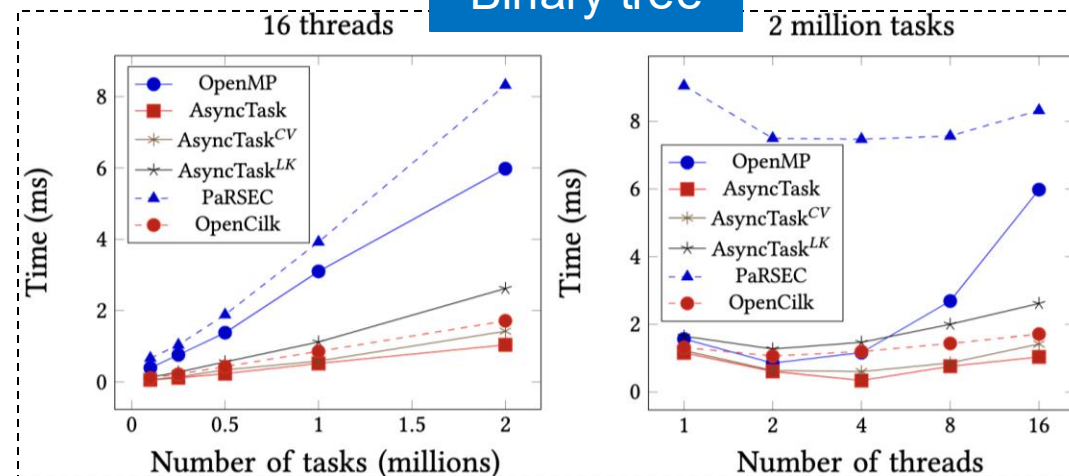
- **We consider the following baselines:**
 - OpenMP tasking: <https://www.openmp.org/spec-html/5.0/openmps99.html>
 - PaRSEC: <https://github.com/ICLDisco/parsec>
 - OpenCilk: <https://github.com/opencilk>
 - AsyncTask^{LK}: replaced AsyncTask's scheduler with OpenMP's task scheduler¹
 - 💡 We want to see how good our lock-free scheduling algorithm is
 - AsyncTask^{CV}: replaced AsyncTask's atomic wait with `std::condition_variable`
 - 💡 We want to see how good our C++20-based notification algorithm is
- **We compiled all programs using g++12 with `-std=c++20` and `-O3`**
 - 64-bit Linux machine with 128 GB RAM and 20 Intel i5-13500 CPU cores at 4.8 GHz
 - All data is an average of ten runs to reduce variance

Microbenchmarks

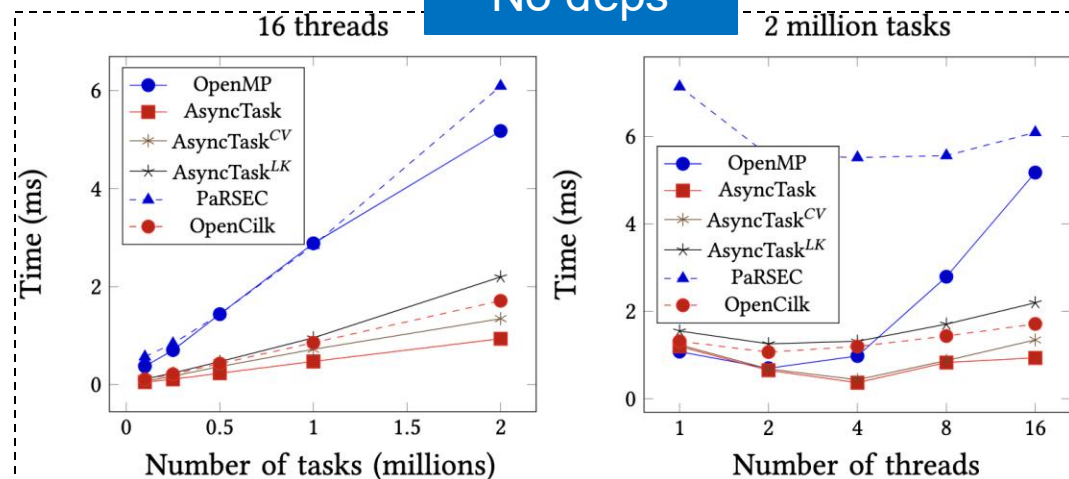
Linear chain



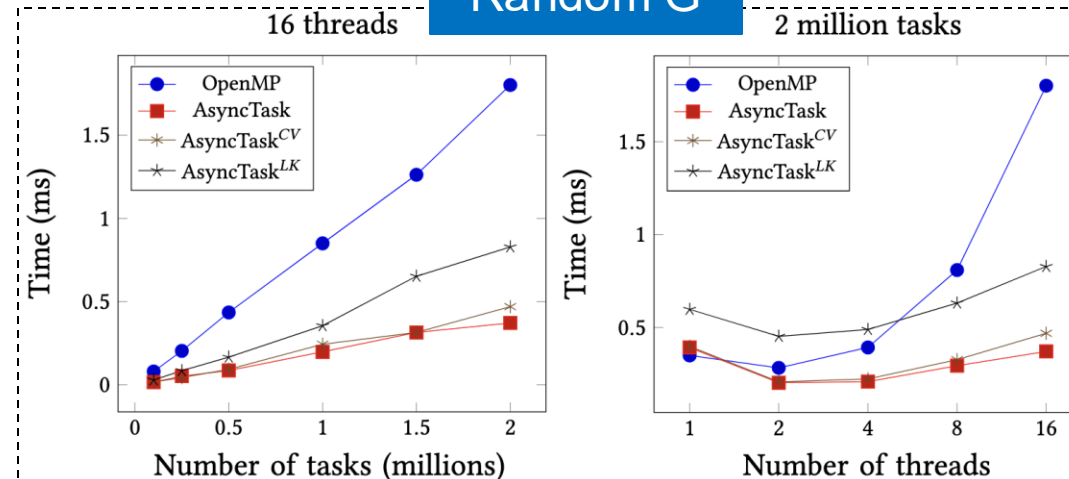
Binary tree



No deps




Random G



Ease of Use of AsyncTask

- **Surveyed 300+ graduate students in my HPC course at UW-Madison**
 - Asked students to finish four microbenchmarks using five DTGP models
 - Rated each library in 1–10 (the higher the better) by the end of this programming assignment



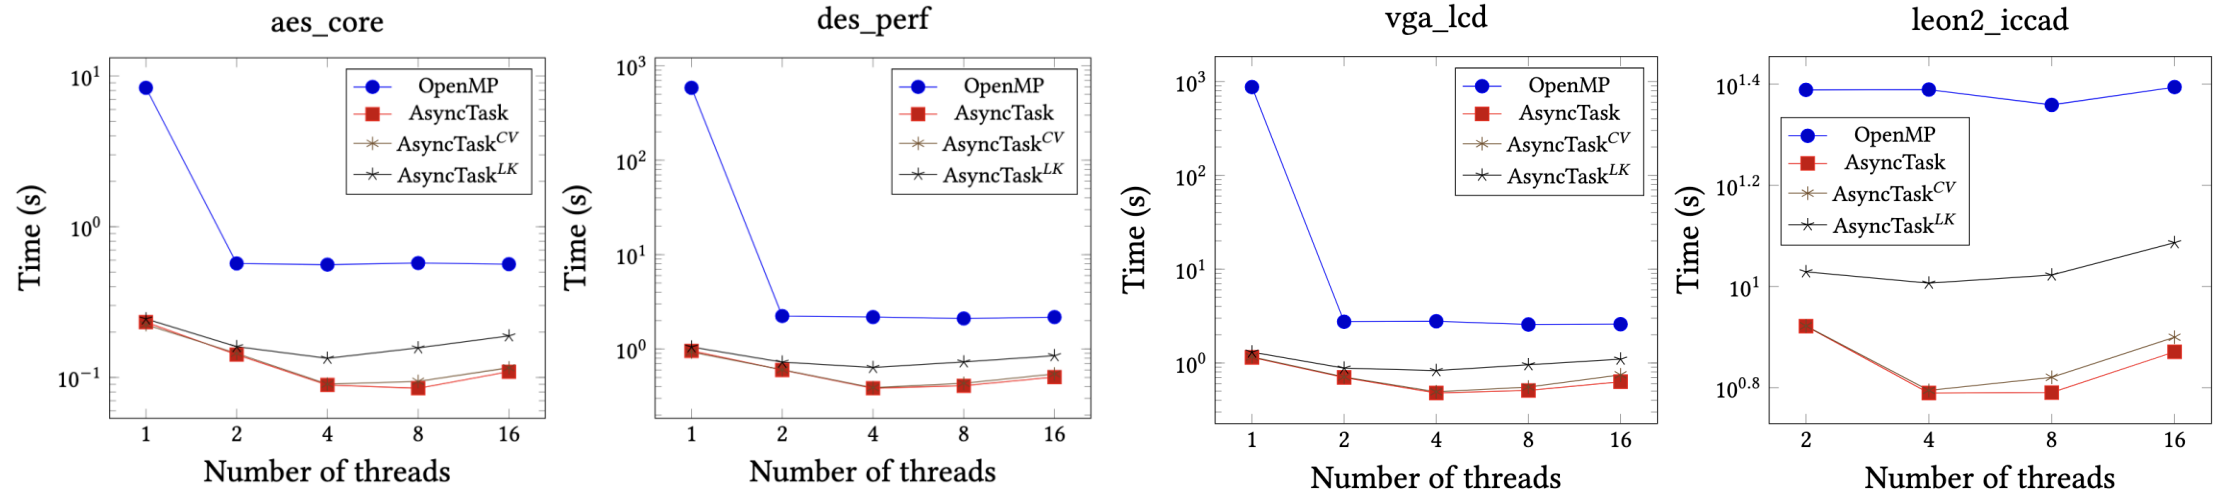

Real-world Application: Static Timing Analysis (STA)

- Implemented task-parallel STA¹ using AsyncTask

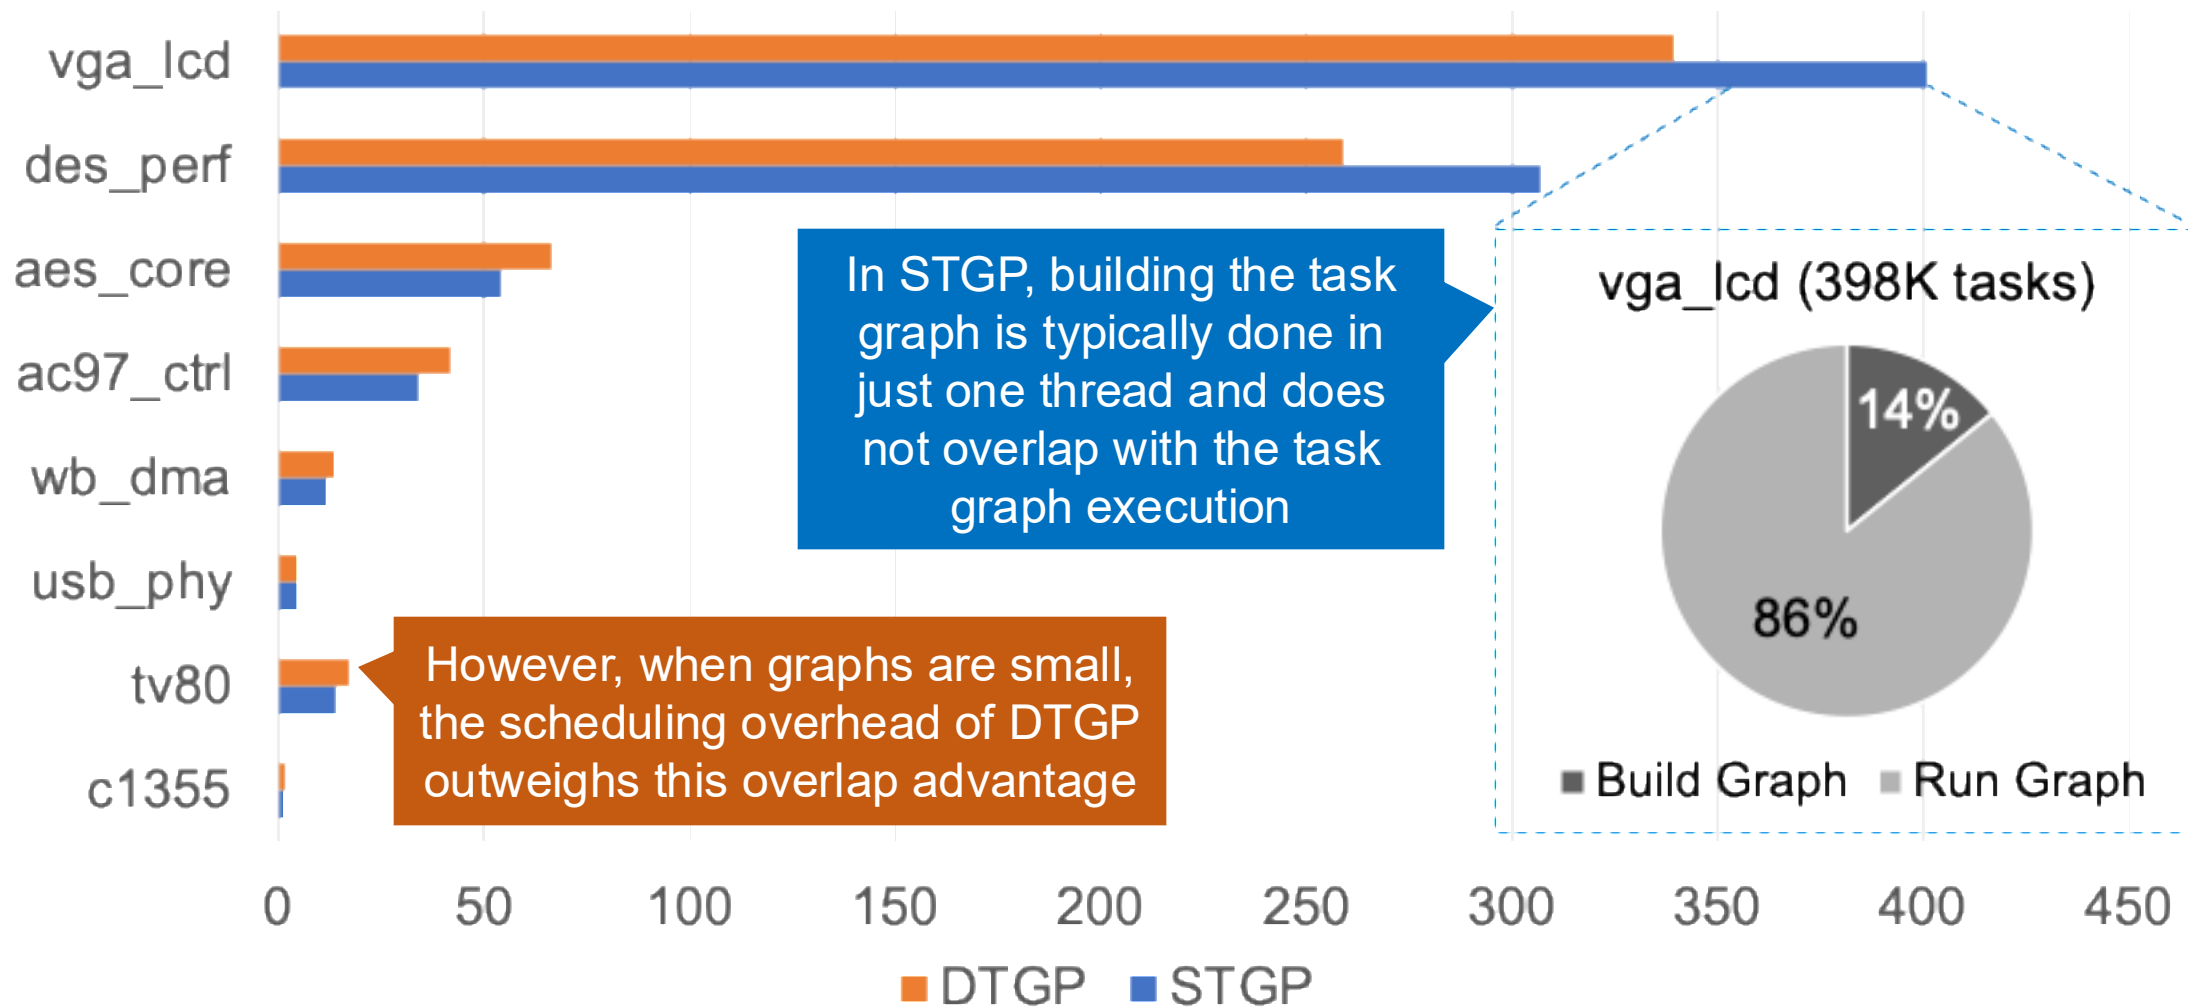
- Models the timing propagation as a dynamic task graph
- Propagates timing data from inputs to outputs

- Evaluated on four industrial circuit graphs

- aes_core: 66,751 tasks and 86,446 dependencies
- des_perf: 303,690 tasks and 387,291 dependencies
- vga_lcd: 397,809 tasks and 498,863 dependencies
- leon2: 4,328,255 tasks and 7,984,262 dependencies



Runtime Comparison: STGP¹ vs DTGP





Takeaways

- Understand the importance of asynchronous tasking with dependencies
- Recognize the limitations of existing asynchronous tasking models
- Introduce a new dynamic task graph programming model called AsyncTask
- Overcome the scheduling challenges to support the model
- Demonstrate the efficiency of AsyncTask
- **Conclude the talk**

Everything has been Integrated to Taskflow¹

- **Taskflow is a header-only C++ library for task-parallel programming**
 - Started in 2018 to help DARPA parallelize critical design automation workloads
- **Using AsyncTask is very easy**

```
# clone the Taskflow project
```

```
~$ git clone https://github.com/taskflow/taskflow.git
```

```
~$ cd taskflow
```

```
# compile your program and tell your compiler where to find Taskflow header files
```

```
~$ g++ -std=c++20 examples/simple.cpp -I ./ -O2 -pthread -o simple
```

```
~$ ./simple
```

```
TaskA
```

```
TaskC
```

```
TaskB
```


```
TaskD
```



Thank you for using Taskflow!



Thank you for Sponsoring Taskflow!



Google Summer of Code



Questions?



Taskflow: <https://taskflow.github.io>



Static task graph parallelism

// Live: <https://godbolt.org/z/j8hx3xnnx>

```
tf::Taskflow taskflow;
tf::Executor executor;
auto [A, B, C, D] = taskflow.emplace(
    [](){ std::cout << "TaskA\n"; },
    [](){ std::cout << "TaskB\n"; },
    [](){ std::cout << "TaskC\n"; },
    [](){ std::cout << "TaskD\n"; }
);

A.precede(B, C);
D.succeed(B, C);
executor.run(taskflow).wait();
```

Dynamic task graph parallelism

// Live: <https://godbolt.org/z/T87PrTarx>

```
tf::Executor executor;
auto A = executor.silent_dependent_async([]{
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([]{
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([]{
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async([]{
    std::cout << "TaskD\n";
}, B, C);
executor.wait_for_all();
```