



Taskflow: A General-purpose Task-parallel Programming System

Dr. Tsung-Wei (TW) Huang, Assistant Professor
Department of Electrical and Computer Engineering
University of Wisconsin at Madison, Madison, WI

<https://tsung-wei-huang.github.io/>

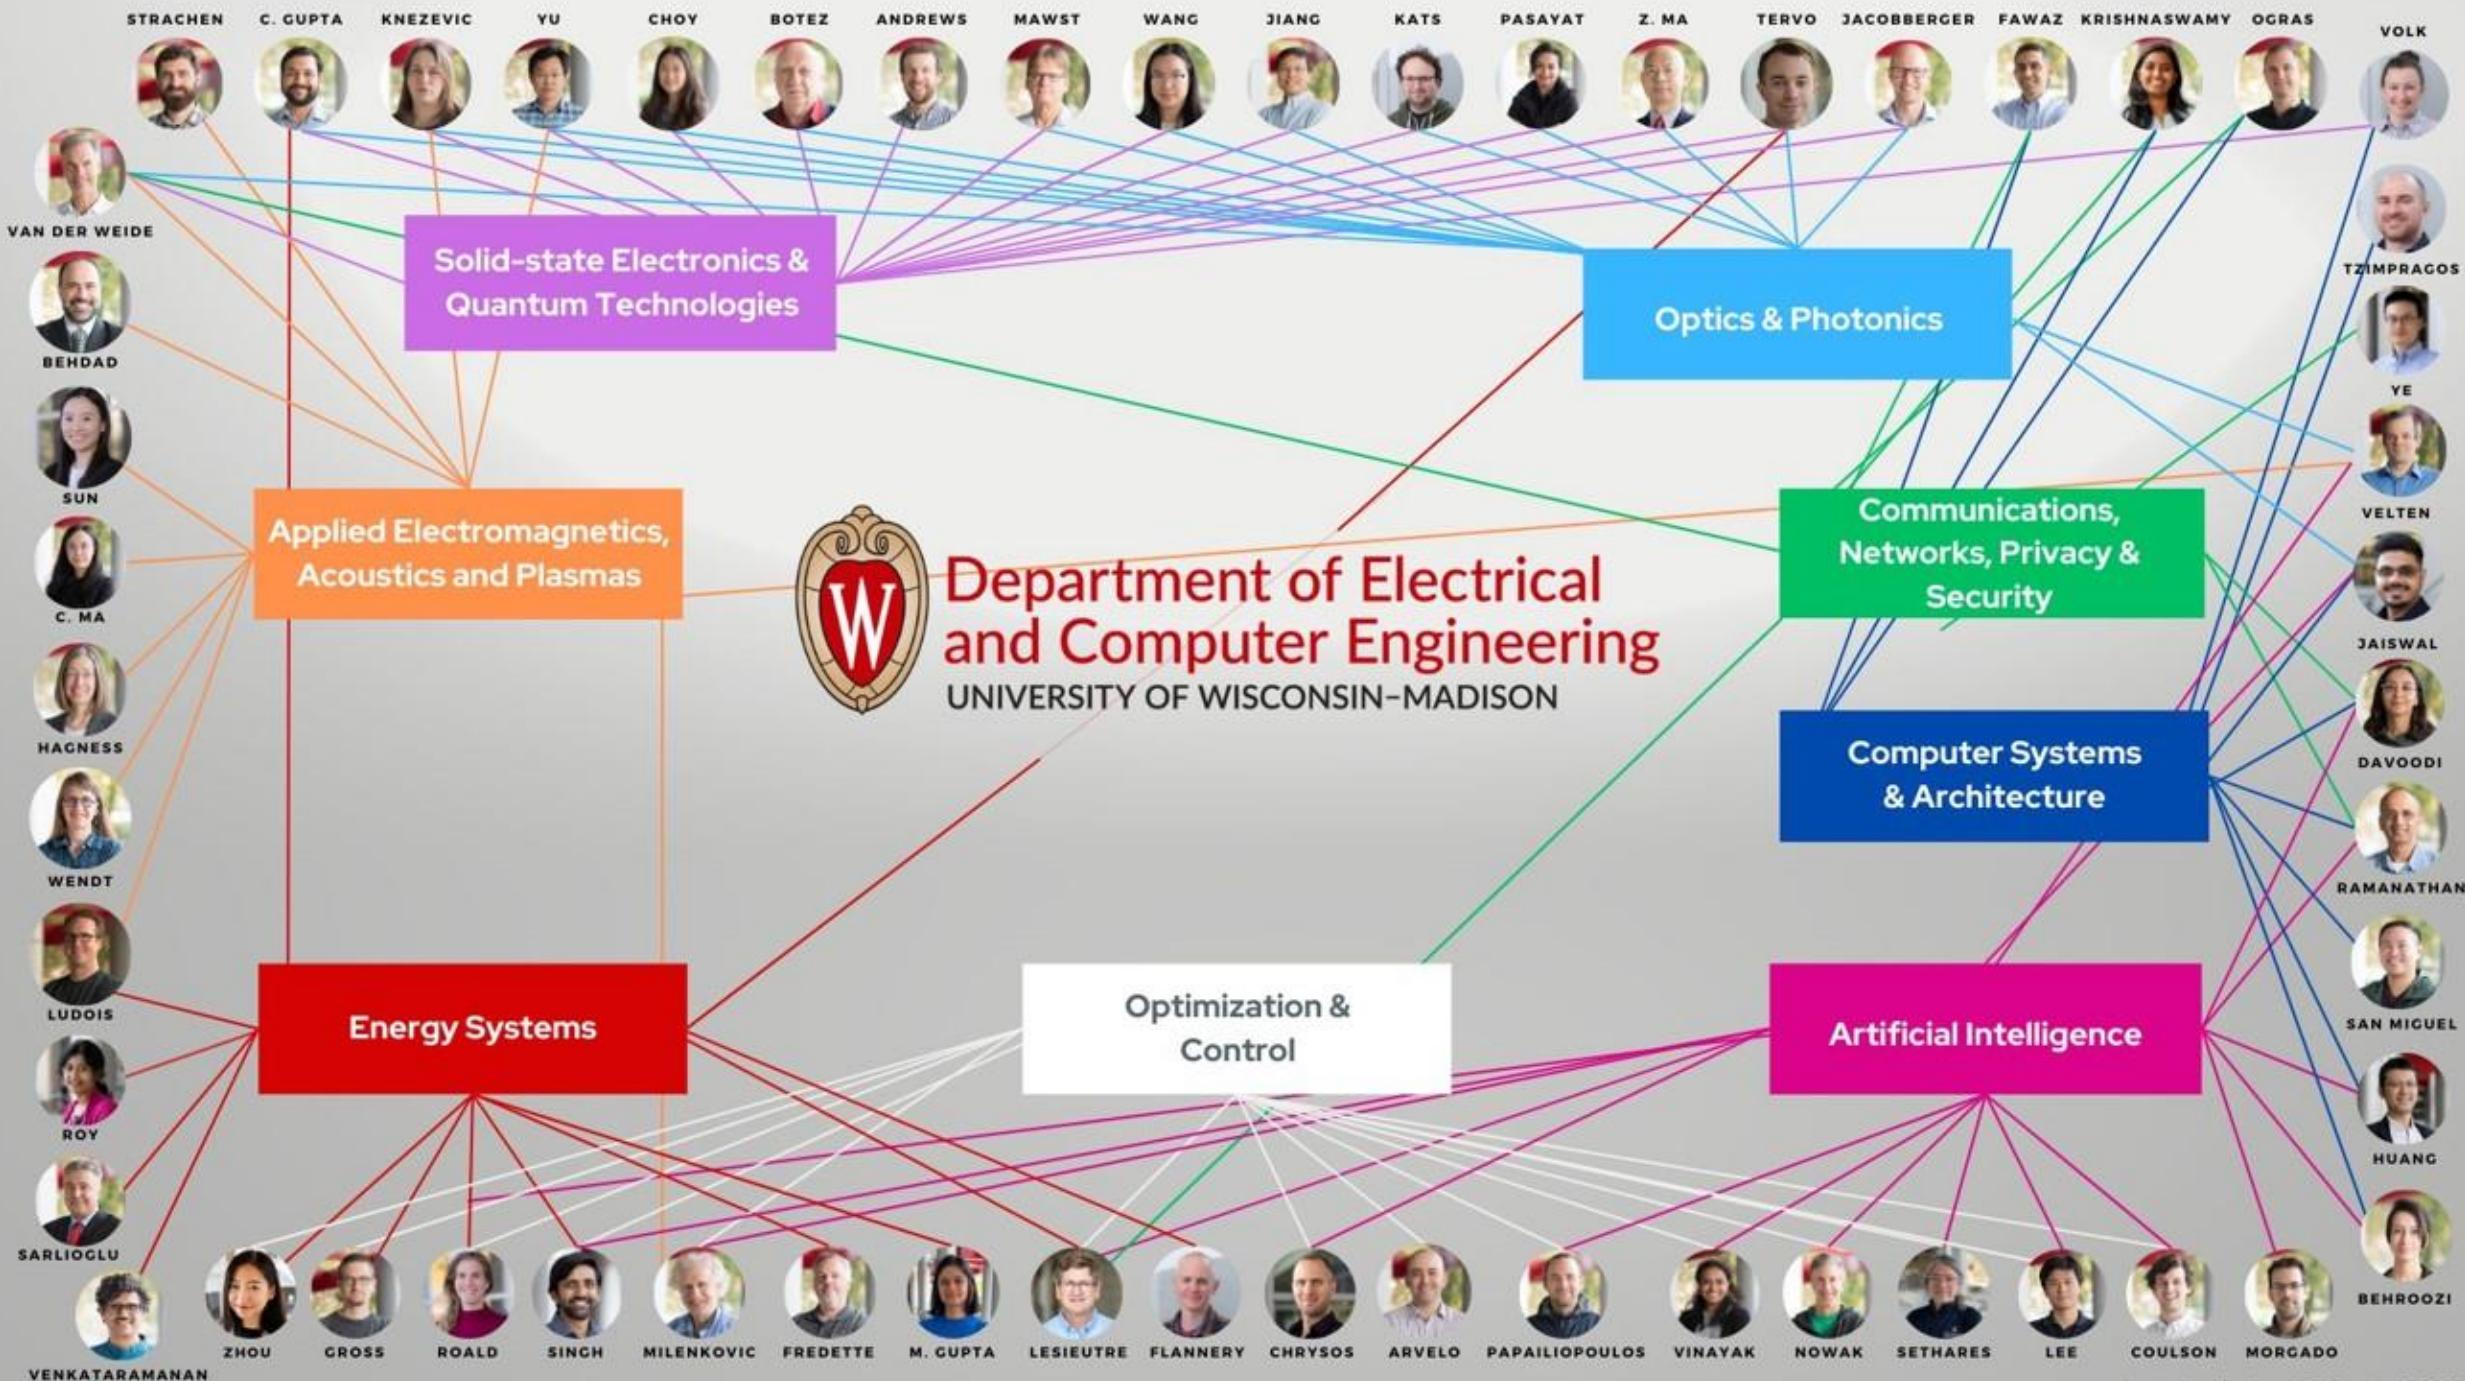




About UW-Madison ECE Department

- Located at the heart of the campus in the beautiful Madison city
 - Surrounded by lakes, consistently ranked top places to live in the US
- Highly ranked undergraduate and graduate ECE programs
 - #9 (EE) and #9 (CE) graduate ranking among public universities – USNEWS'25







Connect with UW-Madison ECE



- **Host intern, co-op, and other job opportunities**
 - Engineering Career Services team helps match top talent with internship, co-op, and full-time opportunities
- **Sponsor research for badger engineers**
 - Philanthropic support
 - Student, faculty or research support
 - Brand reputation & talent pipeline
 - No indirect costs & Tax benefits
 - Sponsored research
 - Drives innovation in mutually beneficial areas
 - Access to broad expertise
 - IP rights & early tech access
- **Give guest lecturers to our classes**
 - Share new ideas of your companies with students



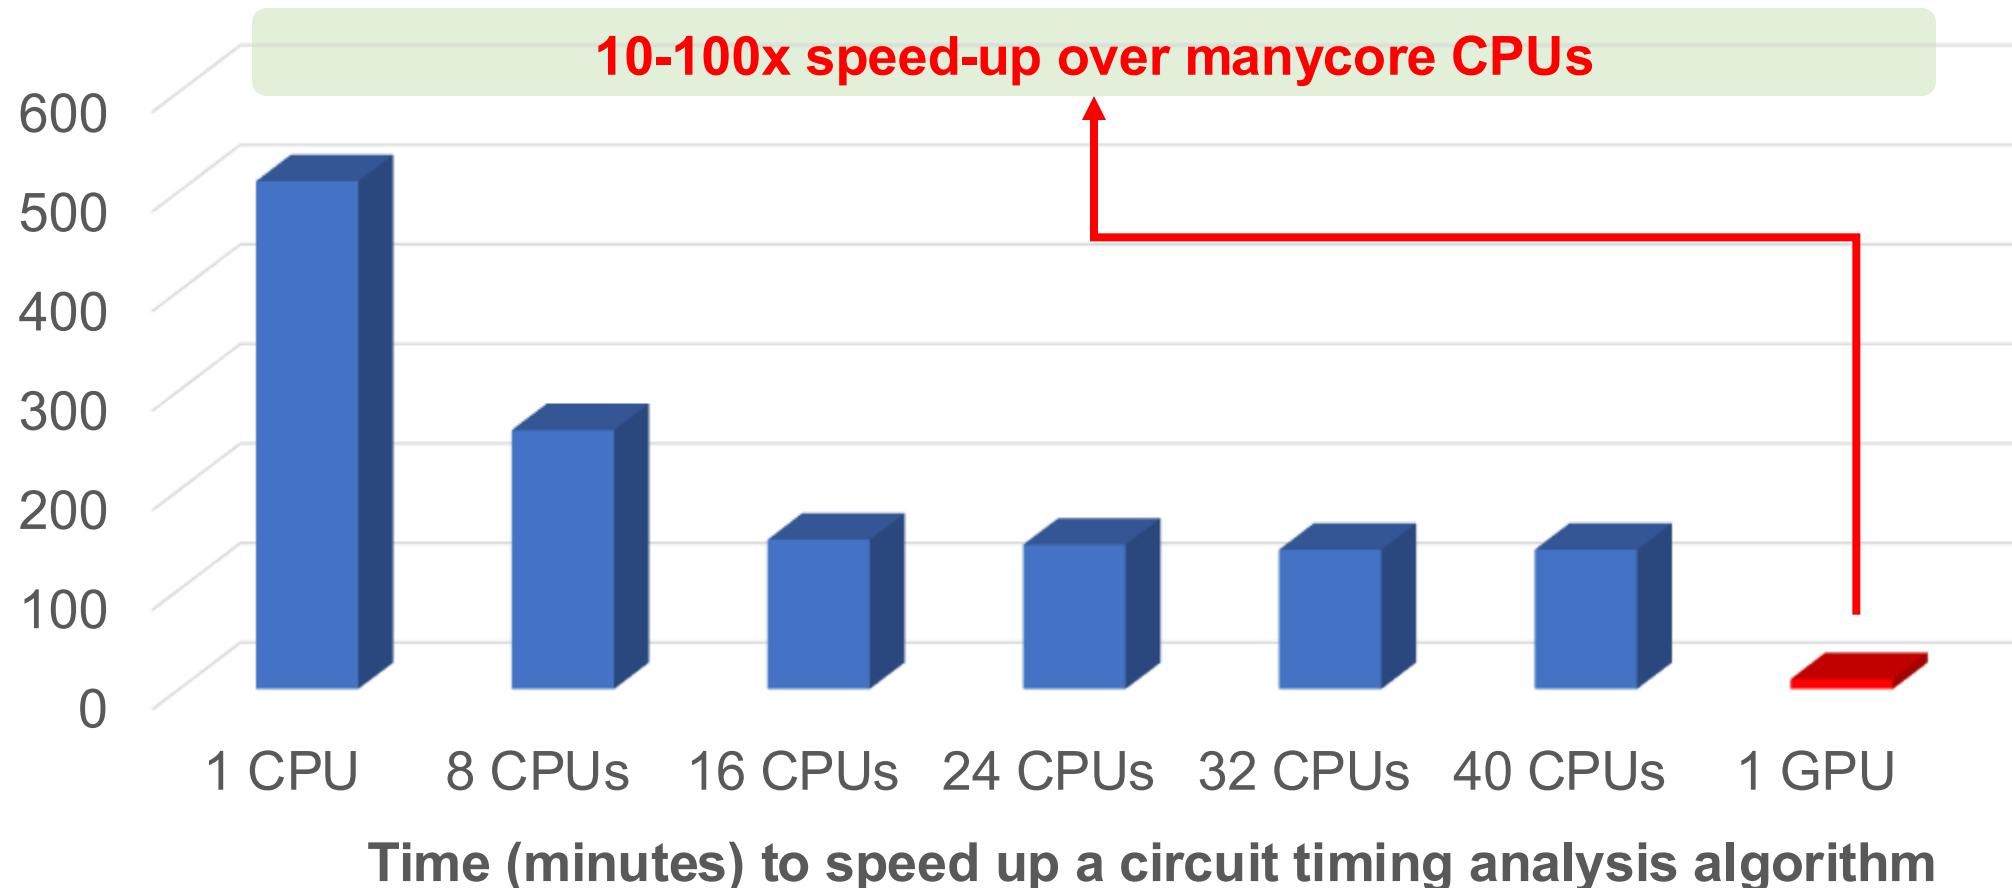


Takeaways

- Express your parallelism in the right way
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Overcome the scheduling challenges
- Demonstrate the efficiency of Taskflow
- Conclude the talk

Why Parallel Computing?

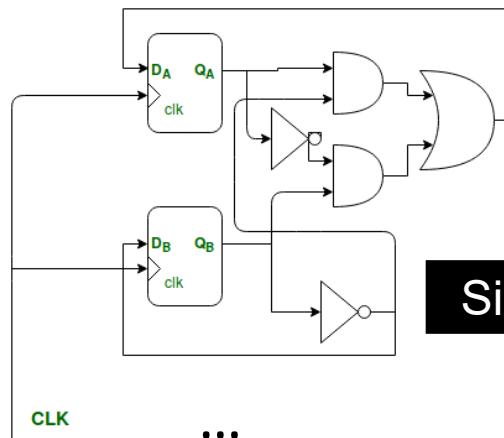
- Advances performance to a new level previously out of reach



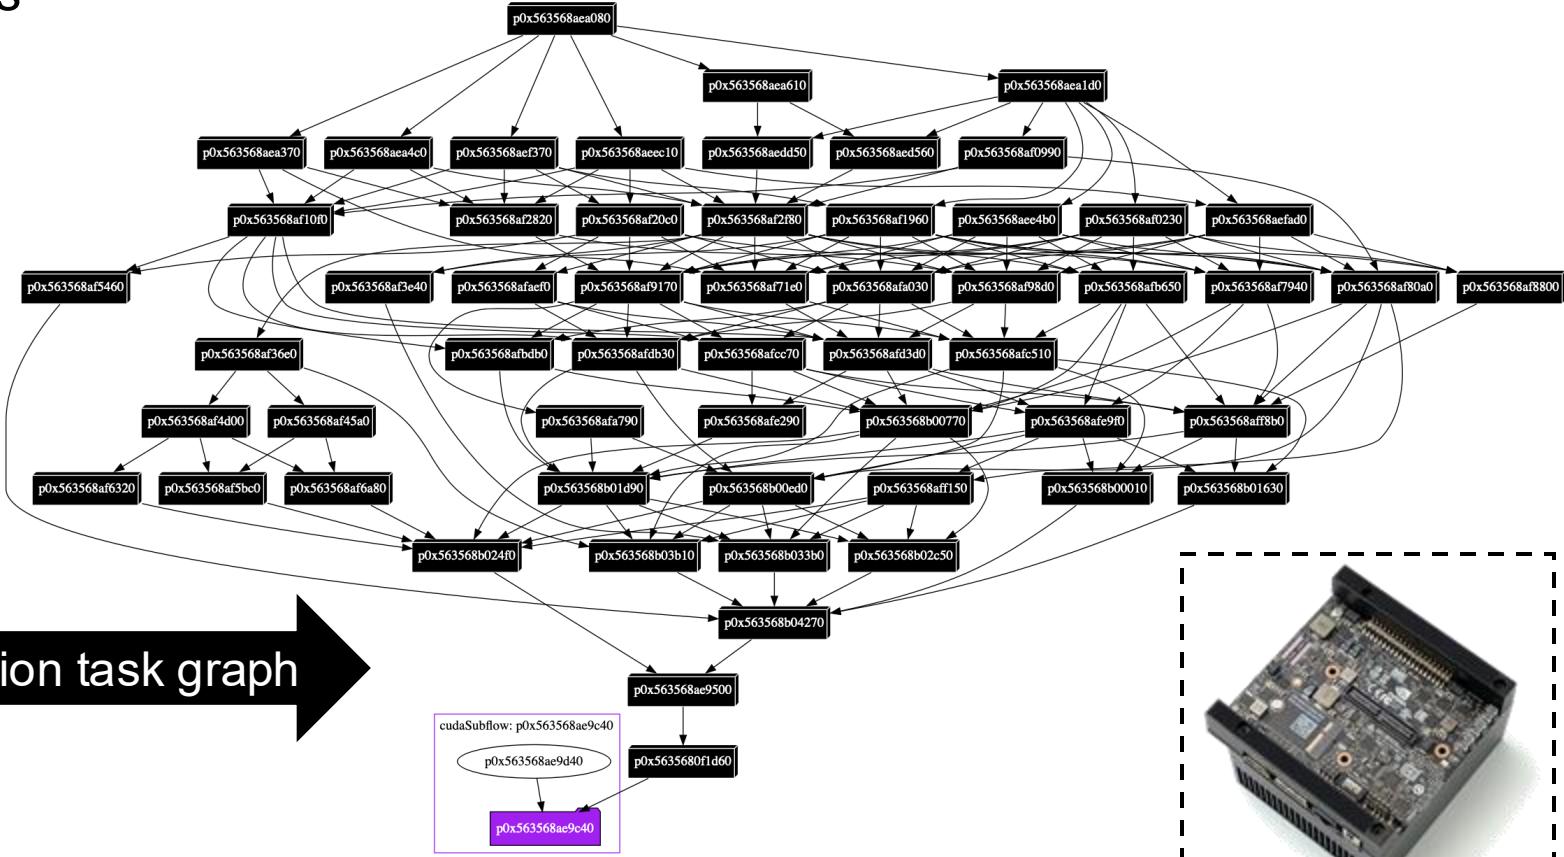


Today's Parallel Workload is Very Complex

- GPU-parallel circuit simulation task graph of Nvidia's NVDLA design¹
 - > 500M gates and nets
 - > 1000 kernels
 - > 1000 dependencies
 - > 2 hours to finish



Simulation task graph





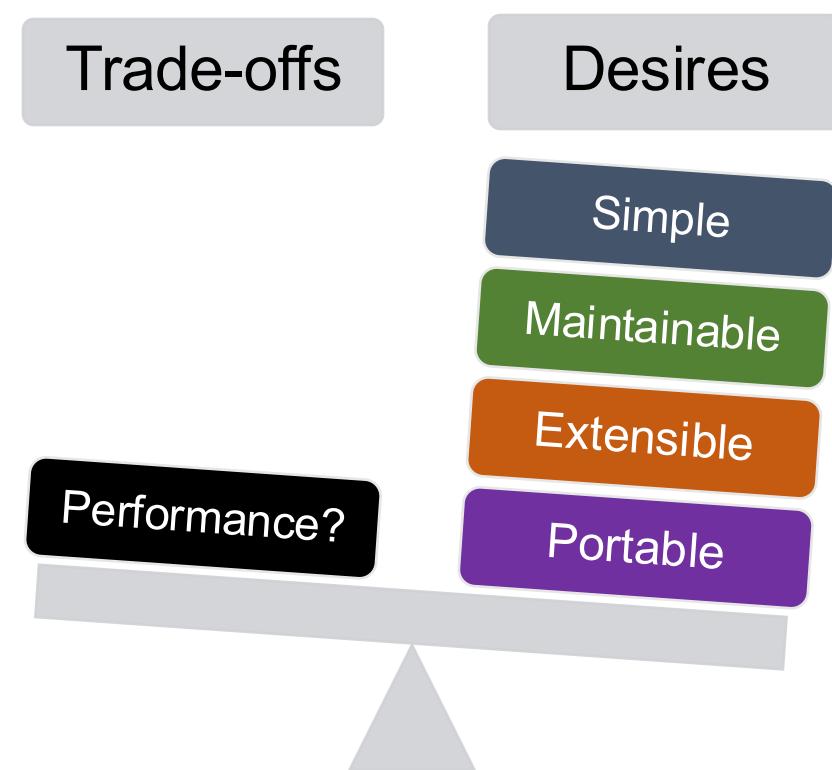
Parallel Programming is Not Easy

- You need to deal with A LOT OF technical details

- Parallelism abstraction (software + hardware)
- Concurrency control
- Task and data race avoidance
- Dependency constraints
- Scheduling efficiencies (load balancing)
- Performance portability
- ...

- And, don't forget about trade-offs

- Performance vs Desires



Need a Good Programming Abstraction

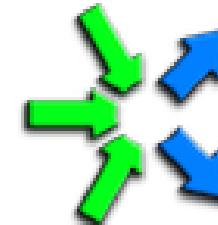
- From user's perspective, the biggest challenge is *transparency*
 - Programming abstraction, runtime optimization, load balancing, etc.
- Observing from the evolution of parallel programming standards:
 - **Task graph parallelism** (TGP) is the best model for future parallel architectures
 - Capture programmers' intention in decomposing a heterogeneous algorithm into a top-down task graph
 - Runtime can schedule dependent tasks across many processing units
- Increasing numbers of task-parallel programming systems

OpenMP

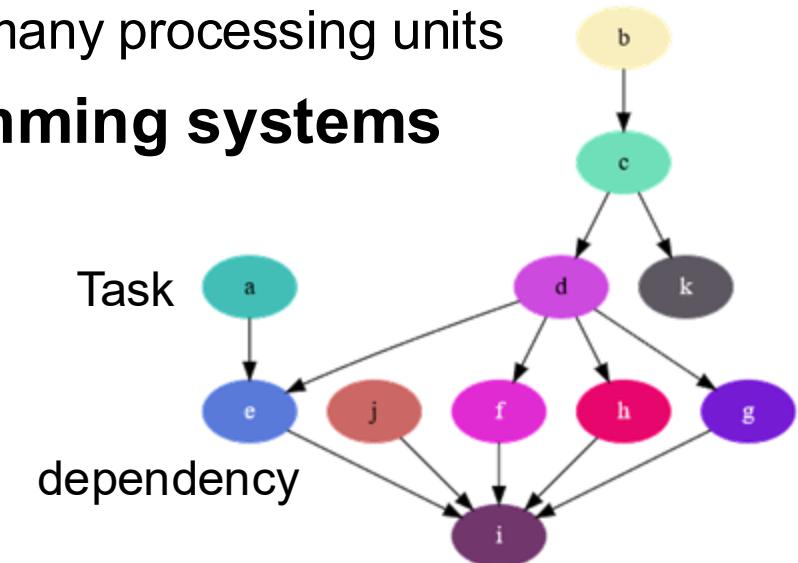


StarPU

PaRSEC



kokkos





Two Problems of Existing Tools for EDA ...

- **EDA has very complex task dependencies**
 - **Example:** analysis algorithms compute the circuit network of multi-millions of nodes and dependencies
 - **Problem:** existing tools are often good at loop parallelism (e.g., embarrassingly-parallel loops) but weak in expressing task graphs at this large scale
- **EDA has very complex control flow**
 - **Example:** synthesis algorithms make essential use of *dynamic control flow* to implement various patterns
 - Combinatorial optimization (e.g., graph algorithms, discrete math)
 - Analytical methods (e.g., physical synthesis)
 - **Problem:** existing tools are often limited to *direct acyclic graph* (DAC) models, requiring users to manually partition their workloads around control-flow or decision-making points



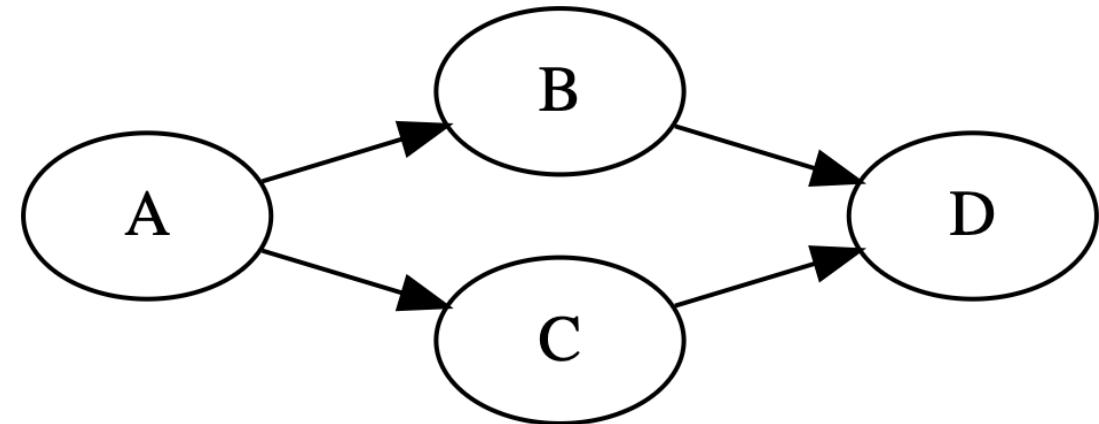
Takeaways

- Express your parallelism in the right way
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Overcome the scheduling challenges
- Demonstrate the efficiency of Taskflow
- Conclude the talk

“Hello World” in Taskflow¹

```
#include <taskflow/taskflow.hpp>
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B, C);
    D.succeed(B, C);
    executor.run(taskflow).wait();
    return 0;
}
```

// live: <https://godbolt.org/z/j8hx3xnnx>





Taskflow Supports Drop-in Integration

- Taskflow is header-only – *no wrangle with installation*

```
# clone the Taskflow project
~$ git clone https://github.com/taskflow/taskflow.git
~$ cd taskflow
# compile your program and tell it where to find Taskflow header files
~$ g++ -std=c++20 examples/simple.cpp -I ./ -O2 -pthread -o simple
~$ ./simple
```

TaskA

TaskC

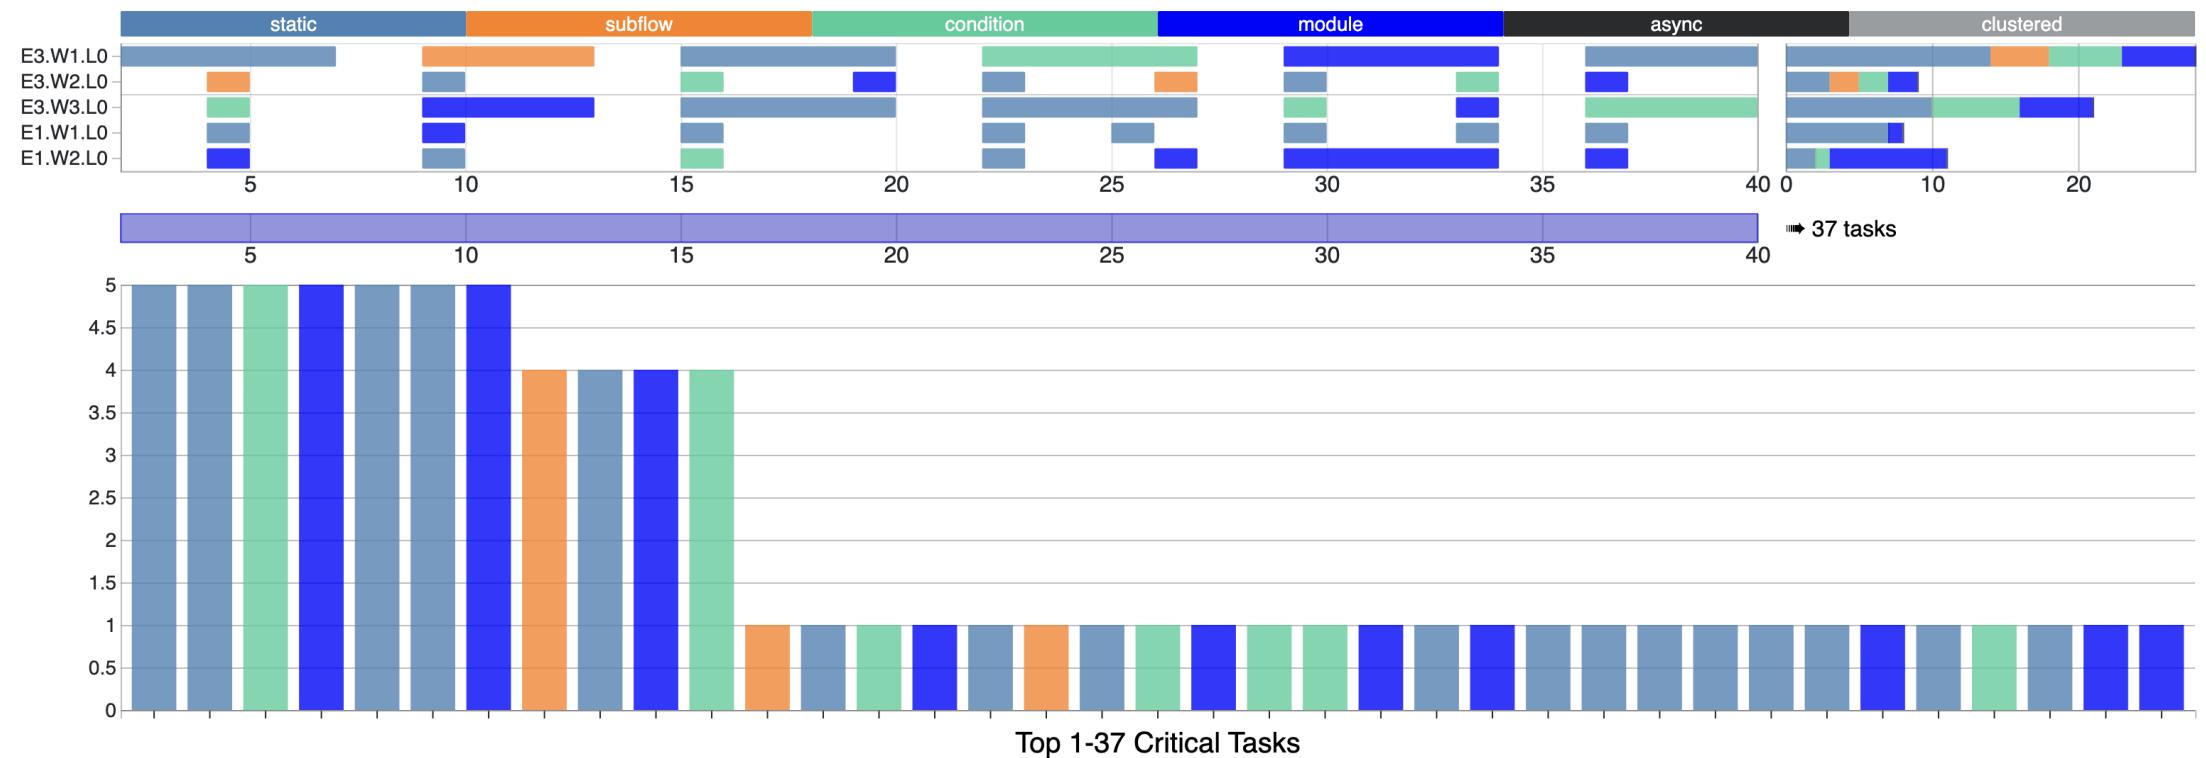
TaskB

TaskD



Built-in Task Execution Visualizer

```
# run your program with the env variable TF_ENABLE_PROFILER enabled  
# and paste the JSON content to https://taskflow.github.io/tfprof/  
~$ TF_ENABLE_PROFILER=simple.json ./simple
```

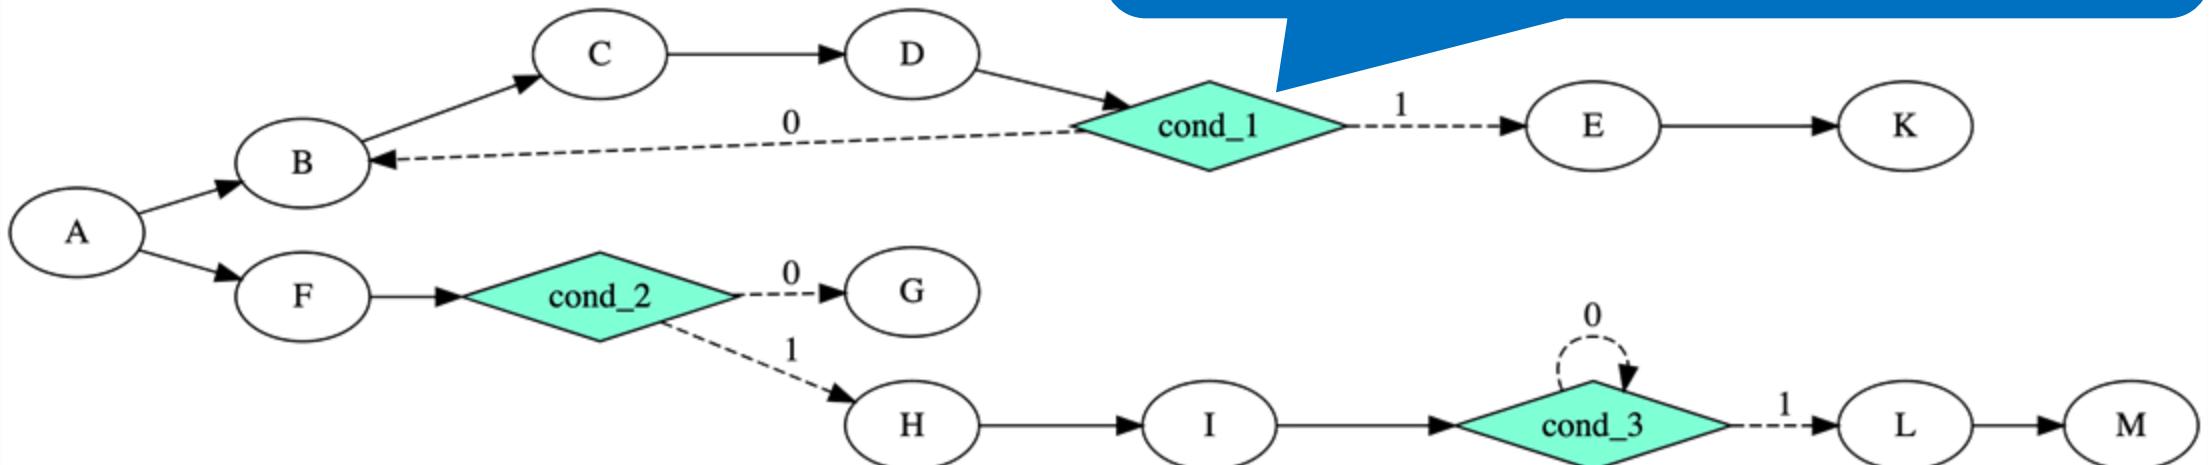


Control Taskflow Graph Programming (CTFG)

// CTFG goes beyond the limitation of traditional DAG-based models

```
auto cond_1 = taskflow.emplace([](){ return run_B() ? 0 : 1; }); // 0: is the index of B
auto cond_2 = taskflow.emplace([](){ return run_G() ? 0 : 1; }); // 0: is the index of G
auto cond_3 = taskflow.emplace([](){ return loop() ? 0 : 1; }); // 0: is the index of cond_3
cond_1.precede(B, E);           // cycle
cond_2.precede(G, H);           // if-else
cond_3.precede(cond_3, L);      // loop
```

Very difficult for existing DAG-based systems to express an efficient overlap between tasks and control flow ...

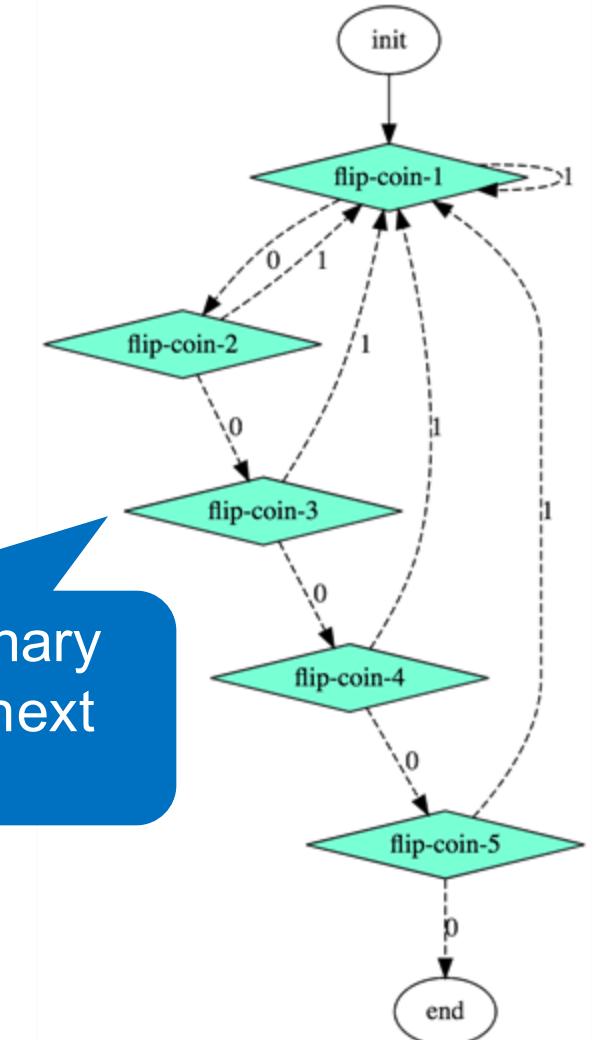


Non-deterministic Control Flow with CTFG

```

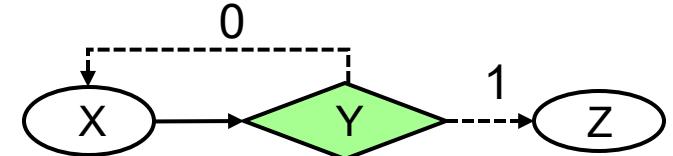
auto A = taskflow.emplace([&](){ } );
auto B = taskflow.emplace([&](){ return rand()%2; } );
auto C = taskflow.emplace([&](){ return rand()%2; } );
auto D = taskflow.emplace([&](){ return rand()%2; } );
auto E = taskflow.emplace([&](){ return rand()%2; } );
auto F = taskflow.emplace([&](){ return rand()%2; } );
auto G = taskflow.emplace([&]());
A.precede(B).name("init");
B.precede(C, B).name("flip-coin-1");
C.precede(D, B).name("flip-coin-2");
D.precede(E, B).name("flip-coin-3");
E.precede(F, B).name("flip-coin-4");
F.precede(G, B).name("flip-coin-5");
G.name("end");
    
```

Each task flips a binary coin to decide the next task to run



Existing Frameworks on Control Flow?

- **Most existing libraries are DAG-based**
 - Do not anticipate conditional execution ...
- **Unroll a task graph over fixed iterations**
 - Task graph size becomes very large ...
- **What about dynamic control flow?**
 - Resort to client-side partitions of the task graph around each decision-making points
 - Synchronize the execution of partitioned task graphs around decision-making points
 - Lack end-to-end parallelism



```
tf::Taskflow G;  
auto X = G.emplace([](){});  
auto Y = G.emplace([](){  
    return converged() ? 1 : 0;  
});  
cond.precede(Z, X);  
executor.run(G).wait();
```

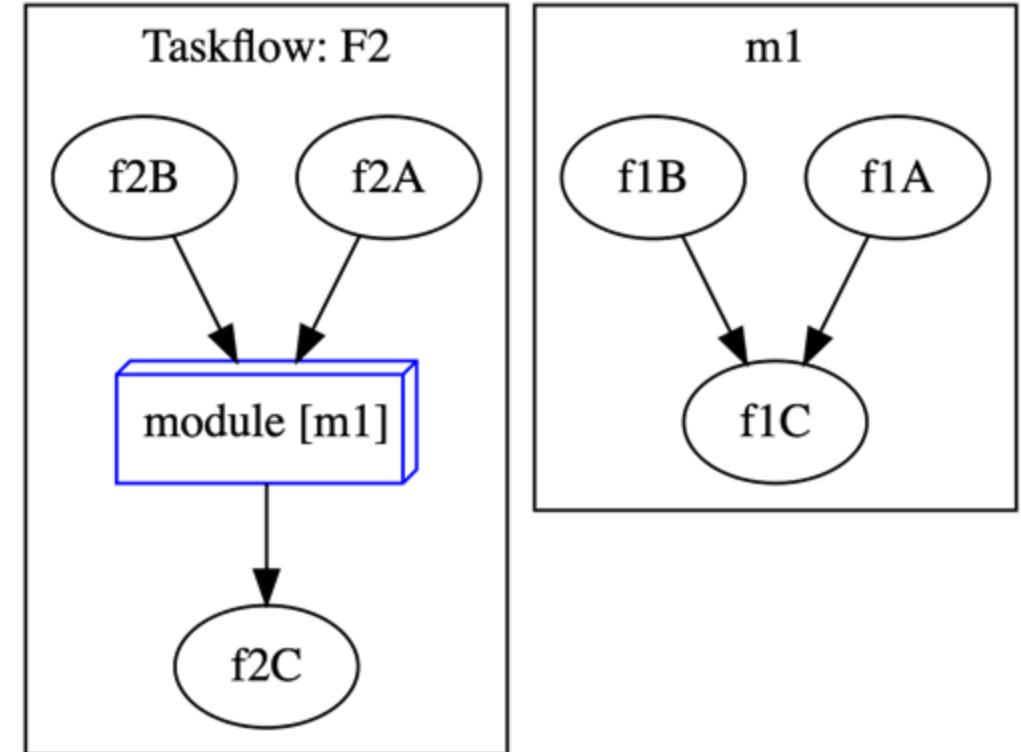
```
tbb::flow::graph1 X, Y, Z;  
do {  
    X.run();  
    Y.run();  
} while (!converged());  
Z.run();
```

Composable Tasking

```

tf::Taskflow f1, f2;
auto [f1A, f1B] = f1.emplace(
    []() { std::cout << "Task f1A\n"; },
    []() { std::cout << "Task f1B\n"; }
);
auto [f2A, f2B, f2C] = f2.emplace(
    []() { std::cout << "Task f2A\n"; },
    []() { std::cout << "Task f2B\n"; },
    []() { std::cout << "Task f2C\n"; }
);
auto f1_module_task = f2.composed_of(f1);
f1_module_task.succeed(f2A, f2B)
    .precede(f2C);

```



GPU Tasking with CUDA Graph¹

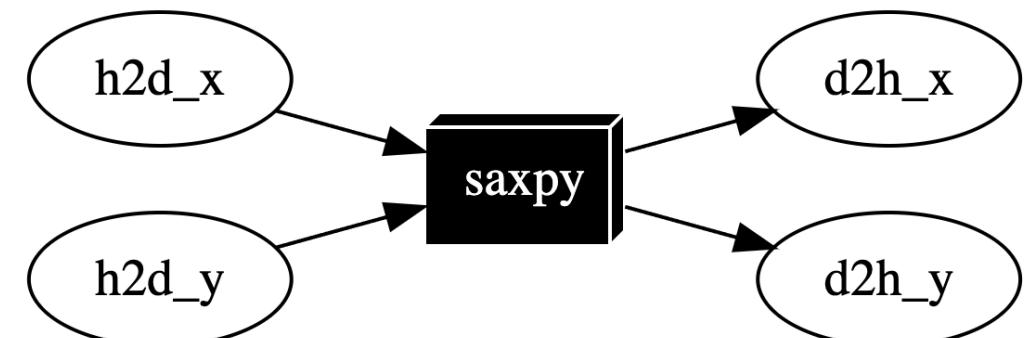
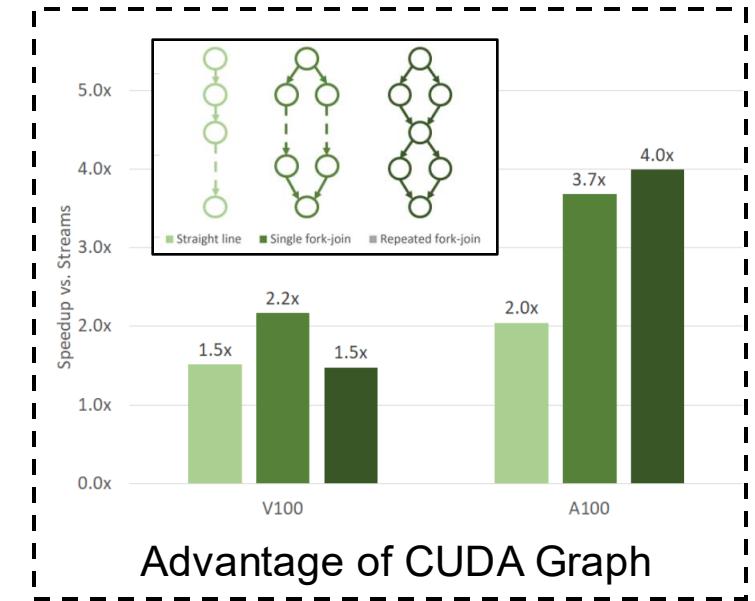
```

tf::cudaGraph cg;
auto h2d_x = cg.copy(dx, hx.data(), N);
auto h2d_y = cg.copy(dy, hy.data(), N);
auto d2h_x = cg.copy(hx.data(), dx, N);
auto d2h_y = cg.copy(hy.data(), dy, N);

// saxpy kernel with 4 blocks each of 512 threads
auto kernel = cg.kernel(4, 512, 0, saxpy, N, 2f, dx, dy);
kernel.succeed(h2d_x, h2d_y).precede(d2h_x, d2h_y);

// create an executable and run it 10 times
tf::cudaGraphExec exec(cg);
tf::cudaStream stream;
stream.run_n(exec, 10).synchronize();

```





Everything is Composable in Taskflow

- **End-to-end parallelism in one graph**
 - Task, dependency, control flow all together
 - Scheduling with whole-graph optimization
 - Efficient overlap among heterogeneous tasks
- **Largely improved productivity!**

Industrial use-case of productivity improvement using Taskflow¹

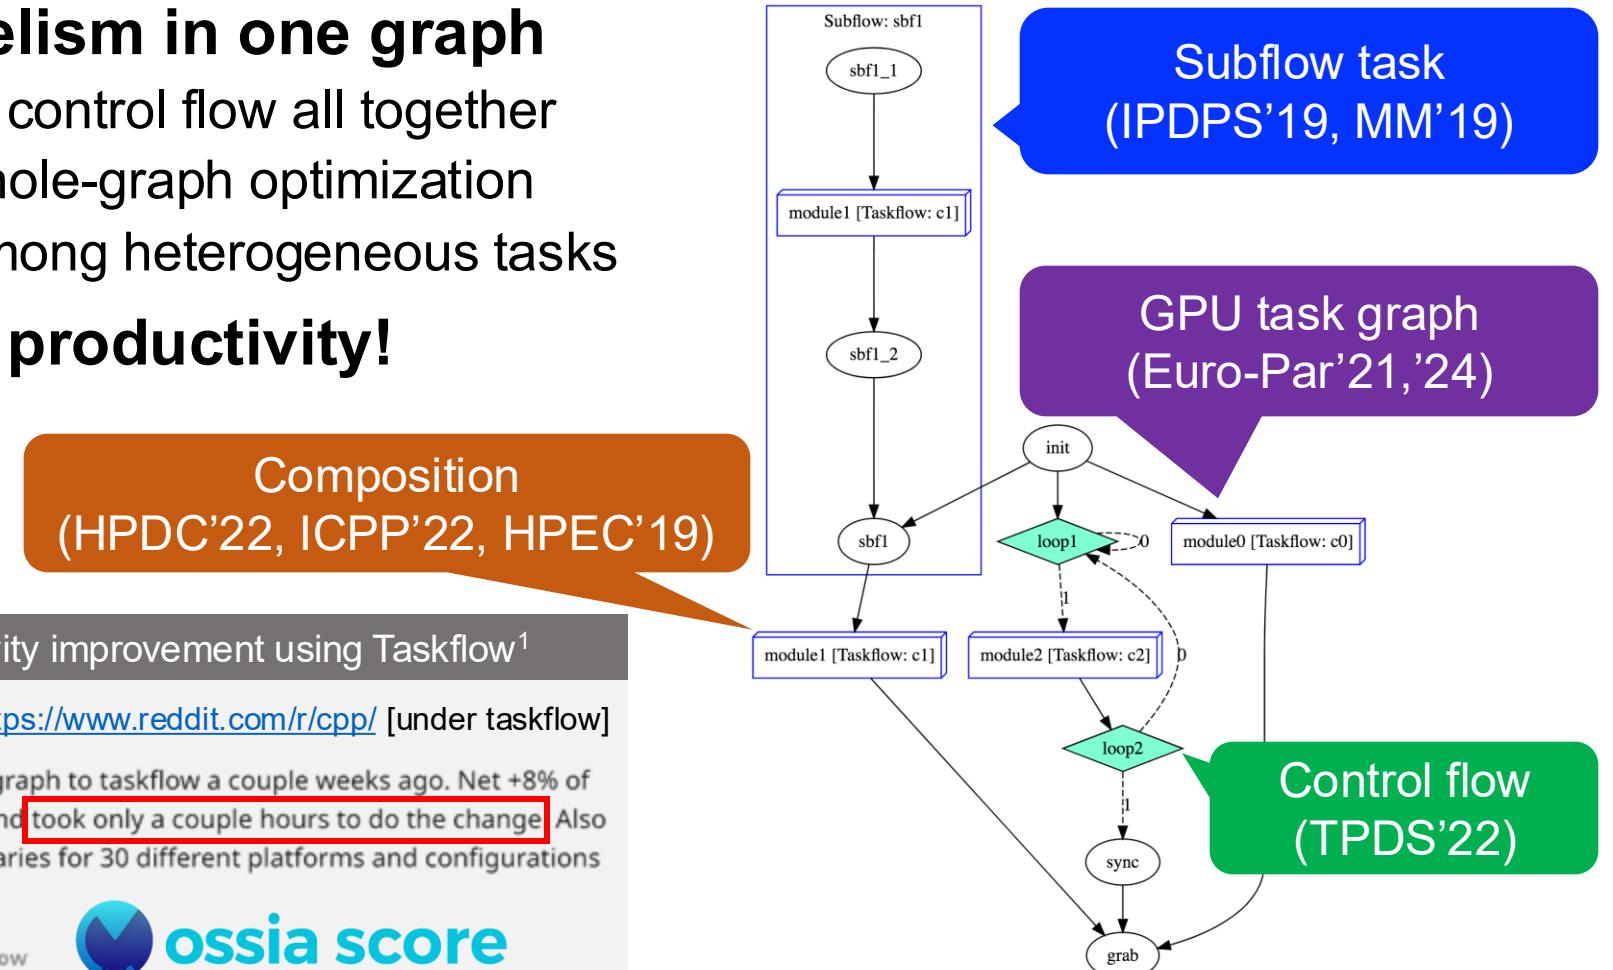
jcelerier ossia score

Reddit: <https://www.reddit.com/r/cpp/> [under taskflow]

I've migrated <https://ossia.io> from TBB flow graph to taskflow a couple weeks ago. Net +8% of throughput on the graph processing itself, and took only a couple hours to do the change! Also don't have to fight with building the TBB libraries for 30 different platforms and configurations since it's header only.

↑ 8 ↓ Reply Share Report Save Follow

 ossia score



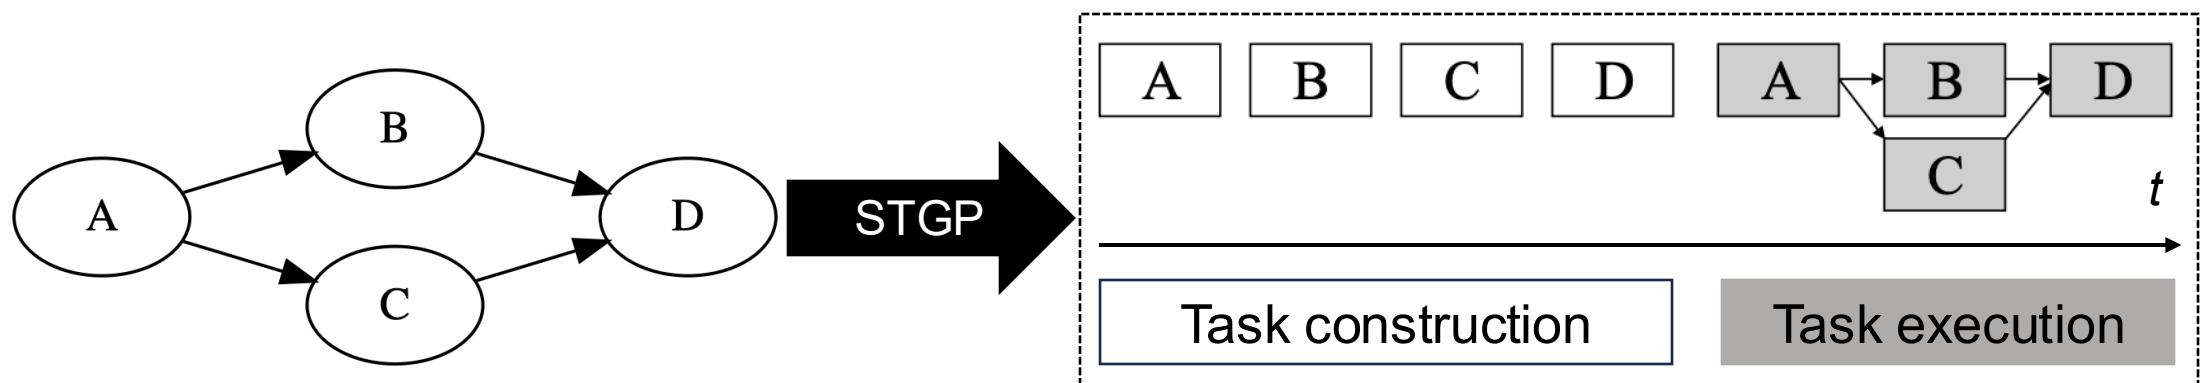


Takeaways

- Express your parallelism in the right way
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Overcome the scheduling challenges
- Demonstrate the efficiency of Taskflow
- Conclude the talk

Static Task Graph Parallelism (STGP)

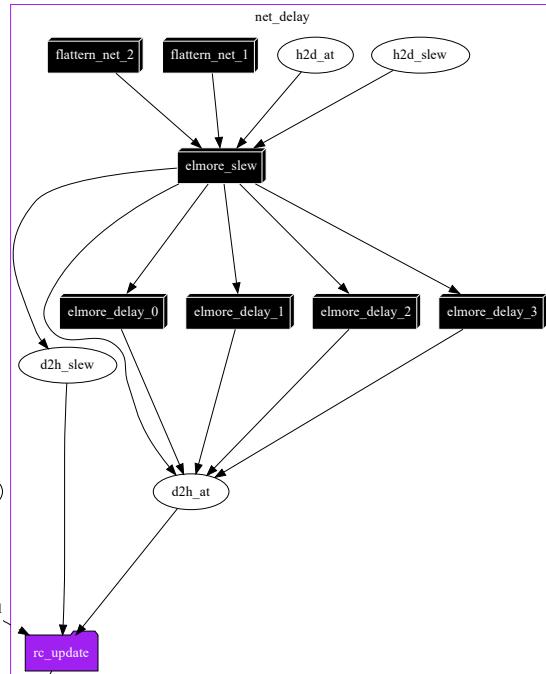
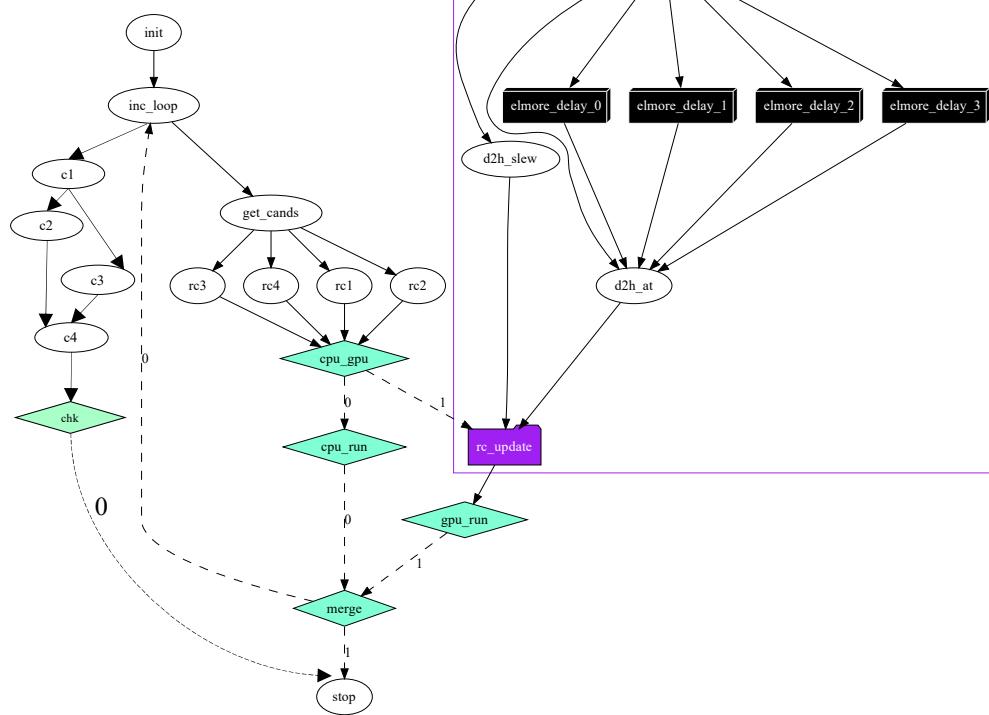
- In STGP, the graph structure must be known up front
 - Execution of STGP is based on the *construct-and-run* model
- Lack of overlap between task construction and task execution
 - For large task graphs (e.g., multi-million tasks and dependencies), such an overlap can bring a significant performance advantage
- Lack of flexible and dynamic expression of TGP
 - Task graph structure cannot depend on runtime values or control-flow results



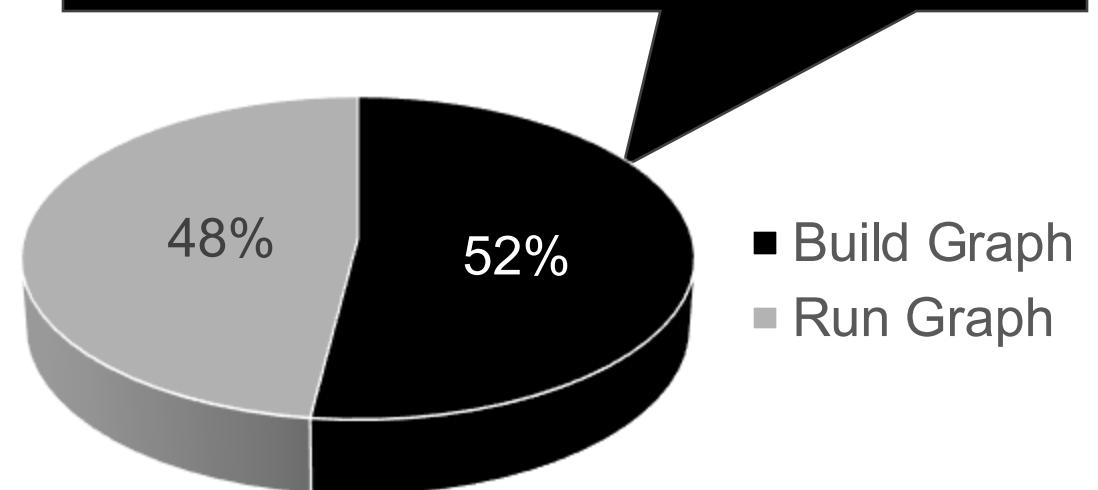
Problem of STGP: Example #1

- Runtime breakdown of a task-parallel circuit timing analyzer¹

- > 10M tasks
- > 10M edges



Task graph construction time takes over 50% of the entire runtime (typically done in one thread)



- Build Graph
- Run Graph



Problem of STGP: Example #2

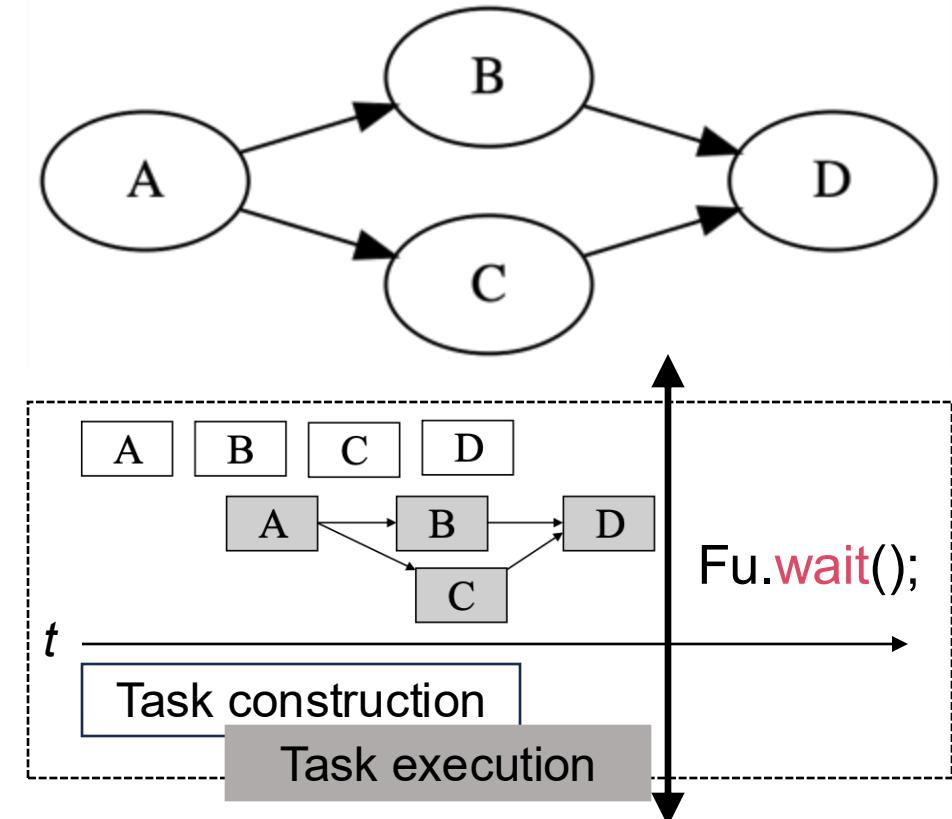
- Express TGP that depends on runtime variables...?

```
if (a == true) {  
    G1 = build_task_graph1();  
    if (b == true) {  
        G2 = build_task_graph2();  
        G1.precede(G2);  
        if (c == true) {  
            ... // need another different TGP  
        }  
    }  
    else {  
        G3 = build_task_graph3();  
        G3.precede(G1);  
    }  
}
```

```
G1 = build_task_graph1();  
G2 = build_task_graph2();  
if (G1.num_tasks() == 100) {  
    G1.precede(G2);  
}  
else {  
    G3 = build_task_graph3();  
    G2.precede(G1, G3);  
    if (G2.num_dependencies() >= 10) {  
        ... // define dependencies on the fly  
    }  
}
```

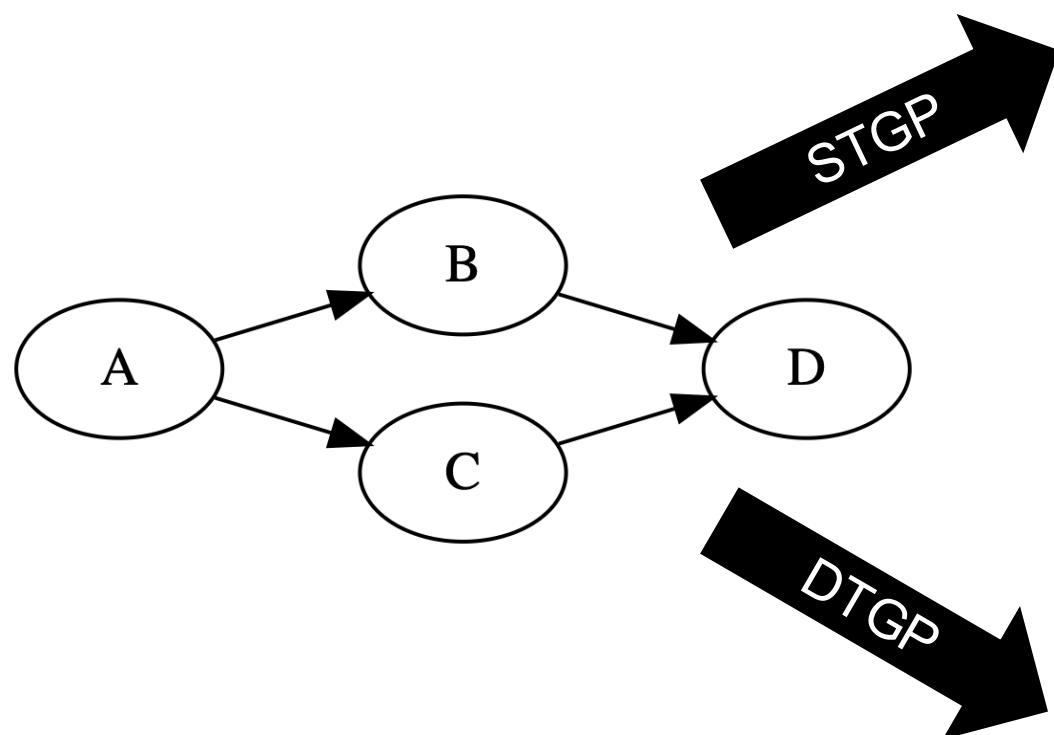
Dynamic TGP (DTGP) in Taskflow

```
// Live: https://godbolt.org/z/j76ThGbWK
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
auto [D, Fu] = executor.dependent_async([](){
    std::cout << "TaskD\n";
}, B, C); ←
Fu.wait();
```



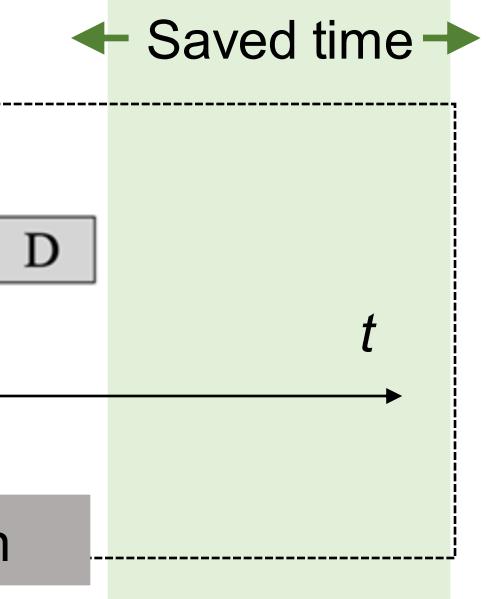
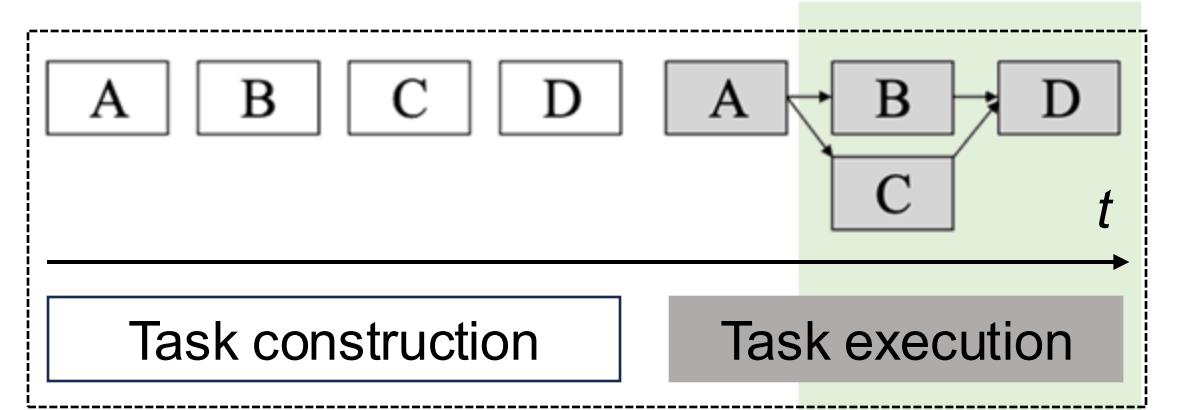
Specify arbitrary task dependencies using
C++ variadic parameter pack

Comparison between STGP and DTGP



STGP

DTGP



DTGP Needs a Correct Topological Order

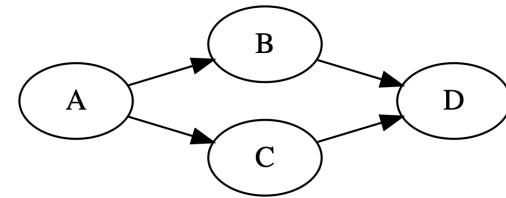
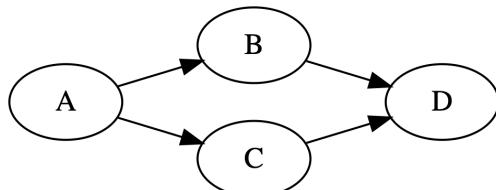
```
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});

auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);

auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);

auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
```

Topological order #1: A → B → C → D



Topological order #2: A → C → B → D

```
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});

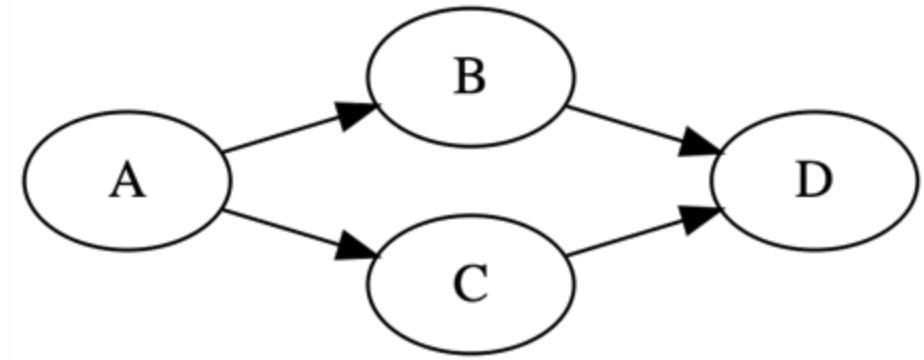
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);

auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);

auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
```

Incorrect Topological Order ...

```
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B-is-unavailable-yet, C-is-unavailable-yet);
auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
executor.wait_for_all();
```



An incorrect topological order ($A \rightarrow D \rightarrow B \rightarrow C$) disallows us from expressing correct DTGP



Variable Range of Task Dependencies

- Both methods can take a range of dependent-async tasks
 - useful when the task dependencies come as a runtime variable

// Live: <https://godbolt.org/z/6Pvco4KeE>

```
std::vector<tf::AsyncTask> tasks = {  
    executor.silent_dependent_async([](){ std::cout << "TaskA\n"; }),  
    executor.silent_dependent_async([](){ std::cout << "TaskB\n"; }),  
    executor.silent_dependent_async([](){ std::cout << "TaskC\n"; }),  
    executor.silent_dependent_async([](){ std::cout << "TaskD\n"; })  
};
```

```
// create a dependent-async tasks that depends on tasks, A, B, C, and D  
executor.dependent_async([](){}, tasks.begin(), tasks.end());
```

```
// create a silent-dependent-async tasks that depends on tasks, A, B, C, and D  
executor.silent_dependent_async([](){}, tasks.begin(), tasks.end());
```

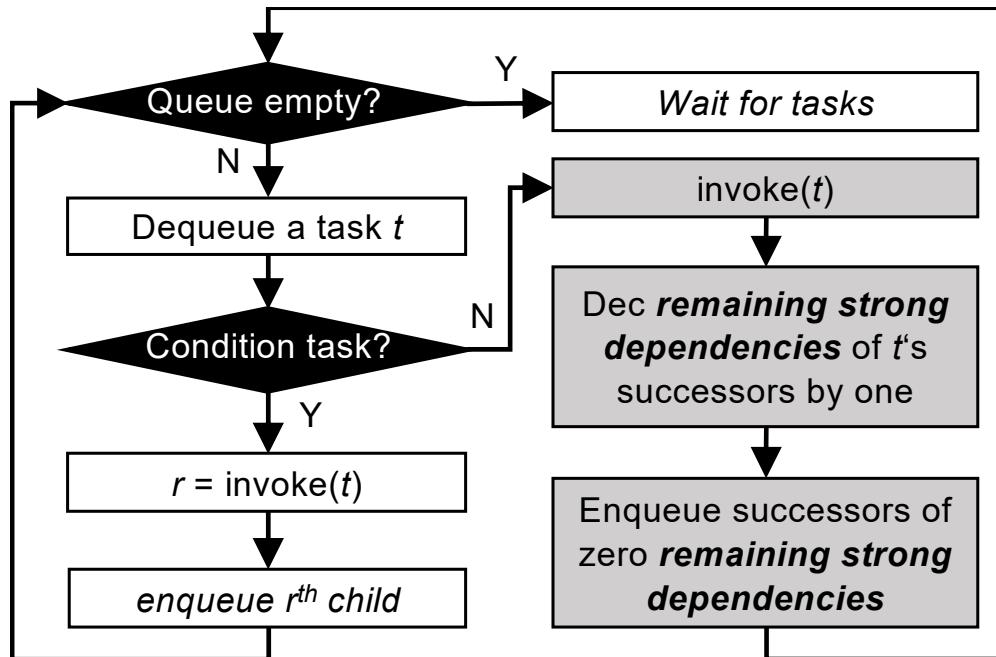


Takeaways

- Express your parallelism in the right way
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Overcome the scheduling challenges
- Demonstrate the efficiency of Taskflow
- Conclude the talk

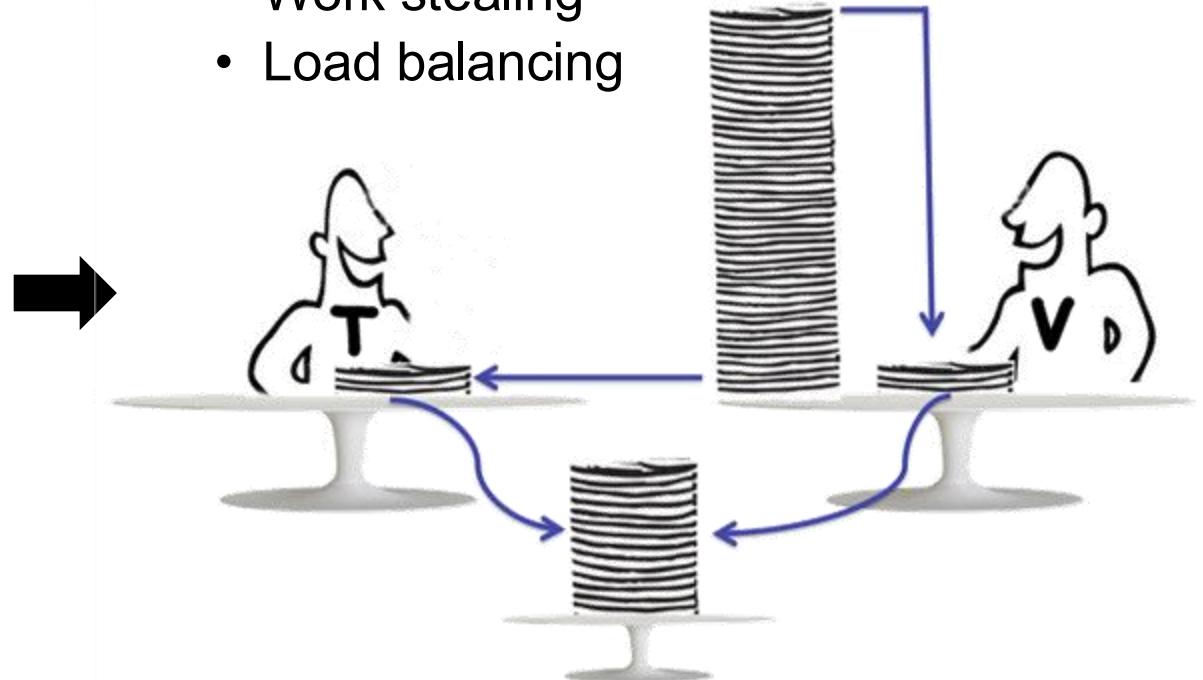
STGP Scheduling Algorithm

- Task-level scheduling



- Worker-level scheduling

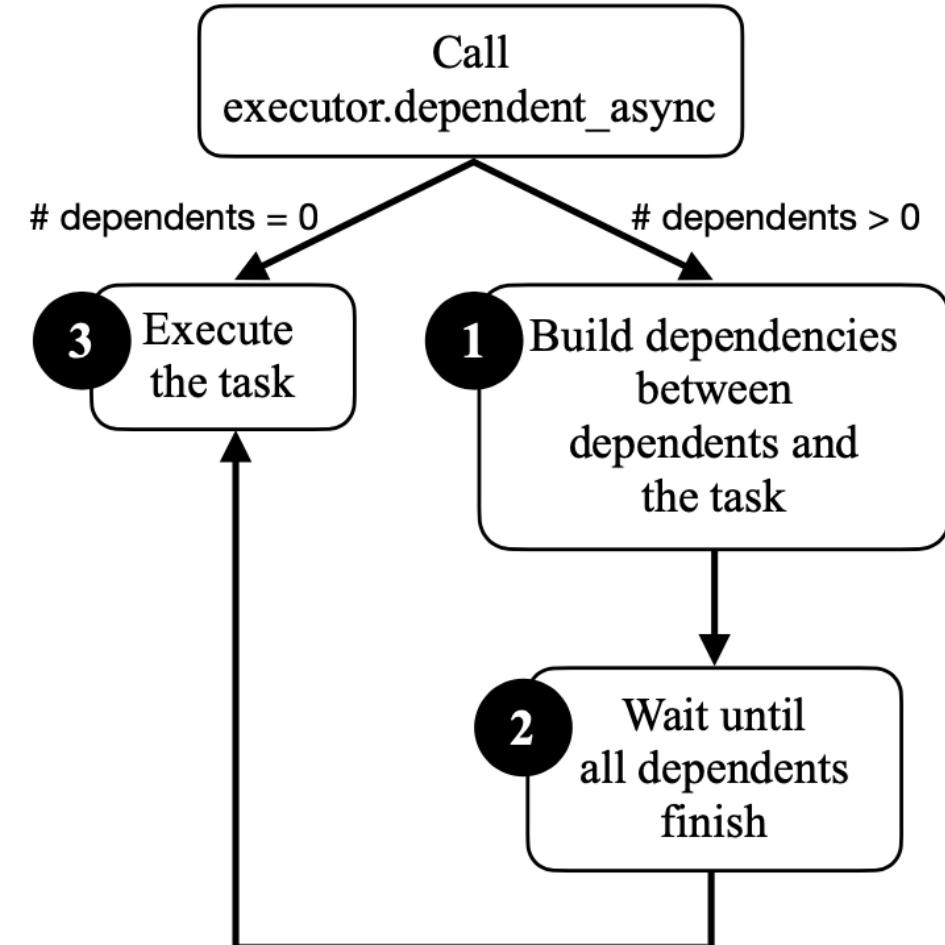
- Work stealing
- Load balancing



Key results: schedule tasks with in-graph control flow with a **strong balance** between the number of active workers and dynamically generated tasks – *low latency, energy efficient, and high throughput*

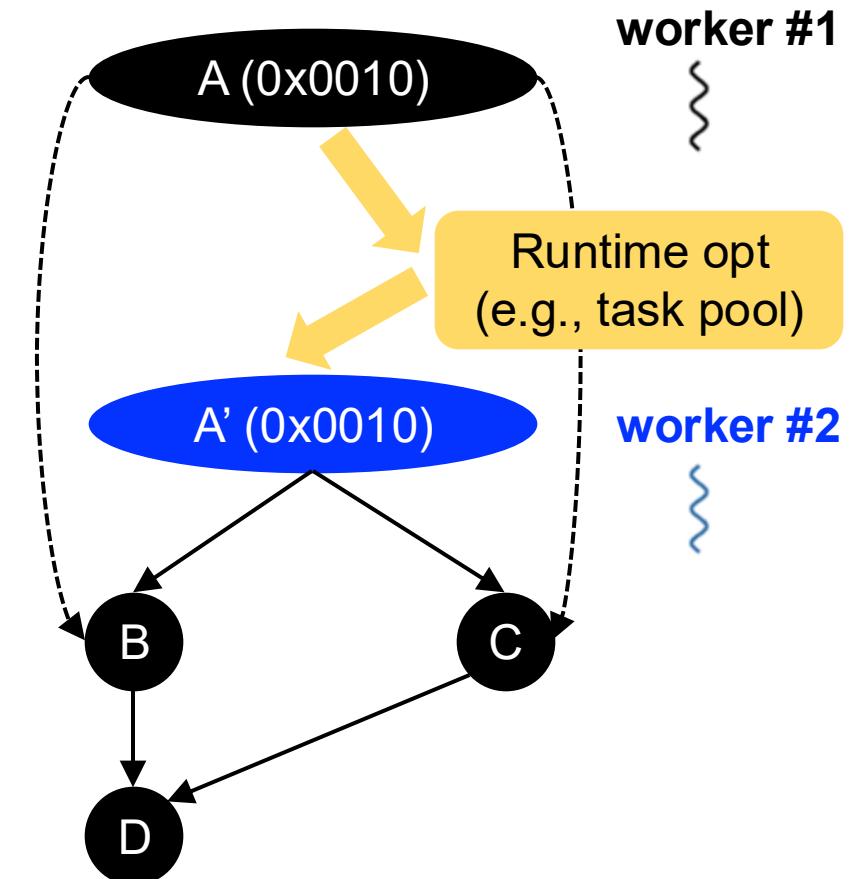
DTGP Scheduling Algorithm

- The algorithm has three parts:
 - Build dependencies
 - Wait for dependents to finish
 - Execute the task
- Three key scheduling challenges:
 1. **ABA** – a specified dependent task must exist correctly
 2. **Data race** – multiple threads may simultaneously add and remove successors to and from a task
 3. **Synchronization** – application can issue a global synchronization at anytime to wait for all tasks to finish



Solving Challenge #1: ABA Problem¹

```
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
executor.wait_for_all();
```





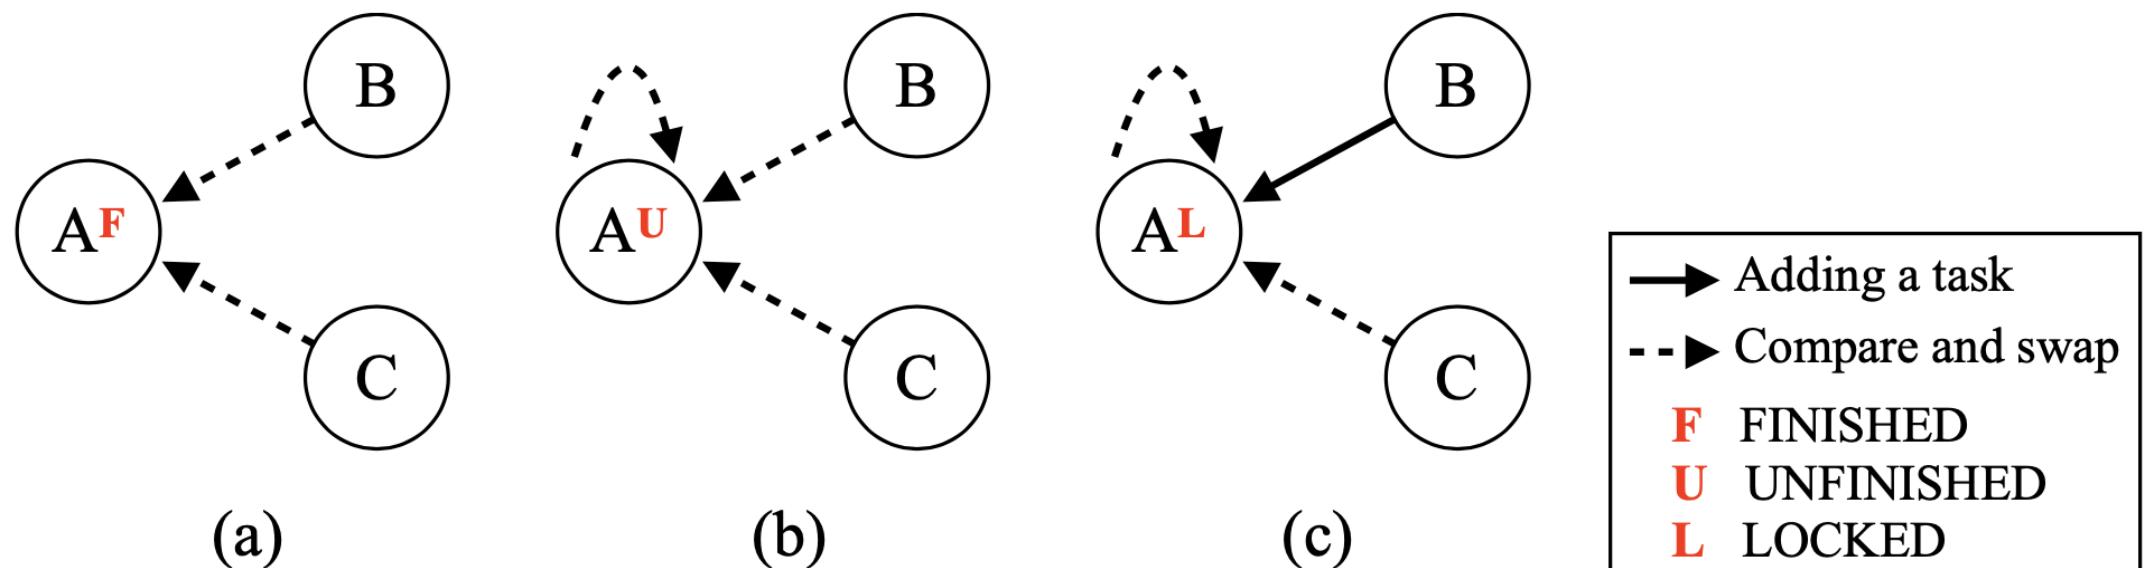
Retain Shared Ownership of Every Task

```
tf::Executor executor;
tf::AsyncTask A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
tf::AsyncTask B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A); ←
tf::AsyncTask C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
tf::AsyncTask D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
executor.wait_for_all();
```

tf::AsyncTask acts like
a std::shared_ptr to
ensure tasks stay alive
when they are used

Solving Challenge #2: Data Race

- Both B and C want to add themselves to the successors of A
 - In the meantime, A may want to remove its successor
- Apply compare-and-swap (CAS) to enable exclusive access
 - As a result, constructing a dynamic task graph can be completely thread-safe





Solving Challenge #3: Synchronization

- Application can issue a global synchronization at any time

- Ex: `executor.wait_for_all()`, `future.get()`, etc.

```
tf::Executor executor;
auto A = executor.silent_dependent_async([](){}); // Task A
auto B = executor.silent_dependent_async([](){}, A); // Task B depends on A
executor.wait_for_all(); // wait for A and B to finish
```

```
auto C = executor.silent_dependent_async([](){}, A); // Task C depends on A
auto D = executor.silent_dependent_async([](){}, B, C); // Task D depends on B and C
executor.wait_for_all(); // wait for C and D to finish
```

Taskflow uses C++20 atomic variables to perform waiting/notifying operations; many synchronizations can happen at user space instead of kernel

```
// lock-based solution
std::unique_lock lock(mutex);
cv.wait(lock, [&](){
    return num_tasks == 0;
});
```

C++17

```
// atomic wait-based solution
auto n = num_tasks.load();
while(n != 0) {
    num_tasks.wait(n);
    n = num_tasks.load();
};
```

C++20



Lock-free Scheduling Algorithm¹

Algorithm 1 dependent_async(callable, deps)

```
1: Create a future
2: num_deps  $\leftarrow$  sizeof(deps)
3: task  $\leftarrow$  initialize_task(callable, num_deps, future)
4: for all dep  $\in$  deps do
5:   process_dependent(task, dep, num_deps)
6: end for
7: if num_deps == 0 then
8:   schedule_async_task(task)
9: end if
10: return (task, future)
```

Algorithm 2 process_dependent(task, dep, num_deps)

```
1: dep_state  $\leftarrow$  dep.state
2: target_state  $\leftarrow$  UNFINISHED
3: if dep_state.CAS(target_state, LOCKED) then
4:   dep.successors.push(task)
5:   dep_state  $\leftarrow$  UNFINISHED
6: else if target_state == FINISHED then
7:   num_deps  $\leftarrow$  AtomDec(task.join_counter)
8: else
9:   goto line 2
10: end if
```

Algorithm 3 schedule_async_task(task)

```
1: target_state  $\leftarrow$  UNFINISHED
2: while not task.state.CAS(target_state, FINISHED)
   do
3:   target_state  $\leftarrow$  UNFINISHED
4: end while
5: Invoke(task.callable)
6: for all successor  $\in$  task.successors do
7:   if AtomDec(successor.join_counter) == 0 then
8:     schedule_async_task(successor)
9:   end if
10: end for
11: if AtomDec(task.ref_count) == 0 then
12:   Delete task
13: end if
```

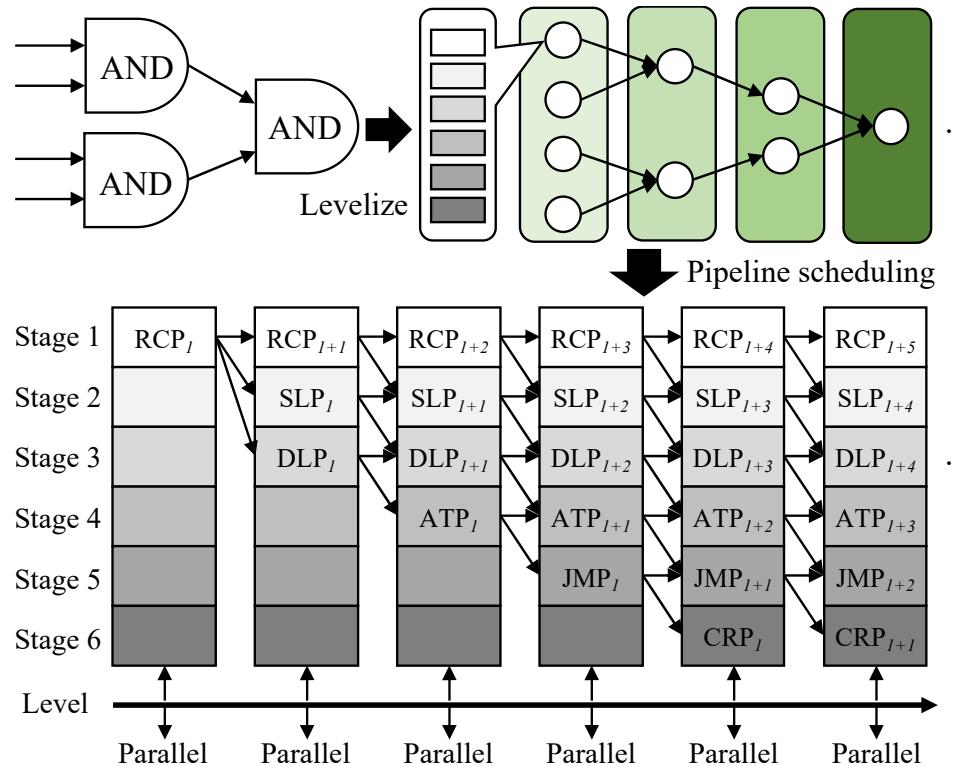


Takeaways

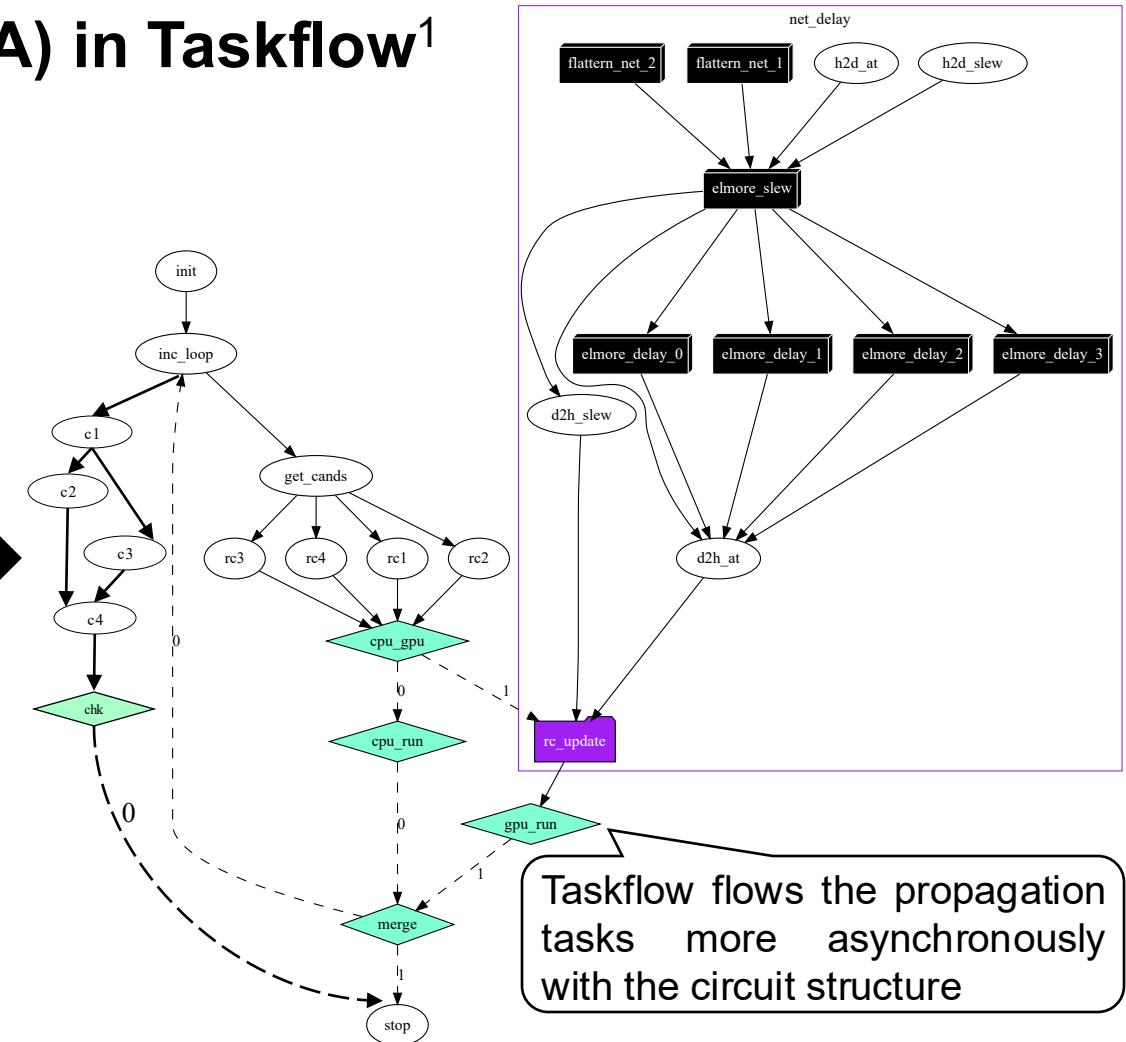
- Express your parallelism in the right way
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Overcome the scheduling challenges
- Demonstrate the efficiency of Taskflow
- Conclude the talk

Case Study 1: Task-parallel STA

- Model graph-based analysis (GBA) in Taskflow¹

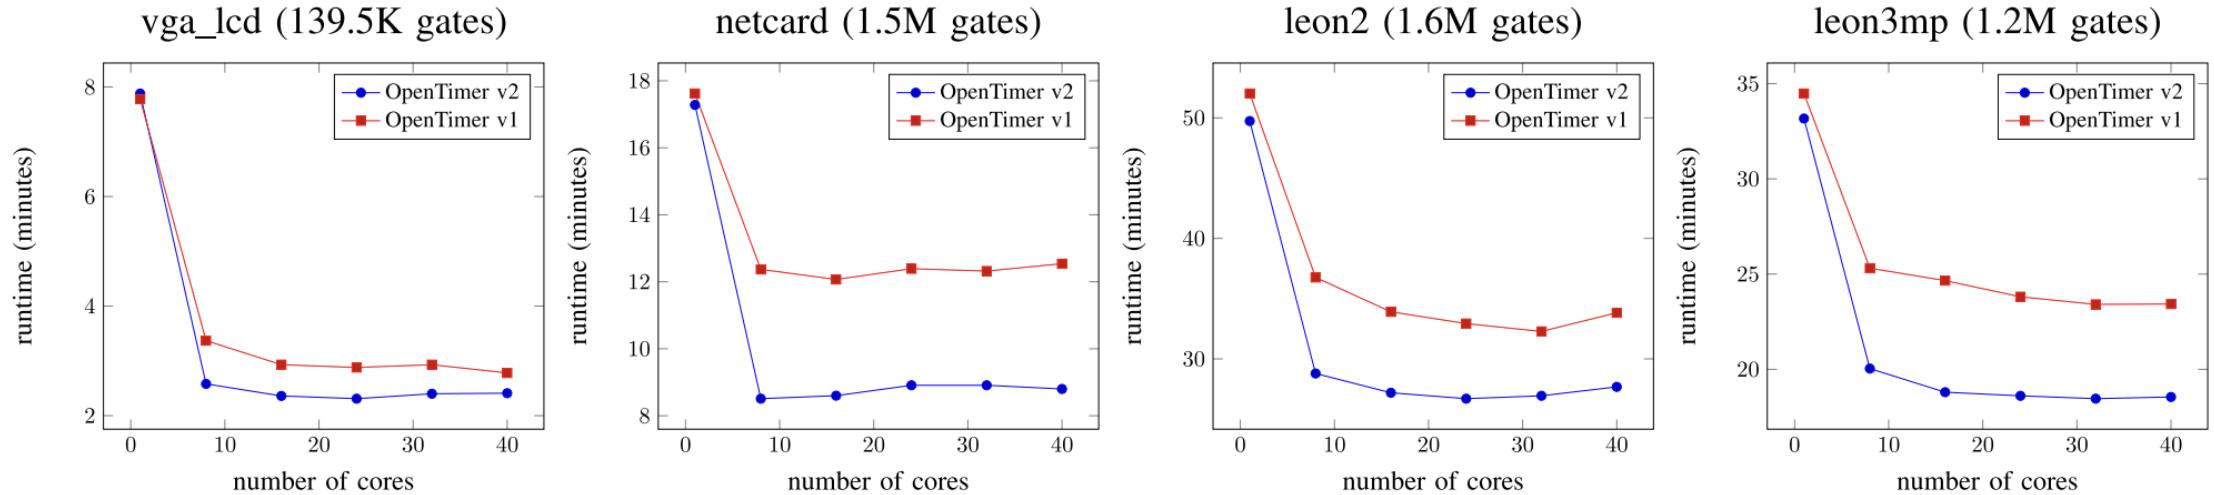


Loop-parallel level-by-level timing propagation in GBA



Levelization-based vs Task-parallel GBA

- **OpenTimer v1: levelization-based parallel timing propagation¹**
 - Implemented using OpenMP “parallel_for” primitive
- **OpenTimer v2: task-parallel timing propagation²**
 - Implemented using Taskflow (<https://taskflow.github.io/>)

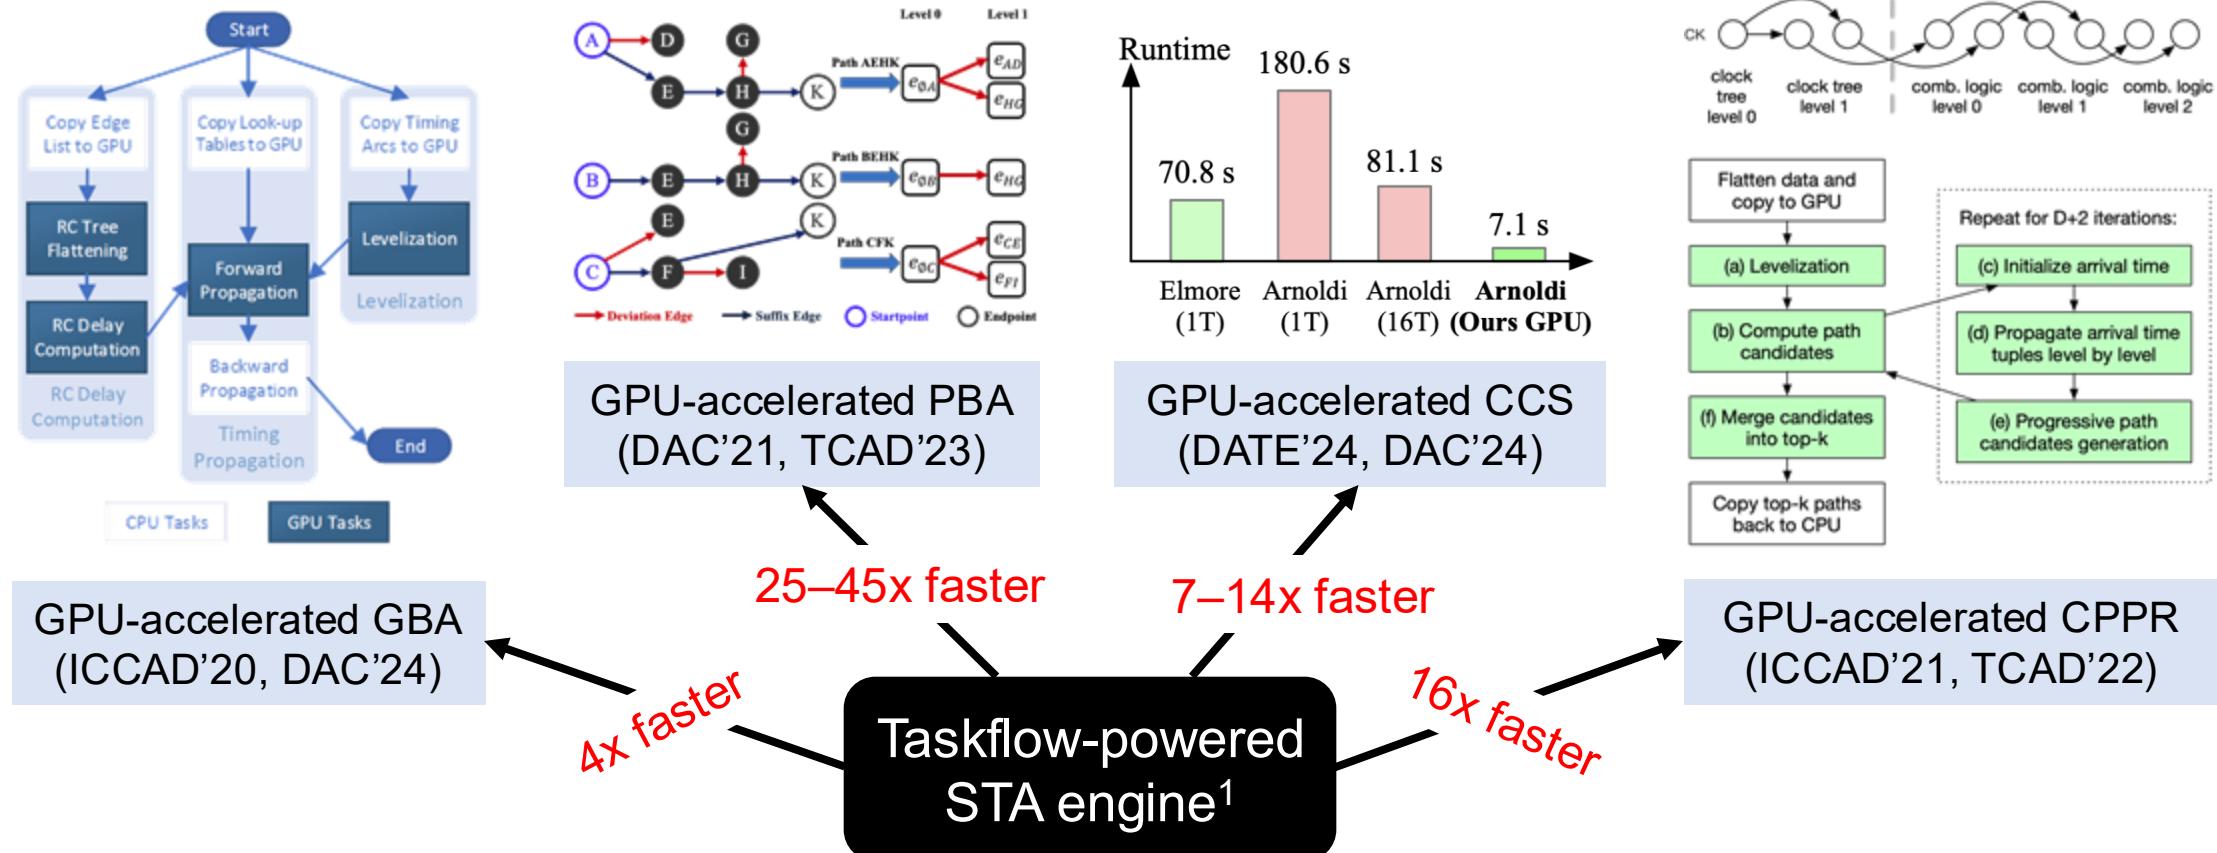


Taskflow allows us to more asynchronously parallelize the timing propagation

¹: Tsung-Wei Huang and Martin Wong, “OpenTimer: A High-Performance Timing Analysis Tool,” *IEEE/ACM ICCAD*, 2015

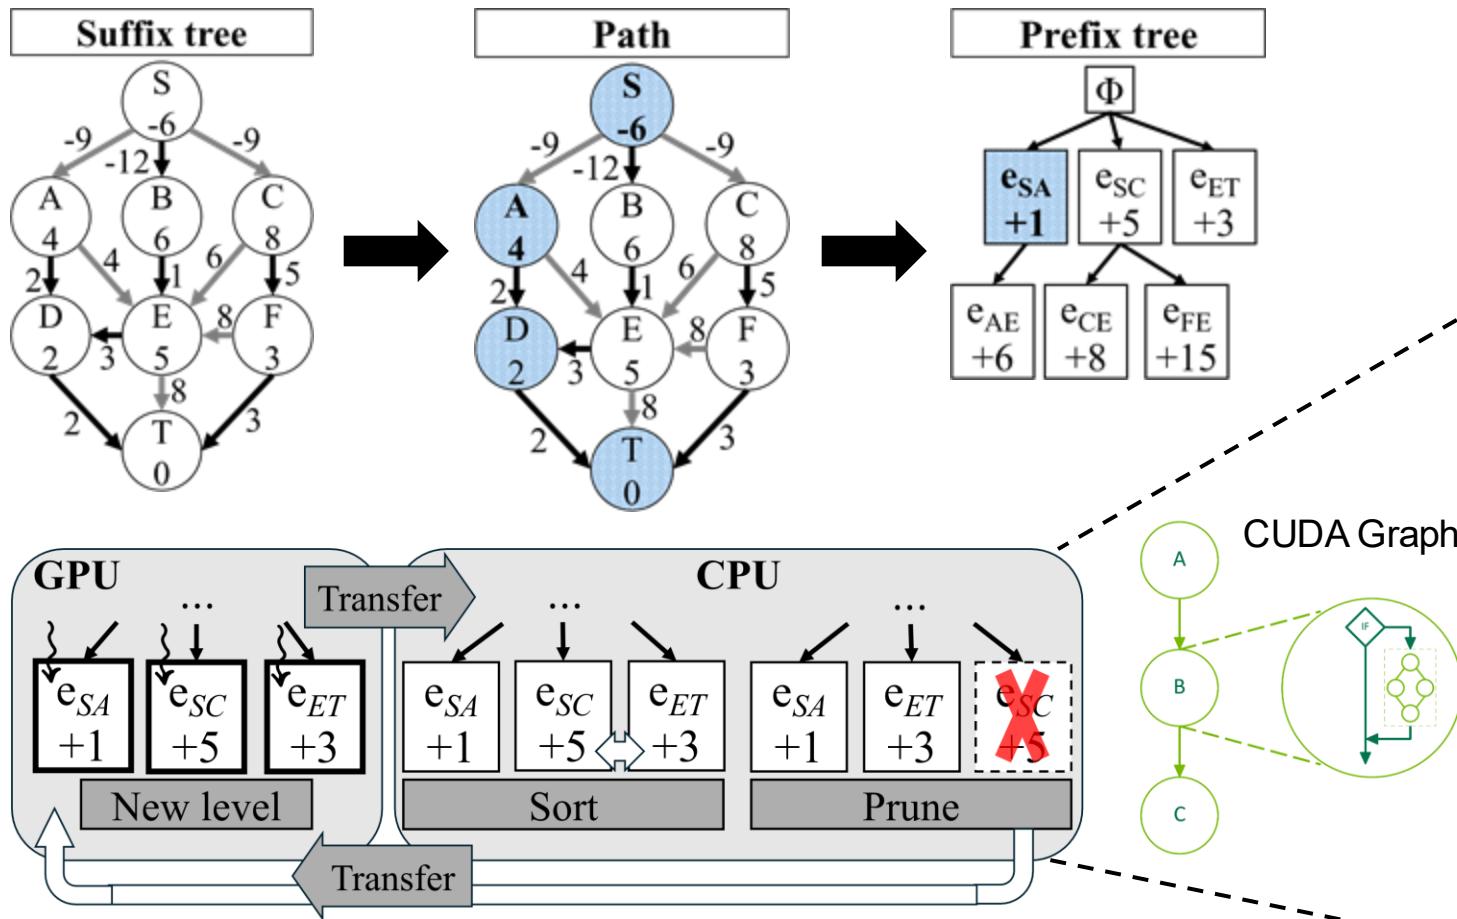
²: Tsung-Wei Huang, et al, “OpenTimer v2: A New Parallel Incremental Timing Analysis Engine,” *IEEE TCAD*, 2022

Our Research atop Task-parallel STA



GPU-accelerated Path-based Analysis (PBA)

- A GPU-parallel path generation algorithm¹

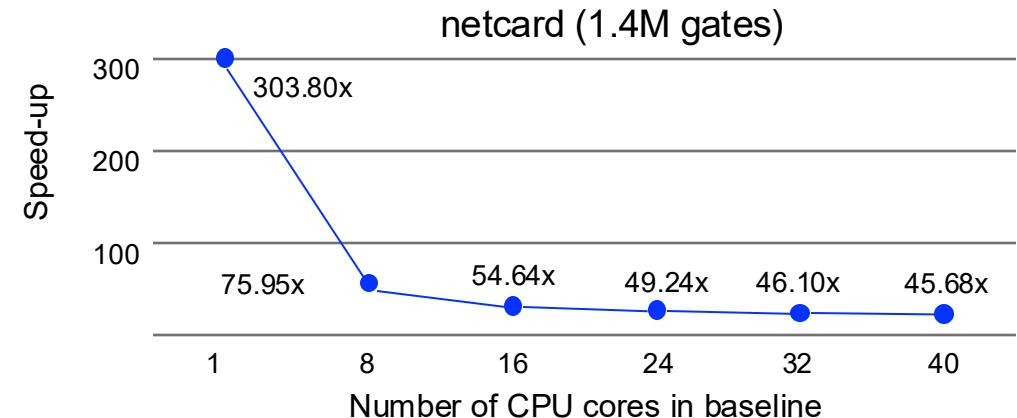
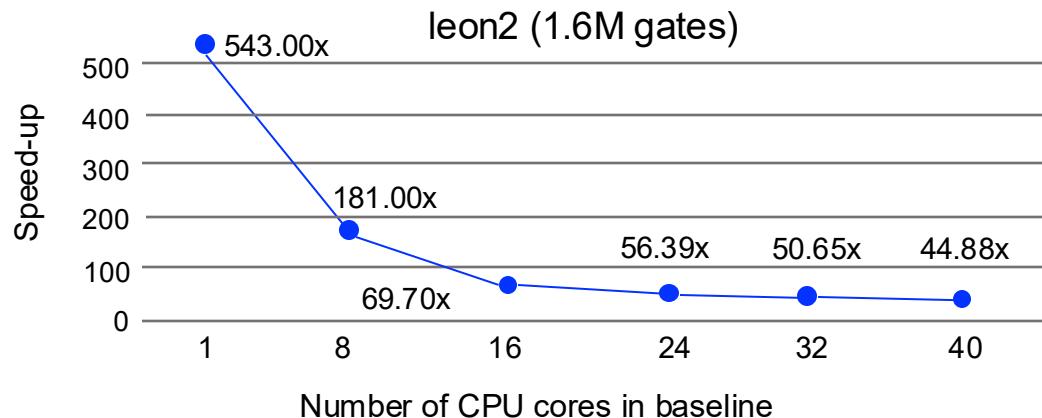


```

Input :  $G^+$  in CSR format,  $N$  as #vertices,  $M$  as #edges,
       vertices[N], edges[M], weights[M]
Input : Shortest path forest, forest[N] as edge array,
       distances[N] as distance array
Input : currLevel as current level
Input : levelSize as the number of entries in current level
Result: Explore critical path candidates in next level
1 tid ← blockIdx.x * blockDim.x + threadIdx.x;
2 if tid ≥ levelSize then
3   | return;
4 end
5 offset ← (tid == 0) ? 0 : currLevel[tid-1].childOffset;
6 level ← currLevel[tid].level;
7 slack ← currLevel[tid].slack;
8 v ← currLevel[tid].to;
9 while v is not endpoint do
10  edgeStart ← vertices[v];
11  edgeEnd ← (v == N-1) ? M : vertices[v+1];
12  for eid ← edgeStart to edgeEnd do
13    neighbor ← edges[eid];
14    weight ← weights[eid];
15    if eid is deviation edge then
16      /* Fill out child path data */
17      newPath ← nextLevel[offset];
18      newPath.level ← level+1;
19      newPath.from ← v;
20      newPath.to ← neighbor;
21      newPath.parent ← tid;
22      newPath.childOffset ← 0;
23      newPath.slack ← slack + distances[neighbor] +
         weight - distances[v];
24      offset ← offset + 1;
25  end
26 end
27 /* Traverse along the shortest path
   forest
   v = forest[v];
28 end
29 return;
30
*/
```

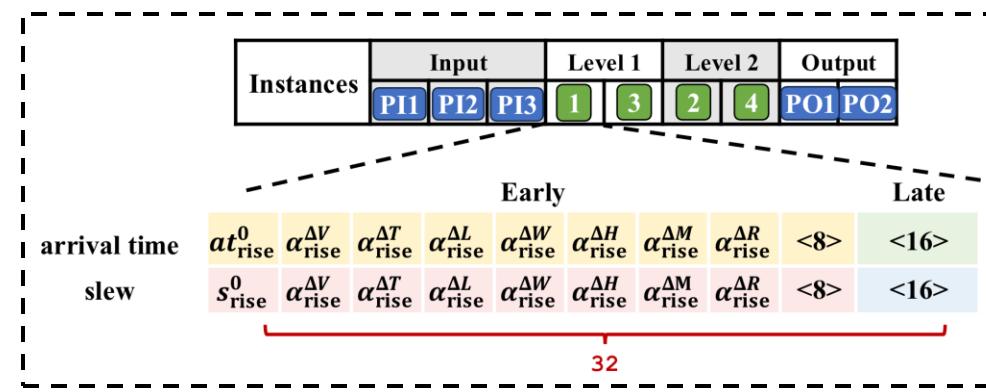
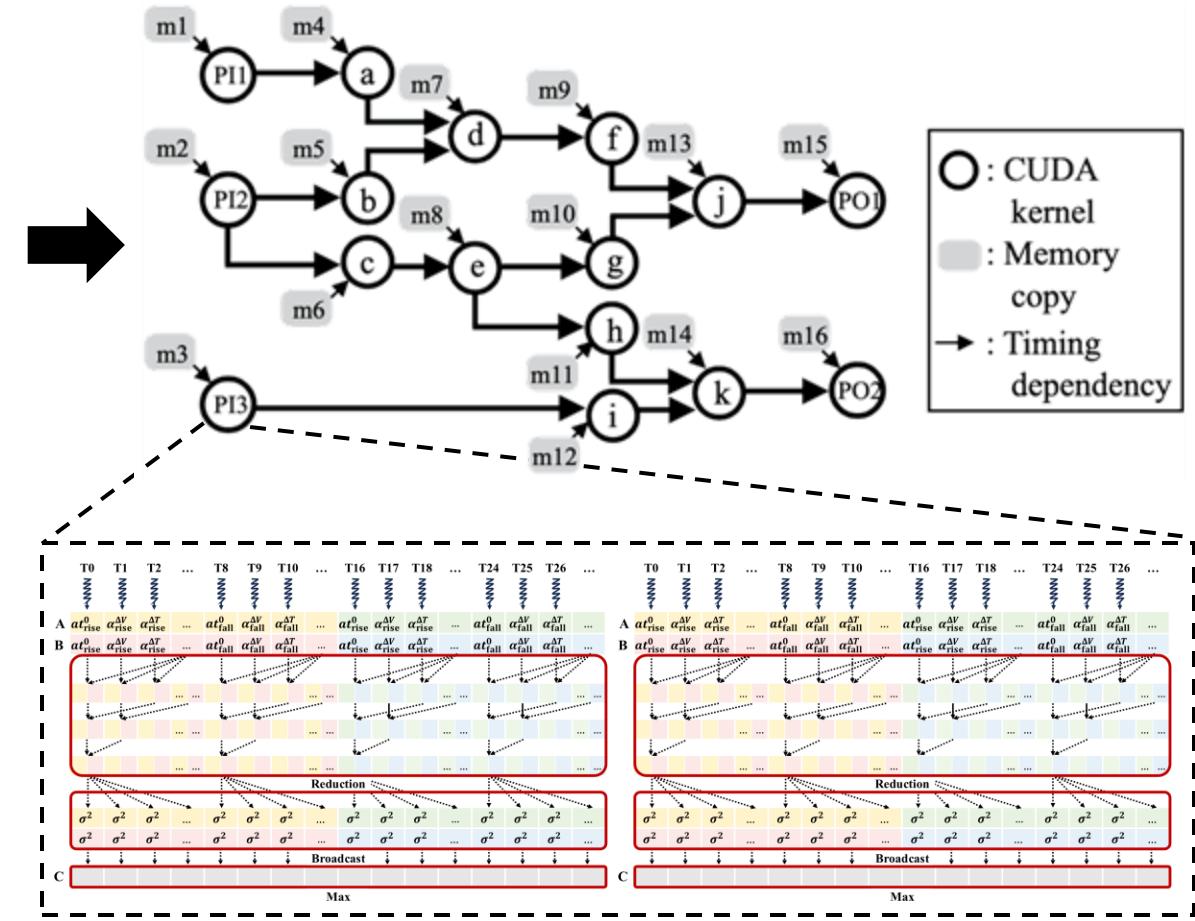
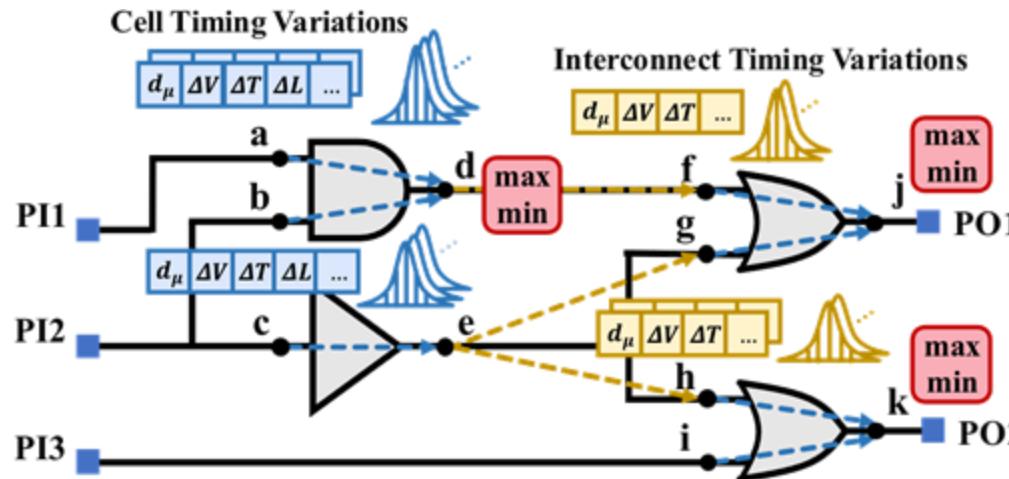
Performance of GPU-based Path Generation

Benchmark	#Pins	#Gates	#Arcs	OpenTimer Runtime	Our Algorithm #MDL=10		Our Algorithm #MDL=15		Our Algorithm #MDL=20	
					Runtime	Speed-up	Runtime	Speed-up	Runtime	Speed-up
leon2	4328255	1616399	7984262	2875783	4708.36	611×	5295.49ms	543×	5413.84	531×
leon3mp	3376821	1247725	6277562	1217886	5520.85	221×	7091.79ms	172×	8182.84	149×
netcard	3999174	1496719	7404006	752188	2050.60	367×	2475.90ms	304×	2484.08	303×
vga_lcd	397809	139529	756631	53204	682.94	77.9×	683.04ms	77.9×	706.16	75.3×
vga_lcd_iccad	679258	259067	1243041	66582	720.40	92.4×	754.35ms	88.3×	766.29	86.9×
b19_iccad	782914	255278	1576198	402645	2144.67	188×	2948.94ms	137×	3483.05	116×
des_perf_ispd	371587	138878	697145	24120	763.79	31.6×	766.31ms	31.5×	780.56	30.9×
edit_dist_ispd	416609	147650	799167	614043	1818.49	338×	2475.12ms	248×	2900.14	212×
mgc_edit_dist	450354	161692	852615	694014	1463.61	474×	1485.65ms	467×	1493.90	465×
mgc_matric_mult	492568	171282	948154	214980	994.67	216×	1075.90ms	200×	1113.26	193×



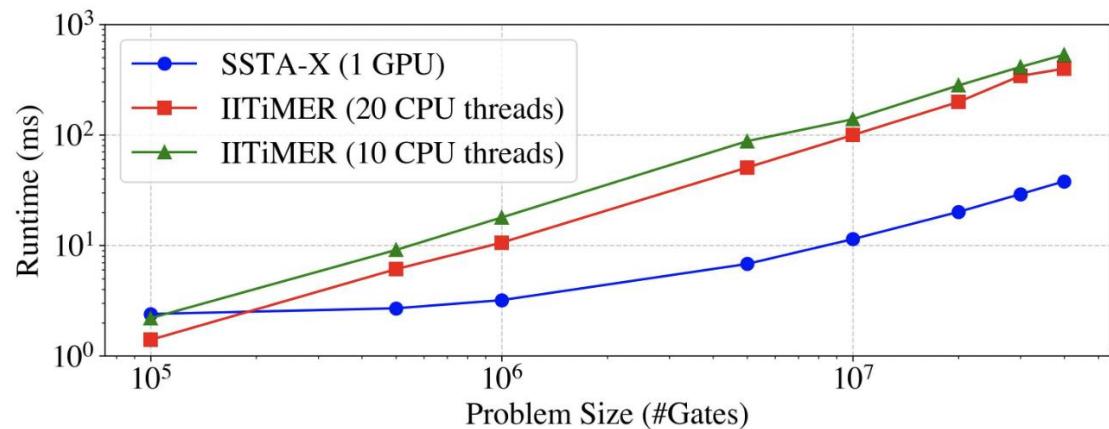
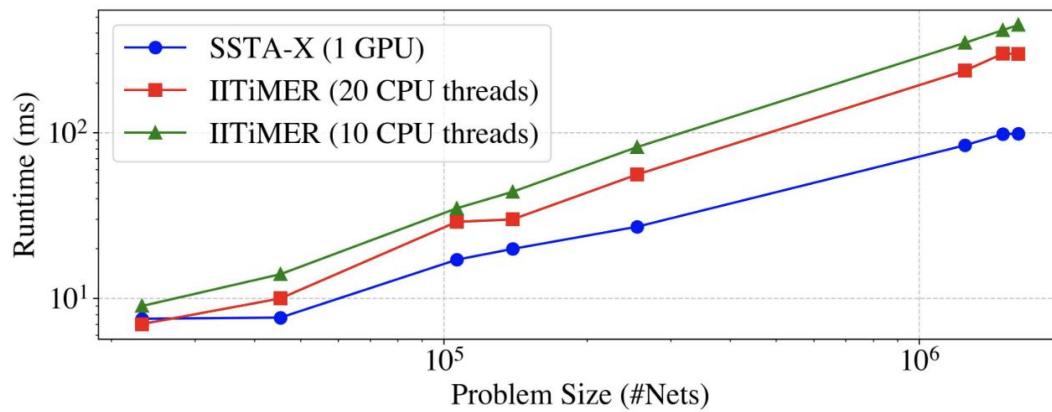
Case Study 2: GPU-accelerated SSTA

- Model the SSTA propagation workload into a GPU task graph



GPU-accelerated STA vs CPU-parallel SSTA¹

Circuit	#Gates	#Nets	#Pins	Runtime (ms)				Speedup of SSTA-X over			Error Rate vs. MC-SSTA μ and σ^2 (%)
				MC-SSTA 20 CPUs	IITiMER 1 CPUs	IITiMER 20 CPUs	SSTA-X 1 GPU	MC-SSTA 20 CPUs	IITiMER 1 CPUs	IITiMER 20 CPUs	
aes_core	22938	23199	66751	3483	129	26	20	174.15×	6.45×	1.30×	0.29
b19	255278	255300	782914	54989	1411	335	134	410.36×	10.50×	2.50×	0.38
cordic	45359	45393	127993	7286	226	70	30	242.86×	7.53×	2.33×	0.07
des_perf	138878	139112	371587	11403	580	115	52	219.28×	11.15×	2.21×	0.36
leon2	1616369	1616984	4328255	178620	6912	1667	554	322.41×	12.47×	3.01×	0.32
leon3mp	1247725	1247979	3376832	167752	5455	1487	476	352.42×	11.46×	3.10×	0.40
netcard	1496719	1498555	3999174	209258	6668	1799	590	354.67×	11.30×	3.04×	0.34
vga_lcd	139529	139635	397809	20233	674	157	72	281.01×	9.36×	2.18×	2.11
vga_lcd_iccad	259067	259152	679258	27455	1068	295	103	266.55×	10.36×	2.86×	0.41
Average								327.96×	11.32×	2.81×	0.58



Other Industrial Applications of Taskflow

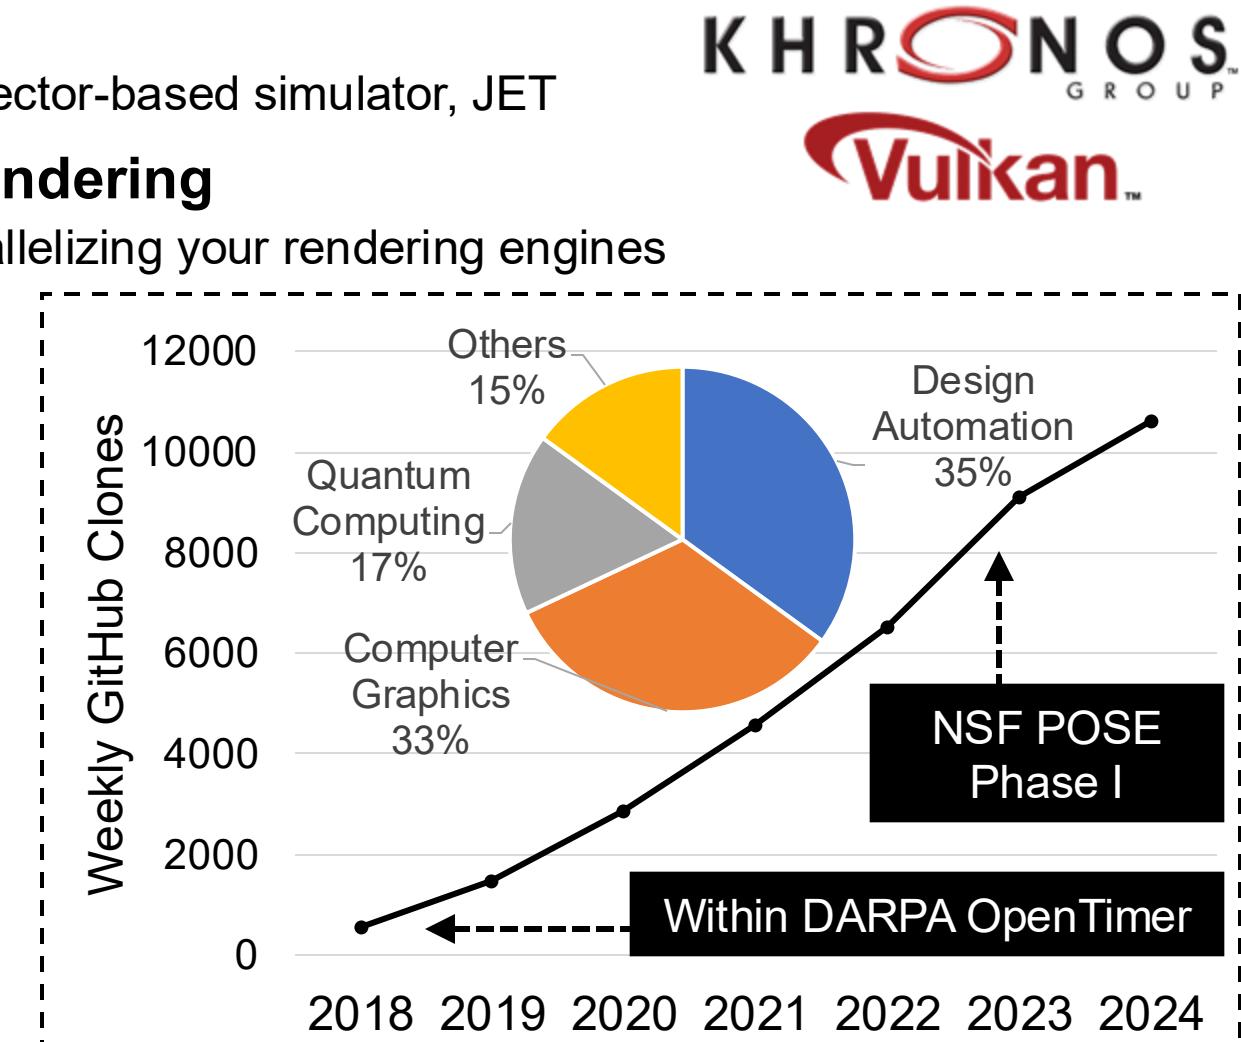
- **Quantum computing**
 - Xanadu used Taskflow in their state vector-based simulator, JET

- **Computer graphics and game rendering**
 - Vulkan recommends Taskflow for parallelizing your rendering engines

- **FPGA physical design**
 - Vivado uses Taskflow for synthesis

- **Embedded/edge computing**
 - Tesseract (robotics planning)
 - Cruise (autonomous car)
 - Reveal.Tech (drone vision)
 - Tesseract Robotic (planning tool)
 - ...

- **C++26 std::exec (coming soon)**
 - Benefit millions of C++ developers ☺



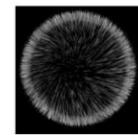


Conclusion

- Expressed your parallelism in the right way
- Programmed static task graph parallelism using Taskflow
- Programmed dynamic task graph parallelism using Taskflow
- Overcame the scheduling challenges
- Demonstrated the efficiency of Taskflow
- **Concluding the talk**



Thank you for using Taskflow!



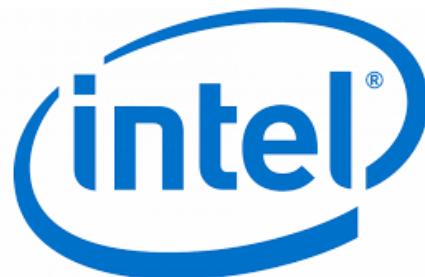
Explosion



...



Thank you for Sponsoring Taskflow!



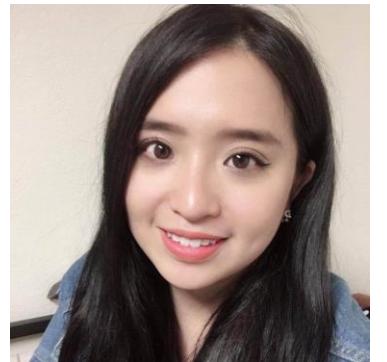
Google Summer of Code





Acknowledgment: Excellent PhD Students

- Please contact me if you have any intern/full-time opportunities!
 - We specialize in CAD, HPC, and GPU heterogeneous programming!
 - <https://tsung-wei-huang.github.io/team/> (or tsung-wei.huang@wisc.edu)





Questions?



Taskflow: <https://taskflow.github.io>

Static task graph parallelism

```
// Live: https://godbolt.org/z/j8hx3xnnx
tf::Taskflow taskflow;
tf::Executor executor;
auto [A, B, C, D] = taskflow.emplace(
    []() { std::cout << "TaskA\n"; },
    []() { std::cout << "TaskB\n"; },
    []() { std::cout << "TaskC\n"; },
    []() { std::cout << "TaskD\n"; });
A.precede(B, C);
D.succeed(B, C);
executor.run(taskflow).wait();
return 0;
```

Dynamic task graph parallelism

```
// Live: https://godbolt.org/z/T87PrTax
tf::Executor executor;
auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});
auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);
executor.wait_for_all();
```