

# **Taskflow: A General-purpose Parallel and Heterogeneous Task Computing System**

---

Dr. Tsung-Wei (TW) Huang


Department of Electrical and Computer Engineering

University of Utah, Salt Lake City, UT




# Why Parallel Computing?

- It's critical to advance your application performance





# Your Computer is Already Parallel

- Intel i7-377K CPU (four cores to run your jobs in parallel)



# Parallel programming is crucial but very challenging ...





*Taskflow offers a solution*

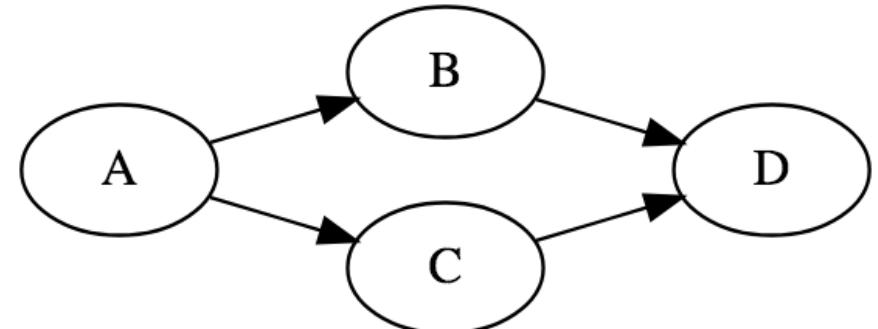
*How can we make it easier for C++  
developers to quickly write parallel and  
heterogeneous programs with **high  
performance scalability** and **simultaneous  
high productivity?***



# “Hello World” in Taskflow

---

```
#include <taskflow/taskflow.hpp> // Taskflow is header-only
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B, C); // A runs before B and C
    D.succeed(B, C); // D runs after B and C
    executor.run(taskflow).wait(); // submit the taskflow to the executor
    return 0;
}
```



# Drop-in Integration



---

- Taskflow is header-only – *no wrangle with installation*

```
~$ git clone https://github.com/taskflow/taskflow.git # clone it only once
~$ g++ -std=c++17 simple.cpp -I taskflow/taskflow -O2 -pthread -o simple
~$ ./simple
TaskA
TaskC
TaskB
TaskD
```

# Built-in Profiler/Visualizer

```
# run the program with the environment variable TF_ENABLE_PROFILER enabled
~$ TF_ENABLE_PROFILER=simple.json ./simple
~$ cat simple.json
[
  {"executor": "0", "data": [{"worker": 0, "level": 0, "data": [{"span": [172, 186], "name": "static"}, {"span": [186, 190], "name": "subflow"}, {"span": [190, 194], "name": "cudaflow"}, {"span": [194, 198], "name": "condition"}, {"span": [198, 202], "name": "module"}, {"span": [202, 206], "name": "clustered"}]}]
]
# paste the profiling json data to https://taskflow.github.io/tfprof/
```



# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- Understand our scheduling algorithm
- Boost performance in real applications
- Collaborate on using Taskflow in your applications


# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- Understand our scheduling algorithm
- Boost performance in real applications
- Collaborate on using Taskflow in your applications

# Motivation: Parallelizing VLSI CAD Tools

- Billions of tasks with diverse computational patterns



How can we write efficient C++ parallel programs for this *monster computational task graph* with **millions of CPU-GPU dependent tasks along with algorithmic control flow**?

# We Invested a lot in Existing Tools ...

---



PaRSEC



StarPU



# Two Big Problems of Existing Tools

---

- Our problems define complex task dependencies
  - **Example:** analysis algorithms compute the circuit network of million of node and dependencies
  - **Problem:** existing tools are often good at loop parallelism but weak in expressing heterogeneous task graphs at this large scale
- Our problems define complex control flow
  - **Example:** optimization algorithms make essential use of *dynamic control flow* to implement various patterns
    - Combinatorial optimization, analytical methods
  - **Problem:** existing tools are *directed acyclic graph* (DAG)-based and do not anticipate cycles or conditional dependencies, lacking *end-to-end* parallelism

# Example: An Iterative Placement Optimizer

- 4 computational tasks with dynamic control flow


#1: starts with init task

#2: enters the optimizer task (e.g., GPU solving linear system)

#3: checks if the optimization converged

- No: loops back to optimizer
- Yes: proceeds to stop

#4: outputs the result



# Need a New C++ Parallel Programming System

---

While designing parallel algorithms is non-trivial ...



what makes parallel programming an enormous challenge is the infrastructure work of  
***“how to efficiently express dependent tasks along with an algorithmic control flow and schedule them across heterogeneous computing resources”***

# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- Understand our scheduling algorithm
- Boost performance in real applications
- Collaborate on using Taskflow in your applications



**WARNING**


**Code Ahead**

# #1: Static Task (“Hello World” Revisited)


```
#include <taskflow/taskflow.hpp> // Taskflow is header-only
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B, C); // A runs before B and C
    D.succeed(B, C); // D runs after B and C
    executor.run(taskflow).wait();
    return 0;
}
```

This talk focuses on three task types:

1. static task
2. cudaFlow task
3. condition task



# “Hello World” in OpenMP



```
#include <omp.h> // OpenMP is a lang ext to describe parallelism using compiler directives
int main(){
    #omp parallel num_threads(std::thread::hardware_concurrency())
    {
        int A_B, A_C, B_D, C_D;
        #pragma omp task depend(out: A_B, A_C)
        {
            std::cout<< "TaskA\n";
        }
        #pragma omp task depend(in: A_B; out: B_D)
        {
            std::cout<< " TaskB\n";
        }
        #pragma omp task depend(in: A_C; out: C_D)
        {
            std::cout<< " TaskC\n";
        }
        #pragma omp task depend(in: B_D, C_D)
        {
            std::cout<< "TaskD\n";
        }
    }
    return 0;
}
```

*Task dependency clauses*


*Task dependency clauses*

*Task dependency clauses*

*Task dependency clauses*

*OpenMP task clauses are **static** and **explicit**;  
Programmers are responsible for a **proper order of writing tasks** consistent with sequential execution*

# “Hello World” in TBB



```
#include <tbb.h> // Intel's TBB is a general-purpose parallel programming library in C++
int main(){
    using namespace tbb;
    using namespace tbb::flow;
    int n = task_scheduler_init::default_num_threads();
    task_scheduler_init init(n);
    graph g;
    continue_node<continue_msg> A(g, [] (const continue msg &) {
        std::cout << "TaskA";
    });
    continue_node<continue_msg> B(g, [] (const continue msg &) {
        std::cout << "TaskB";
    });
    continue_node<continue_msg> C(g, [] (const continue msg &) {
        std::cout << "TaskC";
    });
    continue_node<continue_msg> D(g, [] (const continue msg &) {
        std::cout << "TaskD";
    });
    make_edge(A, B);
    make_edge(A, C);
    make_edge(B, D);
    make_edge(C, D);
    A.try_put(continue_msg());
    g.wait_for_all();
}
```

Use TBB's FlowGraph  
for task parallelism


Declare a task as a  
continue\_node

TBB has excellent performance in generic parallel computing. Its drawback is mostly in the ease-of-use standpoint (simplicity, expressivity, and programmability).

TBB FlowGraph: <https://software.intel.com/content/www/us/en/develop/home.html>

# #2: cudaFlow Task

- Single Precision AX + Y (“SAXPY”)
  - Get x and y vectors on CPU (allocate\_x, allocate\_y)
  - Copy x and y to GPU (h2d\_x, h2d\_y)
  - Run saxpy kernel on x and y (saxpy kernel)
  - Copy x and y back to CPU (d2h\_x, d2h\_y)




# #2: cudaFlow Task (cont'd)

```
const unsigned N = 1<<20;
std::vector<float> hx(N, 1.0f), hy(N, 2.0f);
float *dx{nullptr}, *dy{nullptr};
auto allocate_x = taskflow.emplace([&](){ cudaMalloc(&dx, 4*N);});
auto allocate_y = taskflow.emplace([&](){ cudaMalloc(&dy, 4*N);});
```

```
auto cudaflow = taskflow.emplace([&](tf::cudaFlow& cf) {
    auto h2d_x = cf.copy(dx, hx.data(), N); // CPU-GPU data transfer
    auto h2d_y = cf.copy(dy, hy.data(), N);
    auto d2h_x = cf.copy(hx.data(), dx, N); // GPU-CPU data transfer
    auto d2h_y = cf.copy(hy.data(), dy, N);
    auto kernel = cf.kernel((N+255)/256, 256, 0, saxpy, N, 2.0f, dx, dy);
    kernel.succeed(h2d_x, h2d_y).precede(d2h_x, d2h_y);
});
```

```
cudaflow.succeed(allocate_x, allocate_y);
executor.run(taskflow).wait();
```

To Nvidia  
*cudaGraph*



# Three Key Motivations of cudaFlow


- Our closure enables stateful interface
  - Users capture data in reference to marshal data exchange between CPU and GPU tasks
- Our closure hides implementation details judiciously
  - We use cudaGraph (since cuda 10) due to its excellent performance, much faster than streams in large graphs
- Our closure extend to new accelerator types
  - syclFlow, openclFlow, coralFlow, tpuFlow, fpgaFlow, etc.

```
auto cudaflow = taskflow.emplace([&](tf::cudaFlow& cf) {  
    auto h2d_x = cf.copy(dx, hx.data(), N); // CPU-GPU data transfer  
    auto h2d_y = cf.copy(dy, hy.data(), N);  
    auto d2h_x = cf.copy(hx.data(), dx, N); // GPU-CPU data transfer  
    auto d2h_y = cf.copy(hy.data(), dy, N);  
    auto kernel = cf.kernel((N+255)/256, 256, 0, saxpy, N, 2.0f, dx, dy);  
    kernel.succeed(h2d_x, h2d_y).precede(d2h_x, d2h_y);  
});
```

We do not simplify kernel programming but  
**focus on *CPU-GPU tasking* that affects the performance to a large extent!** (same for data abstraction)

# #3: Condition Task (if-else)


```
auto init      = taskflow.emplace([&](){ initialize_data_structure(); } )  
                  .name("init");  
auto optimizer = taskflow.emplace([&](){ matrix_solver(); } )  
                  .name("optimizer");  
auto converged = taskflow.emplace([&](){ return converged() ? 1 : 0; } )  
                  .name("converged");  
auto output    = taskflow.emplace([&](){ std::cout << "done!\n"; } );  
                  .name("output");  
  
init.precede(optimizer);  
optimizer.precede(converged);  
converged.precede(optimizer, output); // return 0 to the optimizer again
```



*Condition task integrates control flow into a task graph to form **end-to-end parallelism***

# #3: Condition Task (While/For Loop)


```
tf::Taskflow taskflow;
int i;
auto [init, cond, body, back, done] = taskflow.emplace(
    [&](){ std::cout << "i=0"; i=0; },
    [&](){ std::cout << "while i<5\n"; return i < 5 ? 0 : 1; },
    [&](){ std::cout << "i++=" << i++ << '\n'; },
    [&](){ std::cout << "back\n"; return 0; },
    [&](){ std::cout << "done\n"; }
);
init.precede(cond);
cond.precede(body, done);
body.precede(back);
back.precede(cond);
```



# #3: Condition Task (Non-deterministic Loops)

```
auto A = taskflow.emplace([&](){ } );
auto B = taskflow.emplace([&](){ return rand()%2; } );
auto C = taskflow.emplace([&](){ return rand()%2; } );
auto D = taskflow.emplace([&](){ return rand()%2; } );
auto E = taskflow.emplace([&](){ return rand()%2; } );
auto F = taskflow.emplace([&](){ return rand()%2; } );
auto G = taskflow.emplace([&]());
A.precede(B).name("init");
B.precede(C, B).name("flip-coin-1");
C.precede(D, B).name("flip-coin-2");
D.precede(E, B).name("flip-coin-3");
E.precede(F, B).name("flip-coin-4");
F.precede(G, B).name("flip-coin-5");
G.name("end");
```


Each task flips a binary coin to decide the next path



You can describe non-deterministic, nested control flow!

# #3: Condition Task (Switch)

```
auto [source, swcond, case1, case2, case3, target] = taskflow.emplace(  
    [](){ std::cout << "source\n"; },  
    [](){ std::cout << "switch\n"; return rand()%3; },  
    [](){ std::cout << "case 1\n"; return 0; },  
    [](){ std::cout << "case 2\n"; return 0; },  
    [](){ std::cout << "case 3\n"; return 0; },  
    [](){ std::cout << "target\n"; }  
);  
source.precede(swcond);  
swcond.precede(case1, case2, case3);  
target.succeed(case1, case2, case3);
```



# Existing Frameworks on Control Flow?

---


- Expand a task graph across fixed-length iterations
  - Graph size is linearly proportional to decision points
- Unknown iterations? Non-deterministic conditions?
  - Complex dynamic tasks executing “if” on the fly
- Dynamic control-flow tasks?
- ... (resort to client-side decision)

*Existing frameworks on expressing conditional tasking or dynamic control flow suffer from exponential growth of code complexity*




# Everything is Unified in Taskflow


- Use “emplace” to create a task
- Use “precede” to add a task dependency
- No need to learn different sets of API
- You can create a really complex graph
  - Subflow(ConditionTask(cudaFlow))
  - ConditionTask(StaticTask(cudaFlow))
  - Composition(Subflow(ConditionTask))
  - Subflow(ConditionTask(cudaFlow))
  - ...
- Scheduler performs end-to-end optimization
  - Runtime, energy efficiency, and throughput



# Example: *k*-means Clustering



$k=3$



# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- **Understand our scheduling algorithm**
- Boost performance in real applications
- Collaborate on using Taskflow in your applications

# Submit Taskflow to Executor

---

- Executor manages a set of threads to run taskflows
  - All execution methods are *non-blocking*
  - All execution methods are *thread-safe*

```
{  
    tf::Taskflow taskflow1, taskflow2, taskflow3;  
    tf::Executor executor;  
    // create tasks and dependencies  
    // ...  
    auto future1 = executor.run(taskflow1);  
    auto future2 = executor.run_n(taskflow2, 1000);  
    auto future3 = executor.run_until(taskflow3, [i=0](){ return i++>5 });  
    executor.async([](){ std::cout << "async task\n"; });  
    executor.wait_for_all(); // wait for all the above tasks to finish  
}
```



# Executor Scheduling Algorithm

---

- Task-level scheduling
  - Decides how tasks are enqueued under control flow
    - Goal #1: ensures a feasible path to carry out control flow
    - Goal #2: avoids task race under cyclic and conditional execution
    - Goal #3: maximizes the capability of conditional tasking
- Worker-level scheduling
  - Decides how tasks are executed by which workers
    - Goal #1: adopts work stealing to dynamically balance load
    - Goal #2: adapts workers to available task parallelism
    - Goal #3: maximizes performance, energy, and throughput





# Task-level Scheduling

- “*Strong dependency*” versus “*Weak dependency*”
  - Weak dependency: dependencies out of condition tasks
  - Strong dependency: others else



# Task-level Scheduling (cont'd)



- Condition task is powerful but prone to mistakes ...



*It is users' responsibility to ensure a taskflow is properly conditioned, i.e., no task race under our task-level scheduling policy*

# Worker-level Scheduling

- Taskflow adopts *work stealing* to gain *dynamic load balancing*
- What is work stealing? Why?
  - I finish my jobs first, and then steal jobs from you
  - So, we can improve performance and balance our loads





# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- Understand our scheduling algorithm
- **Boost performance in real applications**
- Collaborate on using Taskflow in your applications

# Application 1: VLSI Placement







- Optimize cell locations on a chip



VLSI optimization  
makes essential use of  
dynamic control flow

# Application 1: VLSI Placement (cont'd)


- Runtime, memory, power, and throughput



Performance improvement comes from the *end-to-end* expression of CPU-GPU dependent tasks using condition tasks


# Application 2: Machine Learning

- IEEE HPEC/MIT/Amazon Sparse DNN Challenge
  - Compute a 1920-layer DNN each of 65536 neurons



# Application 2: Machine Learning (cont'd)

- Comparison with TBB and StarPU



- Taskflow's runtime is up to 2x faster
  - Adaptive work stealing balances the worker count with task parallelism
- Taskflow's memory is up to 1.6x less
  - Conditional tasking allows efficient reuse of tasks

# Summary of Experiments

---

- Parallel computing infrastructure matters



*Different models give different implementations. The parallel code/algorithm may run fast, yet the parallel computing infrastructure to support that algorithm may dominate the entire performance.*

Taskflow enables *end-to-end* expression of CPU-GPU dependent tasks along with algorithmic control flow

# Agenda

---

- Express your parallelism in the right way
- Parallelize your applications using Taskflow
- Understand our scheduling algorithm
- Boost performance in real applications
- **Collaborate on using Taskflow in your applications**




Parallelism is  
never standalone

# GAMER + Taskflow




---

- GAMER accelerates maze routing using GPU
  - Primary innovation is GPU kernel – iterative parallel sweeping
- Potential collaboration topics for TCAD extension
  - Leverage Taskflow to achieve end-to-end heterogeneous parallelism
    - Condition tasks allow you to integrate control-flow decisions into dependencies
  - Leverage cudaFlow to reduce the kernel call overheads
  - Leverage multiple GPUs to handle a batch of nets



# Prior Collaborative Results with Taskflow

- We collaborate with Peking U to accelerate STA with Taskflow
  - GPU-accelerated GBA [ICCAD'20]
  - GPU-accelerated PBA [DAC'21, ICCAD'21]
  - GPU-accelerated CPPR [ICCAD'21]
  - Taskflow system [TCAD'21, TPDS'21]
  - ...



# Conclusion

---

- Taskflow is a lightweight parallel task programming system
  - Simple, efficient, and transparent tasking models
  - Efficient heterogeneous work-stealing executor
  - Promising performance in large-scale ML and VLSI CAD
- Taskflow is not to replace anyone but to
  - Complement the current state-of-the-art
  - Leverage modern C++ to express task graph parallelism
- Taskflow is very open to collaboration
  - We want to provide more higher-level algorithms
  - We want to broaden real use cases
  - We want to enhance the core functionalities (e.g., pipeline)

# Thank You All Using Taskflow!

---





# Use the right tool for the right job

Taskflow: <https://taskflow.github.io>

Thank You

Dr. Tsung-Wei Huang

[tsung-wei.huang@utah.edu](mailto:tsung-wei.huang@utah.edu)