

# HeteroTime: Accelerating Static Timing Analysis using GPUs

---

Dr. Tsung-Wei Huang

Department of Electrical and Computer Engineering  
University of Utah, UT

<https://tsung-wei-huang.github.io/>



# Agenda

---

- **Introduce the scalability problem of STA**
  - What is STA and its challenges?
  - Why do we need new parallel paradigm for STA?
- **Accelerate graph-based analysis using GPU**
- **Accelerate path-based analysis using GPU**

# Static Timing Analysis

- Static timing analysis (STA)

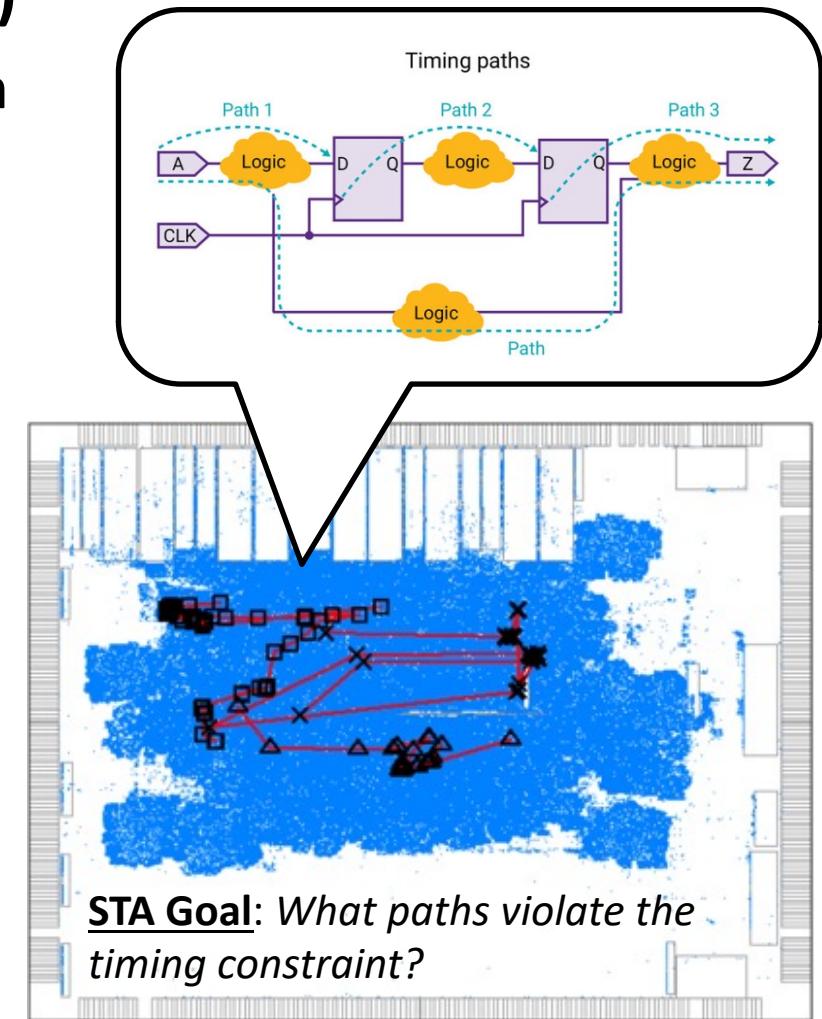
- Key step in the VLSI design
- Verify the circuit timing

- Analyze worst-case timing

- Minimum timing values
- Maximum timing values

- Challenges

- Compute giant graphs
- Analyze millions of paths
- Balance the loads
- ...



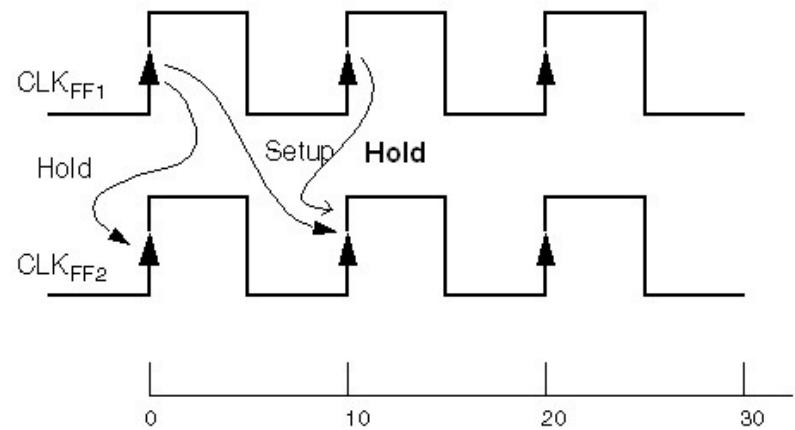
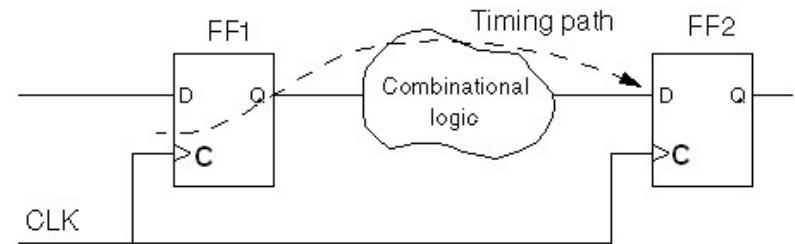
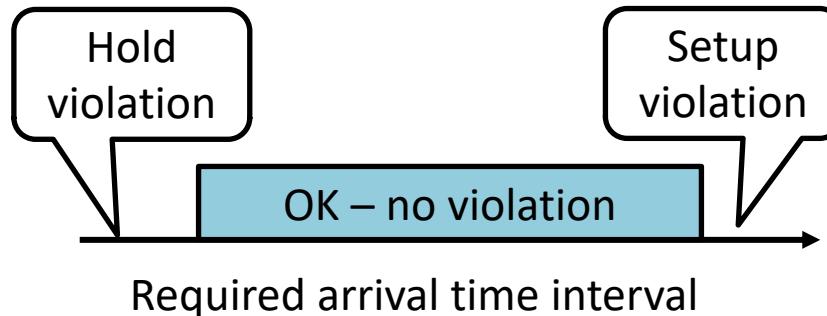
# Timing Checks (Required Arrival Time)

- Modern circuits are sequential

- Drive data signal via clocks
  - Capture data via flip-flops (FF)s

- Timing constraints

- Min required arrival time
    - After clock: hold
  - Max required arrival time
    - Before clock: setup



# “Traffic Light” Analogy

Can I pass the block before the next red light with 40 mph?



# Building a Good Traffic System is Hard

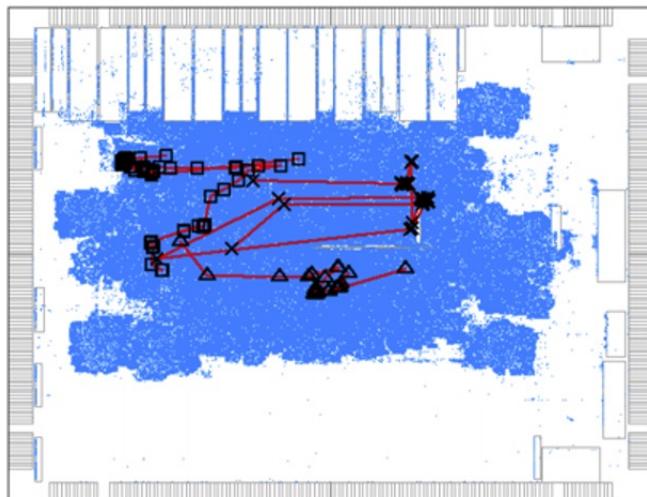
---

Trillions of sections and traffic lights to analyze ...



# STA is Computationally Challenging!

- STA graphs is extremely large and irregular
  - Millions to billions of nodes and edges
  - Propagate timing information along giant graphs

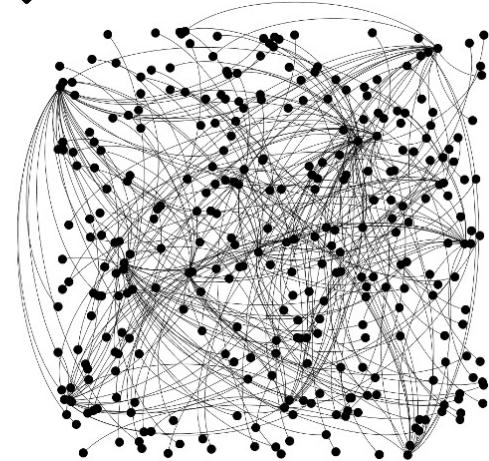
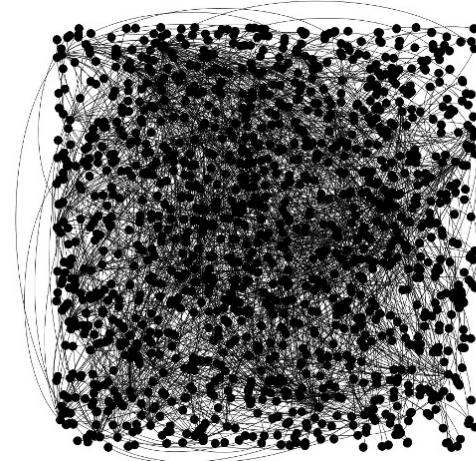


ISPD circuit design (10M gates)

Complete analysis can take **8 hours** and **800 GB RAM**

STA graphs

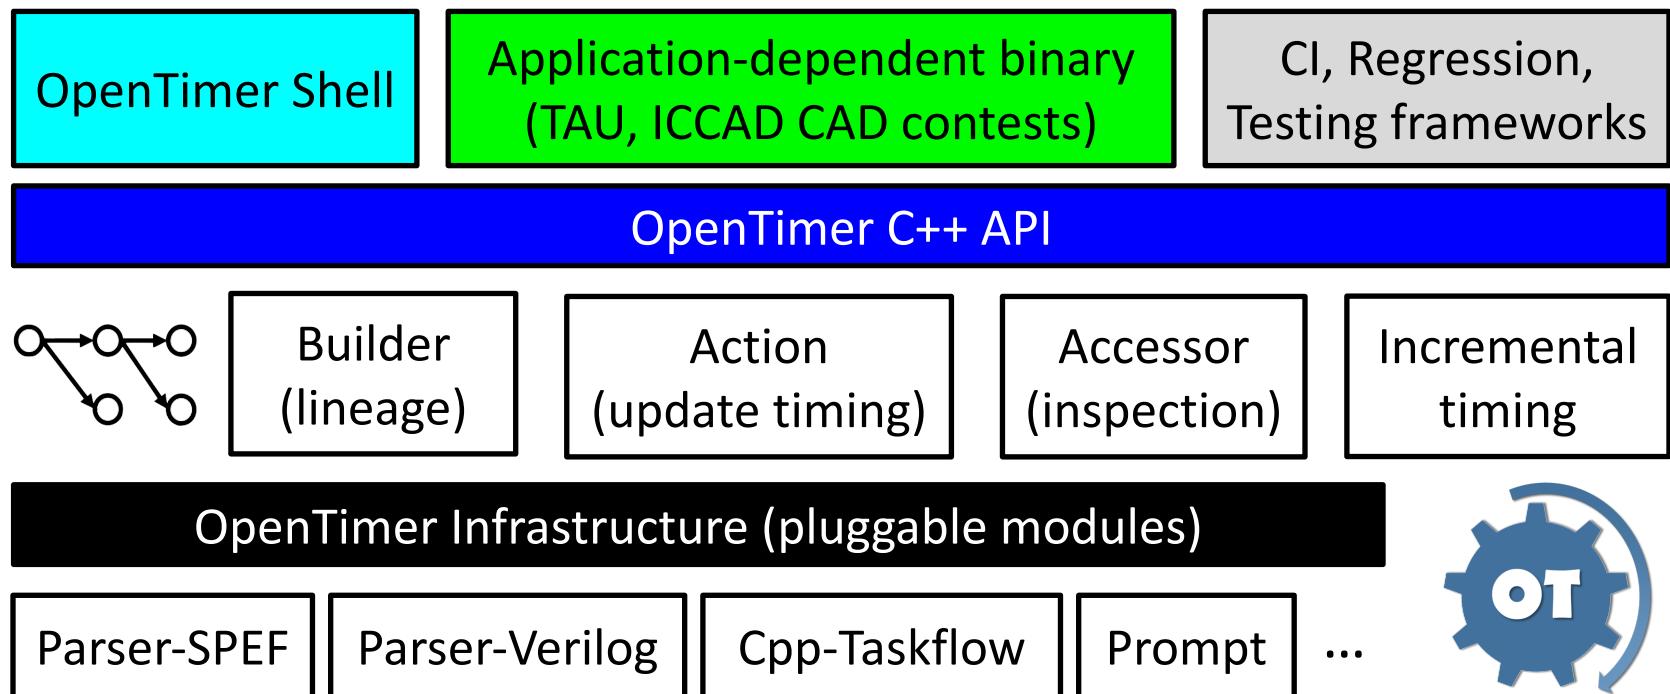
A datapath



STA graphs are extremely large and irregular

# Our STA Solution: OpenTimer

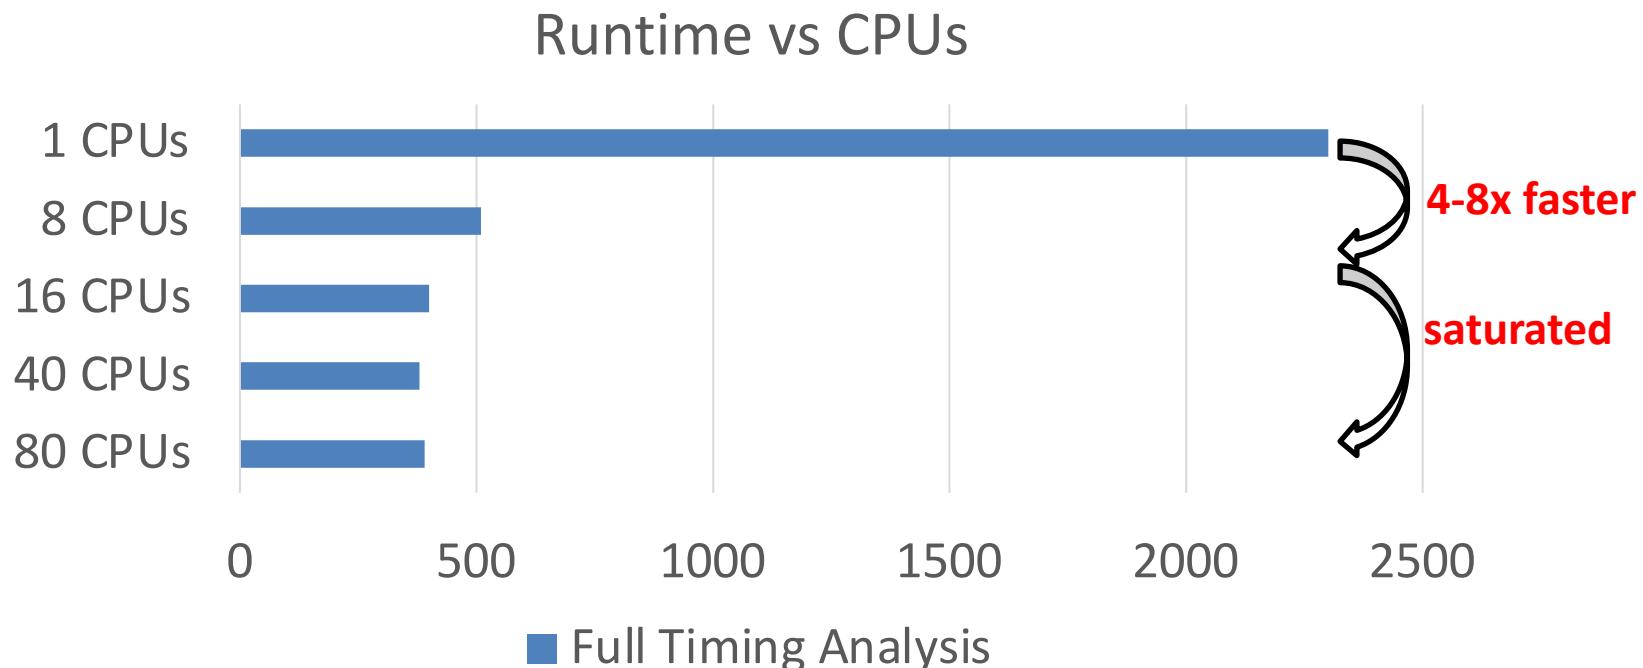
- ❑ Open-source incremental timing analysis software
  - ❑ <https://github.com/OpenTimer/OpenTimer>



T.-W. Huang et al., “OpenTimer v2: A New Parallel Incremental Timing Analysis Engine,” IEEE TCAD21

# How Can We Make STA Run Faster?

- ❑ Leverage many-core CPUs to speed up the runtime
  - ❑ Dramatic speed-up using 8 cores
  - ❑ Yet, scalability *saturates* at about 10–16 cores



# Observed Scalability Bottleneck

---

- ❑ CPU-only parallelism stagnates at about 10 cores
  - ❑ “Amdahl’s Law” limits the strong scalability
    - The maximum speed-up you can achieve from a problem is proportional to the part that can be parallelized
  - ❑ Circuit graph structures limits the maximum parallelism
    - If the graph has only 10 parallel nodes at a level, we won’t achieve 40x speed-up
  - ❑ Irregular computations limits the memory bandwidth
    - STA is graph-oriented, not cache-friendly
- ❑ Need to incorporate new parallel paradigms
  - ❑ GPU opens opportunities for new scalability milestones
    - e.g., 20—80x speed-up reported in placement

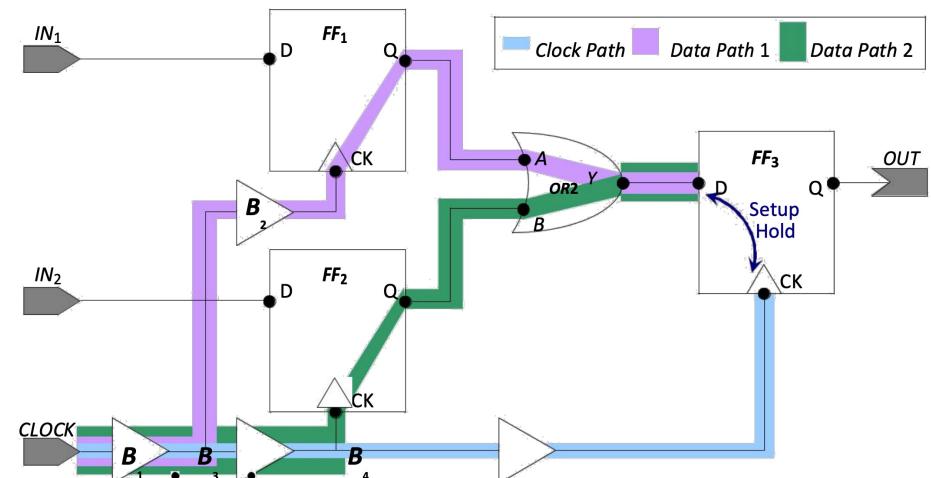
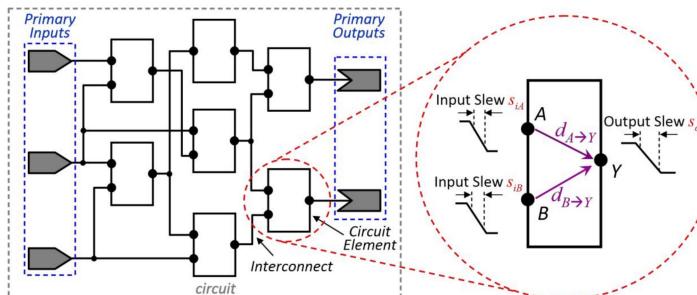


*How can we leverage new GPU parallel paradigms to accelerate static timing analysis algorithms and achieve transformational performance milestones?*

# Leverage GPU to Accelerate STA

- We target two important STA steps:
    - Graph-based analysis (GBA) – *shortest path finding*
    - Path-based analysis (PBA) – *k-shortest path finding*
  - We design new GPU-accelerated STA algorithms:
    - CPU-GPU task decomposition
    - GPU kernels for timing update
- PBA analyzes critical paths one by one on a updated graph

GBA computes the delay, slew, arrival time at each node and edge



# Agenda

---

- Introduce the scalability problem of STA
  - What is STA and its challenges?
  - Why do we need new parallel paradigm for STA?
- Accelerate graph-based analysis using GPU
- Accelerate path-based analysis using GPU

Z Guo, T-W Huang, and Y Lin, “GPU-Accelerated Static Timing Analysis,” *IEEE/ACM International Conference on Computer-aided Design (ICCAD)*, 2020

**Research Question:**

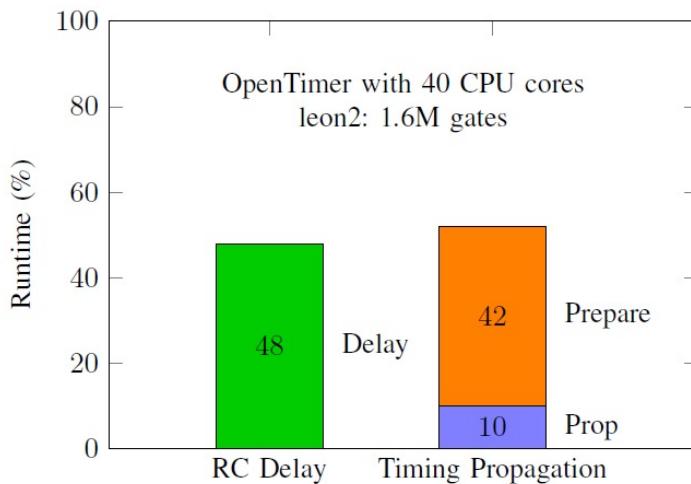
- *How can we use GPU to speed up graph-based analysis (GBA)?*
- *How do we overcome the challenges of running irregular graph computations on GPU?*



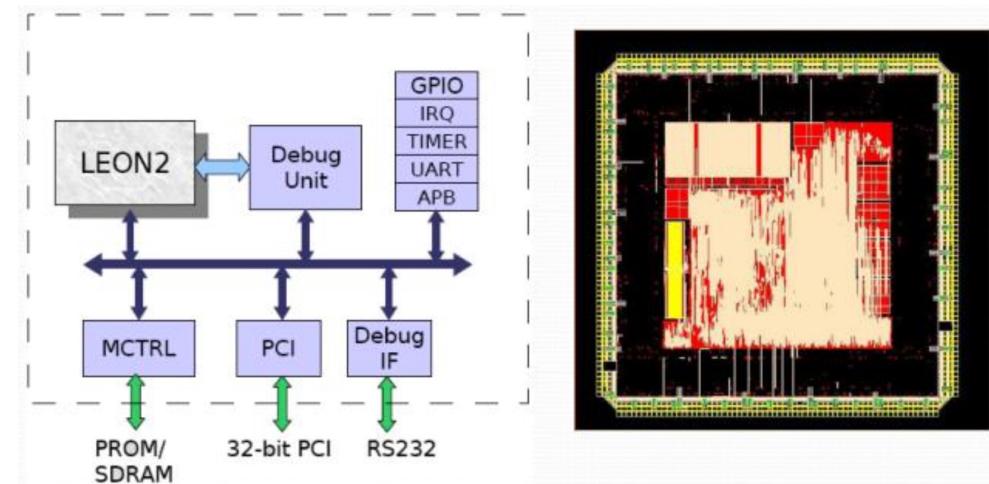
# Runtime Breakdown of GBA

## □ GBA has three time-consuming steps

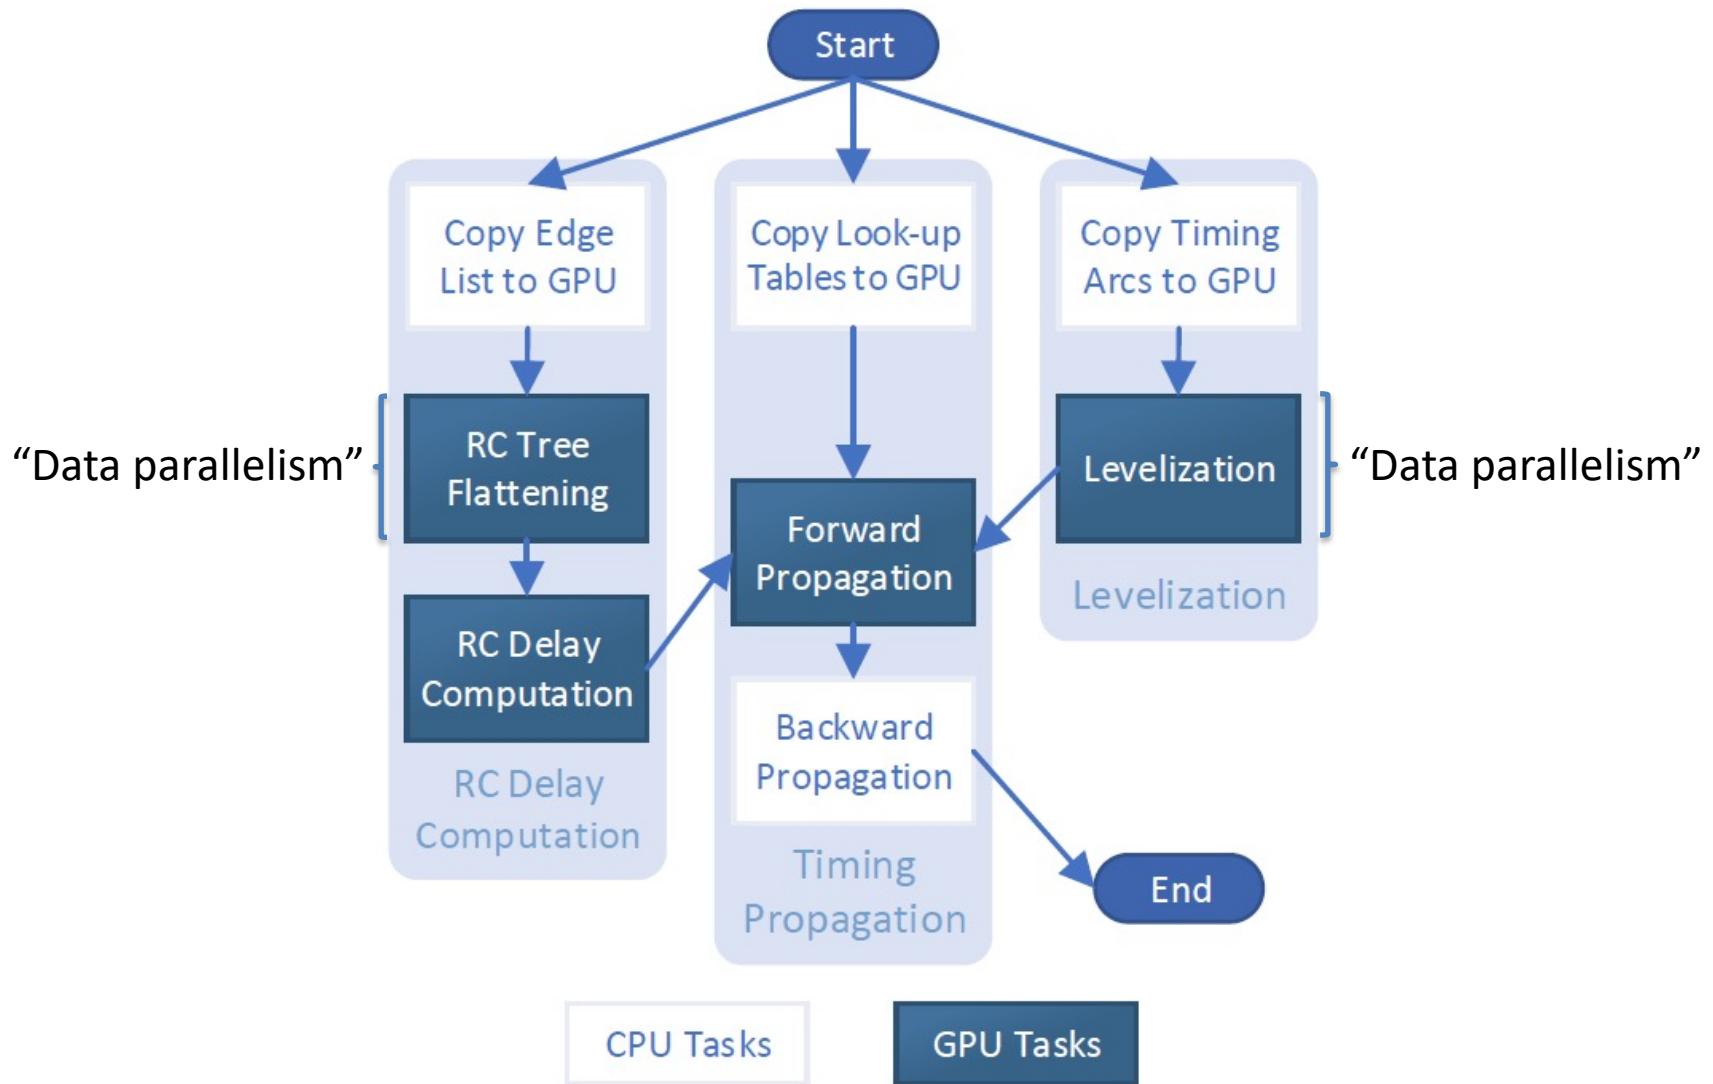
1. Prepare tasks through levelization → 42% runtime
  - All pins at the same level can run in parallel => data parallelism
2. Compute RC delay → 48% runtime
  - RC network (e.g., SPEF) can take GBs of data to compute
3. Propagate timing → 10% runtime



Leon2 GBA runtime breakdown

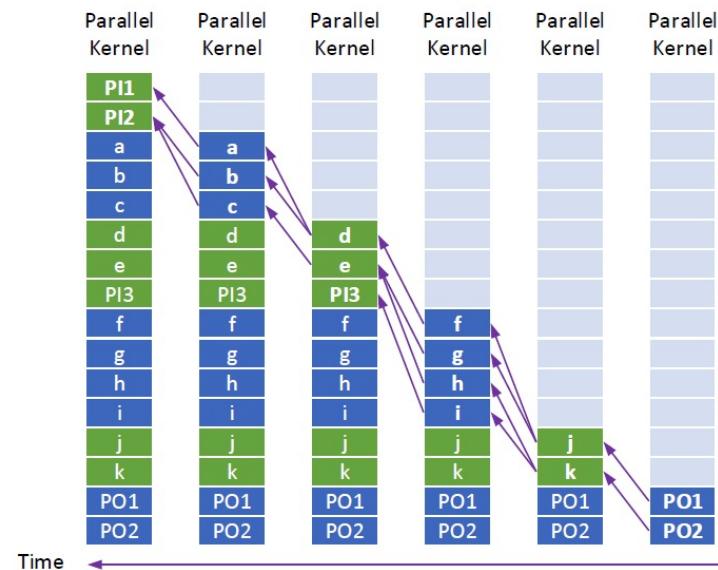
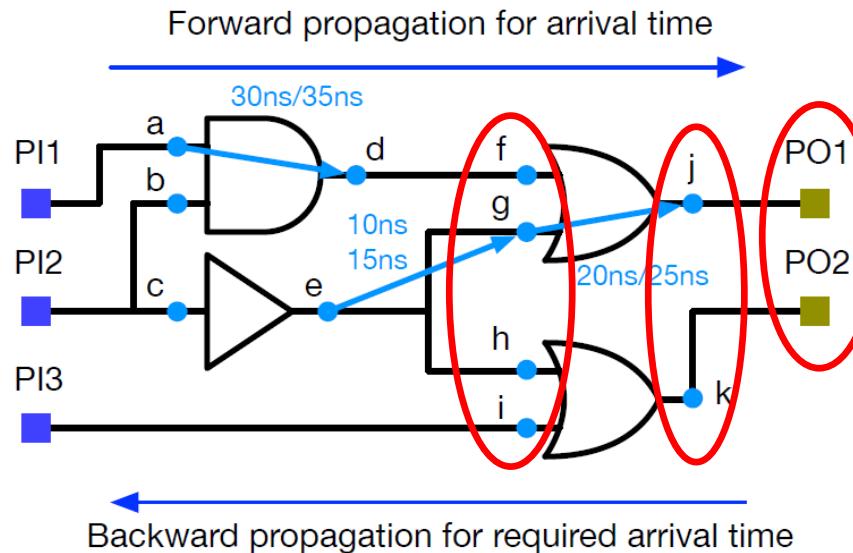


# GPU-Accelerated GBA Algorithm Flow



# Step #1: Levelization

- ☐ Levelize the circuit graph to a 2D levellist
  - ☐ Nodes at the same level can run in parallel (red circle)
  - ☐ Nodes at the same level can be modeled as a batch



- ☐ GPU-accelerated levelization using parallel frontiers

# Step #1: Levelization (cont'd)

## □ Levelization kernels

**Input:** the set of nodes  $nodes$

**Data:** the adjacency list  $out$ , the current in-degree  $in$

**Output:** a level list of nodes

```
1  $F \leftarrow \{f \in nodes : in_f = 0\};$ 
2 while  $F$  is not empty do
3   output  $F$ ;
4    $F' \leftarrow \{\}$ ;
5   Call advanceFrontier on  $F$  and get  $F'$ ;
6    $F \leftarrow F'$ ;
7 end
```

Each GPU thread to expand pin level from  $l$  to  $l+1$

**Input:** the old frontier  $F$

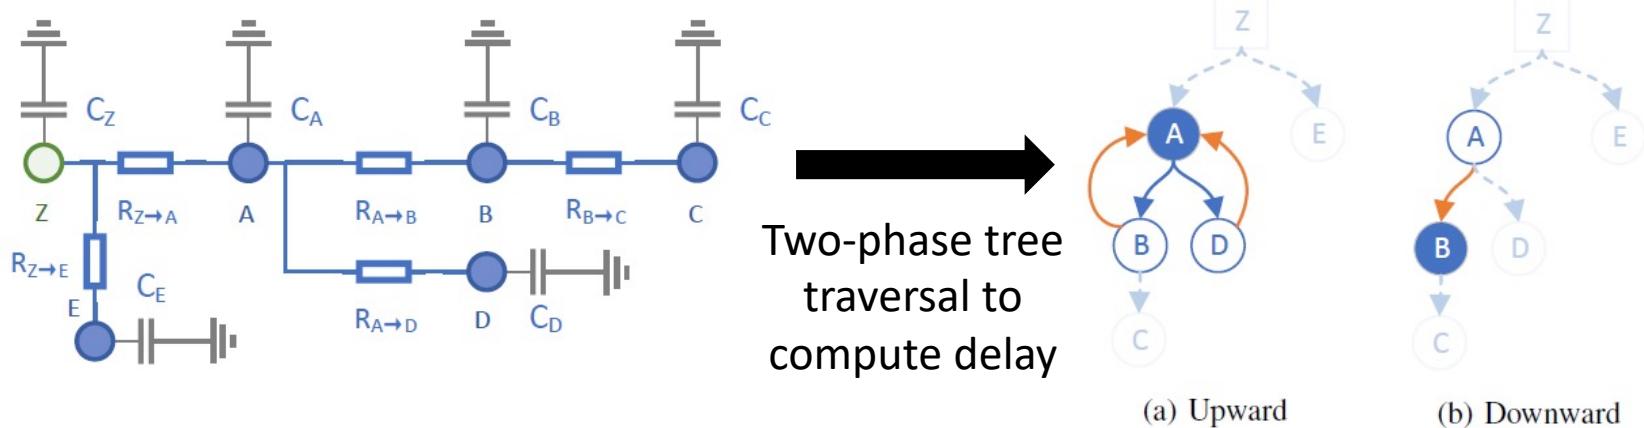
**Data:** the adjacency list  $out$ , in-degree array  $in$

**Output:** the new frontier  $F'$

```
1  $nodeID \leftarrow blockIdx.x \times blockDim.x + threadIdx.x;$ 
2 if  $nodeID \geq \text{size}(F)$  then return;
3 for  $v$  in  $out[nodeID]$  do
4    $oldvalue \leftarrow \text{atomicAdd}(in[v], -1);$ 
5   if  $oldvalue = 1$  then
6     | Add  $v$  to  $F'$ ;
7   end
8 end
9 return  $G$ ;
```

# Step #2: RC Update

- The Elmore delay model
- Phase 1:  $load_u = \sum_{v \text{ is child of } u} cap_v$ 
  - For example,  $load_A = cap_A + cap_B + cap_C + cap_D = cap_A + load_B + load_D$
- Phase 2:  $delay_u = \sum_{v \text{ is any node}} cap_v \times R_{Z \rightarrow LCA(u,v)}$ 
  - For example,  $delay_B = cap_A R_{Z \rightarrow A} + cap_D R_{Z \rightarrow A} + cap_B R_{Z \rightarrow B} + cap_C R_{Z \rightarrow B} = delay_A + R_{A \rightarrow B} load_B$



# Step #2: RC Update Upward Phase

- Store the parent index of each node on GPU
- Perform dynamic programming on trees

DFS\_load(u):

  load[u] = cap[u]

  For child v of u:

    DFS\_load(v)

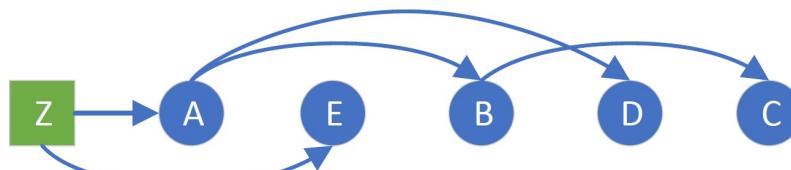
    load[u] += load[v]

GPU\_load:

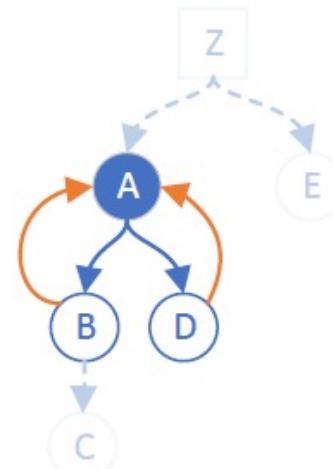
  For u in [C, D, B, E, A]:

    load[u] += cap[u]

    load[u.parent] += load[u]



Parent list representation in memory



(a) Upward

# Step #2: RC Update Downward Phase

- Store the parent index of each node on GPU
- Perform dynamic programming on trees

DFS\_delay(u):

For child v of u:

temp :=  $R[u,v] * \text{load}[v]$

delay[v] = delay[u] + temp

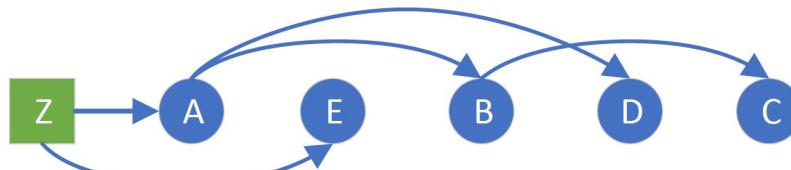
DFS\_delay(v)

GPU\_delay:

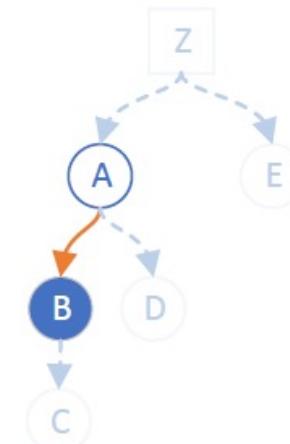
For u in [A, E, B, D, C]:

temp :=  $R[u.\text{parent},u] * \text{load}[u]$

delay[u]=delay[u.parent] + temp



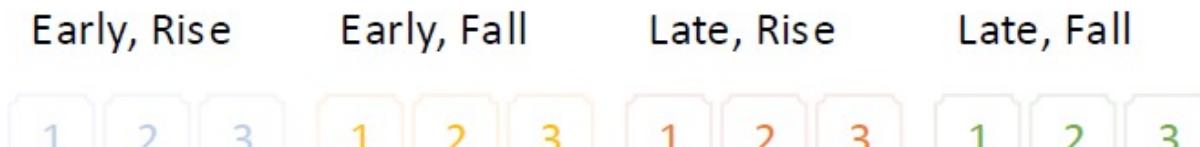
Parent list representation in memory



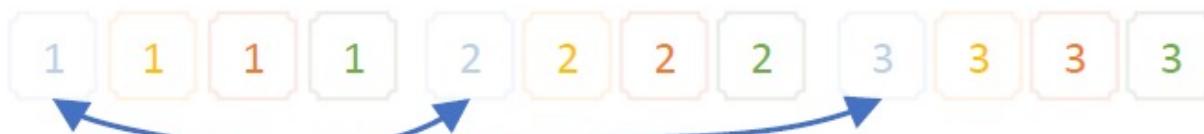
(b) Downward

# Step #2: RC Update Memory Coalesce

- Consecutive threads access consecutive memory
- RC update has four cases: {Rise, Fall} x {Early, Late}



(a)



Early, Rise

(b)

# Step #2: RC Update Kernels (flatten + update)

**Input:**  $N$  as #nets,  $(M, E)$  as (#nodes, #edges) in all nets  
**Input:**  $\text{roots}[0..N - 1]$ , the index of root of each net  
**Input:**  $\text{edges}[0..E - 1]$ , the undirected edges  $\{(a, b)\}$   
**Input:**  $\text{nodedstart}[0..N]$ , the offsets of each net in arrays of nodes, with  $\text{nodedstart}[N] = M$   
**Input:**  $\text{edgestart}[0..N]$ , the offsets of each net in arrays of edges, with  $\text{edgestart}[N] = E$   
**Input:**  $\text{distances}[0..M] = \infty$ ,  $\text{counts}[0..M] = 0$   
**Output:**  $\text{order}[0..M - 1]$ , nodes in BFS order for each net  
/\* Process one net w/ **blockDim.x** threads \*/
1  $\text{netID} = \text{blockIdx.x};$  ▷ **gridDim.x** = #nets  
2  $\text{threadID} = \text{threadIdx.x};$  ▷ **blockDim.x** = 64  
3  $\text{nst} = \text{nodedstart}[\text{netID}];$  ▷ node offset start  
4  $\text{nend} = \text{nodedstart}[\text{netID} + 1];$  ▷ node offset end  
5  $\text{est} = \text{edgestart}[\text{netID}];$  ▷ edge offset start  
6  $\text{eend} = \text{edgestart}[\text{netID} + 1];$  ▷ edge offset end  
7  $\text{distances}[\text{nst} + \text{roots}[\text{netID}]] = 0;$   
8 **for**  $d = 0, 1, 2, \dots, (\text{nend} - \text{nst})$  **do**  
9     **for**  $i = \text{est} + \text{threadID}$  **to**  $\text{eend}$  **step** **blockDim.x** **do**  
10          $(a, b) = \text{edgelist}[i];$   
11         **if**  $\text{distances}[a] == d$  **and**  $\text{distances}[b] > d + 1$  **then**  
12              $\text{distances}[b] = d + 1;$   
13              $\text{atomicAdd}(\text{counts}[d], 1);$   
14         **end**  
15         **else if**  $\text{distances}[b] == d$  **and**  $\text{distances}[a] > d + 1$  **then**  
16              $\text{distances}[a] = d + 1;$   
17              $\text{atomicAdd}(\text{counts}[d], 1);$   
18         **end**  
19     **end**  
20      $\text{__syncthreads();}$  ▷ Sync threads within a block  
21     **break** **when**  $\text{counts}[d] == 0;$   
22 **end**  
23  $\text{countingSort}(\text{distances}, \text{counts}, \text{order}, \text{threadID});$

Flatten the data to array for “**data parallelism**”

Each thread computes one RC tree

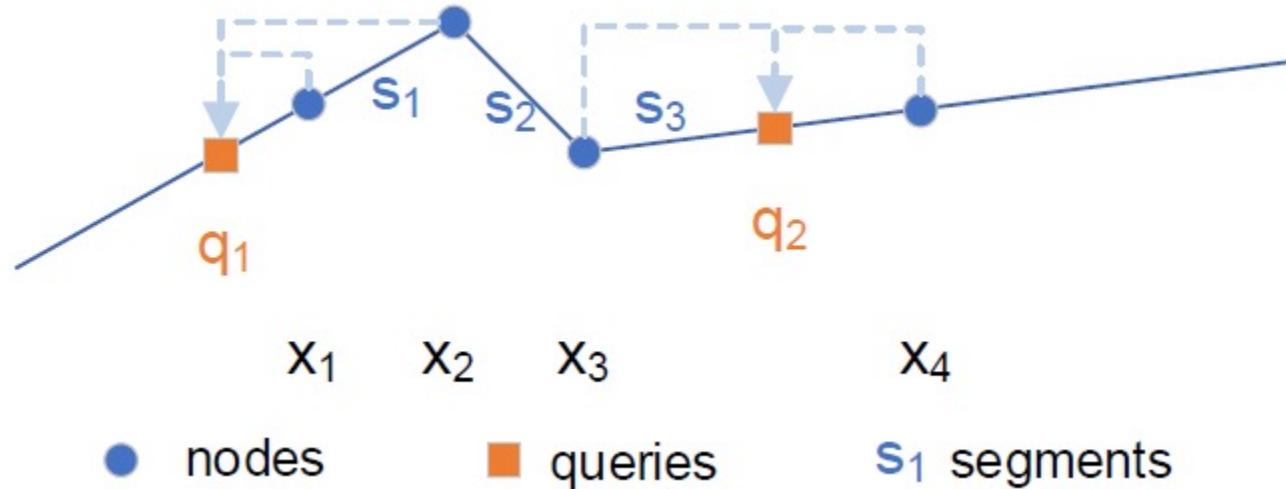
**Input:**  $N$  as #nets,  $M$  as #nodes in all nets  
**Input:**  $\text{start}[0..N]$ , the offsets of each net in arrays of nodes  
**Input:**  $\text{parent}[0..M - 1]$ , the index of parent of every nodes  
**Input:**  $\text{pres}[0..M - 1]$ , the resistance between nodes and their parent  
**Input:**  $\text{cap}[0..4M - 1]$ , the capacitance of nodes, each in 4 different combinations  
**Output:**  $\text{load}[0..4M - 1]$ ,  $\text{delay}[0..4M - 1]$ ,  $\text{impulse}[0..4M - 1]$ : arrays of results of load, delay and impulse, respectively

```

1  $\text{netID} = \text{blockIdx.x} \times \text{blockDim.x} + \text{threadIdx.x};$ 
2  $\text{condID} = \text{threadIdx.y};$ 
3 if  $\text{netID} \geq N$  then return;
4  $\text{offsetL} = \text{start}[\text{netID}];$  ▷ node offset start
5  $\text{offsetR} = \text{start}[\text{netID} + 1];$  ▷ node offset end
6 Initialize  $\text{load}$ ,  $\text{delay}$ ,  $\text{ldelay}$  to zero;
7 Initialize  $\beta = 0$  as an auxiliary array;
8 for  $i = \text{offsetR} - 1$  down to  $\text{offsetL}$  do
9    $\text{load}[4i + \text{condID}] += \text{cap}[4i + \text{condID}];$ 
10   $\text{load}[4\text{parent}[i] + \text{condID}] += \text{load}[4i + \text{condID}];$ 
11 end
12 for  $i = \text{offsetL} + 1$  to  $\text{offsetR} - 1$  do
13    $t = \text{load}[4i + \text{condID}] \times \text{pres}[i];$ 
14    $\text{delay}[4i + \text{condID}] = \text{delay}[4\text{parent}[i] + \text{condID}] + t;$ 
15 end
16 for  $i = \text{offsetR} - 1$  down to  $\text{offsetL}$  do
17    $\text{ldelay}[4i + \text{condID}] +=$ 
      $\text{cap}[4i + \text{condID}] \times \text{delay}[4i + \text{condID}];$ 
18    $\text{ldelay}[4\text{parent}[i] + \text{condID}] += \text{load}[4i + \text{condID}];$ 
19 end
20 for  $i = \text{offsetL} + 1$  to  $\text{offsetR} - 1$  do
21    $t' = \text{ldelay}[4i + \text{condID}] \times \text{pres}[i];$ 
22    $\beta[4i + \text{condID}] = \beta[4\text{parent}[i] + \text{condID}] + t';$ 
23    $\text{impulse}[4i + \text{condID}] =$ 
      $2\beta[4i + \text{condID}] - \text{delay}[4i + \text{condID}]^2;$ 
24 end
```

# Step #3: Cell Delay Update

- ❑ Perform linear inter- and extra-polation in batches
  - ❑ x-axis (input slew) and then y-axis (output load)



# Step #3: Cell Delay Update Kernels

```
/* Input: line (x1,y1)--(x2,y2) */  
/* Input: the x value queried */  
1 Function interpolate(x1,x2,y1,y2,x):  
2   if x1 = x2 then return y1;  
3   else return d1 + (d2 - d1) $\frac{x-x_1}{x_2-x_1}$ ;  
4 end  
/* Input: n × m look-up table */  
/* Input: the point queried (x,y) */  
5 Function lut_lookup(n,m,X,Y,mat,x,y):  
6   i' ← 0;  
7   i ← min(1, n - 1);  
8   while i + 1 < n and X[i] ≤ x do  
9     | i' ← i;  
10    | i ← i + 1;  
11   end  
12   j' ← 0;  
13   j ← min(1, m - 1);  
14   while j + 1 < m and Y[j] ≤ y do  
15     | j' ← j;  
16     | j ← j + 1;  
17   end  
18   ri' ← interpolate(Y[j'], Y[j], mat[i',j'], mat[i',j]);  
19   ri ← interpolate(Y[j'], Y[j], mat[i,j'], mat[i,j]);  
20   r ← interpolate(X[i'], X[i], ri', ri);  
21   return r;  
22 end
```

Each thread performs binary search on x-axis and y-axis to find the slew and capacitance

# Experimental Setting

---

## ❑ Machine configuration

- ❑ Nvidia CUDA, RTX 2080
- ❑ 40 Intel Xeon Gold 6138 CPU cores

## ❑ Execution parameters for GPU kernels

### ❑ RC Tree Flattening

- 64 threads per block with one block for each net

### ❑ Levelization

- 128 threads per block

### ❑ RC delay computation

- 4 threads for each net (one for each Early/Late and Rise/Fall condition) with a block of 64 nets

### ❑ Cell delay computation

- 4 threads for each arc, with a block of 32 arcs

# Overall Performance

- Comparison with OpenTimer of 40 CPUs
  - Run on large TAU15 Benchmarks (>20K gates)

Benchmark	# PIs	# POs	# Gates	# Nets	# Pins	# Nodes	# Edges	OpenTimer Runtime (40 CPUs)	Our Runtime (40 CPUs 1 GPU)	
									Runtime	Speed-up
aes_core	260	129	22938	23199	66751	413588	453508	156 ms	138 ms	1.13×
vga_lcd	85	99	139529	139635	397809	1966411	2185601	829 ms	311 ms	2.67×
vga_lcd_iccad	85	99	259067	259152	679258	3556285	3860916	1480 ms	496 ms	2.98×
b19	22	25	255278	255300	782914	4423074	4961058	1831 ms	585 ms	3.13×
cordic	34	64	45359	45393	127993	7464477	820763	274 ms	167 ms	1.64×
des_perf	234	140	138878	139112	371587	2128130	2314576	832 ms	325 ms	2.56×
edit_dist	2562	12	147650	150212	416609	2638639	2870985	1059 ms	376 ms	2.86×
fft	1026	1984	38158	39184	116139	646992	718566	241 ms	148 ms	1.63×
leon2	615	85	1616369	1616984	4328255	22600317	24639340	10200 ms	2762 ms	3.69×
leon3mp	254	79	1247725	1247979	3376832	17755954	19408705	7810 ms	2585 ms	3.02×
netcard	1836	10	1496719	1498555	3999174	21121256	23027533	9225 ms	2571 ms	3.60×
mgc_edit_dist	2562	12	161692	164254	450354	2436927	2674934	1021 ms	368 ms	2.77×
mgc_matrix_mult	3202	1600	171282	174484	492568	2713241	2994343	1138 ms	377 ms	3.02×
tip_master	778	857	37715	38493	95524	533690	570154	163 ms	143 ms	1.14×

# PIs: number of primary inputs # POs: number of primary outputs

# Pins: number of pins # Nodes: number of nodes in the STA graph

# Gates: number of gates # Nets: number of nets

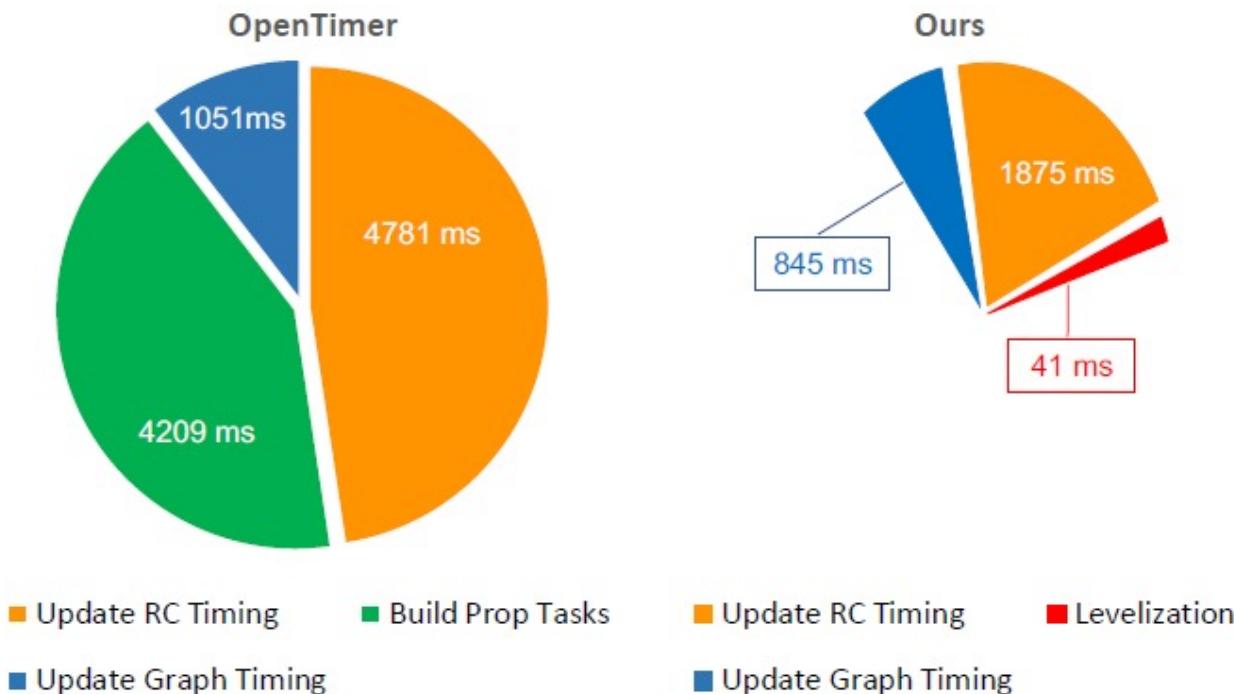
# Edges: number of edges in the STA graph

TAU15 Contest on Incremental Timing Analysis: <https://sites.google.com/site/taucontest2015>

T.-W. Huang et al., "OpenTimer v2: A New Parallel Incremental Timing Analysis Engine," IEEE TCAD21

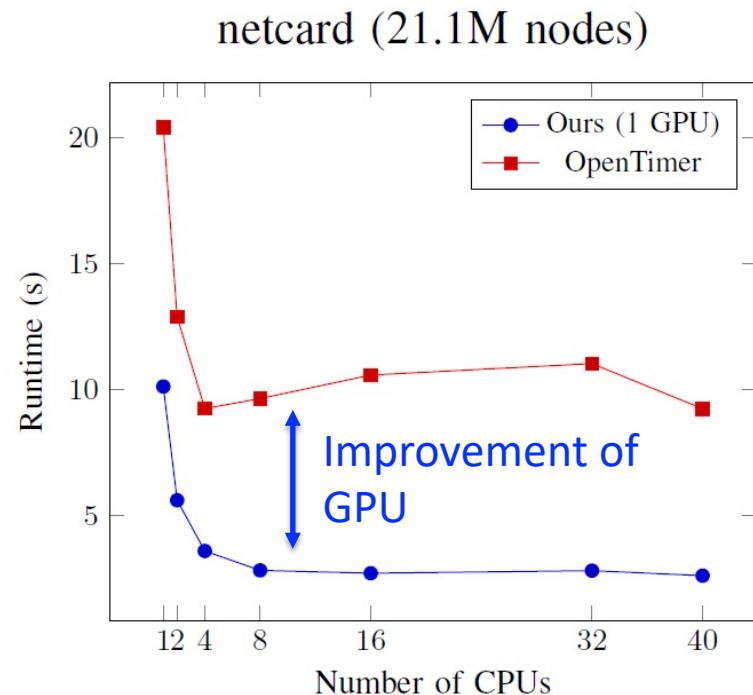
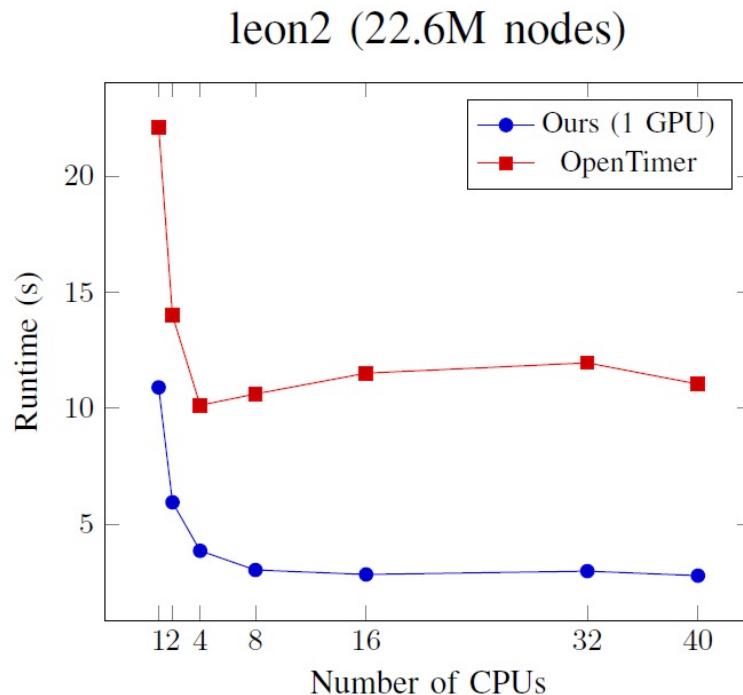
# Runtime Breakdown

## □ Circuit leon2 (21 M nodes)



# Runtime vs CPUs

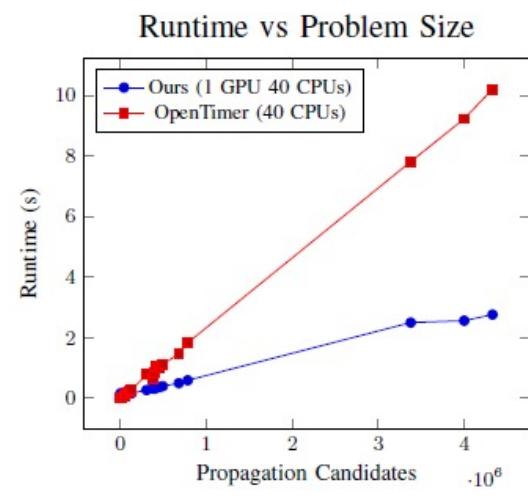
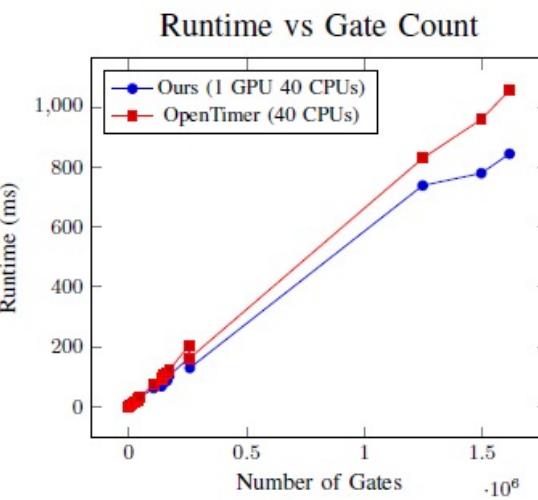
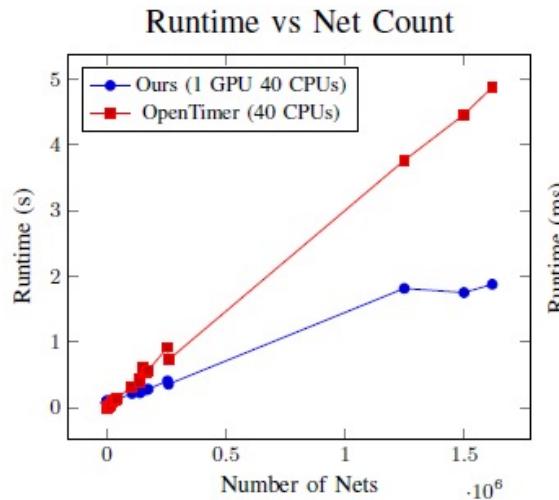
## □ Significant performance gap between CPU and GPU



Our runtime of 1 CPU and 1 GPU is very close to OpenTimer of 40 CPUs

# Runtime vs Problem Sizes

- Problem size matters for GPU acceleration
- When to enable GPU acceleration?
  - Net count > 20K
  - Gate count > 50K
  - Propagation candidate count > 15K



# Agenda

---

- Introduce the scalability problem of STA
  - What is STA and its challenges?
  - Why do we need new parallel paradigm for STA?
- Accelerate graph-based analysis using GPU
- Accelerate path-based analysis using GPU

G Guo, T-W Huang, Y Lin, and M Wong, “GPU-Accelerated Path-based Timing Analysis,”  
*IEEE/ACM Design Automation Conference (DAC)*, 2021

**Research Question:**

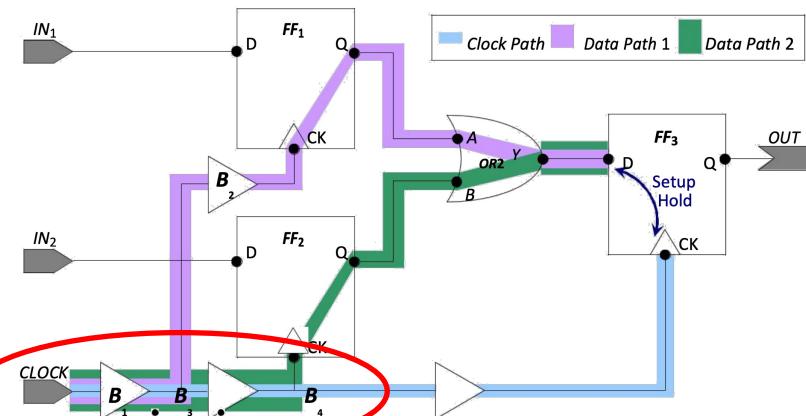
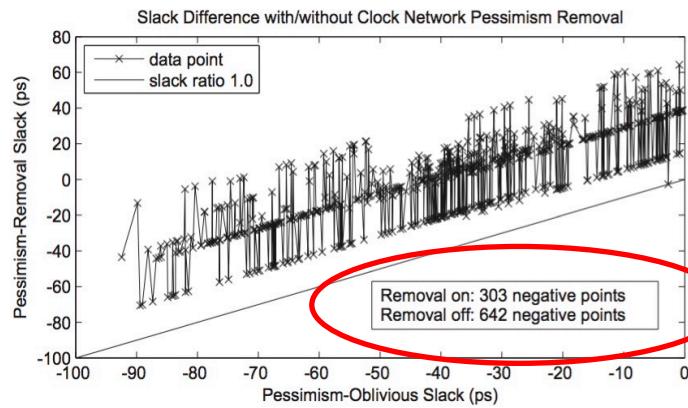
- *How can we use GPU to speed up path-based analysis (PBA)?*
- *How do we overcome the challenges of generating large numbers of critical paths to analyze with GPU?*



# Path-based Analysis (PBA)

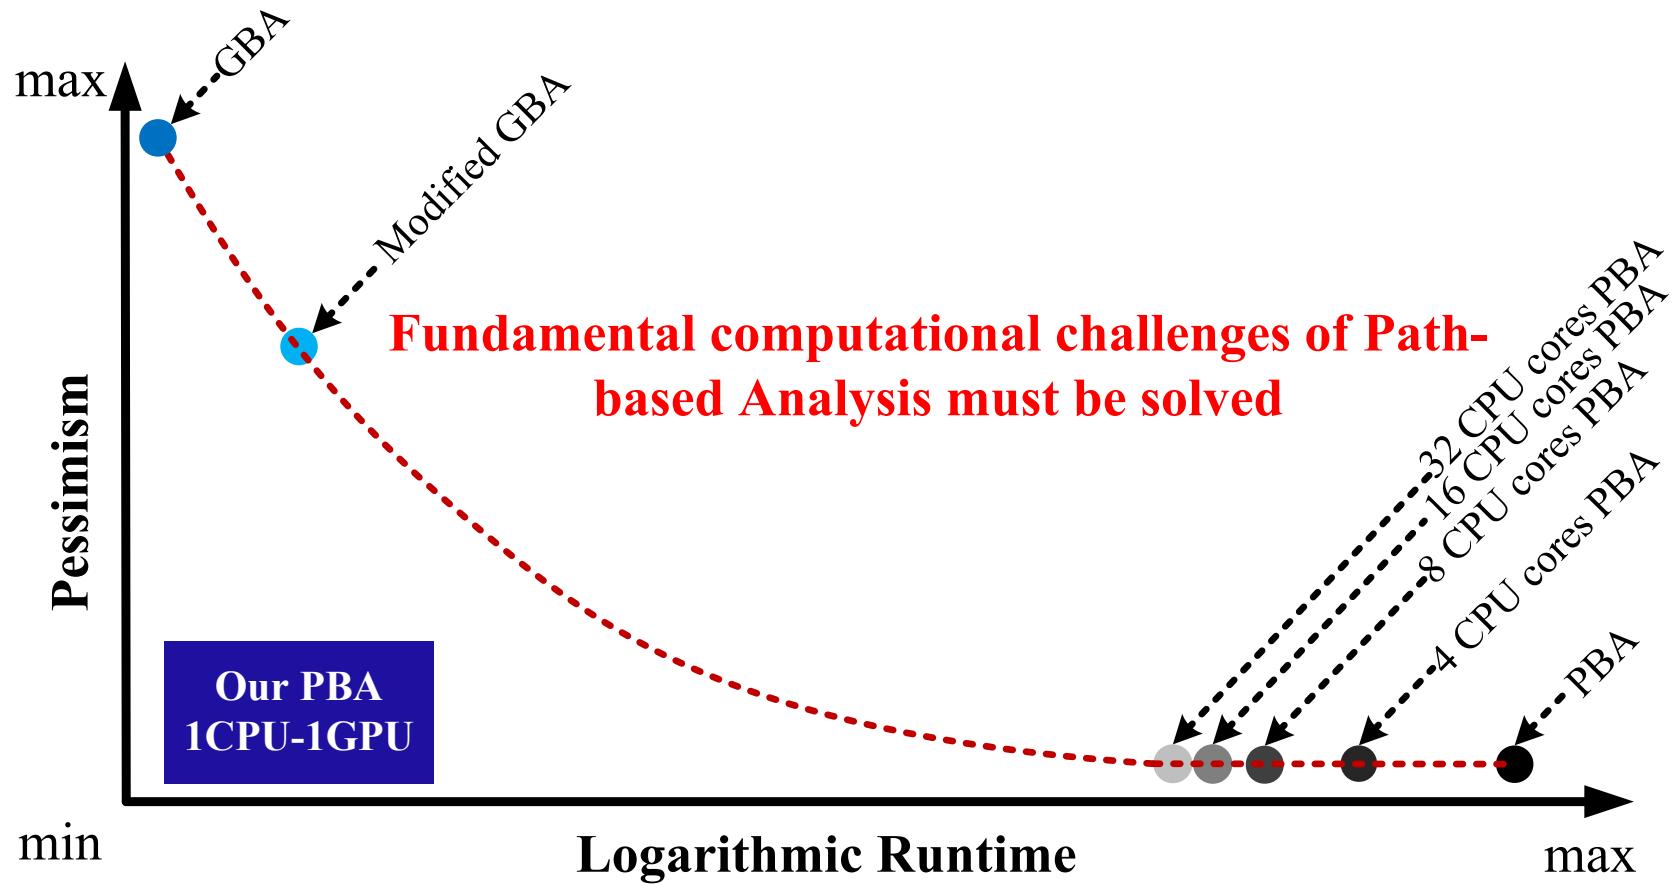
- Identify a set of critical paths from a updated graph
  - Exponential number of paths in the circuit graph
- Re-analyze each path with path-specific update
  - Re-propagate the slew and remove pessimism
  - Advanced on-chip variation (AOCV)
  - Common path pessimism removal (CPPR)
  - ...

*Paths marked failing at GBA may become passing after PBA!*



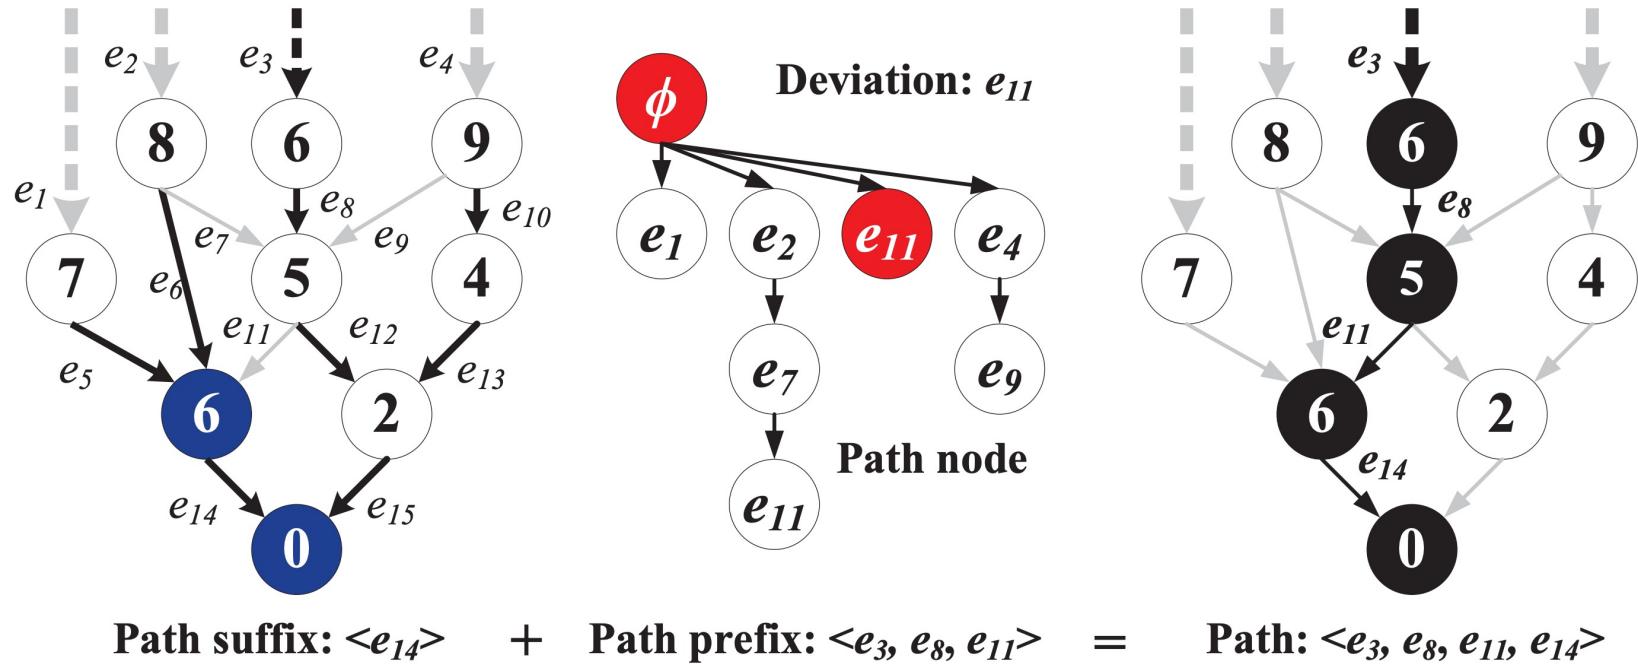
# PBA is Extremely Time-Consuming

- Speed vs Accuracy (pessimism removal) tradeoff

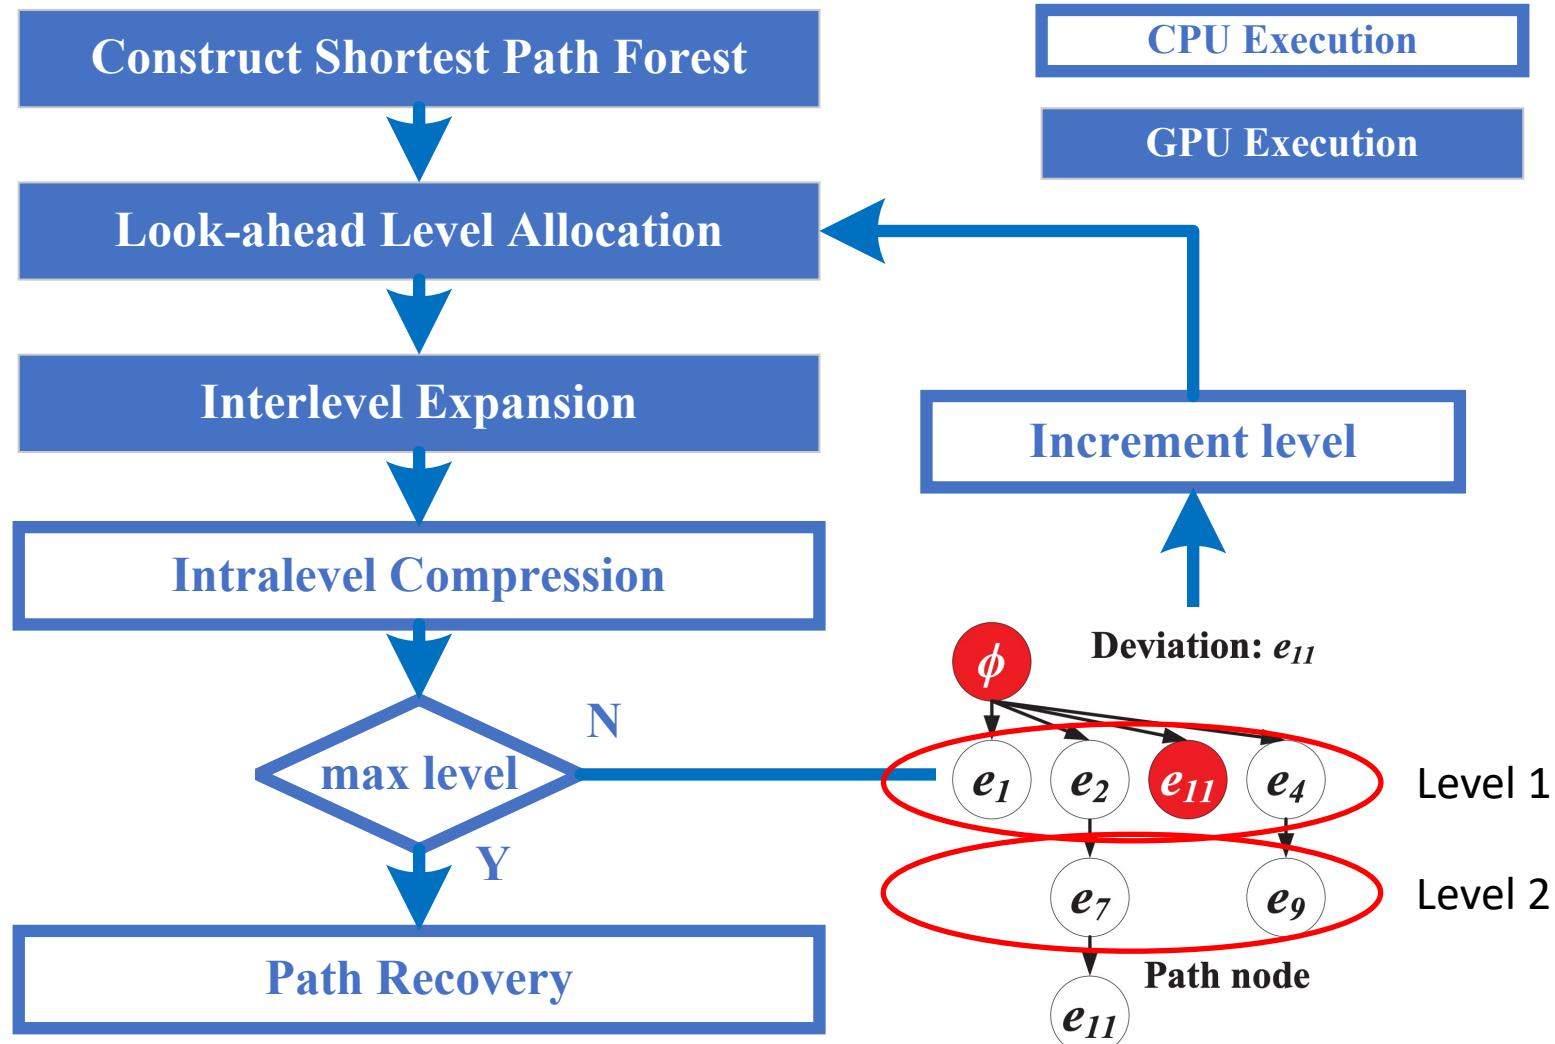


# A Key Step: Generate Critical Paths

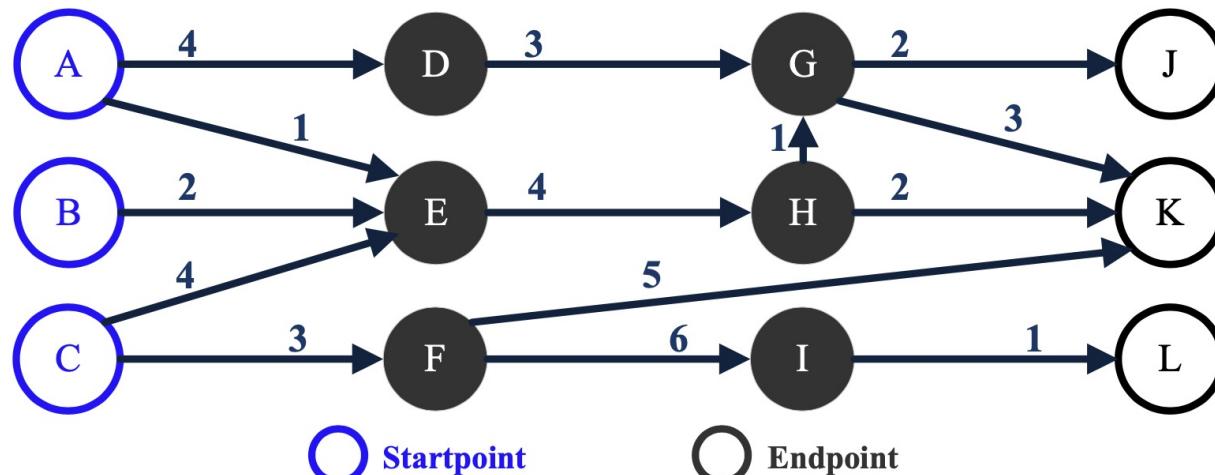
- We introduce implicit path representation
  - Each path is represented using O(1) space and time
  - Each path is ranked through a *prefix tree* & a *suffix tree*



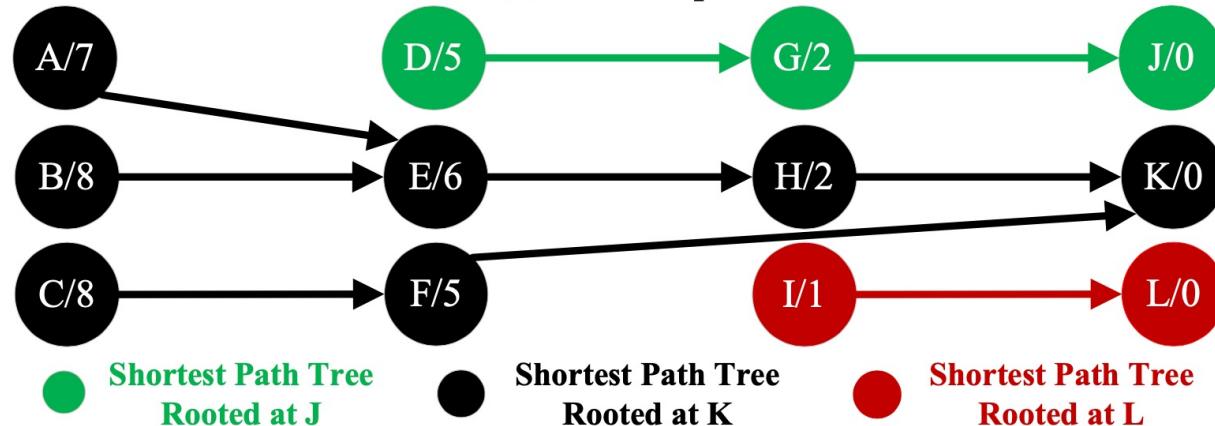
# GPU-Accelerated PBA Algorithm Flow



# Step #1: Generate Suffix Tree on GPU



(a) STA Graph.



(b) Shortest path forest.

# Step #1: Generate Suffix Tree Kernel

**Input** :  $G^-$  in CSR format,  $N$  as #vertices,  $M$  as #edges,  
vertices[N], edges[M], weights[M]

**Input** : Shortest distance cache, distanceCache[N]

**Input** : Array indicating vertices with updated distances,  
distanceUpdated[N]

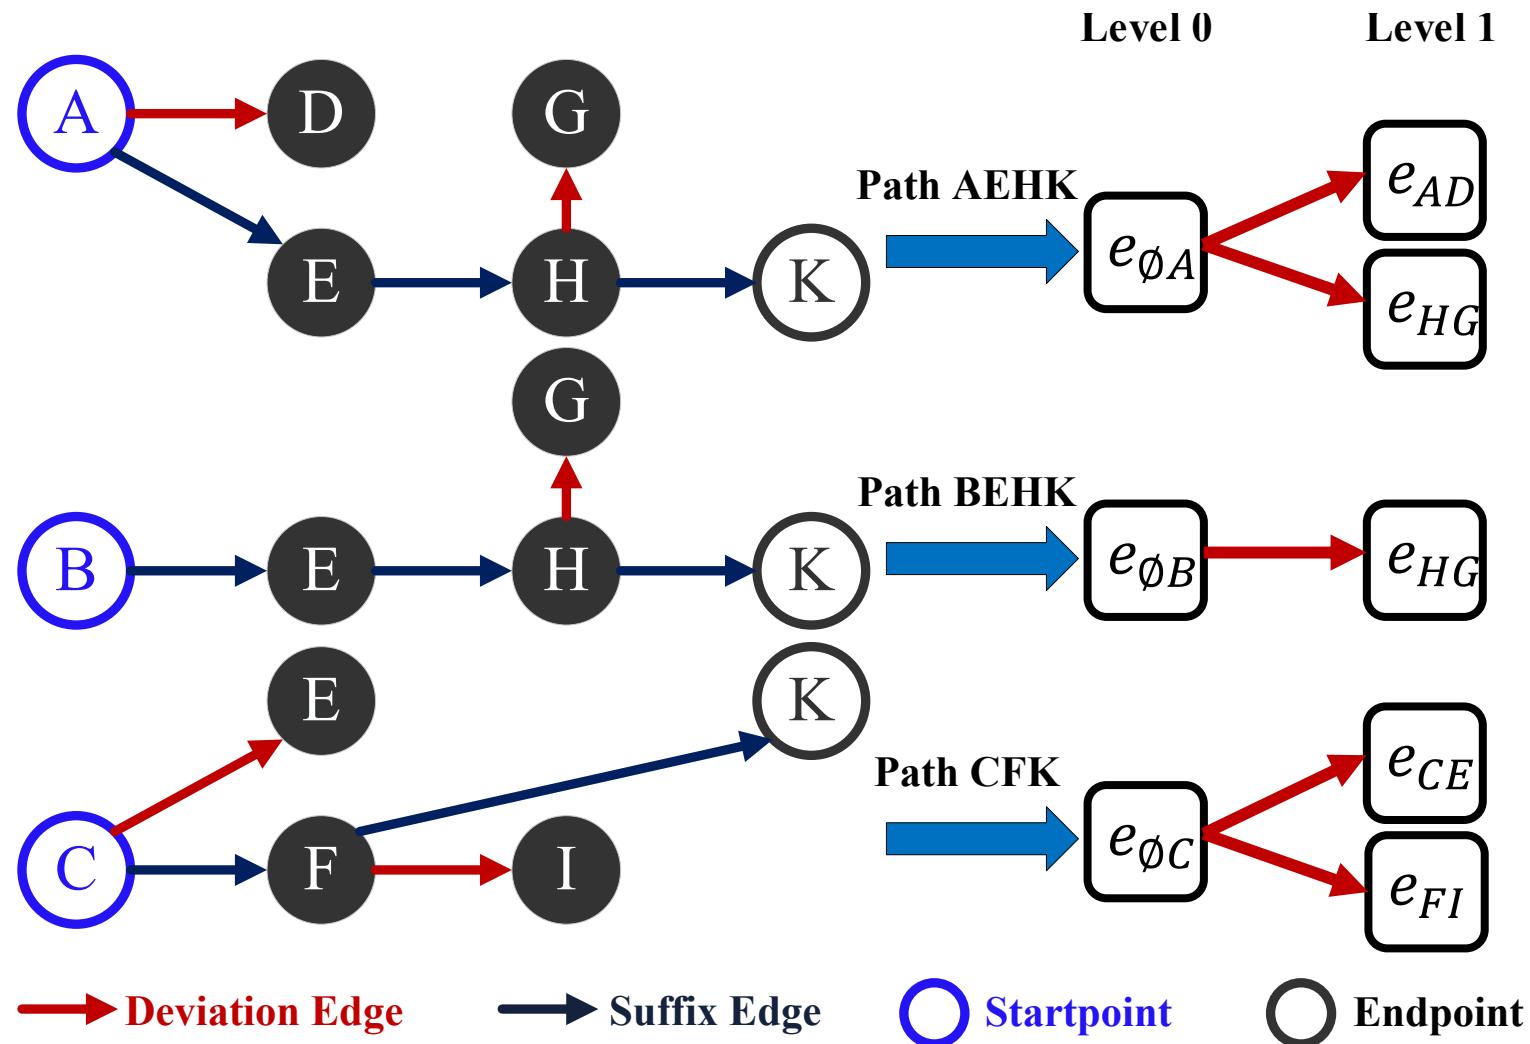
Flattened data  
structure to benefit  
data parallelism

**Result:** Shortest distances array, distances[N]

```
1 tid ← blockIdx.x * blockDim.x + threadIdx.x;
2 if tid ≥ N then
3   | return;
4 end
5 if distanceUpdated[tid] is false then
6   | return;
7 end
8 distanceUpdated[tid] ← false;
9 edgeStart ← vertices[tid];
10 edgeEnd ← (tid == N-1) ? M : vertices[tid+1];
11 for eid ← edgeStart to edgeEnd do
12   neighbor ← edges[eid];
13   weight ← weights[eid];
14   newDis ← distances[tid] + weight;
15   atomicMin (&distanceCache[neighbor], newDis);
16 end
17 return;
```

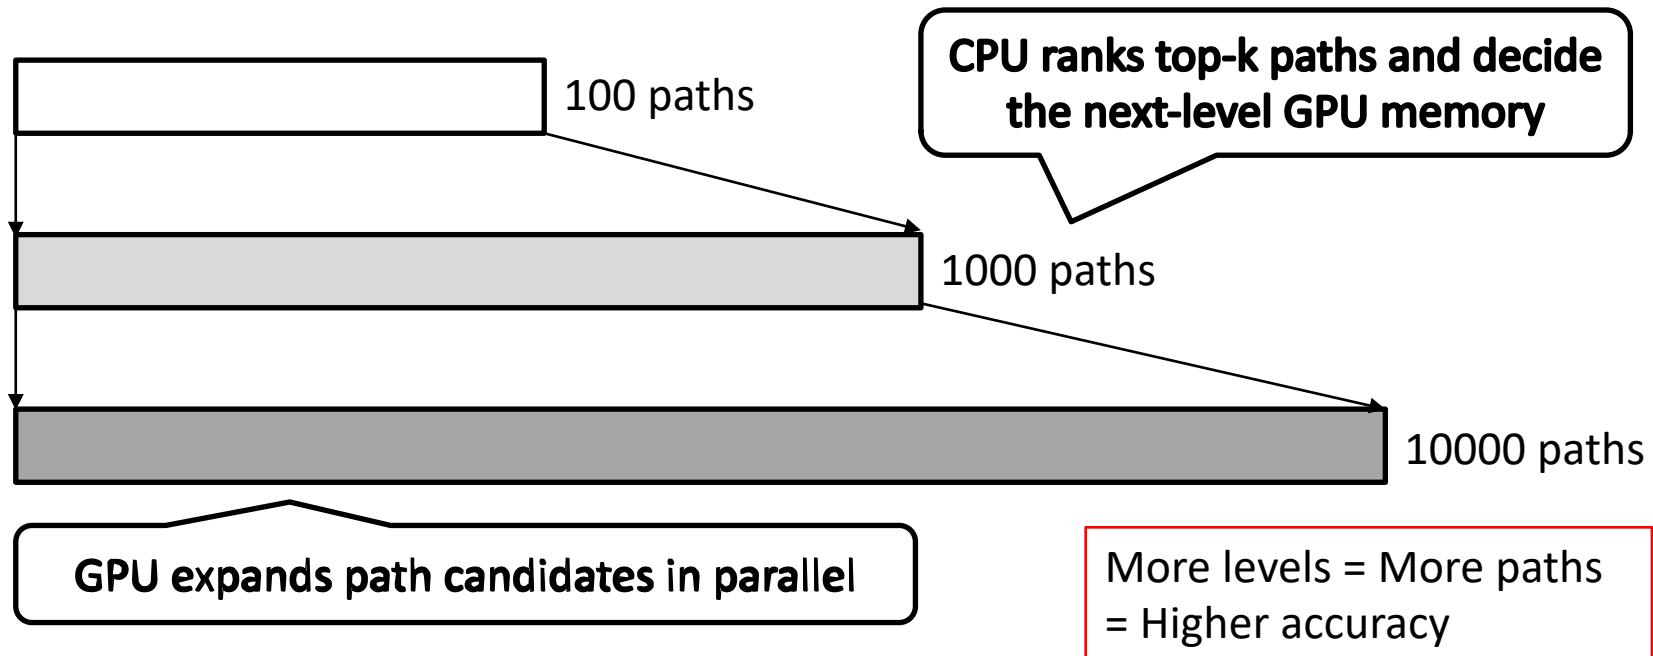
Bellman loop to find  
the shortest path tree  
(each thread per edge)

# Step #2: Expand Prefix Tree on GPU



# Step #2: Expand Prefix Tree on GPU (cont'd)

- Iteratively grow GPU memory at each expansion
  - Each iteration uses GPU to decide path candidates
  - Each iteration uses CPU to prune path candidates
  - Each path candidate takes  $O(1)$  space “deviation edge”



# Step #2: Expand Prefix Tree Kernel

**Input :**  $G^+$  in CSR format,  $N$  as #vertices,  $M$  as #edges,  $\text{vertices}[N]$ ,  $\text{edges}[M]$ ,  $\text{weights}[M]$

**Input :** Shortest path forest,  $\text{forest}[N]$  as edge array,  $\text{distances}[N]$  as distance array

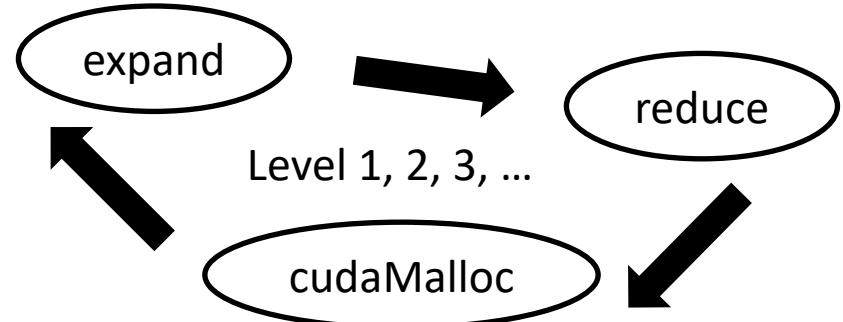
**Input :**  $\text{currLevel}$  as current level

**Input :**  $\text{levelSize}$  as the number of entries in current level

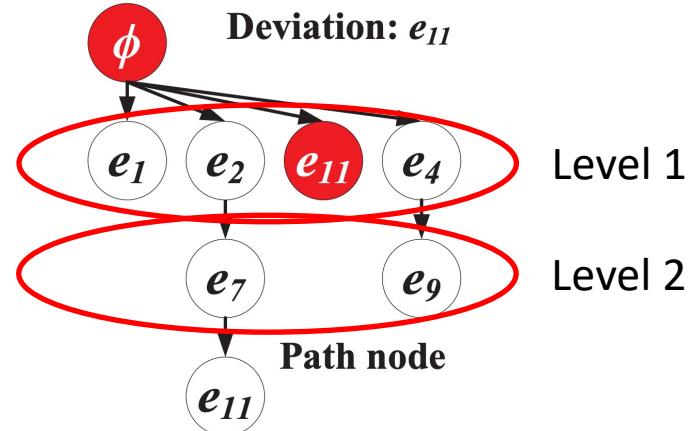
**Result:** Explore critical path candidates in next level

```

1 tid ← blockIdx.x * blockDim.x + threadIdx.x;
2 if tid ≥ levelSize then
3   | return;
4 end
5 offset ← (tid == 0) ? 0 : currLevel[tid-1].childOffset;
6 level ← currLevel[tid].level;
7 slack ← currLevel[tid].slack;
8 v ← currLevel[tid].to;
9 while v is not endpoint do
10   edgeStart ← vertices[v];
11   edgeEnd ← (v == N-1) ? M : vertices[v+1];
12   for eid ← edgeStart to edgeEnd do
13     neighbor ← edges[eid];
14     weight ← weights[eid];
15     if eid is deviation edge then
16       /* Fill out child path data */
17       newPath ← nextLevel[offset];
18       newPath.level ← level+1;
19       newPath.from ← v;
20       newPath.to ← neighbor;
21       newPath.parent ← tid;
22       newPath.childOffset ← 0;
23       newPath.slack ← slack + distances[neighbor] +
24         weight - distances[v];
25       offset ← offset + 1;
26     end
27   end
28   /* Traverse along the shortest path
29   forest
30   v = forest[v];
31 end
32 return;
```



**Key idea: leveled prefix tree**



# Experimental Setting

---

- ❑ Machine configuration
  - ❑ Nvidia CUDA, RTX 2080
  - ❑ 40 Intel Xeon Gold 6138 CPU cores
- ❑ Measure the accuracy-runtime tradeoff
  - ❑ “MDL” stands for maximum deviation level
- ❑ Execution parameters for GPU kernels
  - ❑ Suffix tree kernel
    - 1024 threads per block
  - ❑ Prefix tree kernel
    - 1024 threads per block

# Overall Performance

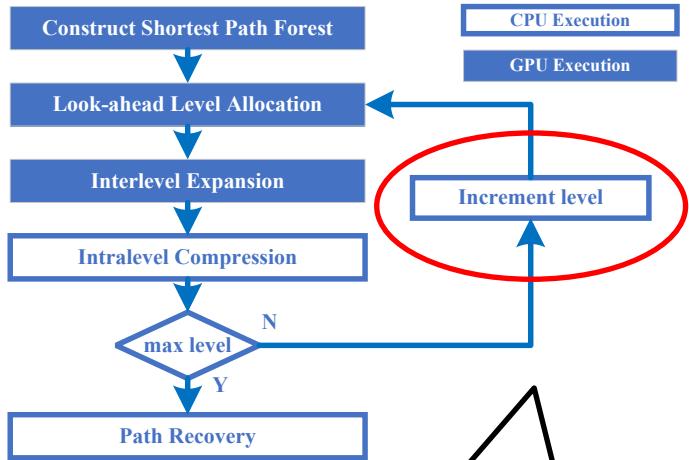
- Compare with OpenTimer's CPU-based PBA
  - Report speed-up at different MDLs

Benchmark	#Pins	#Gates	#Arcs	OpenTimer Runtime	Our Algorithm #MDL=10		Our Algorithm #MDL=15		Our Algorithm #MDL=20	
					Runtime	Speed-up	Runtime	Speed-up	Runtime	Speed-up
leon2	4328255	1616399	7984262	2875783	4708.36	611×	5295.49ms	543×	5413.84	531×
leon3mp	3376821	1247725	6277562	1217886	5520.85	221×	7091.79ms	172×	8182.84	149×
netcard	3999174	1496719	7404006	752188	2050.60	367×	2475.90ms	304×	2484.08	303×
vga_lcd	397809	139529	756631	53204	682.94	77.9×	683.04ms	77.9×	706.16	75.3×
vga_lcd_iccad	679258	259067	1243041	66582	720.40	92.4×	754.35ms	88.3×	766.29	86.9×
b19_iccad	782914	255278	1576198	402645	2144.67	188×	2948.94ms	137×	3483.05	116×
des_perf_ispd	371587	138878	697145	24120	763.79	31.6×	766.31ms	31.5×	780.56	30.9×
edit_dist_ispd	416609	147650	799167	614043	1818.49	338×	2475.12ms	248×	2900.14	212×
mgc_edit_dist	450354	161692	852615	694014	1463.61	474×	1485.65ms	467×	1493.90	465×
mgc_matric_mult	492568	171282	948154	214980	994.67	216×	1075.90ms	200×	1113.26	193×

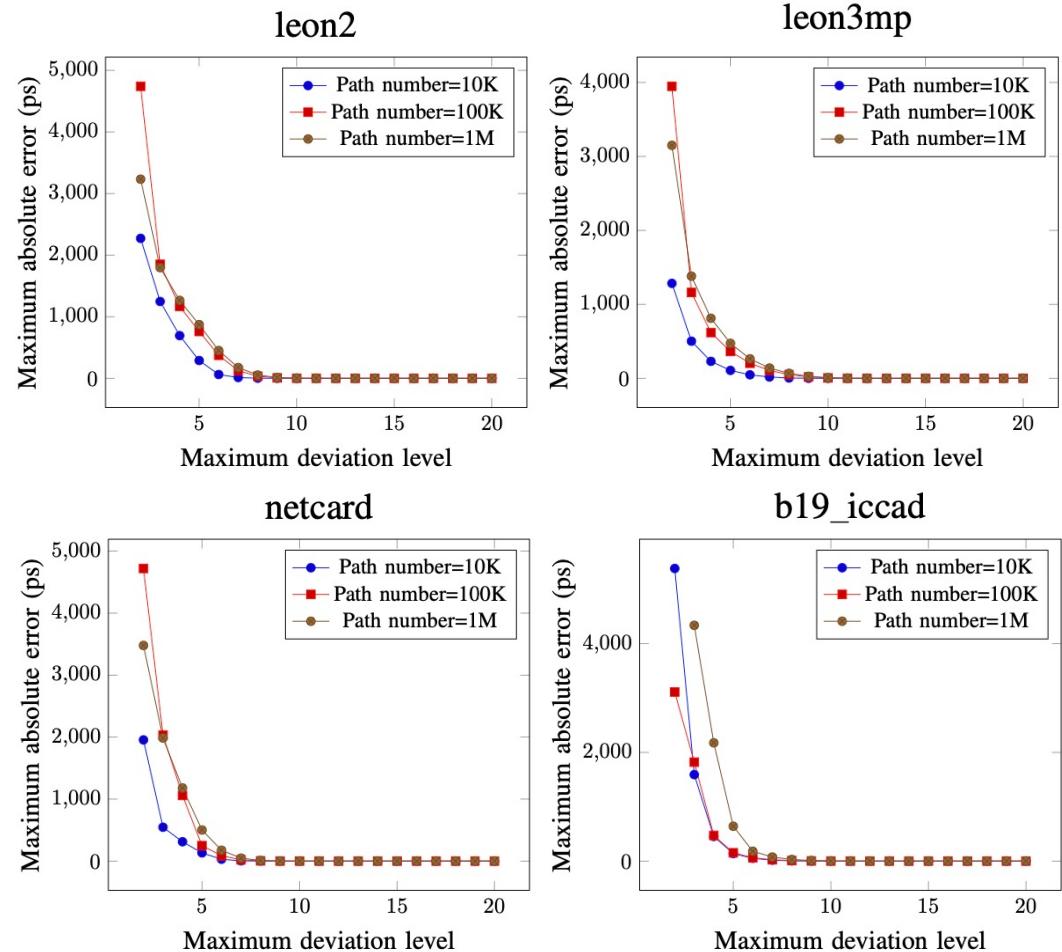
- Achieve significant speed-up at large designs
  - 611x speed-up in leon2 (1.3M gates)
  - 221x speed-up in leon3mp (1.2M gates)

# Path Accuracy vs MDL

## ☐ Achieve decent accuracy at 10–12 GPU iterations



**More GPU expansions  
(iterations) lead to higher  
numbers of paths and thus  
better accuracy**

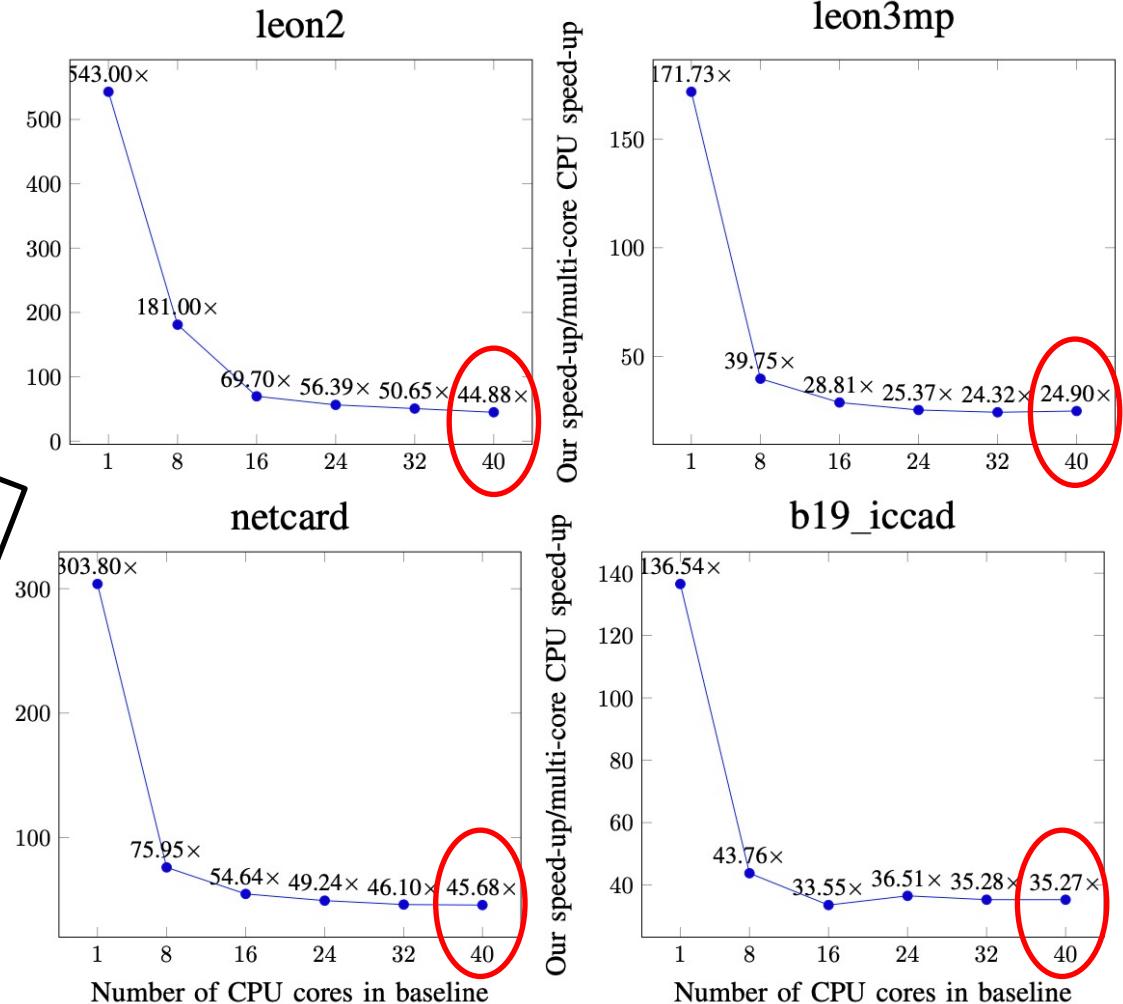


# GPU Speed-up vs CPUs

- ❑ one GPU is even faster than OpenTimer with 40 CPUs

- ❑ 44x on leon2
- ❑ 25x on leon3mp
- ❑ 46x on netcard
- ❑ 35x on b19

In fact, according to our experiments, our GPU-accelerated PBA is always faster than OpenTimer's CPU baseline regardless of the core count



# Conclusion

---

- ❑ **Introduced the runtime challenges of STA**
  - ❑ Introduced graph-based analysis
  - ❑ Introduced path-based analysis
- ❑ **Accelerated the graph-based analysis using GPU**
  - ❑ Achieved 4x speed-up over 40 CPUs on large designs
- ❑ **Accelerated the path-based analysis using GPU**
  - ❑ Achieved 600x speed-up over 40 CPUs on large designs
- ❑ **Future work**
  - ❑ Consider other important STA tasks (CPPR, CCS, etc.)
  - ❑ Leverage cudaMallocAsync to boost memory efficiency
  - ❑ Leverage cudaGraph to reduce kernel launch overheads
  - ❑ Collaborate with experts in Nvidia!

# Thank You

**Use right algorithms for GPU!**

[tsung-wei.huang@utah.edu](mailto:tsung-wei.huang@utah.edu)

